



Step-by-Step Walkthroughs

# Database Migration Guide



# Database Migration Guide: Step-by-Step Walkthroughs

Copyright © 2026 Amazon Web Services, Inc. and/or its affiliates. All rights reserved.

Amazon's trademarks and trade dress may not be used in connection with any product or service that is not Amazon's, in any manner that is likely to cause confusion among customers, or in any manner that disparages or discredits Amazon. All other trademarks not owned by Amazon are the property of their respective owners, who may or may not be affiliated with, connected to, or sponsored by Amazon.

# Table of Contents

|                                                                                  |          |
|----------------------------------------------------------------------------------|----------|
| <b>Database Migration Step-by-Step Walkthroughs .....</b>                        | <b>1</b> |
| Oracle Database .....                                                            | 2        |
| Microsoft SQL Server .....                                                       | 3        |
| MySQL .....                                                                      | 3        |
| BigQuery .....                                                                   | 3        |
| MariaDB .....                                                                    | 4        |
| MongoDB .....                                                                    | 4        |
| PostgreSQL .....                                                                 | 4        |
| SAP ASE .....                                                                    | 4        |
| <b>Migrating Databases to Amazon Web Services Managed Databases .....</b>        | <b>5</b> |
| Migrating a MySQL Database to RDS for MySQL or Aurora MySQL .....                | 5        |
| Full load MySQL database migration .....                                         | 6        |
| Full load MySQL database migration options performance comparison .....          | 12       |
| Migrate MySQL database with AWS DMS ongoing replication .....                    | 12       |
| Migrating PostgreSQL Databases to Amazon RDS for PostgreSQL or Amazon Aurora     |          |
| PostgreSQL .....                                                                 | 13       |
| Summary .....                                                                    | 15       |
| Full load PostgreSQL database migration .....                                    | 15       |
| Full load PostgreSQL database migration options performance comparison .....     | 27       |
| Migrate PostgreSQL database with AWS DMS ongoing replication .....               | 29       |
| Migrating PostgreSQL databases to Amazon RDS for PostgreSQL with DMS homogeneous |          |
| data migrations .....                                                            | 29       |
| Prerequisites for migrating PostgreSQL databases .....                           | 31       |
| PostgreSQL to Amazon RDS migration overview .....                                | 31       |
| Step-by-step PostgreSQL database to Amazon RDS migration walkthrough .....       | 32       |
| PostgreSQL database to Amazon RDS post-migration clean-up .....                  | 46       |
| Migrating an Oracle Database to Amazon RDS for Oracle .....                      | 47       |
| Summary .....                                                                    | 15       |
| Full load Oracle database migration .....                                        | 49       |
| Full load Oracle database migration options performance comparison .....         | 61       |
| Migrate Oracle database with AWS DMS ongoing replication .....                   | 62       |
| Migrating a SQL Server Always On Database to Amazon Web Services .....           | 63       |
| Prerequisites for migrating SQL Server AlwaysOn databases to AWS .....           | 63       |
| SQL Server Always On Availability Groups .....                                   | 64       |

|                                                                                                      |     |
|------------------------------------------------------------------------------------------------------|-----|
| Migrating an Amazon RDS for MySQL Database to an Amazon DynamoDB target .....                        | 67  |
| Why use AWS DMS? .....                                                                               | 68  |
| Example data set .....                                                                               | 68  |
| Solution overview .....                                                                              | 70  |
| Prerequisites .....                                                                                  | 71  |
| Step-by-step Amazon RDS for MySQL database to Amazon DynamoDB migration walkthrough .....            | 72  |
| Migrating an RDS for MySQL database to an S3 data lake .....                                         | 94  |
| Solution overview .....                                                                              | 95  |
| Use case .....                                                                                       | 96  |
| Limitations .....                                                                                    | 97  |
| Choosing an instance class and storage size .....                                                    | 98  |
| Step-By-Step Migration .....                                                                         | 99  |
| Migrating an RDS PostgreSQL database to an S3 data lake .....                                        | 117 |
| Why AWS DMS? .....                                                                                   | 118 |
| Use case .....                                                                                       | 118 |
| Example data set .....                                                                               | 118 |
| Solution overview .....                                                                              | 119 |
| Prerequisites .....                                                                                  | 120 |
| Step-by-step an Amazon RDS PostgreSQL database to an Amazon S3 data lake migration walkthrough ..... | 121 |
| Migrating SQL Server Databases to Amazon RDS for SQL Server .....                                    | 138 |
| Summary .....                                                                                        | 15  |
| Full load SQL Server database migration .....                                                        | 140 |
| Full load SQL Server database migration options performance comparison .....                         | 149 |
| Migrate SQL Server database with AWS DMS ongoing replication .....                                   | 151 |
| Migrating from Amazon RDS for Oracle to Amazon RDS for PostgreSQL and Aurora PostgreSQL .....        | 151 |
| Can My Oracle Database Migrate? .....                                                                | 152 |
| Migration Strategies .....                                                                           | 153 |
| The 12 Step Migration Process .....                                                                  | 154 |
| Automation .....                                                                                     | 157 |
| Oracle application future state architecture design .....                                            | 157 |
| Oracle database schema conversion .....                                                              | 159 |
| Oracle application conversion or remediation .....                                                   | 161 |
| Database migration script/ETL/report conversion .....                                                | 162 |



|                                                                                                                          |            |
|--------------------------------------------------------------------------------------------------------------------------|------------|
| Oracle application migration and integration with third-party applications .....                                         | 164        |
| Amazon RDS for Oracle data migration mechanism .....                                                                     | 164        |
| Oracle database migration testing and bug fixing .....                                                                   | 166        |
| Oracle database migration performance tuning .....                                                                       | 168        |
| Oracle database migration to PostgreSQL setup, DevOps, integration, deployment, and security .....                       | 168        |
| Oracle database migration to PostgreSQL documentation and knowledge transfer .....                                       | 172        |
| Oracle database migration to PostgreSQL project management and version control .....                                     | 172        |
| Oracle database migration to PostgreSQL post-production support .....                                                    | 172        |
| Oracle and PostgreSQL platform differences .....                                                                         | 173        |
| Migrating from SAP ASE to Amazon Aurora MySQL .....                                                                      | 174        |
| Prerequisites for migrating from SAP AWS to Amazon Aurora MySQL .....                                                    | 176        |
| Preparation and assessment for migrating from SAP ASE to Amazon Aurora MySQL .....                                       | 176        |
| SAP ASE to Amazon Aurora MySQL database code conversion and data loading .....                                           | 179        |
| Best practices for migrating from SAP ASE to Amazon Aurora MySQL .....                                                   | 184        |
| <b>Migrating Databases to the Amazon Web Services Cloud Using the Database Migration Service .....</b>                   | <b>186</b> |
| Migrating an On-Premises Oracle Database to Amazon Aurora MySQL .....                                                    | 187        |
| Costs .....                                                                                                              | 189        |
| Migration from Oracle to Aurora MySQL using AWS DMS high-level outline .....                                             | 189        |
| Step-by-step Oracle to Aurora MySQL using AWS DMS migration walkthrough .....                                            | 194        |
| Migrating an Amazon RDS for Oracle Database to Amazon Aurora MySQL .....                                                 | 215        |
| Costs .....                                                                                                              | 217        |
| Prerequisites for migrating from Amazon RDS for Oracle to Amazon Aurora MySQL .....                                      | 218        |
| Migration architecture for migrating from Amazon RDS for Oracle database to Amazon Aurora MySQL-Compatible Edition ..... | 219        |
| Step-by-step Amazon Relational Database Service to Amazon Aurora MySQL-Compatible Edition migration walkthrough .....    | 220        |
| AWS DMS migration from Amazon RDS for Oracle next steps .....                                                            | 256        |
| Migrating a SQL Server Database to Amazon Aurora MySQL .....                                                             | 256        |
| Prerequisites for Migrating from a SQL Server database to Amazon Aurora MySQL .....                                      | 257        |
| Step-by-step SQL Server database to Amazon Aurora MySQL migration walkthrough .....                                      | 259        |
| SQL Server database migration to Amazon Aurora MySQL troubleshooting .....                                               | 281        |
| Migrating a SQL Server AlwaysOn Database on Primary Replica to Amazon Aurora PostgreSQL .....                            | 282        |
| Why Amazon Aurora PostgreSQL? .....                                                                                      | 282        |

|                                                                                                                       |     |
|-----------------------------------------------------------------------------------------------------------------------|-----|
| Common database migration challenges .....                                                                            | 282 |
| Why AWS DMS? .....                                                                                                    | 283 |
| Migration overview .....                                                                                              | 283 |
| Prerequisites for migrating SQL Server AlwaysOn databases on primary replica to Amazon Aurora PostgreSQL .....        | 285 |
| Step-by-step SQL Server AlwaysOn databases on primary replica to Amazon Aurora PostgreSQL migration walkthrough ..... | 286 |
| Migrating an Amazon RDS for Oracle Database to an Amazon S3 Data Lake .....                                           | 305 |
| Why use AWS DMS? .....                                                                                                | 306 |
| Example data set .....                                                                                                | 306 |
| Solution overview .....                                                                                               | 308 |
| Prerequisites for migrating an RDS for Oracle database to an Amazon S3 data lake .....                                | 309 |
| Step-by-step Amazon RDS for Oracle database to Amazon S3 data lake migration walkthrough .....                        | 310 |
| Step-by-step Amazon RDS for Oracle database to Amazon S3 data lake migration conclusion .....                         | 329 |
| Migrating an Amazon RDS for SQL Server Database to an Amazon S3 Data Lake .....                                       | 329 |
| Why Amazon S3? .....                                                                                                  | 330 |
| Why AWS DMS? .....                                                                                                    | 330 |
| Solution overview .....                                                                                               | 331 |
| Prerequisites for migrating from an Amazon RDS for SQL Server database to an Amazon S3 data lake .....                | 333 |
| Step-by-step Amazon RDS for SQL Server database to an Amazon S3 data lake migration walkthrough .....                 | 334 |
| Migrating an Oracle Database to PostgreSQL .....                                                                      | 355 |
| Prerequisites for migrating from an Oracle database to PostgreSQL .....                                               | 356 |
| Step-by-step Oracle database to PostgreSQL migration walkthrough .....                                                | 357 |
| Rolling Back the Migration .....                                                                                      | 380 |
| Oracle database migration to PostgreSQL troubleshooting .....                                                         | 381 |
| Migrating Oracle databases to Amazon Aurora MySQL with DMS Schema Conversion .....                                    | 381 |
| Migration overview .....                                                                                              | 382 |
| Prerequisites for Migrating Oracle databases to Amazon RDS for MySQL with DMS schema conversion .....                 | 383 |
| Step-by-step Oracle database to Amazon RDS for MySQL with DMS schema conversion migration walkthrough .....           | 384 |

|                                                                                                                  |     |
|------------------------------------------------------------------------------------------------------------------|-----|
| Migration from Oracle database to Amazon RDS for MySQL with DMS schema conversion next steps .....               | 399 |
| Migrating Oracle databases to Amazon RDS for PostgreSQL with DMS Schema Conversion ...                           | 400 |
| Migration overview .....                                                                                         | 401 |
| Prerequisites for migrating Oracle databases to Amazon Aurora PostgreSQL with DMS schema conversion .....        | 402 |
| Step-by-step Oracle databases to Amazon Aurora PostgreSQL with DMS schema conversion migration walkthrough ..... | 402 |
| Migration from Oracle databases to Amazon Aurora PostgreSQL with DMS schema conversion next steps .....          | 416 |
| Migrating SQL Server databases to Amazon Aurora PostgreSQL with DMS Schema Conversion .....                      | 417 |
| Migration Overview .....                                                                                         | 418 |
| Prerequisites for migrating SQL Server databases to Aurora PostgreSQL with DMS schema conversion .....           | 419 |
| Step-by-step SQL Server database to Aurora PostgreSQL with DMS schema conversion migration walkthrough .....     | 420 |
| Migration from SQL Server databases to Aurora PostgreSQL with DMS schema conversion next steps .....             | 434 |
| Migrating SQL Server databases to Amazon RDS for MySQL with DMS Schema Conversion ...                            | 435 |
| Migration overview .....                                                                                         | 436 |
| Prerequisites for migrating SQL Server databases to Amazon RDS for MySQL with DMS schema conversion .....        | 437 |
| Step-by-step SQL Server databases to Amazon RDS for MySQL with DMS schema conversion .....                       | 438 |
| Migration from SQL Server databases to Amazon RDS for MySQL with DMS schema conversion next steps .....          | 453 |
| Migrating an Amazon RDS for Oracle Database to Amazon Redshift .....                                             | 454 |
| Prerequisites for migrating from Amazon RDS for Oracle to Amazon Redshift .....                                  | 455 |
| Migration architecture for migrating from Amazon RDS for Oracle to Amazon Redshift ....                          | 456 |
| Step-by-step Amazon RDS for Oracle to Amazon Redshift migration walkthrough .....                                | 457 |
| Migration from Amazon RDS for Oracle to Amazon Redshift next steps .....                                         | 488 |
| Migrating a BigQuery Project to Amazon Redshift .....                                                            | 489 |
| Migration overview .....                                                                                         | 490 |
| Prerequisites for migrating a BigQuery project to Amazon Redshift .....                                          | 494 |
| Step-by-Step BigQuery project to Amazon Redshift migration walkthrough .....                                     | 495 |

|                                                                                                   |     |
|---------------------------------------------------------------------------------------------------|-----|
| Migration from a BigQuery project to Amazon Redshift next steps .....                             | 504 |
| Migrating a MySQL-Compatible Database to Amazon Aurora MySQL .....                                | 505 |
| Migrating Data from an External MySQL Database to an Amazon Aurora MySQL Using<br>Amazon S3 ..... | 505 |
| Migrating MySQL to Amazon Aurora MySQL by Using mysqldump .....                                   | 519 |
| Migrating Data from an Amazon RDS MySQL DB Instance to an Amazon Aurora MySQL DB<br>Cluster ..... | 519 |
| Migrating a MariaDB Database to Amazon RDS for MySQL or Amazon Aurora MySQL .....                 | 529 |
| Set up MariaDB as a source database .....                                                         | 530 |
| Set up Aurora MySQL as a target database .....                                                    | 534 |
| Set up an AWS DMS replication instance .....                                                      | 536 |
| Test the endpoints for MariaDB database migration .....                                           | 537 |
| Create a migration task for a MariaDB database .....                                              | 538 |
| Validate the MariaDB database migration .....                                                     | 539 |
| Cut over for the migration from a MariaDB database .....                                          | 539 |
| Migrating from MongoDB to Amazon DocumentDB .....                                                 | 541 |
| Launch an Amazon EC2 instance for MongoDB migration .....                                         | 542 |
| Install and configure MongoDB community edition .....                                             | 543 |
| Create an AWS DMS replication instance for MongoDB migration .....                                | 545 |
| Create source and target endpoints for MongoDB migration .....                                    | 545 |
| Create and run a MongoDB migration task .....                                                     | 548 |

# Database Migration Step-by-Step Walkthroughs

You can use AWS Database Migration Service (AWS DMS) to migrate your data to and from most widely used commercial and open-source databases such as Oracle, PostgreSQL, Microsoft SQL Server, Amazon Redshift, Amazon Aurora, MariaDB, and MySQL. The service supports homogeneous migrations such as Oracle to Oracle, and also heterogeneous migrations between different database platforms, such as Oracle to MySQL or MySQL to Amazon Aurora MySQL-Compatible Edition. Alternatively, you can use AWS DMS to move from existing, self-managed, open-source, and commercial databases to fully managed AWS databases of the same engine.

You can use DMS Schema Conversion to migrate to a different database engine. This service automatically assesses and converts your source schemas to a new target engine. Alternatively, you can download the AWS Schema Conversion Tool AWS SCT to your local PC to convert your source schemas.

In this guide, you can find step-by-step walkthroughs that go through the process of schema conversion and data migration of the following source databases:

Amazon Web Services (AWS) has several services that allow you to run a MySQL-compatible database on AWS. Amazon Relational Database Service (Amazon RDS) supports MySQL-compatible databases including MySQL, MariaDB, and Amazon Aurora MySQL. Amazon Elastic Compute Cloud (Amazon EC2) provides platforms for running MySQL-compatible databases.

| Migrating From                          | Solution                                                                                                                                                                                                                                 |
|-----------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| An RDS for MySQL DB instance            | You can migrate data directly from an Amazon RDS for MySQL DB snapshot to an Amazon Aurora MySQL DB cluster. For details, see <a href="#">Migrating Data from an Amazon RDS MySQL DB Instance to an Amazon Aurora MySQL DB Cluster</a> . |
| A MySQL database external to Amazon RDS | If your database supports the InnoDB or MyISAM tablespaces, you have these options for migrating your data to an Amazon Aurora MySQL DB cluster:                                                                                         |

| Migrating From                          | Solution                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |
|-----------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                                         | <ul style="list-style-type: none"> <li>You can create a dump of your data using the <code>mysqldump</code> utility, and then import that data into an existing Amazon Aurora MySQL DB cluster.</li> <li>You can copy the source files from your database to an Amazon Simple Storage Service (Amazon S3) bucket, and then restore an Amazon Aurora MySQL DB cluster from those files. This option can be considerably faster than migrating data using <code>mysqldump</code>.</li> </ul> <p>For more information, see <a href="#">Migrating MySQL to Amazon Aurora MySQL by Using mysqldump</a>.</p> <p>However, for very large databases, you can significantly reduce the amount of time that it takes to migrate your data by copying the source files for your database and restoring those files to an Amazon Aurora MySQL DB instance as described in <a href="#">Migrating Data from an External MySQL Database to an Amazon Aurora MySQL Using Amazon S3</a>.</p> |
| A database that is not MySQL-compatible | <p>You can also use AWS Database Migration Service (AWS DMS) to migrate data from a not MySQL-compatible database. For more information about AWS DMS, see <a href="https://docs.aws.amazon.com/dms/latest/userguide/Welcome.html">https://docs.aws.amazon.com/dms/latest/userguide/Welcome.html</a></p>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |

## Oracle Database

- [the section called “Migrating an On-Premises Oracle Database to Amazon Aurora MySQL”](#)

- [the section called “Migrating an Amazon RDS for Oracle Database to Amazon Aurora MySQL”](#)
- [the section called “Migrating an Amazon RDS for Oracle Database to an Amazon S3 Data Lake”](#)
- [Migrating an Oracle Database to PostgreSQL](#)
- [the section called “Migrating Oracle databases to Amazon Aurora MySQL with DMS Schema Conversion”](#)
- [the section called “Migrating Oracle databases to Amazon RDS for PostgreSQL with DMS Schema Conversion”](#)
- [Migrating an Amazon RDS for Oracle Database to Amazon Redshift](#)
- [Migrating an Oracle Database to Amazon RDS for Oracle](#)
- [Migrating from Amazon RDS for Oracle to Amazon RDS for PostgreSQL and Aurora PostgreSQL](#)

## Microsoft SQL Server

- [Migrating a SQL Server Database to Amazon Aurora MySQL](#)
- [Migrating an Amazon RDS for SQL Server Database to an Amazon S3 Data Lake](#)
- [the section called “Migrating SQL Server databases to Amazon Aurora PostgreSQL with DMS Schema Conversion”](#)
- [the section called “Migrating SQL Server databases to Amazon RDS for MySQL with DMS Schema Conversion”](#)
- [Migrating a SQL Server Always On Database](#)
- [Migrating SQL Server Databases to Amazon RDS for SQL Server](#)

## MySQL

- [Migrating a MySQL-Compatible Database to Amazon Aurora MySQL](#)
- [Migrating a MySQL Database to Amazon RDS for MySQL or Amazon Aurora MySQL](#)
- [Migrating data from MySQL databases with homogeneous data migrations in AWS DMS](#)

## BigQuery

- [Migrating a BigQuery Project to Amazon Redshift](#)

## MariaDB

- [Migrating a MariaDB Database to Amazon RDS for MySQL or Amazon Aurora MySQL](#)

## MongoDB

- [Migrating from MongoDB to Amazon DocumentDB](#)

## PostgreSQL

- [Migrating PostgreSQL Databases to Amazon RDS for PostgreSQL or Amazon Aurora PostgreSQL](#)
- [the section called “Migrating PostgreSQL databases to Amazon RDS for PostgreSQL with DMS homogeneous data migrations”](#)

## SAP ASE

- [Migrating from SAP ASE to Amazon Aurora MySQL](#)

In the DMS User Guide, you can find additional resources:

- [Migrating large data stores](#)



# Migrating Databases to Amazon Web Services Managed Databases

You can move from existing, self-managed, open-source, and commercial databases to fully managed AWS databases of the same engine. The following walkthroughs show how to move your databases to Amazon Relational Database Service (Amazon RDS) and Amazon Aurora.

## Topics


- [Migrating a MySQL Database to RDS for MySQL or Aurora MySQL](#)
- [Migrating PostgreSQL Databases to Amazon RDS for PostgreSQL or Amazon Aurora PostgreSQL](#)
- [Migrating PostgreSQL databases to Amazon RDS for PostgreSQL with DMS homogeneous data migrations](#)
- [Migrating an Oracle Database to Amazon RDS for Oracle](#)
- [Migrating a SQL Server Always On Database to Amazon Web Services](#)
- [Migrating an Amazon RDS for MySQL Database to an Amazon DynamoDB target](#)
- [Migrating an RDS for MySQL database to an S3 data lake](#)
- [Migrating an RDS PostgreSQL database to an S3 data lake](#)
- [Migrating SQL Server Databases to Amazon RDS for SQL Server](#)
- [Migrating from Amazon RDS for Oracle to Amazon RDS for PostgreSQL and Aurora PostgreSQL](#)
- [Migrating from SAP ASE to Amazon Aurora MySQL](#)

## Migrating a MySQL Database to RDS for MySQL or Aurora MySQL


You can use these two main approaches for migrating a self-managed MySQL database to an Amazon RDS for MySQL or Amazon Aurora MySQL database.

- Use a native or third-party database migration tool such as mysqldump to perform the full load and MySQL replication to perform ongoing replication. Typically this is the simplest option.
- Use a managed migration service such as the AWS Database Migration Service (AWS DMS). AWS DMS provides migration-specific services such as data validation that are not available in the native or third-party tools.

The following diagram displays these two approaches.



You can use a hybrid strategy that combines native or third-party tools for full load and AWS DMS for ongoing replication. The following diagram displays the hybrid migration approach.



The hybrid option delivers the simplicity of the native or third-party tools along with the additional services that AWS DMS provides. For example, in AWS DMS, you can automatically validate your migrated data, row by row and column by column, to ensure the data quality in the target database. Or, if you are only migrating a subset of the tables, it will be simpler to use AWS DMS to filter your tables than the equivalent configuration in the native or third-party tools.

## Topics

- [Full load MySQL database migration](#)
- [Full load MySQL database migration options performance comparison](#)
- [Migrate MySQL database with AWS DMS ongoing replication](#)

## Full load MySQL database migration

You can use one of these three tools to move data from your MySQL database to Amazon RDS for MySQL or Amazon Aurora MySQL. Follow the steps described in this document to perform the full data load.

## mysqldump

This native MySQL client utility installs by default with the engine that performs logical backups, producing a set of SQL statements that you can execute to reproduce the original database object definitions and table data. `mysqldump` dumps one or more MySQL databases for backup or transfer to another MySQL server. For more information, see the [mysqldump documentation](#).

`mysqldump` is appropriate when the following conditions are met:

- The data set is smaller than 10 GB.
- The network connection between source and target databases is fast and stable.
- Migration time is not critical, and the cost of re-trying the migration is very low.
- You don't need to do any intermediate schema or data transformations.

You can decide not to use this tool if any of the following conditions are true:

- You migrate from an Amazon RDS for MySQL DB instance or a self-managed MySQL 5.5 or 5.6 database. In that case, you can get better performance results with Percona XtraBackup.
- It is impossible to establish a network connection from a single client instance to source and target databases due to network architecture or security considerations.
- The network connection between the source and target databases is unstable or very slow.
- The data set is larger than 10 GB.
- An intermediate dump file is required to perform schema or data manipulations before you can import the schema or data.

For details and step-by-step instructions, see [Importing data to an Amazon RDS for MySQL or MariaDB DB instance with reduced downtime](#) in the *Amazon RDS User Guide*.

Follow these three steps to perform full data load using `mysqldump`.

1. Produce a dump file containing source data.
2. Restore this dump file on the target database.
3. Retrieve the binary log position for ongoing replication.

For example, the following command creates the dump file. The `--master-data=2` parameter creates a backup file, which you can use to start the replication in AWS DMS.

```
sudo mysqldump \  
  --databases <database_name> \  
  --master-data=2 \  
  --single-transaction \  
  --order-by-primary \  
  -r <backup_file>.sql \  
  -u local_user \  
  -p <local_password>
```

For example, the following command restores the dump file on the target host.

```
mysql -h host_name -P 3306 -u db_master_user -p < backup_file.sql
```

For example, the following command retrieves the binary log file name and position from the dump file. Save this information for later when you configure AWS DMS for ongoing replication.

```
head mysqldump.sql -n80 | grep "MASTER_LOG_POS"  
  
-- Will Get output similar to  
-- CHANGE MASTER TO MASTER_LOG_FILE='mysql-bin.000125', MASTER_LOG_POS=150;
```

## Percona XtraBackup

Amazon RDS for MySQL and Amazon Aurora MySQL support migration from Percona XtraBackup files that are stored in an Amazon S3 bucket. Percona XtraBackup produces a binary backup files which can be significantly faster than migrating from logical schema and data dumps using tools such as mysqldump. The tool can be used for small-scale to large-scale migrations.

Percona XtraBackup is appropriate when the following conditions are met:

- You have administrative, system-level access to the source database.
- You migrate database servers in a 1-to-1 fashion: one source MySQL server becomes one new Amazon RDS for MySQL or Aurora DB cluster.

You can decide not to use this tool if any of the following conditions are true:

- You can't use third-party software because of operating system limitations.
- You migrate into existing Aurora DB clusters.

- You migrate multiple source MySQL servers into a single Aurora DB cluster.
- For more information, see [Limitations and recommendations for importing backup files from Amazon S3 to Amazon RDS](#).

For details and step-by-step instructions, see [Migrating data from MySQL by using an Amazon S3 Bucket](#) in the *Amazon RDS User Guide*.

Follow these three steps to perform full data load using Percona XtraBackup.

1. Produce a backup file containing source data.
2. Restore this backup file from Amazon S3 while launching a new target database.
3. Retrieve the binary log position for ongoing replication.

For example, the following command creates the backup file and streams it directly to Amazon S3.

```
xtrabackup --user=<myuser> --backup --parallel=4 \  
--stream=xbstream --compress | \  
aws s3 cp - s3://<bucket_name>/<backup_file>.xbstream
```

Use the Amazon RDS console to restore the backup files from the Amazon S3 bucket and create a new Amazon Aurora MySQL DB cluster. For more information, see [Restoring an Aurora MySQL DB cluster from an Amazon S3 bucket](#).

For example, the following command prints the binary log (binlog) information after you finish the creation of a compressed backup.

```
MySQL binlog position: filename 'mysql-bin.000001', position '481'
```

For example, the following command retrieves the binary log file name and position from the `xtrabackup_binlog_info` file. This file is located in the main backup directory of an uncompressed backup.

```
$ cat </on-premises/backup>/xtrabackup_binlog_info  
// Output  
mysql-bin.000001      481
```

## mydumper

mydumper and myloader are third-party utilities that perform a multithreaded schema and data migration without the need to manually invoke any SQL commands or design custom migration scripts. mydumper functions similarly to mysqldump, but offers many improvements such as parallel backups, consistent reads, and built-in compression. Another benefit to mydumper is that each individual table gets dumped into a separate file. The tools are highly flexible and have reasonable configuration defaults. You can adjust the default configuration to satisfy the requirements of both small-scale and large-scale migrations.

mydumper is appropriate when the following conditions are met:

- Migration time is critical.
- You can't use Percona XtraBackup.

You can decide not to use this tool if any of the following conditions are true:

- You migrate from an Amazon RDS for MySQL DB instance or a self-managed MySQL 5.5 or 5.6 database. In that case, you might get better results Percona XtraBackup.
- You can't use third-party software because of operating system limitations.
- Your data transformation processes require intermediate dump files in a flat-file format and not an SQL format.

For details and step-by-step instructions, see the [mydumper project](#).

Follow these three steps to perform full data load using mydumper.

1. Produce a dump file containing source data.
2. Restore this dump file on the target database using myloader.
3. Retrieve the binary log position for ongoing replication.

For example, the following command creates the backup of dbName1 and dbName2 databases using mydumper.

```
mydumper \
--host=<db-server-address> \
--user=<mydumper-username> --password=<mydumper-password> \
```

```
--outputdir=/db-dump/mydumper-files/ \  
-G -E -R --compress --build-empty-files \  
--threads=4 --compress-protocol \  
--regex '^(DbName1\.|DbName2\..)' \  
-L /<mydumper-logs-dir>/mydumper-logs.txt
```

For example, the following command restores the backup to the Amazon RDS instance using `myloader`.

```
myloader \  
--host=<rds-instance-endpoint> \  
--user=<db-username> --password=<db-password> \  
--directory=<mydumper-output-dir> \  
--queries-per-transaction=50000 --threads=4 \  
--compress-protocol --verbose=3 -e 2><myload-output-logs-path>
```

For example, the following command retrieves the binary log information from the `mydumper` metadata file.

```
cat <mydumper-output-dir>/metadata  
# It should display data similar to the following:  
SHOW MASTER STATUS:SHOW MASTER STATUS:  
  Log: mysql-bin.000129  
  Pos: 150  
  GTID:
```


### Note

1. To ensure a valid dump file of logical backups in `mysqldump` and `mydumper`, don't run data definition language (DDL) statements while the dump process is running. It is recommended to schedule a maintenance window for these operations. For details, see the [single-transaction documentation](#).
2. While exporting the data with logical backups, it is recommended to exclude MySQL default schemas (`mysql`, `performance_schema`, and `information_schema`), functions, stored procedures, and triggers.
3. Remove definers from schema files before uploading extracted data to Amazon RDS. For more information, see [How can I resolve definer errors](#).
4. Any backup operation acquires a global read lock on all tables (using `FLUSH TABLES WITH READ LOCK`). As soon as this lock has been acquired, the binary log coordinates

are read and the lock is released. For more information, see [Establishing a Backup Policy](#). For logical backups this step done at the beginning of the logical dump, however for physical backup (Percona XtraBackup) this step done at the end of backup.

## Full load MySQL database migration options performance comparison

We tested these three full load options using a MySQL 5.7 database on EC2 as the source and Aurora MySQL 5.7 as the target. The source database contained the AWS DMS [sample database](#) with a total of 9 GB of data. The following image shows the performance results.



Percona XtraBackup performed 4x faster than mysqldump and 2x faster than mydumper backups. We tested larger datasets, for example with a total of 400 GB of data, and found that the performance scaled proportionally to the dataset size.

Percona XtraBackup creates a physical backup of the database files whereas the other tools create logical backups. Percona XtraBackup is the best option for full load if your use case conforms to the restrictions listed in the Percona XtraBackup section above. If Percona XtraBackup isn't compatible with your use case, mydumper is the next best option. For more information about physical and logical backups, see [Backup and Recovery Types](#).

## Migrate MySQL database with AWS DMS ongoing replication

To configure the ongoing replication in AWS DMS, enter the native start point for MySQL, which you have retrieved at the end of the full load process as described for each tool. The native start point will be similar to `mysql-bin-changelog.000024:373`.

In the **Create database migration task** page, follow these three steps to create the migration task.

1. For **Migration type**, choose **Replicate ongoing changes**.
2. Under **CDC start mode for source transactions**, choose **Enable custom CDC start mode**.
3. Under **Custom CDC start point**, paste the native start point you saved earlier.



For more information, see [Creating tasks for ongoing replication](#) and [Migrate from MySQL to Amazon RDS](#).

 **Note**

The AWS DMS CDC replication uses plain SQL statements from the binary log to apply data changes in the target database. Therefore, it is slower and more resource-intensive than the native Primary/Replica binary log replication in MySQL. For more information, see [Replication with a MySQL or MariaDB instance running external to Amazon RDS](#).

You should always remove triggers from the target during the AWS DMS CDC replication. For example, the following command generates the script to remove triggers.

```
# In case required to generate drop triggers script
SELECT Concat('DROP TRIGGER ', Trigger_Name, ';') FROM information_schema.TRIGGERS
WHERE TRIGGER_SCHEMA not in ('sys','mysql');
```


## Migrating PostgreSQL Databases to Amazon RDS for PostgreSQL or Amazon Aurora PostgreSQL

This walkthrough gets you started with homogeneous database migration from PostgreSQL to Amazon Relational Database Service (Amazon RDS) for PostgreSQL or Amazon Aurora PostgreSQL-Compatible Edition. This guide provides a quick overview of the data migration process and provides suggestions on how to select the best option to use.

Customers looking to migrate self-managed PostgreSQL databases to Amazon RDS for PostgreSQL or Aurora PostgreSQL, can use one of the three main approaches.

- Use a native or third-party database migration method such as `pg_dump` and `pg_restore` for full load only migrations.
- Use a managed service such as AWS Database Migration Service (AWS DMS) for full load and ongoing replication.
- Use a native tool for full load and a managed AWS DMS service for ongoing replication. We call this strategy the *hybrid approach*.

This document describes the hybrid approach. The following diagram shows the components of the hybrid approach.



The hybrid approach provides the following advantages.

- Automation of the creation of secondary database objects such as views, indexes, and constraints.
- AWS DMS data validation to ensure that your target data matches with the source, row by row and column by column.
- Other capabilities provided by AWS DMS, for example CloudWatch monitoring and table statistics. It may be simpler to use AWS DMS to track migration progress, transactional workload, receive and transmit throughput, source and target latency, and so on.

This document describes the native options for the full load. It also includes a comparison so that you can evaluate the options for your migration requirements. In conclusion, you can find a brief description of how to use AWS DMS for ongoing replication.

## Topics

- [Summary](#)

- [Full load PostgreSQL database migration](#)
- [Full load PostgreSQL database migration options performance comparison](#)
- [Migrate PostgreSQL database with AWS DMS ongoing replication](#)

## Summary

This document describes the hybrid approach for migrating between PostgreSQL databases. We analyzed three options for full load and demonstrated the relative performance of each for a test database.

If you need to create secondary database objects, then `pg_dump` and `pg_restore` is the most appropriate option. However, this option incurs a performance tradeoff compared to other options.

`pglogical` has a slight performance advantage over publisher and subscriber. However, you need to install the `pglogical` extension on your source database server.

You can use these guidelines to choose the option that best matches your migration goal.

For more information about the performance of these tools, see [Performance Comparison](#).

## Full load PostgreSQL database migration

The full load migration phase populates the target database with a copy of the source data. This chapter describes the following native methods to help you choose the one that best matches your migration scenario.

- `pg_dump` and `pg_restore`
- Publisher and Subscriber
- `pglogical`

We recommend that you begin by reviewing the following table to understand the tools suitable for your use case.

| Method                   | Supported versions         | Support of metadata migration | Suitable database sizes | Performance |
|--------------------------|----------------------------|-------------------------------|-------------------------|-------------|
| pg_dump and pg_restore   | All versions of PostgreSQL | Yes                           | 100 GB or less          | Medium      |
| Publisher and Subscriber | PostgreSQL 10.0 and higher | No                            | Any size                | High        |
| pglogical                | PostgreSQL 9.4 and higher  | Yes                           | Any size                | High        |

The suitable database sizes provided in the preceding table are the AWS DMS recommendations. These recommendations are based on customer migration experiences and aren't the limitation of the native tools.

## Topics

- [Preparing for Ongoing Replication](#)
- [PostgreSQL pg\\_dump and pg\\_restore utility](#)
- [PostgreSQL publisher and subscriber model](#)
- [PostgreSQL pglogical extension](#)

## Preparing for Ongoing Replication

Before you start full load, make sure that you record the current log sequence number (LSN) as the starting position for ongoing replication. Use this LSN when you configure the ongoing replication task in AWS DMS.

To avoid data loss or duplication with hybrid approach, make sure of the following:

- You create a replication slot on the source database before you start the full load using the following command:

```
SELECT * FROM pg_create_logical_replication_slot('test_slot', 'test_decoding');
```

- When you create a replication slot, your source database doesn't have open transactions. To confirm, use the following command:

```
SELECT * FROM pg_stat_activity where state <> 'idle';
```

- If you have open transactions, wait for them to complete or cancel them.

To capture the LSN, use the following command.

```
SELECT slot_name, confirmed_flush_lsn from pg_replication_slots where slot_name like 'test_slot';
```

```
slot_name | confirmed_flush_lsn  
test_slot | 12/68000000
```

From the output of the preceding command, copy the `confirmed_flush_lsn` value. In the example preceding, this value is set to `12/68000000`. After you complete the full load, you can use this value as the start position for the AWS DMS task.

## PostgreSQL `pg_dump` and `pg_restore` utility

`pg_dump` and `pg_restore` is a native PostgreSQL client utility. You can find this utility as part of the database installation. It produces a set of SQL statements that you can run to reproduce the original database object definitions and table data.

The `pg_dump` and `pg_restore` utility is suitable for the following use cases if:

- Your database size is less than 100 GB.
- You plan to migrate database metadata as well as table data.
- You have a relatively large number of tables to migrate.

The `pg_dump` and `pg_restore` utility may not be suitable for the following use cases if:

- Your database size is greater than 100 GB.
- You want to avoid downtime.

## Example

At a high level, you can use the following steps to migrate the [dms\\_sample](#) database.

1. Export data to one or more dump files.
2. Create a target database.
3. Import the dump file or files.
4. (Optional) Migrate database roles and users.

## Export Data

You can use the following command to create dump files for your source database.

```
pg_dump -h <hostname> -p 5432 -U <username> -Fc -b -v -f <dumpfilelocation.sql> -d  
<database_name>
```

-h is the name of source server where you would like to migrate your database.  
-U is the name of the user present on the source server  
-Fc: Sets the output as a custom-format archive suitable for input into pg\_restore.  
-b: Include large objects in the dump.  
-v: Specifies verbose mode  
-f: Dump file path

## Create a Database on Your Target Instance

First, login to your target database server.

```
psql -h <hostname> -p 5432 -U <username> -d <database_name>
```

-h is the name of target server where you would like to migrate your database.  
-U is the name of the user present on the target server.  
-d is the name of database name present on target already.

Then, use the following command to create a database.

```
create database migrated_database;
```

## Import Dump Files

You can use the following command to import the dump file into your Amazon RDS instance.

```
pg_restore -v -h <hostname> -U <username> -d <database_name> -j 2  
<dumpfilelocation.sql>
```

```
-h is the name of target server where you would like to migrate your database.  
-U is the name of the user present on the target server.  
-d is the name of database name that was created in step 2.  
<dumpfilelocation.sql> is the dump file that was created to generate the script of the  
database using pg_dump
```

## Migrate Database Roles and Users

To export such database objects as roles and users, you can use the `pg_dumpall` utility.

To generate a script for users and roles, run the following command on the source database.

```
pg_dumpall -U <username> -h <hostname> -f <dumpfilelocation.sql> --no-role-passwords -g
```

```
-h is the name of source server where you would like to migrate your database.  
-U is the name of the user present on the source server.  
-f: Dump file path.  
-g: Dump only global objects (roles and tablespaces), no databases.
```

To restore users and roles, run the following command on your target database.

```
psql -h <hostname> -U <username> -f <dumpfilelocation.sql>
```

```
-h is the name of target server where you would like to migrate your database.  
-U is the name of the user present on the target server.  
-f: Dump file path.
```

To complete the export and import operations, the `pg_dump` and `pg_restore` requires some time. This time depends on the following parameters.

- The size of your source database.
- The number of jobs.
- The resources that you provision for your instance used to invoke `pg_dump` and `pg_restore`.

## PostgreSQL publisher and subscriber model

In PostgreSQL, logical replication uses a publisher and subscriber model. In this model, one or more subscribers subscribe to one or more publications on a publisher node. Subscribers pull data

from the publications they subscribe to and may subsequently re-publish data to allow cascading replication or more complex configurations. PostgreSQL version 10.0 and higher supports the native publisher and subscriber model.

The publisher and subscriber model is suitable for the following use cases if:

- Your database size is greater than 100 GB.
- You have an existing schema in your target database and migrate only data.
- You want to capture ongoing changes.
- You want to minimize downtime.

The publisher and subscriber may not be suitable for the following use cases if:

- You plan to migrate database metadata.
- Your source PostgreSQL database version is lower than 10.x.
- You want to replicate the schema, DDL, and sequences.

## Example

The migration process involves copying a snapshot of the publishing database to the subscriber. This step is also called the table synchronization phase. To reduce the time spent in this phase, you can spawn multiple table synchronization workers. However, you can only have one synchronization worker for each table.

The following example shows how to migrate all tables from a public schema.

## Configure the Source Database

To configure the source database with built-in logical replication, complete the following steps.

1. In the source database, edit the `postgresql.conf` file to add the following parameters.

```
wal_level = 'logical'  
max_replication_slots = 10  
max_wal_senders = 10
```

Make sure that you set `wal_level` to `logical`. This adds information necessary to support logical decoding. Also, make sure that you set `max_replication_slots` to



at least the number of subscriptions expected to connect, plus some reserve for table synchronization. Finally, make sure that you set `max_wal_senders` to at least the same value as `max_replication_slots` plus the number of physical replicas that are connected at the same time.

2. Restart the source PostgreSQL instance for these parameters to take effect.
3. To make sure that you configured parameters on your source database correctly, run the following command.

```
psql -h <hostname> -p 5432 -U <username> -d <database_name>
-c "select name, setting from pg_settings where name in
('wal_level', 'max_worker_processes', 'max_replication_slots', 'max_wal_senders', 'shared_preload_libraries');"
```

-h is the name of source server where you would like to migrate your database.  
-U is the name of the user present on the source server.  
-d is the name of database name present on source already.

## Set Up the Logical Replication

Now, you can configure built-in logical replication between your self-managed PostgreSQL databases and Amazon RDS for PostgreSQL or Aurora PostgreSQL.

To create the publication on the source database server, run the following command.

```
CREATE PUBLICATION my_publication FOR ALL TABLES;
```

You can specify only the tables that you want to publish. You can also limit the changes that will be published.

To replicate DELETE and UPDATE operations, make sure that the published table has a replica identity, which can be a primary key. This makes it possible for the subscriber to identify the modified rows. You can replicate INSERT operations without a replica identity. After you created the publications, you can create subscriptions in the subscriber node.

Before you create a subscriber node, make sure that you created the target Amazon RDS for PostgreSQL or Aurora PostgreSQL database. For more information, see [Creating a PostgreSQL DB instance and connecting to a database on a PostgreSQL DB instance](#).

To create a database on Amazon RDS for PostgreSQL or Aurora PostgreSQL, run the following command.

```
psql -h <hostname> -p 5432 -U <username> -d <database_name> -c "create database
migrated_database;"
```

-h is the name of target server where you would like to migrate your database.

-U is the name of the user present on the target server.

-d is the name of database name present on target already.

To create the subscription on your target database, run the following command.

```
CREATE SUBSCRIPTION <subscription_name>
CONNECTION 'host=<host> port=<port_number> dbname=<database_name> user=<username>
password=<password>' PUBLICATION <publication_name> WITH (copy_data=true);
```

-subscription\_name: Provide the name of the subscription created at target.

-host: Provide the hostname of the source database.

-port\_number: Provide the port on which the source database is running.

-database\_name: Provide the name of the database where the publication is created.

-publication\_name: Provide the name of the publication created at source.

-copy\_data: Specifies whether the existing data in the publications that are being subscribed to should be copied once the replication starts. The default is true.

## Verify that the Data Replication Is Running

Make sure that no active transactions or data changes are happening on your source database.

Then, check the status of your replication by running the following statement on your source database. Make sure that the WAL locations are the same for the sent\_location, write\_location, and replay\_location. This indicates that the target database is at the same LSN position as the source database.

```
SELECT * FROM pg_stat_replication;
```

## Stop the Replication

When the data is in sync between your source and target databases, stop the subscriber on your target database.

```
ALTER SUBSCRIPTION <subscription_name> DISABLE;
```

-subscription\_name: Provide the name of the subscription created at target.

Capture the slot name created by the publisher and subscriber on the target database.

```
select subslotname from pg_subscription where subname like 'subscription_name';
```

-subscription\_name: Provide the name of the subscription created at target.

Capture the `confirmed_flush_lsn` value from the replication slot fetched. You can use this value as the start position for the AWS DMS task.

```
SELECT slot_name, confirmed_flush_lsn from pg_replication_slots where slot_name like 'replication_slot_name';
```

## Drop Publication and Subscription Artifacts

To drop the subscription on your target database, run the following command.

```
DROP SUBSCRIPTION <subscription_name>;
```

-subscription\_name: Provide the name of the subscription created at target.

To drop the publication on your source database, run the following command.

```
DROP PUBLICATION <publication_name>;
```

-publication\_name: Provide the name of the publication created at source.

## PostgreSQL pglogical extension

The pglogical extension for PostgreSQL implements logical streaming replication, using a similar publish and subscribe built-in approach.

The pglogical extension is suitable for the following use cases if:

- Your database size is greater than 100 GB.
- You want to replicate the schema, DDL, sequences, and table data.
- You want to capture ongoing changes.
- You want to avoid downtime.

The pglogical extension may not be suitable for the following use cases if:

- You have UNLOGGED and TEMPORARY tables.

- You plan to migrate database metadata.

## Example

The following example shows how to migrate the public schema.

### Configure the Source Database

To configure the source database with logical replication, complete the following steps.

1. In the source database, edit the `postgresql.conf` file to add the following parameters.

```
wal_level = 'logical'  
max_worker_processes = 10  
max_replication_slots = 10  
max_wal_senders = 10  
shared_preload_libraries = 'pglogical'
```

2. Restart the source PostgreSQL instance for these parameters to take effect.
3. To make sure that you configured parameters on your source database correctly, run the following command.

```
psql -h <hostname> -p 5432 -U <username> -d <database_name>  
-c "select name, setting from pg_settings where name in  
( 'rds.logical_replication', 'shared_preload_libraries' );"
```

-h is the name of source server where you would like to migrate your database.  
-U is the name of the user present on the source server.  
-d is the name of database name present on source already.

### Configure the Target Database

By default, Amazon RDS for PostgreSQL and Aurora PostgreSQL have the `pglogical` extension. To configure the target DB parameter group, complete the following steps.

1. To turn on the logical replication in the target database, set the following parameters in the database parameter group. For more information, see [Working with parameter groups](#).

```
rds.logical_replication=1  
shared_preload_libraries = 'pglogical'
```

2. Reboot your Amazon RDS instance for these parameters to take effect.
3. To make sure that you configured parameters on your target database correctly, run the following command.

```
psql -h <hostname> -p 5432 -U <username> -d <database_name>  
-c "select name, setting from pg_settings where name in  
( 'rds.logical_replication', 'shared_preload_libraries' );"
```

-h is the name of target server where you would like to migrate your database.  
-U is the name of the user present on the target server.  
-d is the name of database name present on target already.

## Set Up the Logical Replication

Now, you can configure the logical replication between your self-managed PostgreSQL databases and Amazon RDS for PostgreSQL or Aurora PostgreSQL.

1. Download the `pglogical` rpm and install it on your source database.

- For PostgreSQL 9.6, run the following command.

```
curl https://access.2ndquadrant.com/api/repository/dl/default/release/9.6/rpm |  
bash  
yum install postgresql96-pglogical
```

- For PostgreSQL 10, run the following command.

```
curl https://access.2ndquadrant.com/api/repository/dl/default/release/10/rpm | bash  
yum install postgresql10-pglogical
```

2. Create the `pglogical` extension on your provider and subscriber.

```
CREATE EXTENSION pglogical;
```

3. Create the publisher node on your source database.

```
SELECT pglogical.create_node(  
    node_name := 'publisher_name',  
    dsn := 'host=<publisher_hostname> port=port_number dbname=<database_name>'  
);
```

```
-publisher_name: Provide the name of the publication created at source.  
-publisher_hostname: Provide the hostname of the source database.  
-port_number: Provide the port on which the source database is running.  
-database_name: Provide the name of the database where the publication is created.
```

#### 4. Add tables in public schema to the default replication set.

```
SELECT pglogical.replication_set_add_all_tables('default', ARRAY['public']);
```

#### 5. Create the subscriber node on target database.

```
SELECT pglogical.create_node(  
    node_name := 'subscriber_name',  
    dsn := 'host=<subscriber_hostname> port=port_number dbname=<database_name>'  
);  
  
-subscriber_hostname: Provide the hostname of the target database.  
-port_number: Provide the port on which the target database is running.  
-database_name: Provide the name of the database where the subscription is created.  
-subscriber_name: Provide the name of the subscription created at target.
```

#### 6. Create the subscription on the subscriber node. This subscription starts synchronization and replication processes in background.

```
SELECT pglogical.create_subscription(  
    subscription_name := 'subscription_name',  
    provider_dsn := 'host=<publisher_hostname> port=port_number  
    dbname=<database_name>'  
);  
  
SELECT pglogical.wait_for_subscription_sync_complete('subscription_name');  
  
-publisher_hostname: Provide the hostname of the source database.  
-port_number: Provide the port on which the target database is running.  
-database_name: Provide the name of the database where the subscription is created.  
-subscription_name: Provide the name of the subscription created at target.
```

### Verify that the Data Replication Is Running

Make sure that no active transactions or data changes are happening on your source database. Then, check the status of your replication by running the following statement on your

source database. Make sure that the WAL locations are the same for the `sent_location`, `write_location`, and `replay_location`. This indicates that the target database is at the same LSN position as the source database.

```
SELECT * FROM pg_stat_replication;
```

## Stop the Replication

When the data is in sync between your source and target databases, stop the subscriber on your target database.

```
select pglogical.alter_subscription_disable('subscriber_name');  
  
-subscriber_name: Provide the name of the subscription created at target.
```

Capture the `confirmed_flush_lsn` value from the replication slot created by the `pglogical` setup. You can use this value as the start position for the AWS DMS task.

```
SELECT slot_name, confirmed_flush_lsn from pg_replication_slots where slot_name like  
'replication_slot_name';
```

## Drop the Subscription


To drop the subscription on your target database, run the following command.

```
select pglogical.drop_subscription('subscriber_name');  
  
-subscriber_name: Provide the name of the subscription created at target.
```

## Full load PostgreSQL database migration options performance comparison

We analyzed the performance of `pg_dump` and `pg_restore`, publisher and subscriber, and `pglogical` in a full load migration. We migrated a 70 GB database that includes 6 tables and LOB data. We lifted and shifted this database from the source to the target.

The following image represents the performance comparison of the three migration methods. We expect similar performance trends for larger datasets.



We performed this test to provide a basic overview of the full load performance. This performance may vary because it depends on such factors as network bandwidth, data structure, data size, and so on.

The elapsed time shown in the diagram is the actual migration time. It doesn't include the time spent on implementing prerequisites.

You can compare the results of all three methods.

- 35 minutes is the total elapsed time for pglogical.
- 37 minutes is the total elapsed time for publisher and subscriber.
- 46 minutes is the total elapsed time for pg\_dump and pg\_restore. This time includes:
  - 19 minutes to unload data using pg\_dump.
  - 27 minutes to load data using pg\_restore.

From the comparison, you can see that pglogical has the best performance among the three full load options. Consider this approach if you don't need to migrate secondary database objects such as views, stored procedures, triggers, and so on. This is the preferred approach where the database size is greater than 100 GB and when you don't have transformation or filtering requirements.



Publisher and subscriber may be appropriate if you don't need to migrate secondary database objects such as views, stored procedures, triggers, and so on. You can use publisher and subscriber for smaller migrations where ease of use considerations override the minor performance gains provided by pglogical.

Using `pg_dump` and `pg_restore` is slower than both pglogical and publisher and subscriber. `pg_dump` is the only option that migrates your secondary database objects. Additionally, data files created by `pg_dump` may be orders of magnitude larger than the original table size.

## Migrate PostgreSQL database with AWS DMS ongoing replication

After you complete the full load, make sure that you perform ongoing replication using AWS DMS to keep the source and target databases in sync. To configure the ongoing replication task, open the Database Migration Service console. On the **Create database migration task** page, follow these steps to create a migration task.

1. For **Migration type**, choose **Replicate ongoing changes**.
2. Under **CDC start mode for source transactions**, choose **Enable custom CDC start mode**.
3. Under **Custom CDC start point**, paste the native start point you captured when you prepared for ongoing replication. For more information, see [Preparing for Ongoing Replication](#).

### Note

PostgreSQL as a source doesn't support a custom CDC start time. This is because the PostgreSQL database engine doesn't have a way to map a timestamp to an LSN or SCN as Oracle and SQL Server do.

For more information, see [Creating tasks for ongoing replication](#) and [Migrate from PostgreSQL to Amazon RDS](#).

## Migrating PostgreSQL databases to Amazon RDS for PostgreSQL with DMS homogeneous data migrations

This walkthrough gets you started with a homogeneous database migration from PostgreSQL to Amazon RDS for PostgreSQL. To automate the migration, we use homogeneous data migrations in AWS DMS. For homogeneous data migrations, AWS DMS uses native database tools to provide

easy and performant like-to-like migrations. This approach helps you effectively set up and run the migration from the source PostgreSQL database to its equivalent target.

With [homogeneous data migrations](#), you can migrate data, table partitions, data types, and secondary objects such as functions, stored procedures, and so on. Homogeneous data migrations in AWS DMS precisely map your source database to its equivalent Amazon RDS or Amazon Aurora target. You can also use homogeneous data migrations to replicate ongoing changes from your source database to your compatible target.

Note that when using homogeneous data migrations, AWS DMS migrates your source views as tables to the target database. Otherwise, the schema and data of the target matches the schema and data of the source. This typically results in a substantially faster migration from start to finish than using AWS DMS migration tasks.

This introductory exercise shows how you can use homogeneous data migrations in AWS DMS to migrate your self-managed PostgreSQL database to the AWS Cloud.

At a high level, this migration includes the following steps:

- Use the AWS Management Console to create the required resources:
  - Create a VPC in the Amazon VPC console.
  - Create IAM roles in the IAM console.
  - Create your target Amazon RDS for PostgreSQL database in the Amazon RDS console.
  - Store database credentials in AWS Secrets Manager.
- Use the AWS DMS console to configure your migration resources:
  - Create a subnet group and an instance profile for your migration project.
  - Create data providers for your source and target databases.
  - Create a migration project.
  - Create and run a data migration.

Watch [this video](#) to learn how to use homogeneous data migrations in AWS DMS.

This walkthrough takes approximately three hours to complete. Make sure that you delete resources at the end of this walkthrough to avoid additional charges.

## Topics

- [Prerequisites for migrating PostgreSQL databases](#)

- [PostgreSQL to Amazon RDS migration overview](#)
- [Step-by-step PostgreSQL database to Amazon RDS migration walkthrough](#)
- [PostgreSQL database to Amazon RDS post-migration clean-up](#)

## Prerequisites for migrating PostgreSQL databases

The following prerequisites are also required to complete this walkthrough:

- Familiarity with the Amazon Relational Database Service (Amazon RDS), AWS Database Migration Service (AWS DMS), and SQL.
- Create an AWS account with an AWS Identity and Access Management (IAM) credentials. This account should allow you to launch Amazon RDS instances and run AWS DMS data migrations in your and AWS Region. For more information, see [Create an IAM User](#).
- Basic knowledge of the Amazon Virtual Private Cloud (Amazon VPC) service and of security groups. For information about using Amazon VPC with Amazon RDS, see [Amazon Virtual Private Cloud \(VPCs\) and Amazon RDS](#). For information about Amazon RDS security groups, see [Controlling access with security groups](#).
- An understanding of the supported features and limitations of homogeneous data migrations in AWS DMS. For example, you can't apply table mapping rules to your homogeneous data migration. For more information, see [Limitations for homogeneous data migrations](#).

We recommend that you don't use your production workloads for the migration in this walkthrough. After you get familiar with migration tools and AWS services, you can migrate your production workloads. Also, make sure that you use a source PostgreSQL database that is version 10.5 or later.

Make sure that you create all your resources in the AWS Regions that support homogeneous data migrations in AWS DMS. For more information, see the [list of supported Regions](#).

For more information about migrating self-managed PostgreSQL databases to the AWS Cloud, [Migrating PostgreSQL Databases to Amazon RDS for PostgreSQL or Amazon Aurora PostgreSQL](#).

## PostgreSQL to Amazon RDS migration overview

This section provides high-level guidance for customers looking to migrate their PostgreSQL database to Amazon RDS for PostgreSQL using homogeneous data migrations in AWS DMS.


AWS DMS creates a serverless environment for your data migration. Depending on the type of your data migration, AWS DMS automatically chooses an appropriate native PostgreSQL database tool.

For full load migrations, AWS DMS uses `pg_dump` and `pg_restore`.

For full load and change data capture (CDC) migrations, AWS DMS uses `pg_dump`, `pg_restore`, and a publisher and subscriber model for logical replication.

For homogeneous data migrations of the change data capture type, AWS DMS configures the data replication from the start point that you provide in settings.

The following diagram illustrates how AWS DMS migrates data from PostgreSQL databases with homogeneous data migrations.



Start the walkthrough by [creating the required resources](#).

## Step-by-step PostgreSQL database to Amazon RDS migration walkthrough

In the following sections, you can find step-by-step instructions for migrating your PostgreSQL database to Amazon RDS for PostgreSQL using homogeneous data migrations in AWS DMS.

### Topics

- [Step 1: Create AWS Resources](#)
- [Step 2: Configure Your Source Database](#)

- [Step 3: Create Your Target Amazon RDS for PostgreSQL Database](#)
- [Step 4: Store Database Credentials in AWS Secrets Manager](#)
- [Step 5: Create an Instance Profile](#)
- [Step 6: Configure Data Providers](#)
- [Step 7: Create a Migration Project](#)
- [Step 8: Configure a Data Migration](#)
- [Step 9: Running and Monitoring a Data Migration](#)

## Step 1: Create AWS Resources

In this step, you create and configure the required AWS resources for homogeneous data migrations in AWS DMS.

### Topics

- [Creating a VPC](#)
- [Creating an IAM policy](#)
- [Creating an IAM role](#)

### Creating a VPC

In this section, you create a virtual private cloud (VPC). This VPC is based on the Amazon Virtual Private Cloud (Amazon VPC) service and contains your AWS resources. Make sure that you create this VPC in one of the AWS Regions that support homogeneous data migrations in AWS DMS. For more information, see the [list of supported Regions](#).

To migrate your on-premises source database, make sure that you configure a private network to connect to your target database. For more information, see the [Using an on-premises source data provider](#).

### To create a VPC for homogeneous data migrations

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. Choose your AWS Region.
3. Choose **Create VPC**.
4. On the **Create VPC** page, enter the following settings:

- **Resources to create — VPC and more**
- **Name tag auto-generation** — Choose **Auto-generate** and enter a globally unique name. For example, enter `dm-vpc`.
- **IPv4 CIDR block** — `10.0.1.0/24`
- **NAT gateways** — **In 1 AZ**
- **VPC endpoints** — **None**

5. Keep the rest of the settings as they are, and choose **Create VPC**.

Use this VPC when you create your target Amazon RDS database in [Step 3](#) and your subnet group in [Step 5](#).

## Creating an IAM policy

In this section, you create an AWS Identity and Access Management (IAM) policy that AWS DMS requires to run homogeneous data migrations.

### To create an IAM policy for homogeneous data migrations

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the navigation pane, choose **Policies**.
3. Choose **Create policy**.
4. On the **Create policy** page, choose the **JSON** tab.
5. Paste the following JSON into the editor.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "ec2:DescribeRouteTables",
        "ec2:DescribeSecurityGroups",
        "ec2:DescribeVpcPeeringConnections",
        "ec2:DescribeVpcs",
        "ec2:DescribePrefixLists",
        "logs:DescribeLogGroups"
      ]
    }
  ]
}
```

```

    ],
    "Resource": "*"
  },
  {
    "Effect": "Allow",
    "Action": [
      "servicequotas:GetServiceQuota"
    ],
    "Resource": "arn:aws:servicequotas:*:*:vpc/L-0EA8095F"
  },
  {
    "Effect": "Allow",
    "Action": [
      "logs:CreateLogGroup",
      "logs:DescribeLogStreams"
    ],
    "Resource": "arn:aws:logs:*:*:log-group:dms-data-migration-*"
  },
  {
    "Effect": "Allow",
    "Action": [
      "logs:CreateLogStream",
      "logs:PutLogEvents"
    ],
    "Resource": "arn:aws:logs:*:*:log-group:dms-data-migration-*:log-
stream:dms-data-migration-*"
  },
  {
    "Effect": "Allow",
    "Action": "cloudwatch:PutMetricData",
    "Resource": "*"
  },
  {
    "Effect": "Allow",
    "Action": [
      "ec2:CreateRoute",
      "ec2:DeleteRoute"
    ],
    "Resource": "arn:aws:ec2:*:*:route-table/*"
  },
  {
    "Effect": "Allow",
    "Action": [
      "ec2:CreateTags"

```

```

    ],
    "Resource": [
        "arn:aws:ec2:*:*:security-group/*",
        "arn:aws:ec2:*:*:security-group-rule/*",
        "arn:aws:ec2:*:*:route-table/*",
        "arn:aws:ec2:*:*:vpc-peering-connection/*",
        "arn:aws:ec2:*:*:vpc/*"
    ]
},
{
    "Effect": "Allow",
    "Action": [
        "ec2:AuthorizeSecurityGroupEgress",
        "ec2:AuthorizeSecurityGroupIngress"
    ],
    "Resource": "arn:aws:ec2:*:*:security-group-rule/*"
},
{
    "Effect": "Allow",
    "Action": [
        "ec2:AuthorizeSecurityGroupEgress",
        "ec2:AuthorizeSecurityGroupIngress",
        "ec2:RevokeSecurityGroupEgress",
        "ec2:RevokeSecurityGroupIngress"
    ],
    "Resource": "arn:aws:ec2:*:*:security-group/*"
},
{
    "Effect": "Allow",
    "Action": [
        "ec2:AcceptVpcPeeringConnection",
        "ec2:ModifyVpcPeeringConnectionOptions"
    ],
    "Resource": "arn:aws:ec2:*:*:vpc-peering-connection/*"
},
{
    "Effect": "Allow",
    "Action": "ec2:AcceptVpcPeeringConnection",
    "Resource": "arn:aws:ec2:*:*:vpc/*"
}
]
}

```



1. Choose **Next** The **Review, and create** page opens.
2. For **Name**, enter `HomogeneousDataMigrationsPolicy`, and choose **Create policy**.

Use this IAM policy when you create the IAM role.

## Creating an IAM role

In this section, you create an IAM role for homogeneous data migrations. AWS DMS uses this IAM role to access database credentials stored in AWS Secrets Manager, store log files in Amazon CloudWatch, and interact with Amazon EC2.

### To create an IAM role that provides access to AWS Secrets Manager

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the navigation pane, choose **Roles**.
3. Choose **Create role**.
4. On the **Select trusted entity** page, choose **AWS service**. For **Use case**, Choose **DMS**.
5. Choose **Next**. The **Add permissions** page opens.
6. Choose **HomogeneousDataMigrationsPolicy** that you created before. Also, choose **SecretsManagerReadWrite**.
7. Choose **Next**. The **Name, review, and create** page opens.
8. For **Role name**, enter `HomogeneousDataMigrationsRole` and choose **Create role**.
9. On the **Roles** page, enter `HomogeneousDataMigrationsRole` for **Role name**. Choose **HomogeneousDataMigrationsRole**.
10. Choose the **Trust relationships** tab and choose **Edit trust policy**.
11. On the **Edit trust policy** page, paste the following JSON into the editor, replacing the existing text.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "",
      "Effect": "Allow",
      "Principal": {
```

```
        "Service": [
            "dms-data-migrations.amazonaws.com",
            "dms.your_region.amazonaws.com"
        ],
        "Action": "sts:AssumeRole"
    }
}
```

+ Replace *your\_region* with the name of your Region, such as `us-east-1`. . Choose **Update policy**.

Use this IAM role when you create your instance profile in [Step 5](#) and your migration project in [Step 7](#).

## Step 2: Configure Your Source Database

In this step, you create a new database user on your source PostgreSQL database and configure the data replication.

Use the following script to create a database user with the required permissions in your PostgreSQL source database.

```
CREATE USER your_user WITH LOGIN PASSWORD 'your_password';
ALTER USER your_user WITH SUPERUSER;
GRANT SELECT ON ALL TABLES IN SCHEMA schema_name TO your_user;
```

In the preceding example, replace *your\_user* with the name of your user. Next, replace *your\_password* with a secure password. Finally, replace *schema\_name* with the name of your database schema. Run the GRANT query for each schema that you migrate to AWS.

To replicate ongoing changes in your source database after the data migration, configure the logical replication. To turn on logical replication, set the following parameters and values in the `postgresql.conf` configuration file.

- Set `wal_level` to `logical`.
- Set `max_replication_slots` to a value greater than 1. Set the `max_replication_slots` value according to the number of tasks that you want to run. For example, to run five tasks you set a minimum of five slots. Slots open automatically as soon as a migration starts and remain open even when the migration is no longer running. Make sure to manually delete open slots.

- Set `max_wal_senders` to a value greater than 1. The `max_wal_senders` parameter sets the number of concurrent tasks that can run.
- The `wal_sender_timeout` parameter ends replication connections that are inactive longer than the specified number of milliseconds. The default is 60000 milliseconds (60 seconds). Setting the value to 0 (zero) disables the timeout mechanism.

After you edit the `postgresql.conf` configuration file, restart your PostgreSQL database server to apply new values of static parameters.

### Step 3: Create Your Target Amazon RDS for PostgreSQL Database

In this step, you create a new Amazon RDS for PostgreSQL database to use as a migration target. Also, you configure a new database user on your target Amazon RDS for PostgreSQL database.

If you already created the target database, skip this step and proceed with the configuration of your database user.

#### To create an Amazon RDS for PostgreSQL database for homogeneous data migrations

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. Choose your AWS Region.
3. Choose **Create database**.
4. For **Engine type**, choose **PostgreSQL**.
5. For **Templates**, choose **Free tier**.
6. For **DB instance identifier**, enter a unique name for your PostgreSQL database.
7. For **Master password** and **Confirm master password**, enter a secure password that includes at least 8 printable characters.
8. For **Virtual private cloud (VPC)** under **Connectivity**, choose `dm-vpc`. You created this VPC in [Step 1](#).
9. For **Public access**, choose **Yes**.
10. Keep the rest of the settings as they are, and then choose **Create database**.

After you create your Amazon RDS for PostgreSQL database, configure a new database user. Then, store the credentials of this user in AWS Secrets Manager.

You can use the following code example to create a database user with the required permissions.

```
CREATE USER your_user WITH LOGIN PASSWORD 'your_password';
GRANT USAGE ON SCHEMA schema_name TO your_user;
GRANT CONNECT ON DATABASE db_name TO your_user;
GRANT CREATE ON DATABASE db_name TO your_user;
GRANT CREATE ON SCHEMA schema_name TO your_user;
GRANT UPDATE, INSERT, SELECT, DELETE, TRUNCATE ON ALL TABLES IN SCHEMA schema_name TO
your_user;
```

In the preceding example, replace *your\_user* with the name of your user. Next, replace *your\_password* with a secure password. Finally, replace *db\_name* and *schema\_name* with your values.

To turn on logical replication for your RDS for PostgreSQL target, set the `rds.logical_replication` parameter in your DB parameter group to 1. This static parameter requires a reboot of the DB instance or DB cluster to take effect. Some parameters are static, and you can only set them at server start. AWS DMS ignores changes to their entries in the DB parameter group until you restart the server.

## Step 4: Store Database Credentials in AWS Secrets Manager

To connect to your source and target databases in an AWS DMS migration project, store your database credentials in AWS Secrets Manager. Make sure that you replicate these secrets to your AWS Region.

### To store your source database credentials in AWS Secrets Manager

1. Sign in to the AWS Management Console and open the AWS Secrets Manager console at <https://console.aws.amazon.com/secretsmanager/>.
2. Choose your AWS Region.
3. Choose **Store a new secret**. The **Choose secret type** page opens.
4. For **Secret type**, choose **Credentials for other database**.
5. For **User name** and **Password**, enter the credentials of the database user that you created for your source database in [Step 2](#).
6. For **Database**, choose **PostgreSQL**.
7. For **Server address**, **Database name**, and **Port**, enter your PostgreSQL database connection information.

8. Choose **Next**. The **Configure secret** page opens.
9. For **Secret name**, enter `dm-postgresql-source`.
10. Choose **Next**. The **Configure rotation** page opens.
11. Choose **Next**. The **Review** page opens.
12. Choose **Store**.

### To store your target database credentials in AWS Secrets Manager

1. Sign in to the AWS Management Console and open the AWS Secrets Manager console at <https://console.aws.amazon.com/secretsmanager/>.
2. Choose your AWS Region.
3. Choose **Store a new secret**. The **Choose secret type** page opens.
4. For **Secret type**, choose **Credentials for Amazon RDS database**.
5. For **User name** and **Password**, enter the credentials of the database user that you created for your target database in [Step 3](#).
6. For **Database**, choose your Amazon RDS for PostgreSQL DB instance.
7. Choose **Next**. The **Configure secret** page opens.
8. For **Secret name**, enter `dm-postgresql-target`.
9. Choose **Next**. The **Configure rotation** page opens.
10. Choose **Next**. The **Review** page opens.
11. Choose **Store**.

Use these secrets when you create your migration project in [Step 7](#).

## Step 5: Create an Instance Profile

Before you create an instance profile, configure a subnet group for your instance profile.

### To create a subnet group

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.

3. In the navigation pane, choose **Subnet groups**, and then choose **Create subnet group**.
4. For **Name**, enter DataMigrationSubnetGroup.
5. For **Description**, enter A group of private subnets.
6. For **VPC**, choose dm-vpc. You created this VPC in [Step 1](#).
7. For **Add subnets**, choose two private subnet IDs.
8. Choose **Create subnet group**.

Before you create your migration project, you set up an instance profile. An instance profile specifies network and security settings for your migration project.

### To create an instance profile

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. In the navigation pane, choose **Instance profiles**, and then choose **Create instance profile**.
4. For **Name**, enter a unique name for your instance profile. For example, enter dm-instance-profile.
5. For **Virtual private cloud (VPC)**, choose dm-vpc. You created this VPC in [Step 1](#).
6. For **Subnet group**, choose the DataMigrationSubnetGroup subnet group that you created before.
7. Choose **Create instance profile**.

Use this instance profile when you create your migration project in [Step 7](#).

## Step 6: Configure Data Providers

In this step, you create data providers that describe your source and target databases. A data provider stores a data store type and the location information about your database. Data providers don't include database credentials. You store database credentials in AWS Secrets Manager. Make sure that you include data providers and database secrets in your migration project.

You can create only one data provider for a single database. If you try to create a second data provider for the same database, AWS DMS displays an error message. However, you can use one data provider in multiple migration projects.

## To create a data provider for your source database

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. In the navigation pane, choose **Data providers**, and then choose **Create data provider**.
4. For **Configuration**, choose **Enter manually**.
5. For **Name**, enter a unique name for your source data provider. For example, enter dm-postgresql-source-provider.
6. For **Engine type**, choose **PostgreSQL**.
7. For **Server name**, enter the Domain Name Service (DNS) name or IP address of your database server.
8. For **Port**, enter the port used to connect to your database server.
9. For **Database name**, enter the name of your source database.
10. For **Secure Socket Layer (SSL) mode**, choose **none**. Optionally, choose the type of your SSL enforcement, and provide the certificate information.
11. Choose **Create data provider**.

## To create a data provider for your Amazon RDS for PostgreSQL database

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. In the navigation pane, choose **Data providers**, and then choose **Create data provider**.
4. For **Configuration**, choose **RDS database instance**.
5. For **Database from RDS**, choose the Amazon RDS for PostgreSQL database that you created in [Step 3](#).
6. For **Name**, enter a unique name for your target data provider. For example, enter dm-postgresql-target-provider.
7. Choose **Create data provider**.

Use these data providers when you create your migration project in [Step 7](#).

## Step 7: Create a Migration Project

Now you can create a migration project. A migration project describes your instance profile, source and target data providers, and secrets from AWS Secrets Manager.

To create a migration project

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. Choose **Migration projects**, and then choose **Create migration project**.
4. For **Name**, enter a unique name for your migration project. For example, enter dm-project.
5. For **Instance profile**, choose dm-instance-profile. You created this instance profile in [Step 5](#).
6. For **Source**, choose **Browse**, and then choose dm-postgresql-source-provider. You created this data provider in [Step 6](#).
7. For **Secret ID**, choose dm-postgresql-source. You created this secret in [Step 4](#).
8. For **IAM role**, choose HomogeneousDataMigrationsRole. You created this role in [Step 1](#).
9. For **Target**, choose **Browse**, and then choose dm-postgresql-target-provider. You created this data provider in [Step 6](#).
10. For **Secret ID**, choose dm-postgresql-target. You created this secret in [Step 4](#).
11. For **IAM role**, choose HomogeneousDataMigrationsRole. You created this role in [Step 1](#).
12. Choose **Create migration project**.

Use this migration project to migrate your source PostgreSQL database to your Amazon RDS for PostgreSQL database.

## Step 8: Configure a Data Migration

After you create the migration project with two PostgreSQL data providers, you can use this project for homogeneous data migrations.

To create a data migration

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.



2. Choose your AWS Region.
3. Choose **Migration projects**. The **Migration projects** page opens.
4. Choose dm-project, and then choose **Data migrations**.
5. Choose **Create data migration**.
6. For **Name**, enter a unique name for your data migration. For example, enter postgresql-replication.
7. For **Replication type**, choose **Full load and change data capture (CDC)** to migrate your existing source data and replicate ongoing changes. For this replication type, AWS DMS deletes all data, tables, and other database objects on your target database. Make sure you create a backup of your target database before you start your data migration.
8. Select the check box for **Turn on CloudWatch logs** to store data migration logs in Amazon CloudWatch.
9. For **IAM service role**, choose the IAM role that you created in [Step 1](#).
- 10 For **Stop mode**, choose **Don't stop CDC**.
- 11 Choose **Create data migration**.

AWS DMS creates your data migration and sets its status to **Ready**. To migrate your data, you must start the data migration manually. For more information, see [Step 9](#).

## Step 9: Running and Monitoring a Data Migration

After you create a data migration, you can run it and monitor its status.

### To start a data migration

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose **Migration projects**. The **Migration projects** page opens.
3. Choose the migration project that you created in [Step 7](#).
4. On the **Data migrations** tab, choose the data migration that you created in [Step 7](#).
5. For **Actions**, choose **Start**.

The first launch of a homogeneous data migration requires some setup. AWS DMS creates a serverless environment for your data migration. This process takes up to 15 minutes.

## To monitor a data migration

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose **Migration projects**. The **Migration projects** page opens.
3. Choose the migration project that you created in [Step 7](#).
4. On the **Data migrations** tab, see the **Status** column for your data migration. For more information about values in this column, see [Statuses of homogeneous data migrations](#).
5. For a running data migration, the **Migration progress** column displays the percentage of migrated data.
6. Choose your data migration. On the **Details** tab, you can see the progress of your homogeneous data migration.

After AWS DMS completes the full load process, your data migration starts the replication of ongoing changes.

## PostgreSQL database to Amazon RDS post-migration clean-up

After you migrate your PostgreSQL database to Amazon RDS for PostgreSQL using homogeneous data migrations in AWS DMS, you can explore several other resources:

- Use DMS Fleet Advisor to inventory your source databases and discover other candidates to move to the cloud. For more information, see the [DMS Fleet Advisor User Guide](#).
- Learn more about Amazon RDS for PostgreSQL. For more information, see the [Amazon Relational Database Service User Guide](#).

After you've finished using your migration project, clean up your resources.

### To clean up your AWS DMS resources

- Sign in to the AWS Management Console and open the AWS DMS console.
- In the navigation pane, choose **Migration projects**, and choose dm-project. On the **Data migrations** tab, choose **Stop** for **Actions**.
- After AWS DMS stops your data migration, choose **Delete** for **Actions** and confirm your choice.
- Choose **Migration projects**, and choose dm-project. Choose **Delete** for **Actions** and confirm your choice.

- Choose **Instance profiles**, and choose `dm-instance-profile`. Choose **Delete** and confirm your choice.
- Choose **Data providers**, and then select the check boxes for `dm-postgresql-source-provider` and `dm-postgresql-target-provider`. Choose **Delete** and confirm your choice.
- Delete your database users that you created in [Step 2](#) and [Step 3](#).
- Drop the replication slot and the publisher in the source database by using the following code example.

```
SELECT pg_drop_replication_slot('migration_subscriber_{ARN}');  
DROP PUBLICATION publication_{ARN};
```


Also, make sure that you delete your database secrets in AWS Secrets Manager, IAM role, IAM policy, and the virtual private cloud (VPC).

## Migrating an Oracle Database to Amazon RDS for Oracle

You can use these three main approaches to migrate self-managed Oracle databases to Amazon Relational Database Service (Amazon RDS) for Oracle.

- Using a native database tool such as Oracle Data Pump.
- Using a managed service such as the AWS Database Migration Service (AWS DMS).
- Using a native tool for the full load phase and AWS DMS for ongoing replication.

This document describes the third strategy — we call this the hybrid approach. The following diagram shows the components of the hybrid approach.



The hybrid approach provides the following advantages.

- To automate the creation of secondary database objects such as views, indexes, and constraints.
- To use AWS DMS data validation to ensure your target data matches with the source, row by row and column by column.
- To use some of the other capabilities that AWS DMS provides, for example data filtering or renaming tables and columns.

This document describes the native options for the full load. It also includes a comparison so you can evaluate the options against your migration requirements. In conclusion, you can find a brief description of how to use AWS DMS for ongoing replication.

## Topics

- [Summary](#)
- [Full load Oracle database migration](#)
- [Full load Oracle database migration options performance comparison](#)
- [Migrate Oracle database with AWS DMS ongoing replication](#)

## Summary

To migrate database objects and data, use either Oracle Export/Import or Oracle Data Pump. Oracle Export/Import and Oracle Data Pump automate schema object creation. Oracle Data Pump has better performance than Oracle Export/Import and it's a newer version of Oracle Export/Import.

You can still choose to work with Oracle Export/Import for relatively small data sets which are less than 10 GB because of the ease of use. To migrate table data only, choose any native full load option described in the full load and performance comparison sections.

For full load and ongoing replication, use the hybrid approach. AWS DMS recommends using Oracle Data Pump for full load because it's faster than other tools, and it automates target object creation.

## Full load Oracle database migration

The full load phase populates the target database with a copy of the source. This chapter describes the following methods to help you choose the one that best matches your migration scenario.

- Oracle Export/Import.
- Oracle Data Pump.
- Database link.
- Oracle SQL Developer database copy.
- Oracle materialized views.
- Oracle SQL\*Loader.

We recommend that you begin by reviewing the following table to understand the tools suitable for your use case.

| Oracle native tools  | Supports metadata migration | Suitable for large number of tables | Suitable database sizes | Performance |
|----------------------|-----------------------------|-------------------------------------|-------------------------|-------------|
| Oracle Export/Import | Yes                         | Yes                                 | Less than 10 GB         | Medium      |

| Oracle native tools         | Supports metadata migration | Suitable for large number of tables | Suitable database sizes | Performance |
|-----------------------------|-----------------------------|-------------------------------------|-------------------------|-------------|
| Oracle Data Pump            | Yes                         | Yes                                 | Any size                | High        |
| Database link               | No                          | No                                  | Less than 10 GB         | Medium      |
| Oracle SQL*Loader           | No                          | No                                  | Less than 10 GB         | High        |
| Oracle materialized views   | No                          | No                                  | Less than 500 MB        | Low         |
| SQL Developer database copy | Yes                         | Yes                                 | Less than 200 MB        | Low         |

#### Note

The sizes that are provided in the table are AWS DMS recommendations based on customer migration experiences and not the limitations of the native tools.

This document doesn't discuss Oracle external tables because their use is similar to Oracle Data Pump. AWS DMS recommends using Oracle Data Pump instead.

Oracle Data Guard and RMAN are the excellent options to migrate an Oracle instance into EC2, however, Amazon RDS doesn't support these options.

You can use the [Amazon DMS Sample Database for Oracle: version 1.0](#) to run the following migration examples.

The [Performance Comparison](#) section to know the performance of these native tools.

The AWS DMS ongoing replication task requires a start position such as system change number (SCN) from the source database. AWS DMS ongoing replication task replicates changes from a position that you specify during the task configuration. If all exported objects are consistent as of same SCN or if there is no data modification after this SCN, we call it as consistent full load.

To avoid data loss and/or duplication with hybrid approach, make sure of the following

- Full load is consistent.
- You captured the SCN from the source database before you start the full load when there are no open transactions in the database.

To capture the SCN, use the following command before you start the full load.

```
select current_scn from v$database;
```

## Oracle Export/Import

Oracle Export/Import is a native database migration tool set that is provided as part of the database installation. Oracle replaced Export/Import with Oracle Data Pump, but many DBAs are familiar with Export/Import because of its long history and usage.

Export/Import may be suitable for your use case if:

- Your database size is less than 10 GB.
- You plan to migrate your database metadata, as well as table data.
- You have a relatively large number of tables to migrate.

Export/Import may not be suitable for your use case if:

- Your database size is greater than 10 GB.
- You have LOBs or other binary data values.

One limitation of Export/Import is that it performs a serial migration. If you have a large data volume and/or large objects like LOB or CLOB values, then using Export/Import may be slower than the other options.

### Example

You use Oracle `exp` to unload the source database into a dump file, and Oracle `imp` to load the dump file into the target database. You can run `exp` and `imp` on the same host to avoid copying the dump file between database servers.

For example, you can perform the following steps on the source database to migrate the `dms_sample` schema using `exp` and `imp`. The schema includes tables, views, indexes, packages, stored procedures, and other database objects.

First, capture the current SCN on the source database. Use this value to get a consistent image in the database export.

```
select current_scn from v$database;  
  
CURRENT_SCN  
86409924
```

Then, export the `dms_sample` schema using the `exp` utility. Set `flashback_scn` to the SCN value that you obtained in the previous step to get a consistent export.

```
exp <user>/<password> file=export_file.dmp owner=dms_sample log=explog_file.log  
flashback_scn=86409924 statistics=None
```

Finally, import the `export_file.dmp` dump file into the target Amazon RDS database

```
imp <user>/<password>@targetdb fromuser=dms_sample touser=dms_sample  
file=export_file.dmp log=implog_file.log
```

This command imports the dump file into the target `dms_user` schema.

## Oracle Data Pump

Oracle Data Pump is a native database migration tool set that is provided as part of the database installation. Such as `exp` and `imp`, Oracle Data Pump has separate utilities for export and import. These utilities are `expdp` and `impdp`. You can use Oracle Data Pump to migrate data and database objects such as tables, indexes, views, procedure, packages, and functions, and so on.

Oracle Data Pump may be suitable for your use case if:

- You want to automate creation of secondary database objects such as views, indexes, and constraints.
- You have a large number of tables to migrate. Because Oracle Data Pump works at different levels such as table, schema, and database level, it doesn't add additional steps or complexity with an increase of the number of tables.



Oracle Data Pump may not be suitable for your use case if you have a small database, which is less than 10 GB. Though Oracle Data Pump works well for databases of any size, for smaller data sets you have many native tools available as well.

## Example

At a high level, we can use the following steps to migrate the `dms_sample` database.

- Export data to one or more dump files.
- Move the dump file or files.
- Import the dump file or files.

## Export data

To perform a consistent export, make sure that you use the `flashback_scn` parameter in `expdp`. Run the following commands on the source database to get the current SCN, and create a database directory to store the dump and log files.

```
SQL> select current_scn from v$database;

CURRENT_SCN
60045321

SQL> create directory expdp_sample_bkp as '/u001/oraarch';
Directory created.
```

Now, use the 60045321 SCN and the `expdp_sample_bkp` directory in the following command to dump the `dms_sample` schema.

```
expdp userid=<user>/<password> directory=expdp_sample_bkp dumpfile=exp_dms_sample.dmp
logfile=exp_dms_sample.log schemas=dms_sample flashback_scn=60045321
```

## Move the dump file or files

After you export your data, you have the `exp_dms_sample.dmp` dump file. You can move this file to Amazon RDS using different methods. In this example, we use the Amazon RDS and Amazon S3 integration feature to transfer the dump file to the Amazon RDS instance. For more information, see [Amazon S3 integration](#) in the *Amazon Relational Database Service User Guide*.

The first step is to copy the dump file into the Amazon S3 bucket in the same account and region as the target database.

```
aws s3 cp exp_dms_sample.dmp s3://YourBucket
```

Now, use the `rdsadmin_s3_tasks` package to copy the file from Amazon S3 to your Amazon RDS instance. The file will be copied to the `DATA_PUMP_DIR` directory which was created as part of your Amazon RDS instance.

```
select rdsadmin.rdsadmin_s3_tasks.download_from_s3(  
    p_bucket_name => 'YourBucket',  
    p_directory_name => 'DATA_PUMP_DIR', p_s3_prefix => 'exp%')  
as task_id from dual;
```

The preceding statement returns the ID of the task. You can view the result by displaying the task's output file.

```
SELECT text FROM table(rdsadmin.rds_file_util.read_text_file('BDUMP','dbtask-*task-id*.log'));
```

Replace *task-id* with the task ID returned by the `download_from_s3`.

## Import the dump file or files

You can use the following command to import the dump file into your Amazon RDS instance.

```
DECLARE  
    v_hdn1 NUMBER;  
BEGIN  
    v_hdn1 := DBMS_DATAPUMP.OPEN( operation => 'IMPORT', job_mode => 'SCHEMA', job_name  
=> null);  
    DBMS_DATAPUMP.ADD_FILE(  
        handle      => v_hdn1,  
        filename    => 'exp_dms_sample.dmp',  
        directory   => 'DATA_PUMP_DIR',  
        filetype    => dbms_datapump.ku$_file_type_dump_file);  
    DBMS_DATAPUMP.ADD_FILE(  
        handle      => v_hdn1,  
        filename    => 'sample_imp.log',  
        directory   => 'DATA_PUMP_DIR',
```

```
filetype => dbms_datapump.ku$_file_type_log_file);  
DBMS_DATAPUMP.METADATA_FILTER(v_hdn1, 'SCHEMA_EXPR', 'IN ('DMS_SAMPLE')');  
DBMS_DATAPUMP.START_JOB(v_hdn1);  
END;
```

## Database Link

In Oracle, a database link enables you to access objects in another database. To migrate tables using a database link, you first create the database link and then run the insert and select statements for individual tables.

Using a database link may suit your use case if:

- You have a relatively small database size, which is less than 10 GB.
- you need to migrate table data only.
- You have a relatively small number of tables to migrate. For large number of tables you need to create a script to perform the migration.

Database link may not suit your use case if:

- Your database size is greater than 10 GB.
- You need to migrate schema objects other than table.

## Example

The following example migrates the `sporting_event_ticket` table from the `dms_sample` schema.

First, create a database link to the source on the target database.

```
SQL>create database link rdsmigration_link  
CONNECT TO <user> identified by <password>  
USING '(description=(address=(protocol=tcp) (host=<Self Managed Database Hostname>  
(port=<Listener Port number>)) (connect_data=(sid=<sourcedb sid>)))';
```

Use the following query to verify that the database link works correctly.

```
SQL> select sysdate from dual@rdsmigration_link;
```

Run the following query on the source to generate the DDL statement for the table. You can use the `dbms_metadata` system package to extract the DDL. After generating the DDL, create the table on the target database.

```
select dbms_metadata.get_ddl('TABLE','SPORTING_EVENT_TICKET') from dual;
```

Run the following query on target database to transfer data using the insert and select statements.

```
insert into dms_sample.sporting_event_ticket select * from  
  sporting_event_ticket@rdsmigration_link;  
Commit;
```

## Oracle SQL\*Loader

Oracle SQL\*Loader or `sqlldr` is a native database utility. Oracle provides this utility as part of the Oracle installation. Oracle SQL\*Loader loads data from flat files into an Oracle database.

Oracle SQL\*Loader may be suitable for your use case if:

- Your database size is less than 10 GB.
- You need to migrate data only.
- You have a small number of tables to migrate, the tool requires data export and control file creation for each table.

Oracle SQL\*Loader may not be suitable for your use case if:

- Your database size is greater than 10 GB.
- You need to migrate database objects along with data.

## Example

The migration process includes the following steps.

- Export the table data and create configuration files.
- Create the table on the target database.
- Load the exported files into the target database.

The following example migrates the `sporting_event_ticket` table from the `dms_sample` schema.

### Export the table data and create configuration files

Use the SQL Developer to connect to your source database. Choose **Tools**, and then choose **Database Export** to open the **Export Wizard**. Choose the following option to generate the data file, control file, and table DDL for the target database.

- Select your connection from the **Connect Panel**.
- Turn on **Export DDL**.
- Turn on **Export Data**.
- Make sure that the data format is set to **loader**.
- Save as separate files.

To export the `sporting_event_ticket` table, open the **Types of Export** page and do the following:

- Choose **Table**.
- Search for the `sporting_event_ticket` table.
- Use default options for other pages.
- Choose **Finish**.

This step exports the `sporting_event_ticket` data, generates the control file, and the table DDL in separate operating system files.

### Create the table on the target database

On the target Amazon RDS database, run the DDL script that was generated by Oracle SQL Developer to create the table.

### Load the exported files into the target database

Use SQL\*Loader from same host to import data. If required, you can copy the dump files to an Amazon EC2 instance to perform the import.

```
sqlldr userid=<user>/<password>@targetdb control=sporting_event_ticket.ctl
```

```
log=load.log bad=load.bad discard=load.dsc direct=y skip_index_maintenance=true
```

This example uses Oracle SQL Developer to generate the configuration file. If you prefer, you can use SQL\*Plus to generate the configuration file. For more information, see [Oracle SQL\\*Loader](#) in the *Amazon Relational Database Service User Guide*.

## Oracle SQL Developer Database Copy

Oracle SQL Developer is a graphical SQL client available from Oracle and can be installed on Windows, Linux, or macOS. You can use the Database Copy option in SQL Developer to transfer the data from one database to another; you can choose to copy individual objects or an entire schema.


Oracle SQL Developer database copy may be suitable for your use case if:

- You have a relatively small database size, which is less than 200 MB.
- You need to migrate both data and metadata.


### Example

Use the following steps to migrate the `sporting_event_ticket` table to the target database.

1. Install Oracle SQL Developer.
2. Open the Database Copy wizard on the Tools menu.
3. Connect to your source and target databases.
4. Select the appropriate options to migrate the tables in your source database.



5. On the next page, choose the `sporting_event_ticket` table, and then choose **Finish** to start the migration.



## Oracle Materialized Views

A *materialized view* is an object that contains the results of a query. You can use Oracle materialized views to migrate data over a database link. With Oracle materialized views, you can perform a full load migration and keep your target tables continuously in-sync with the source.

Oracle materialized views may be suitable for your use case if:

- You have a relatively small database size, which is less than 500 MB.
- You need to migrate data only.
- You have a few tables to migrate as it requires preparation for each table.



Oracle materialized views may not be suitable for your use case if You need to migrate objects other than table data.

## Example

The following example shows how to migrate the `sporting_event_ticket` table.

On the target Amazon RDS instance, create a database link to the source database.

```
create database link rdsmigration_link
CONNECT TO <user> identified by <password>
USING '(description=(address=(protocol=tcp) (host=<Self Managed Database Hostname>)
(port=<Listener Port number>)) (connect_data=(sid=<sourcedb sid>)))';
```

Test the database link to make sure you can access source database.

```
Select sysdate from dual@rdsmigration_link;
```

On source database, create a materialized view log.

```
create materialized view log on sporting_event_ticket;
```

On the target Amazon RDS instance, create a materialized view.

```
create materialized view sporting_event_ticket
build immediate refresh fast
as (select *
from dms_sample.sporting_event_ticket@rdsmigration);
```

When you're ready to switch to the new database, drop the materialized view using the `PRESERVE TABLE` clause to retain the underlying table and its contents.


```
drop materialized view sporting_event_ticket preserve table;
```

For more information, see [CREATE MATERIALIZED VIEW](#) in the Oracle documentation.

## Full load Oracle database migration options performance comparison

We analyzed the Oracle Export/Import, Oracle Data Pump, database link, and SQL\*Loader tools for their performance in a full load migration. We populated the `sporting_event_ticket` table with 10 GB of data to use as a test environment.

We expect the similar trend for larger data sets too. We didn't include Oracle materialized views or SQL Developer database copy because those tools aren't recommended for data sets larger than 1 GB.



- 18:08 minutes is the total elapsed time for Oracle Data Pump. This time includes:
  - 3:07 minutes to unload data and metadata using expdp.
  - 2:00 minutes to upload the data dump to Amazon S3 from Amazon EC2.
  - 3:01 minutes to download the data dump to Amazon RDS instance.
  - 7:00 Minutes to load data and metadata into Amazon RDS using impdp.
- 22 minutes is the total elapsed time for Oracle SQL\*Loader. This time includes:
  - 15 minutes to unload data.
  - 7 minutes to load data into the target using direct=y.
- 29 minutes is the total elapsed time for database link.
- 52 minutes is the total elapsed time for Oracle Export/Import. This time includes:
  - 14 minutes to unload data using exp.
  - 38 minutes to load data using imp.

## Migrate Oracle database with AWS DMS ongoing replication

After you complete the full load, make sure that you perform ongoing replication using AWS DMS to keep the source and target databases in sync. To configure the ongoing replication task, sign in to the AWS Management Console and follow these steps.

1. Choose **Database Migration Service**, and then choose **Database migration tasks**.
2. Choose **Create task**.
3. For **Migration type**, choose **Replicate data changes only**.
4. For **CDC start mode for source transactions**, choose **Enable custom CDC start mode**.
5. For **Custom CDC start point**, choose **Specify a log sequence number** and enter the SCN that you captured before starting the full load.

For more information, see [Continuous replication tasks](#) and [Migrate from Oracle to Amazon RDS](#).

## Migrating a SQL Server Always On Database to Amazon Web Services

Microsoft SQL Server Always On is a high-availability feature for Microsoft SQL Server databases. With the synchronous-commit secondary replica, your application remains transparent to a failover. If the primary node in the Always On Availability Group (AAG) fails due to unforeseen circumstances or due to maintenance, your applications remain unaware of the failure, and automatically redirect to a functional node. You can use AWS Database Migration Service (AWS DMS) to migrate a SQL Server Always On database to all supported target engines. AWS DMS has the flexibility to adapt to your Always On configuration, but it may be unclear how to set up the optimal AWS DMS configuration.

Using this guide, you can learn how to configure AWS DMS to migrate a SQL Server Always On database to AWS. This guide also describes specific configuration and troubleshooting issues and best practices to resolve them.

The guide includes a customer use case and covers the issues that the customer encountered when configuring AWS DMS, along with the solutions employed.

### Topics

- [Prerequisites for migrating SQL Server AlwaysOn databases to AWS](#)
- [SQL Server Always On Availability Groups](#)

## Prerequisites for migrating SQL Server AlwaysOn databases to AWS

The following prerequisites are required to complete this walkthrough:

- Understand how to work with Microsoft SQL Server as a source for AWS DMS. For information about working with SQL Server as a source, see [Using a SQL Server Database as a Source](#).
- Understand how to work with SQL Server Always On availability groups. For more information about working with SQL Server Always On availability groups, see [Working with SQL Server Always On availability groups](#)
- Understand how to run prerequisite tasks for AWS DMS, such as setting up your source and target databases. For information about prerequisites for AWS DMS, see [Prerequisites](#).
- Understand the supported features and limitations of AWS DMS. For information about AWS DMS, see [What Is Database Migration Service?](#).

For more information about AWS DMS, see the [Database Migration Service user guide](#).


## SQL Server Always On Availability Groups

Always On availability groups provide high availability, disaster recovery, and read-scale balancing. These availability groups require a cluster manager. The Always On availability groups feature provides an enterprise-level alternative to database mirroring. Introduced in SQL Server 2012 (11.x), Always On availability groups maximizes the availability of a set of user databases for an enterprise. An availability group supports a fail-over environment for a discrete set of user databases, known as availability databases, that fail over together. An availability group supports a set of read-write primary databases and sets of corresponding secondary databases. Optionally, secondary databases can be made available for read-only access and/or some backup operations.

### AWS DMS Use Case

A customer used AWS DMS to migrate data from a SQL Server 2017 source database. This database was clustered in a 4-node Always On Availability Group (AAG) configuration. The customer configured the AWS DMS source endpoint to connect directly to the IP address of the primary node of the AAG by using an IP address. With this setup, the customer used the AAG HA/DR functionality for internal applications. In this case, AWS DMS can't use the secondary database if a failover happens. The customer used the target endpoint to populate an Operational Data Store (ODS) of the Amazon RDS for SQL Server database and an Amazon Simple Storage Service (Amazon S3) data lake.

The following diagram displays the customer's existing architecture.



## Issues with This Approach

Maintenance activities (operating system patching, RDBMS patching) can cause a server failover and AWS DMS will not be able to connect to the source.

Activity and transactions continue to occur on the failover database as shown in the preceding image. Because of this, the change data capture task becomes out of sync when the cluster fails back to the primary node.


At the start of the task, AWS DMS polls all the nodes in Always On cluster for transaction backups. The AWS DMS task can also fail if transaction backup happens from any other node than the primary.

## The Solution Recommended by AWS DMS

To address connectivity design deficiencies, AWS DMS recommended to configure the AWS DMS source endpoint to connect to the AAG listener IP address or a canonical name record instead of connecting directly to the IP address of the primary node. In case of a failover, AWS DMS will

interact with the secondary databases, like any other application. Without using the AAG listener IP address, AWS DMS will not be aware of the secondary replica to connect in case of a failover.

The following diagram displays the proposed architecture.



AWS DMS recommended to set the extra connection attribute `MultiSubnetFailover=Yes` in the customer's AWS DMS endpoint. This ODBC driver attribute helps AWS DMS connect to the new primary in case of an Availability Group failover. This attribute is designed for situations when the connection is broken. In these situations, AWS DMS attempts to connect to all IP addresses associated with the AAG listener. For more information, see [Multi-subnet failovers](#).

Also, AWS DMS recommended to set the extra connection attribute `alwaysOnSharedSynchedBackupIsEnabled=false` to poll all the nodes in Always On cluster for transaction backups.

For more information on extra connection attributes for SQL Server as source, see [Extra connection attributes when using SQL Server as a source](#).

# Migrating an Amazon RDS for MySQL Database to an Amazon DynamoDB target

This walkthrough helps you to understand the process of migrating data from Amazon Relational Database Service (Amazon RDS) for MySQL to Amazon DynamoDB using AWS Database Migration Service (AWS DMS).

Amazon DynamoDB is a key-value and document database that delivers single-digit millisecond performance at any scale for modern applications. It's a fully managed, multi-region, multi-master, durable database with built-in security, backup and restore, and in-memory caching for internet-scale applications. DynamoDB can handle more than 10 trillion requests per day and can support peaks of more than 20 million requests per second. Many of the world's fastest growing businesses depend on the scale and performance of DynamoDB to support their mission-critical workloads.

Customers use DynamoDB for banking/finance, gaming, ad-tech, retail, media & entertainment workloads to build internet-scale applications supporting user-content metadata and caches. It requires high concurrency and connections for millions of users and requests, where there is a requirement for a very stringent response time. With DynamoDB, you can use design patterns for deploying shopping carts, workflow engines, inventory tracking, customer profiles, fraud detection, and leader boards, to name a few.

In this document, we will talk about a use case where a customer is running an application that handles a COVID-19 vaccination drive and stores this information in a data store. Currently, they use RDS MySQL to store vaccine data, but because of the sheer scale where data of millions of people can be getting stored at the same time, MySQL poses scalability challenges vis-à-vis response time. As business and application requirements are sensitive enough for response time in both writing data and reading it back, a relational database like MySQL cannot meet the SLA requirements. So, the customer decides to migrate to DynamoDB, which is purpose built to be performant at scale and is specifically designed to handle such use cases. The business also requires that the initial transfer of data from RDS MySQL to Amazon DynamoDB must complete within a 15-hour window.

To illustrate the process, we use AWS DMS to migrate data from an example database. AWS DMS is a managed service that helps migrate between heterogeneous sources and targets. In our case, we migrate an RDS MySQL database to Amazon DynamoDB. AWS DMS supports not only the migration of your existing data, but also ensures that the source and target are synchronized for ongoing transactions.

## Topics

- [Why use AWS DMS?](#)
- [Example data set](#)
- [Solution overview](#)
- [Prerequisites](#)
- [Step-by-step Amazon RDS for MySQL database to Amazon DynamoDB migration walkthrough](#)

## Why use AWS DMS?

When migrating from a relational database like MySQL to Dynamo DB, there are multiple approaches that you can take. One can be dumping your data using a CSV dump and loading that into Amazon DynamoDB Tables from S3. However, it comes with its own challenges in regard to size and requires taking extended downtime. AWS DMS supports binary log-based replication between MySQL based engines and Dynamo DB which can help achieve such migrations with minimal downtime. Also, Relational Database Management System (RDBMS) tables store the data in a normalized way across multiple tables. However, using DMS, you can customize the target table using the object mapping feature to denormalize the data into a single target table.

In this document, we guide you through the steps that you take to migrate the example MySQL database into Amazon DynamoDB. In the next sections, we describe the characteristics of the database. Then, we build the replication resources in AWS DMS that we use to migrate the database, paying close attention to matching the AWS DMS configuration with our particular use case.

## Example data set

In this walkthrough, the following is the table information that is used to store the vaccine drive data. As it can be noted that the schema does not completely play out the relational model of normalization, and all data are stored in a single table in a de-normalized way.



| vaccine_drive_stats ▼ |                               |
|-----------------------|-------------------------------|
| 🔑                     | user_id INT(11)               |
| ◇                     | patient_name VARCHAR(1000)    |
| ◇                     | phone_num INT(11)             |
| ◇                     | date_of_birth DATE            |
| ◇                     | age TINYINT(4)                |
| ◇                     | date_vacc1 DATE               |
| ◇                     | date_vacc2 DATE               |
| ◇                     | date_booster DATE             |
| ◇                     | fully_vaccinated BIT(64)      |
| ◇                     | age_group VARCHAR(50)         |
| ◇                     | state VARCHAR(1000)           |
| ◇                     | zipcode INT(11)               |
| ◇                     | gender VARCHAR(50)            |
| ◇                     | city VARCHAR(50)              |
| 🔑                     | area_code VARCHAR(200)        |
| ◇                     | vaccine_type VARCHAR(300)     |
| ◇                     | vaccine_name VARCHAR(100)     |
| ◇                     | rural_or_urban VARCHAR(100)   |
| ◇                     | certificate_link VARCHAR(100) |
| ◇                     | vaccinated_by VARCHAR(200)    |
| ◇                     | vaccinated_at VARCHAR(100)    |
| ◇                     | next_due_date DATE            |
| Indexes ▼             |                               |
| PRIMARY               |                               |

Generally, relational tables are used to fetch a fixed data set based on the table definition. However, in this use case, we define the tables in a de-normalized manner, and going forward based on the business requirement schema, growth can be exponential in rate and dynamic in nature. Services like Amazon DynamoDB help application developers and architects to rethink the data model in a key-value format for such use cases, and plan to move the data store on DynamoDB.

The “vaccine\_drive\_stats” table contains 1022 million records with a size of 210 GB. This table mainly collects the information for people who participated in the vaccination program, including their vaccine status and user details.

Note that the table contains composite primary keys for the “user\_id” and “area\_code” columns. In MySQL, the application and admin user accesses the data using composite keys for reporting and manipulating the records in the tables.

There are additional use cases to get aggregate data , such as the total number of people who have received the first or second vaccination, state-wise vaccine numbers, total percentage of the population receiving vaccination, etc. All of these aggregate use cases can be handled using a DynamoDB schema designed to cater to aggregations.


Migration of this use case can be handled using one-to-one mapping from RDBMS MySQL to a DynamoDB table.

Similarly, if you have the following types of tables, you can consider migrating to a DynamoDB target using DMS with less downtime.

1. Table with non-relational data
2. Logging tables
3. User preference tables
4. Application Session state tables

## Solution overview

The following diagram displays a high-level architecture of the solution, where we use AWS DMS to move data from a MySQL database hosted on RDS to Amazon DynamoDB.



To connect to the source database where your data resides and target Amazon DynamoDB, you will create two endpoint resources in AWS DMS. An “endpoint” is a resource for storing connection information such as hostname, username, and password. For DynamoDB, it stores an IAM role name that provides access to resources. Endpoint resources also store unique settings for each endpoint to configure the endpoint behavior.

The endpoint itself does not have a mechanism to connect to the source or target. A resource called a “replication task” connects to the source and target to migrate data. One source and target endpoint can be associated with single replication task. Tasks can use source and target endpoints, which are used by other tasks.

A replication instance is a resource where your replication task is running. It has a network interface connected to your VPC, through which AWS DMS tasks communicate with sources and targets.

In summary, in this walkthrough you will set up the following resources in AWS DMS

- **Replication Instance** — An AWS managed instance that hosts the AWS DMS engine. You control the type or size of the instance based on your workload.
- **Source Endpoint** — A resource that provides connection details, data store type, and credentials to connect to a source database. For this use case, we will configure the source endpoint to point to the Amazon RDS for MySQL database.
- **Target table** - A DynamoDB table used on this scenario to consume the data from the Source database. We will create a DynamoDB table with customized settings for migration.
- **Target Endpoint** — AWS DMS supports several target systems including Amazon RDS, Amazon Aurora, Amazon Redshift, Amazon Kinesis Data Streams, Amazon S3, and more. For this use case, we will configure Amazon Dynamo DB as the target endpoint.
- **Replication Task** — A resource that runs on the replication instance and connects to endpoints to replicate data from the source to the target.

## Prerequisites

The following prerequisites are required to complete this walkthrough:

- An understanding of Amazon Relational Database Service (Amazon RDS), the applicable database technologies, and SQL.
- A user with AWS Identity and Access Management (IAM) credentials that allows you to launch Amazon RDS and AWS Database Migration Service (AWS DMS) instances in your AWS Region. For information about IAM credentials, see [Create an IAM user](#).

- An understanding of the Amazon Virtual Private Cloud (Amazon VPC) service and security groups. For information about using Amazon VPC with Amazon RDS, see [Amazon Virtual Private Cloud \(VPCs\) and Amazon RDS](#). For information about Amazon RDS security groups, see [Controlling access with security groups](#).
- An understanding of the supported features and limitations of AWS DMS. For information about AWS DMS, see [What is Database Migration Service?](#)
- An understanding of how to work with MySQL as a source and Amazon DynamoDB as a target. For information about working with MySQL as a source, see [Using an MySQL database as a source](#). For information about working with Amazon DynamoDB as a target, see [Using Amazon DynamoDB as a target](#).
- An understanding of the supported data type conversion options for MySQL and Amazon DynamoDB. For information about data types for MySQL as a source, see [Source data types for MySQL](#). For information about data types for Amazon DynamoDB as a target, see [Target data types for Amazon DynamoDB](#).

For more information about AWS DMS, see [Getting started with Database Migration Service](#).

## Step-by-step Amazon RDS for MySQL database to Amazon DynamoDB migration walkthrough

The following steps provide instructions for migrating an Amazon RDS for MySQL database to DynamoDB. These steps assume that you have already prepared your source database as described previously.

Step 1: [Create an AWS DMS Replication Instance](#)

Step 2: [Configure a Source Amazon RDS for MySQL Database](#)

Step 3: [Create an AWS DMS Source Endpoint](#)

Step 4: [Configure a Target Amazon DynamoDB table](#)

Step 5: [Configure an AWS DMS Target Endpoint](#)

Step 6: [Create an AWS DMS Task](#)

Step 7: [Run the AWS DMS Task](#)

## Step 1: Create replication instance

An AWS DMS replication instance hosts the software that migrates data between the source and target. The replication instance also caches the transaction logs during the migration. The CPU and memory capacity of the replication instance influences the overall time needed for the migration. Make sure that you consider the specifics of your particular use case when you determine the size of your replication instance. A full load task consumes a lot of memory if it is run multithreaded. For more information, see [Choosing the right replication instance for your migration](#).

For our use case, we have a limited time window of 15 hours to complete the full load, and the table that includes 210 GB of data. Our goal is to fit into the 10-hour window. Therefore, we scale the replication instance to accommodate these requirements.

Each type of instance class has a different CPU, memory, and I/O capacity. Sizing the replication instance should be based on factors such as data volume, transaction frequency, large objects (LOBs) within storage of the data migration, and so on. We initially chose a DMS `dms.c5.large` instance running the latest AWS DMS engine version and default task configuration. We then upgraded to a `dms.c5.12xlarge` instance with a customized task configuration to see the performance differences. We will discuss the performance and configuration details in an upcoming section.

We also upgraded the storage of the replication instance to 200 GB, and as a result, 600 IOPS were available for our replication instance. By default, DMS allocates 50 GB of storage to a replication instance. This may not be sufficient for use cases where more tasks are running on same replication instance or when running tasks with parallel load for large tables. With 600 IOPS, we saved several minutes of migration time. For more information about storage volume performance and burst I/O credits, see [General Purpose SSD \(gp2\) volumes](#).

Because we replicate production data in this walkthrough, we use the Multi-AZ deployment option for our replication instance for high availability. Also, we didn't make this replication instance publicly accessible for additional security. For information about best practices for using AWS DMS, see [Database Migration Service Best Practices](#).

To create an AWS DMS replication instance, do the following:

1. Sign in to the AWS Management Console, and open the [AWS DMS console](#).
2. If you are signed in as an AWS Identity and Access Management (IAM) user, you must have the appropriate permissions to access AWS DMS. For more information about the permissions required, see [IAM permissions](#).

3. On the Welcome page, choose **Create replication instance** to start a database migration.
4. On the **Create replication instance** page, specify your replication instance information.

| For this parameter             | Do the following                                                                                                                                                                                                     |
|--------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Name</b>                    | Enter <b>mysql-to-ddb-migration-ri</b> . If you are using multiple replication servers or sharing an account, choose a name that helps you quickly differentiate between the different servers.                      |
| <b>Description</b>             | Enter <b>Migrate MySQL to Amazon DynamoDB</b> .                                                                                                                                                                      |
| <b>Instance class</b>          | Choose <b>dms.c5.12xlarge</b> . Each size and type of instance class has increasing CPU, memory, and I/O capacity.                                                                                                   |
| <b>Engine version</b>          | Leave the default value, which is the latest stable version of the AWS DMS replication engine.                                                                                                                       |
| <b>Allocated storage (GiB)</b> | Choose <b>200 GiB</b> .                                                                                                                                                                                              |
| <b>VPC</b>                     | Choose the virtual private cloud (VPC) in which your replication instance will launch. Select the same VPC in which your source is placed.                                                                           |
| <b>Multi AZ</b>                | In this scenario, choose <b>No</b> . If you choose <b>Yes</b> , AWS DMS creates a second replication server in a different Availability Zone for failover if there is a problem with the primary replication server. |
| <b>Publicly accessible</b>     | Choose <b>Yes</b> . If either your source or target database resides outside of the VPC in which your replication server resides, you must make                                                                      |

your replication server policy publicly accessible.

## 5. Choose **Create**.

## Step 2: Configure a Source Amazon RDS for MySQL Database

Before setting up AWS DMS resource, there are some setups are required to configure your MySQL DB instances as a source for AWS DMS. As you know, in this walkthrough we are using a MySQL database on Amazon RDS, so DMS MySQL required prerequisites has to be enabled at the instance parameter group.

### Binary logging and its retention

To use AWS DMS change data capture (CDC), enable binary logging on the source MySQL RDS instance. To enable binary logs for RDS for MySQL and for RDS for MariaDB, enable automatic backups at the instance level. For more information about setting up automatic backups, see [Working with automated backups](#) in the *Amazon RDS User Guide*.

Next, the following parameters must be configured on the parameter group used by the source database. You can't modify a default parameter group. If the database instance is using a default parameter group, create a new parameter group and associate it with the database instance. After you perform these steps, you must reboot the database instance for your changes to apply.

The following parameters are dynamic types, so a custom parameter group with the below values doesn't require an instance reboot.

```
binlog_format=ROW
```

```
binlog_checksum=NONE
```

```
binlog_row_image=FULL
```

To ensure that binary logs are available to AWS DMS, you should increase the length of time that the logs remain available in the database instance host. For example, to increase the log retention to 24 hours, execute the following procedure call on the source database.

```
call mysql.rds_set_configuration('binlog retention hours', 24);
```

Also, it is recommended to retain the binary logs until the task completes the full load phase, and runs the CDC phase with less latency. In the planning phase, test your workload, and based on that, retain the logs for the production migration.

## Source User Permission

You must have an account for AWS DMS that has the Replication Admin role. The role needs the following privileges to run the CDC task.

**REPLICATION CLIENT** – This privilege is required for CDC tasks only. In other words, full-load-only tasks don't require this privilege.

**REPLICATION SLAVE** – This privilege is required for CDC tasks only. In other words, full-load-only tasks don't require this privilege.

The AWS DMS user must also have SELECT privileges for the source tables designated for replication.

## Network configuration

In this walkthrough, the DB instance and the replication instance are placed in the same VPC and the same subnet, so all you need to do is configure security groups, network ACLs, and route tables so that these Amazon RDS for MySQL DB instances and AWS DMS replication instances can communicate within the subnet. If you have source databases in different subnets or different VPCs, you need to configure your network to allow communication between the Amazon RDS for MySQL DB instance and the AWS DMS replication instance. For more information about network setup, see [Network configurations for database migration](#) in the *DMS user Guide*.

## Inbound connection rule

To ensure that the replication instance can access the server and the port for the database, you need to make changes to the relevant security groups and network access control lists. AWS DMS only requires access to the MySQL database listener port (3306). Also, the connection is always from the AWS DMS replication instance to MySQL. Therefore, allow connections from the replication instance to the ingress rule of the security group attached to the DB instance. We recommend you to add the complete subnet group range in the ingress rule, because the AWS DMS replication instance is a managed service and the IP address may change automatically.



## Step 3: Create an AWS DMS Source Endpoint

After you complete the network configurations, you can create a source endpoint. To create a source endpoint, do the following:

1. Open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose **Endpoints**.
3. Choose **Create endpoint**.
4. On the **Create endpoint** page, enter the following information.

| Parameter                             | Value                                                      |
|---------------------------------------|------------------------------------------------------------|
| <b>Endpoint type</b>                  | Choose <b>Source endpoint</b>                              |
| <b>Endpoint identifier</b>            | Enter <b>mysql-source-dms-datastore</b>                    |
| <b>Source engine</b>                  | Choose <b>MySQL</b> .                                      |
| <b>Access to endpoint database</b>    | Choose <b>Provide access information manually</b> .        |
| <b>Server name</b>                    | Enter the MySQL Database host amazon ec2 instance IP       |
| <b>Port</b>                           | Enter <b>3306</b> .                                        |
| <b>Secure Socket Layer (SSL) mode</b> | Choose <b>none</b> .                                       |
| <b>User name</b>                      | Enter <b>dms_user</b> .                                    |
| <b>Password</b>                       | Enter the password that you created for the dms_user user. |

## Step 4: Configure a Target Amazon DynamoDB table

A DMS task can create a target DynamoDB table based on the source table definition. When AWS DMS sets DynamoDB parameter values for a migration task, the default Read Capacity Units (RCU) parameter value is set to 200. The Write Capacity Units (WCU) parameter value is also set, but its value depends on several other settings:

- The default value for the WCU parameter is 200.
- If the `ParallelLoadThreads` task setting is set greater than 1 (the default is 0), then the WCU parameter is set to 200 times the `ParallelLoadThreads` value.

In this case, DMS uses the default provisioned capacity, which will not be sufficient to handle the workload from Source database. To avoid a resource constraint issue, and to customize the key usage, consider creating the target DynamoDB table with the configuration required by your workload.

In this walkthrough, use the below source MySQL table definition to create the target DynamoDB table. As you can see below, the source table contains composite primary keys (`user_id`, `area_code`), so you can use these fields to create a DynamoDB table with a partition key and a sort key.

```
Table: vaccine_drive_stats
Create Table: CREATE TABLE `vaccine_drive_stats` (
  'user_id' int(11) NOT NULL AUTO_INCREMENT,
  'patient_name' varchar(1000) DEFAULT NULL,
  'phone_num' int(11) DEFAULT NULL,
  'date_of_birth' date DEFAULT NULL,
  'age' tinyint(4) DEFAULT NULL,
  'date_vacc1' date DEFAULT NULL,
  'date_vacc2' date DEFAULT NULL,
  'date_booster' date DEFAULT NULL,
  'fully_vaccinated' bit(64) DEFAULT NULL,
  'age_group' varchar(50) DEFAULT NULL,
  'state' varchar(1000) DEFAULT NULL,
  'zipcode' int(11) DEFAULT NULL,
  'gender' varchar(50) DEFAULT NULL,
  'city' varchar(50) DEFAULT NULL,
  'area_code' varchar(200) NOT NULL,
  'vaccine_type' varchar(300) DEFAULT NULL,
  'vaccine_name' varchar(100) DEFAULT NULL,
  'rural_or_urban' varchar(100) DEFAULT NULL,
  'certificate_link' varchar(100) DEFAULT NULL,
  'vaccinated_by' varchar(200) DEFAULT NULL,
  'vaccinated_at' varchar(100) DEFAULT NULL,
  'next_due_date' date DEFAULT NULL,
  PRIMARY KEY ('user_id','area_code')
) ENGINE=InnoDB AUTO_INCREMENT=13462359 DEFAULT CHARSET=utf8mb4;
```

To create an Amazon DynamoDB table, do the following.

1. Open the DynamoDB console at <https://console.aws.amazon.com/dynamodb/>.
2. Choose **Create Table**. In the **Create DynamoDB table** screen, do the following:
3. On the Table name box, enter the name of the table as “vaccine\_drive\_stats\_tab”.

 **Note**

The target table can be renamed as per your requirements, but make sure to map the table name using a DMS object mapping rule.

The Dynamo DB sort/partition key for a table should be picked based on the table access patterns. DMS has the limitation in the CDC phase that DynamoDB doesn't allow updates to the primary key attributes. This restriction is important when using ongoing replication with change data capture (CDC) because it can result in unwanted data on the target.

Depending on how you have the object mapping, a CDC operation that updates the primary key can do one of two things: It can either fail, or insert a new item with the updated primary key and incomplete data. So, choose the partition key and sort key carefully to avoid issues in the migration.

For the Primary key, do the following:

DynamoDB query performance depends on the partition key and sort key selection for a table. So, choose a high cardinality column as the partition key to distribute the data across partitions in a DDB table. The sort key is used to sort and order items in a partition internally at the DDB table level. So, choose a sort key that collects related information together in one partition area, so that query performance can be improved. In this use case, we have chosen user\_id as the partition key and "area\_code" as the sort key to distribute and organize the data based on the application access pattern. Refer [Choosing the Right DynamoDB Partition Key](#) for more details.

4. In the Partition key box, enter column name as “user\_id” and set the data type to String.
5. Choose To add sort key.
6. In the Sort key box, enter column name as “area\_code” and set the data type to String.
7. In table settings, choose Customize Settings and then select On-Demand Read/Write capacity
8. When the settings are as you want them, choose Create.

In this Walkthrough, we are pre-creating the target table with On-demand capacity mode for migration. Later, based on the traffic flow, you can change the capacity mode on the target to save costs after the migration completes. For more information, see [Amazon DynamoDB create table](#).

## Step 5: Configure an AWS DMS Target Endpoint

Before you begin to work with a DynamoDB database as a target for AWS DMS, make sure that you create an IAM role. This IAM role should allow AWS DMS to assume the application role, and grants access to the DynamoDB tables that are being migrated into. The minimum set of access permissions is shown in the following IAM policy.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "",
      "Effect": "Allow",
      "Principal": {
        "Service": "dms.amazonaws.com"
      },
      "Action": "sts:AssumeRole"
    }
  ]
}
```

DMS creates the control tables “awsdms\_apply\_exceptions” and “awsdms\_full\_load\_exceptions” on the DynamoDB target to record the failures in loading/applying the records in the migration. So, the role that you use for the migration to DynamoDB must have the following permissions, including for control tables.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "dynamodb:PutItem",
        "dynamodb:CreateTable",
        "dynamodb:DescribeTable",
        "dynamodb>DeleteTable",
        "dynamodb>DeleteItem",
        "dynamodb:UpdateItem"
      ]
    }
  ]
}
```

```

    ],
    "Resource": [
      "arn:aws:dynamodb:us-west-2:111122223333:table/name1",
      "arn:aws:dynamodb:us-west-2:111122223333:table/OtherName*",
      "arn:aws:dynamodb:us-west-2:111122223333:table/awsdms_apply_exceptions",
      "arn:aws:dynamodb:us-west-2:111122223333:table/awsdms_full_load_exceptions"
    ],
  },
  {
    "Effect": "Allow",
    "Action": [
      "dynamodb:ListTables"
    ],
    "Resource": "*"
  }
]
}

```

To create a target endpoint for Amazon DynamoDB, do the following:

1. Open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose **Endpoints**.
3. Choose **Create endpoint**.
4. On the **Create endpoint** page, enter the following information.

| Parameter                      | Value                                      |
|--------------------------------|--------------------------------------------|
| <b>Endpoint type</b>           | Choose <b>Target endpoint</b>              |
| <b>Endpoint identifier</b>     | Enter <b>dynamodb-target-dms-datastore</b> |
| <b>Target engine</b>           | Choose <b>Amazon DynamoDB</b> .            |
| <b>Service access role ARN</b> | Provide the IAM role ARN created above     |

## Step 6: Create DMS Task

Before you create the replication task, it is important to understand the workload on the source database, and the usage pattern of the tables being replicated. This helps plan an effective

migration approach, and minimizes any configuration or workload related issues. In this section, we first review the important considerations, and then learn how to configure our walkthrough DMS task accordingly by applying table mappings and task settings.

## Considerations Before Creating an AWS DMS Task

### Size and number of records

The volume of migrated records affects the full load completion time. It is difficult to predict the full load time upfront, but testing with a replica of a production instance should provide a baseline. Use this estimate to decide whether you should parallelize full load by using multiple tasks or by using the parallel load option.

DMS supports parallel load threads for a target DynamoDB endpoint. However, other features such as parallel-load table level mapping aren't supported for a target Dynamo DB endpoint.

**ParallelLoadThreads** – Use this option to specify the number of threads that AWS DMS uses to load each table into its DynamoDB target table. The default value is 0 (single-threaded). The maximum value is 200. You can contact support to have this maximum limit increased.

**ParallelLoadBufferSize** – Use this option to specify the maximum number of records to store in the buffer that the parallel load threads use to load data to the DynamoDB target. The default value is 50. The maximum value is 1,000. Use this setting with **ParallelLoadThreads**.

**ParallelLoadBufferSize** is valid only when there is more than one thread. **ParallelLoadThreads** related settings responsible for only loading the data to target table using multiple threads. However, it doesn't help to unload the source data in parallel.

To speed up the full load of large tables such as “vaccine\_drive\_stats” table in our use case, we can increase the number of parallel load threads in a task.

### Transactions per second

While full load is affected by the number of records, the ongoing replication performance relies on the number of transactions on the source MySQL database. Performance issues during change data capture (CDC) generally stem from resource constraints on the source database, replication instance, target database, and network bandwidth or throughput. Knowing average and peak TPS on the source and recording CDC throughput and latency metrics helps baseline AWS DMS performance and identify an optimal task configuration. For more information, see [Replication task metrics](#).

In this walkthrough, the source database is an RDS MySQL database where transaction volume depends on number of people attending the vaccination drive. So, a considerable amount of read and write traffic is expected during the day on the Source RDS MySQL database during the migration.

This approach requires a replication instance with higher compute capacity if the data volume is huge. We chose the compute-intensive c5 class replication instance to speed up the process.

If you are not sure about your data volumes or performance expectations from the migration task, start with general t3-class instances, and then migrate to c5-class instances for compute-intensive tasks, or r5-class instances for memory intensive tasks. You should monitor the task metrics continuously, and choose the appropriate instance class that best suits your needs.

### Unsupported data types

Identify data types used in tables and check that AWS DMS supports these data types. For more information, see [Source data types for MySQL](#).

Validate that the target DynamoDB has the corresponding data types. For more information, see [Target data types for DynamoDB](#).

After you run the initial load test, validate that AWS DMS converted data as you expected. You can also initiate a pre-migration assessment to identify any unsupported data types in the migration scope. For more information, see [Specifying individual assessments](#).

### Source filtering in full load phase

Running AWS DMS replication tasks for large tables can add to the workload on the source database especially during the full load phase when AWS DMS reads whole tables from source database without any filters to restrict rows. When you use filters in AWS DMS task table mapping, confirm that appropriate indexes exist on the source tables and indexes are actually being used during full load. Regularly monitor the source database to identify any workload related issues. For more information, see [Using table mapping to specify task settings](#).

In this walkthrough, we migrate one large source table (a table with 1 billion records and 210 GB in size) with the DMS default configuration to migrate the existing data to check the performance. Based on the full load run time and resource utilization metrics on the source MySQL database instance and replication instance, we used the AWS DMS parallel load option to further improve full load performance.

## Task configuration

In this walkthrough, we migrate the existing and incremental changes to the target DynamoDB. To do so, we use the Full Load + CDC option. For more information about the AWS DMS task creation steps and available configuration options, see [Creating a task](#).

We will first focus on the following settings.

### LOB Settings

DMS considers source MySQL data types such as JSON, LONGTEXT, MEDIUMTEXT as LOB fields during migration. AWS DMS handles large binary object (LOB) columns differently compared to other data types. For more information, see [Migrating large binary objects \(Lobs\)](#).

A detailed explanation of LOB handling by AWS DMS is out of scope for this walkthrough. However, remember that increasing the LOB Max Size increases the task's memory utilization. Because of that, we recommended that you don't set LOB Max Size to a large value. For more information about LOB settings, see [Task Configuration](#).

For this use case, the source table doesn't contain any large object data types, so we decided to disable LOB settings in the task "TargetMetadata" configuration. Refer to the below task setting for more details.

```
{
  "TargetMetadata": {
    "TargetSchema": "",
    "SupportLobs": false,
    "FullLobMode": false,
    "LobChunkSize": 0,
    "LimitedSizeLobMode": false,
    "LobMaxSize": 0,
    "InlineLobMaxSize": 0,
    "LoadMaxFileSize": 0,
    "ParallelLoadThreads": 0,
    "ParallelLoadBufferSize": 0,
    "BatchApplyEnabled": false,
    "TaskRecoveryTableEnabled": false,
    "ParallelLoadQueuesPerThread": 0,
    "ParallelApplyThreads": 0,
    "ParallelApplyBufferSize": 0,
    "ParallelApplyQueuesPerThread": 0
  },
}
```



```
}
```

DMS has the following limitations in migrating large objects. If you have source table with large objects, check the respective source database DMS documentation for support scope, and based on that, configure the migration task.

- AWS DMS doesn't support LOB data unless it is a CLOB. AWS DMS converts CLOB data into a DynamoDB string when migrating the data.

## Table Object mappings

DMS has the following limitations for a target DynamoDB endpoint.

- AWS DMS only supports replication of tables with non-composite primary keys. The exception is if you specify an object mapping for the target table with a custom partition key or sort key, or both.

For this use case, the source MySQL table contains a composite primary key. Initially, we tried migrating the composite primary key table with a target prep mode of "DROP and CREATE" with only a DMS selection mapping rule. However, the table got suspended from the migration with following error, as mentioned in the limitations section prior:

```
00019383: 2023-03-14T08:48:33 [TARGET_LOAD ]E: Table 'vaccine_drive_stats' has
  composite primary key [1025900] (dynamodb_imp.c:368)
00019383: 2023-03-14T08:48:33 [TARGET_LOAD ]E: Unable to determine hash key for table
  'vaccine_drive_stats' [1025900] (dynamodb_table_requests.c:399)
00019383: 2023-03-14T08:48:33 [TARGET_LOAD ]E: Failed to initialize create table
  request. [1020413] (dynamodb_table_requests.c:92)
00019383: 2023-03-14T08:48:33 [TARGET_LOAD ]E: Handling new table
  'valis'. 'vaccine_drive_stats' failed [1020413] (endpointshell.c:2712)
00019382: 2023-03-14T08:48:33 [SOURCE_UNLOAD ]I: Unload finished for table
  'valis'. 'vaccine_drive_stats' (Id = 1). 20970 rows sent. (streamcomponent.c:3543)
00019374: 2023-03-14T08:48:33 [TASK_MANAGER ]W: Table
  'valis'. 'vaccine_drive_stats' (subtask 1 thread 1) is suspended
  (replicationtask.c:2550)
```

To mitigate this issue, we created the target table as mentioned in Step 4, and then configured the task with the following object mapping rule. In our case, we used the "map-record-to-record"

option to restructure the target table and its data storing method. Refer to the source table "vaccine\_drive\_stats" definition with the following object mapping for more clarity.

```
{
  "rules": [
    {
      "rule-type": "selection",
      "rule-id": "1",
      "rule-name": "1",
      "object-locator": {
        "schema-name": "valis",
        "table-name": "vaccine_drive_stats"
      },
      "rule-action": "include"
    },
    {
      "rule-type": "object-mapping",
      "rule-id": "2",
      "rule-name": "2",
      "rule-action": "map-record-to-record",
      "object-locator": {
        "schema-name": "valis",
        "table-name": "vaccine_drive_stats"
      },
      "target-table-name": "vaccine_drive_stats_tab",
      "mapping-parameters": {
        "partition-key-name": "user_id",
        "sort-key-name": "area_code",
        "exclude-columns": [
          "patient_name",
          "phone_num",
          "date_of_birth",
          "age",
          "date_vacc1",
          "date_vacc2",
          "date_booster",
          "fully_vaccinated",
          "age_group",
          "state",
          "zipcode",
          "gender",
          "city",
          "vaccine_type",

```

```

        "vaccine_name",
        "rural_or_urban",
        "certificate_link",
        "vaccinated_by",
        "vaccinated_at",
        "next_due_date"
    ],
    "attribute-mappings": [
        {
            "target-attribute-name": "user_id",
            "attribute-type": "scalar",
            "attribute-sub-type": "string",
            "value": "${user_id}"
        },
        {
            "target-attribute-name": "area_code",
            "attribute-type": "scalar",
            "attribute-sub-type": "string",
            "value": "${area_code}"
        },
        {
            "target-attribute-name": "rural_or_urban",
            "attribute-type": "scalar",
            "attribute-sub-type": "string",
            "value": "${rural_or_urban}"
        },
        {
            "target-attribute-name": "PatientDetails",
            "attribute-type": "scalar",
            "attribute-sub-type": "string",
            "value": "{\"patient_name\": \"${patient_name}\", \"phone_num\": \"${phone_num}\", \"date_of_birth\": \"${date_of_birth}\", \"age\": \"${age}\", \"gender\": \"${gender}\", \"state\": \"${state}\", \"zipcode\": \"${zipcode}\"}"
        },
        {
            "target-attribute-name": "PatientVaccineinfo",
            "attribute-type": "scalar",
            "attribute-sub-type": "string",
            "value": "{\"date_vacc1\": \"${date_vacc1}\", \"date_vacc2\": \"${date_vacc2}\", \"date_booster\": \"${date_booster}\", \"fully_vaccinated\": \"${fully_vaccinated}\", \"vaccine_type\": \"${vaccine_type}\", \"vaccine_name\": \"${vaccine_name}\", \"certificate_link\": \"${certificate_link}\"}"
        }
    ]
}

```

```

        "target-attribute-name": "PatientVaclocation",
        "attribute-type": "scalar",
        "attribute-sub-type": "string",
        "value": "{\"vaccinated_by\": \"${vaccinated_by}\", \"vaccinated_at\": \"${vaccinated_at}\", \"next_due_date\": \"${next_due_date}\"}"
    }
  ]
}
]
}

```

In this case, the source table contains 22 columns in total, but by using object mapping, we restructured the total number of columns to 6, and concatenated other fields into new columns, as mentioned following. Similarly, you can restructure the target based on your requirements using the object mapping feature. For more information, see [Using object mapping to migrate data to DynamoDB](#). The following DynamoDB console screenshot shows the records in the table. As you can see, DMS migrated the records based on object mapping configuration.

| Items returned (50)      |         |              |                                                                                         |                                                                             |                                                                           |                | Actions | Create Item |
|--------------------------|---------|--------------|-----------------------------------------------------------------------------------------|-----------------------------------------------------------------------------|---------------------------------------------------------------------------|----------------|---------|-------------|
|                          | user_id | area_code    | PatientDetails                                                                          | PatientVaccineInfo                                                          | PatientVaclocation                                                        | rural_or_urban |         |             |
| <input type="checkbox"/> | 3635    | areacdareacd | { "patient_name": "ramsedramsedramsed", "phone_num": "83886", "date_of_birth": "202..." | { "date_vacc1": "2023-03-13", "date_vacc2": "2023-03-13", "date_booster..." | { "vaccinated_by": "GOVT", "vaccinated_at": "HOSPITAL", "next_due_dat..." | URBEN          |         |             |
| <input type="checkbox"/> | 228     | areacdareacd | { "patient_name": "ramsedramsedramsed", "phone_num": "83886", "date_of_birth": "202..." | { "date_vacc1": "2023-03-13", "date_vacc2": "2023-03-13", "date_booster..." | { "vaccinated_by": "GOVT", "vaccinated_at": "HOSPITAL", "next_due_dat..." | URBEN          |         |             |
| <input type="checkbox"/> | 9990    | areacdareacd | { "patient_name": "ramsedramsedramsed", "phone_num": "83886", "date_of_birth": "202..." | { "date_vacc1": "2023-03-13", "date_vacc2": "2023-03-13", "date_booster..." | { "vaccinated_by": "GOVT", "vaccinated_at": "HOSPITAL", "next_due_dat..." | URBEN          |         |             |
| <input type="checkbox"/> | 1668    | areacdareacd | { "patient_name": "ramsedramsedramsed", "phone_num": "83886", "date_of_birth": "202..." | { "date_vacc1": "2023-03-13", "date_vacc2": "2023-03-13", "date_booster..." | { "vaccinated_by": "GOVT", "vaccinated_at": "HOSPITAL", "next_due_dat..." | URBEN          |         |             |

## Parallel load configuration

High values for `ParallelLoadThreads` cause heavy write traffic on the target DynamoDB tables. In such a scenario, you might find an increase in throttling events even in On-demand capacity mode. While increasing the setting value, monitor the target table's monitoring graph and make sure that no throttling events occur.

In our use case, the task is initially configured to use 200 for the `ParallelLoadThreads` setting. However, the task experienced the following DynamoDB throttling error. To avoid this DynamoDB error, we reduced the values from 200 to 150 to avoid having high throttling write events on the target table. After this change, the number of throttling events was reduced to zero on target table. For more information about throttling, see [Why is my on-demand DynamoDB table being throttled?](#)

```
00143766: 2023-03-15T05:26:07 [SOURCE_CAPTURE ]E: PutItem failed with error: Throughput
exceeds the current capacity of your table or index. DynamoDB is automatically scaling
your table or index so please try again shortly. If exceptions persist, check if you
have a hot key: https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/bp-
partition-key-design.html. [1001788] (ddb_item_actions.cpp:78)
00143766: 2023-03-15T05:26:07 [TARGET_LOAD ]E: Encountered a non-data error. Thread is
exiting. [1025906] (dynamodb_load.c:83)
```

Task setting used for parallel load configuration:

```
{
  "TargetMetadata": {
    "TargetSchema": "",
    "SupportLobs": false,
    "FullLobMode": false,
    "LobChunkSize": 0,
    "LimitedSizeLobMode": false,
    "LobMaxSize": 0,
    "InlineLobMaxSize": 0,
    "LoadMaxFileSize": 0,
    "ParallelLoadThreads": 150,
    "ParallelLoadBufferSize": 1000,
    "BatchApplyEnabled": false,
    "TaskRecoveryTableEnabled": false,
    "ParallelLoadQueuesPerThread": 0,
    "ParallelApplyThreads": 0,
    "ParallelApplyBufferSize": 0,
    "ParallelApplyQueuesPerThread": 0
  },
}
```

## Other task settings

Choose **Enable CloudWatch Logs** to upload the AWS DMS task run log to Amazon CloudWatch. You can use these logs to troubleshoot issues, because they include error and warning messages, start and end times of the run, configuration issues, and so on. To diagnose performance issues, you can change the task logging setting, such as to enable debugging or tracing.

**Note** Cloud Watch log usage is charged at standard rates. For more information, see [Amazon CloudWatch pricing](#).

For **Target table preparation mode**, choose one of the following options: **Do nothing**, **Truncate**, or **Drop**. Use **Truncate** in data pipelines where the downstream systems rely on a fresh dump of clean data and do not rely on historical data. In our use case, the truncate option doesn't support DynamoDB as a target. In this walkthrough, we choose **Do nothing** because we pre-created the target table as per the use case requirements.

For **Maximum number of tables to load in parallel**, enter the number of parallel threads that AWS DMS initiates during the full load. You can increase this value to improve the full load performance and minimize the load time when you have numerous tables. In this walkthrough, we use the default value of 8 because the task is only migrating one source table.

For **Commit rate during full load**, enter a value to indicate the maximum number of records that can be transferred together to the target table. The default value is 10000. In this walkthrough, use 50000 for better performance.

Configuration used for FullLoadSettings :

```
"FullLoadSettings": {  
  "TargetTablePrepMode": "DO_NOTHING",  
  "CreatePkAfterFullLoad": false,  
  "StopTaskCachedChangesApplied": false,  
  "StopTaskCachedChangesNotApplied": false,  
  "MaxFullLoadSubTasks": 8,  
  "TransactionConsistencyTimeout": 600,  
  "CommitRate": 50000  
},
```

## Note

Increasing this parameter induces additional load on the source database, replication instance, and target database.

## To create a database migration task

1. Log in to the AWS Management Console, and open the [AWS DMS console](#).
2. Choose **Database migration tasks**, then choose **Create task**.
3. On the **Create database migration task page**, enter the following information.

| Parameter                                                                                              | Action                                                                             |
|--------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------|
| <b>Task identifier</b>                                                                                 | Enter <b>mysql-to-dynamodb-data-migration</b> .                                    |
| <b>Replication instance</b>                                                                            | Choose <b>mysql-to-ddb-migration-ri</b> . You configured this value in Step 1.     |
| <b>Source database endpoint</b>                                                                        | Choose <b>mysql-source-dms-datastore</b> . You configured this value in Step 3.    |
| <b>Target database endpoint</b>                                                                        | Choose <b>dynamodb-target-dms-datastore</b> . You configured this value in Step 5. |
| <b>Migration type</b>                                                                                  | Choose <b>Migrate existing data and replicate ongoing changes</b> .                |
| <b>Editing mode</b>                                                                                    | Choose <b>Wizard</b> .                                                             |
| <b>Custom CDC stop mode for source transactions</b>                                                    | Choose <b>Disable custom CDC stop mode</b> .                                       |
| <b>Target table preparation mode</b>                                                                   | Choose <b>Do nothing</b> .                                                         |
| <b>Stop task after full load completes</b>                                                             | Choose <b>Don't stop</b> .                                                         |
| <b>Include LOB columns in replication</b>                                                              | Choose <b>Don't include LOB columns</b> .                                          |
| <b>Advanced task settings, Full load tuning settings, Maximum number of tables to load in parallel</b> | Use default value                                                                  |
| <b>Enable validation</b>                                                                               | Turn off because DynamoDB doesn't support validation.                              |
| <b>Enable CloudWatch logs</b>                                                                          | Turn on.                                                                           |

4. Keep the default values for other parameters, and choose Create task.

AWS DMS runs the task immediately. The **Database migration tasks** section displays the status of the migration task.

## Step 7: Run the AWS DMS Task

After you create your AWS Database Migration Service (AWS DMS) task, do a test run to identify the full load run time and ongoing replication performance. You can validate that initial configurations work as expected. You can do this by monitoring and documenting resource utilization on the source database, replication instance, and target database. These details make up the initial baseline and help determine if you need further optimization.

After you started the task, the full load operation starts loading tables. You can see the table load completion status in the **Table Statistics** section and the corresponding records in the target DynamoDB instance.

In our use case, the following image shows table statistics for the **dms.c5.12xlarge** replication instance with parallel-load threads option. The full load for migrating 1 billion records completed in 14 hours. This means that we achieved our goal of completing full load in less than 15 hours. Further, if you still want to reduce the full load time, you can distribute the table workload using multiple tasks with DMS source filter conditions and Parallel load threads configurations. Following this approach, you can migrate the data in parallel with better performance.

**Summary**

|                                              |                                        |                                      |                                         |
|----------------------------------------------|----------------------------------------|--------------------------------------|-----------------------------------------|
| Status<br>Load complete, replication ongoing | Type<br>Full load, ongoing replication | Source<br>mysql-source-dms-datastore | Target<br>dynamodb-target-dms-datastore |
|----------------------------------------------|----------------------------------------|--------------------------------------|-----------------------------------------|

Overview details | **Table statistics** | CloudWatch metrics | Mapping rules | Premigration assessments | Tags

**Table statistics (1)**  
Total rows include loaded source table rows from Inserts, Deletes, Updates, DDLs, and Full load rows.

Find schema

| Schema name         | Table               | Load state      | Elapsed load time | Inserts | Deletes | Updates | DDLs | Applied inserts | Applied deletes | Applied updates | Applied DDLs | Full load rows | Total rows |
|---------------------|---------------------|-----------------|-------------------|---------|---------|---------|------|-----------------|-----------------|-----------------|--------------|----------------|------------|
| vaccine_drive_stats | vaccine_drive_stats | Table completed | 14 h 3 m 12 s     | 50,000  | 0       | 0       | 0    | 44,501          | 0               | 0               | 0            | 1,022,485,341  | 1,022,535  |

A task with instance class “dms.c5.large” and default configuration was able to migrate 1 Billion records in 278 hours. Later, the task moved to the failed state due to unavailability of the source binary log from the full load start time. To avoid this issue, ensure that you are retaining the binary log based on the full load completion time. Using these statistics, you can understand the benefits of using a parallel load configuration to speed up the migration phase. See the following screenshot for details.



mysql-to-dynamodb-data-migration Actions Quick view and compare View logs

**Summary**

Status Failed Type Full load, ongoing replication Source mysql-source-dms-datastore Target dynamodb-target-dms-datastore

Overview details **Table statistics** CloudWatch metrics Mapping rules Premigration assessments Tags

**Table statistics (1)** Export to CSV Validate again Reload table data

Total rows include loaded source table rows from Inserts, Deletes, Updates, DDLs, and Full load rows.

Find schema

| Schema name | Table               | Load state      | Elapsed load time | Full load rows | Total rows  | Validation state | Validation pending | Validation failed | Validation suspended |
|-------------|---------------------|-----------------|-------------------|----------------|-------------|------------------|--------------------|-------------------|----------------------|
| valis       | vaccine_drive_stats | Table completed | 277 h 51 m 24 s   | 224,034,462    | 224,034,462 | Not enabled      | 0                  | 0                 | 0                    |

We also monitored the CloudWatch metrics such as compute, memory, and network to identify the resource usage of the AWS DMS instances. You have to identify the resource constraint and scale-up to the AWS DMS instance class that serves your workloads better. You could also scale-down the AWS DMS instance to a t3 or r5 instance class based on the transaction volume for your ongoing replication task.

Because we turned on the parallel-load option, the I/O load on the replication instance is expected to increase. We described in Step 1 that you should monitor the **Write IOPS** and **Read IOPS** metrics in CloudWatch to make sure that the total IOPS (write + read IOPS) doesn't exceed the total IOPS available for your replication instance. If it does, make sure that you allocate more storage to scale for better I/O performance. For more information, see [Monitoring replication tasks using Amazon CloudWatch](#).

We covered most of the prerequisites that help avoid errors related to configuration. If you observe issues when running the task, then see [Troubleshooting migration tasks in Database Migration Service](#) or [Best practices for Database Migration Service](#), or reach out to AWS Support for further assistance.

Optionally, you could choose to validate the successful completion of the data migration by querying the target DynamoDB table from the console. You can use the "Get live item count" option to get the total table record counts.

**Items summary** Get live item count

DynamoDB updates the following information approximately every six hours.

|                             |                               |                                   |
|-----------------------------|-------------------------------|-----------------------------------|
| Item count<br>1,022,485,341 | Table size<br>564.3 gigabytes | Average item size<br>551.91 bytes |
|-----------------------------|-------------------------------|-----------------------------------|

When you choose "Start scan" from **Get live item count**, you will perform a DynamoDB scan to determine the most-recent item count. This scan might consume additional table read capacity

units. Generally, it is not recommended to perform this action on very large tables or tables that serve critical production traffic. You can pause the action at any time to avoid consuming extra read capacity.

After you complete the migration, validate that your data migrated successfully, and delete the cloud resources that you created.

## Conclusion

We covered all of the steps that you need to migrate a table from RDS MySQL to Amazon DynamoDB, and used the available configuration details to complete the migration in less time. Once the data was completely migrated to the target DB, then you can view the application traffic on the DynamoDB table. In this walkthrough, we achieved the crucial business requirements by using AWS DMS. Try out these steps to migrate your data to DynamoDB and explore how you can centralize your data with a low-cost solution. To learn more about AWS DMS, see the [Database Migration Service User Guide](#).

## Migrating an RDS for MySQL database to an S3 data lake

A data lake is a system architecture that enables you to store data in a centralized repository, allowing for categorization, cataloging security, and analysis by a diverse range of users and tools. In a data lake, you can analyze structured, semi-structured, and unstructured data, as well as transform these raw data assets as necessary.

Thousands of customers are building data lakes in AWS, using the cloud-scale storage provided by Amazon S3. The transformation capabilities of services such as AWS Glue, Amazon EMR, and the analytic capabilities of services such as Amazon Athena, Amazon Redshift, and Amazon SageMaker enable you to utilize data lakes easily and cost efficiently.

When building a data lake, a common concern is how to hydrate your data lake: populating data from upstream systems, and keeping the lake up-to-date as the source data grows and changes. Traditionally, customers have relied on SQL-level solutions to extract changed records from source systems, e.g., filtering on “last updated” timestamps, or performing full-refreshes on a periodic basis. Both solutions have drawbacks: last updated filters rely on the timestamps being accurately populated, and full refresh has performance and timeliness considerations.

A different approach is to use a database replication service like [AWS Database Migration Service \(AWS DMS\)](#). AWS DMS captures source data changes from the database transaction logs and

logically replicates them on the target (Change Data Capture, CDC). It can also perform a "full-load" to populate the data lake with an initial snapshot of your source data. Then, as changes occur on the source, AWS DMS finds and applies those changes to your data lake, ensuring your data is consistent.


In this document, we will describe the process of setting up an AWS data lake using source data from an Amazon RDS for MySQL database. We will host the lake on Amazon S3, and use AWS DMS to hydrate the data. After describing some prerequisites, we will walk through the steps to setup AWS DMS, connect to the source database, and discuss considerations you should know about when using AWS DMS.

## Topics

- [Solution overview](#)
- [Use case](#)
- [Limitations](#)
- [Choosing an instance class and storage size](#)
- [Step-By-Step Migration](#)

## Solution overview

The following diagram displays a high-level architecture of the solution, where we use AWS DMS to move data from two MySQL databases hosted on Amazon RDS to Amazon S3.



This walkthrough assumes that the source data is sharded over two MySQL instances with identical schemas. Note that the only difference from having a single source instance is that you will create an additional endpoint and task. Therefore, this walkthrough can be applied even if the source is single instance. The schema and table structures used in this walkthrough will be explained in further detail later in the use case section.

In this walkthrough, you will set up the following resources in AWS DMS:

- **Replication Instance** — An AWS managed instance that hosts the AWS DMS engine. You control the type and size of the instance based on your workload.
- **Source Endpoint** — A resource that provides connection details, data store type, and credentials to connect to a source database. For this use case, we will configure the source endpoint to point to the Amazon RDS for MySQL database.
- **Target Endpoint** — AWS DMS supports several target systems including Amazon RDS, Amazon Aurora, Amazon Redshift, Amazon Kinesis Data Streams, Amazon S3, and more. For this use case, we will configure Amazon S3 as the target endpoint.
- **Replication Task** — A resource that runs on the replication instance and connects to endpoints to replicate data from the source to the target.

## Use case

The source MySQL engine version that we will use in this walkthrough is 8.0.31. AWS DMS supports Amazon RDS for MySQL 5.6 or higher as a source. There are three tables under the `dms_sample` schema in the two MySQL databases. The total size is about **220 GiB**. We assume a data change amount of about tens of GiB per day. A similar size of data exists in both instances. The primary keys of the `posts` and `post_history` tables are `id` and `creation_date`, and the tables are partitioned with 180 partitions on the `creation_date` column. The `votes` table is not partitioned and the `id` column is the primary key.

| posts                                | post_history                      | votes                      |
|--------------------------------------|-----------------------------------|----------------------------|
| id DECIMAL(11,0)                     | id DECIMAL(11,0)                  | id DECIMAL(11,0)           |
| accepted_answer_id DECIMAL(11,0)     | post_history_type_id DECIMAL(4,0) | post_id DECIMAL(11,0)      |
| answer_count DECIMAL(11,0)           | post_id DECIMAL(11,0)             | user_id DECIMAL(11,0)      |
| body LONGTEXT                        | revision_guid VARCHAR(36)         | bounty_amount DECIMAL(8,0) |
| closed_date DATETIME(3)              | creation_date DATETIME(3)         | vote_type_id DECIMAL(4,0)  |
| comment_count DECIMAL(11,0)          | user_id DECIMAL(11,0)             | creation_date DATETIME(3)  |
| community_owned_date DATETIME(3)     | user_display_name VARCHAR(40)     |                            |
| creation_date DATETIME(3)            | comment LONGTEXT                  |                            |
| favorite_count DECIMAL(8,0)          | text LONGTEXT                     |                            |
| last_activity_date DATETIME(3)       |                                   |                            |
| last_edit_date DATETIME(3)           |                                   |                            |
| last_editor_display_name VARCHAR(40) |                                   |                            |
| last_editor_user_id DECIMAL(11,0)    |                                   |                            |
| owner_user_id DECIMAL(11,0)          |                                   |                            |
| parent_id DECIMAL(11,0)              |                                   |                            |
| post_type_id DECIMAL(4,0)            |                                   |                            |
| score DECIMAL(8,0)                   |                                   |                            |
| tags VARCHAR(150)                    |                                   |                            |
| title VARCHAR(250)                   |                                   |                            |
| view_count DECIMAL(8,0)              |                                   |                            |
| Indexes                              | Indexes                           | Indexes                    |

## Limitations

As a managed service, AWS DMS allows users to start migration in a few steps. However there are some limitations/restrictions depending on the type of source and target endpoints.

There are some data types that are not supported as MySQL source. Before you start your migration, it's a good idea to find out if there are any unsupported data types. Premigration assessments can help you find unsupported data in your source database. For information about datatypes supported in MySQL, see [Data types](#).

For other MySQL source or S3 target endpoint limitations, see the following documents:

- \* [https://docs.aws.amazon.com/dms/latest/userguide/CHAP\\_Source.MySQL.html#CHAP\\_Source.MySQL.Limitations](https://docs.aws.amazon.com/dms/latest/userguide/CHAP_Source.MySQL.html#CHAP_Source.MySQL.Limitations) [Limitations on using a MySQL database as a source for [.shared]`DMS`]

\* [https://docs.aws.amazon.com/dms/latest/userguide/CHAP\\_Target.S3.html#CHAP\\_Target.S3.Limitations](https://docs.aws.amazon.com/dms/latest/userguide/CHAP_Target.S3.html#CHAP_Target.S3.Limitations)[Limitations to using Amazon S3 as a target]

## Choosing an instance class and storage size

Before you start migrating your database, you need to consider your source, target, and replication instance resources such as CPU, memory, disk space, and network bandwidth/latency. How much workload will be placed on the source database, how to determine the sizing of the replication instance, and what instance class should be used for the target database are common questions when starting a migration.

There is no single answer to these questions. It's hard to calculate because the optimal configuration varies depending on the amount of data in your source database, your workload, your AWS DMS task configuration, and the number of tasks running concurrently. One of the benefits of using AWS is the ability to flexibly and easily resize resources as needed. You can change your replication instance class or target database instance class in-place as needed in a few clicks and minutes. Test your migrations using a larger instance class first, then check the resource usage provided by CloudWatch metrics and resize if necessary.

In this walkthrough, we will use the following instance classes.

### Source database

- **Instance class:** db.c5.4xlarge
- **Allocated storage size:** 1024 GiB
- **Storage type:** io1
- **IOPS:** 20000

Full-load typically requires more resources from the source database than CDC because full-load simultaneously transfers data from the source with the number of parallels you specify in the task settings. The default parallelism is 8, which means that all data from the source table will be transferred to the target through the replication instance in 8 parallel threads. In this walkthrough, we will allocate the resources described above to use a maximum possible parallelism of **49** threads for an AWS DMS task.

Note that this setting yields **250-300 MiB/s** read throughput on the source database. If you want to reduce the workload on the source database, you can lower the parallelism number described in later section.

## Replication instance

- **dms.c5.9xlarge**
- **Allocated storage size:** 100 GiB

Because we are performing a heterogeneous migration and using the parallel full-load option with a maximum of 49 parallel threads, we start with the relatively large compute optimized instance **dms.c5.9xlarge** as the replication instance class. This instance class has enough performance to migrate source data to S3 in 49 parallel threads in our use case. It is also possible to use a smaller instance class if it reduces the number of threads. We'll discuss this in a later section.

When Amazon S3 is the target, storage throughput is the primary factor when determining the full-load performance. This is because when AWS DMS outputs a CSV or a Parquet file to Amazon S3, AWS DMS first writes the file to storage on your replication instance, and then AWS DMS uploads the file to the Amazon S3 bucket.

AWS DMS supports GP2 EBS storage. IOPS for GP2 EBS storage depends on storage size. It increases at a rate of 3 IOPS/GiB. This value is the same as the EBS burst credits added per second. A single GP2 volume performs up to 3000 IOPS as long as it has burst credits, but once it runs out of credits, it only performs as much performance as the credits provided at 3 IOPS/GiB. For example, 100 GiB is 300 IOPS.

In this scenario, we will allocate 100 GiB of storage for a temporary maximum throughput of about 20-30 minutes. This is enough with this workload. Find the optimal disk size for your workload by running a test task. The storage size can be changed online even while the task is running. However, the storage performance may be temporarily degraded during the change. Also, the storage size can increase, but cannot decrease unless you recreate the replication instance.

## Step-By-Step Migration

The following steps provide instructions for migrating Amazon RDS for MySQL databases to an Amazon S3 data lake.

### Topics

- [Step 0: Configure the source Amazon RDS for MySQL database](#)

- [Step 1: Create a replication instance](#)
- [Step 2: Create an AWS DMS source endpoint](#)
- [Step 3: Configure a target Amazon S3 bucket](#)
- [Step 4: Create an AWS DMS Task](#)
- [Step 5: Run and monitor your AWS DMS Task](#)
- [Step 6: Monitor your migration](#)
- [Conclusion](#)

## Step 0: Configure the source Amazon RDS for MySQL database

Before setting up AWS DMS resources, you need to configure your Amazon RDS for MySQL database instances as a source for AWS DMS.

### Amazon RDS Backup configuration

Your Amazon RDS for MySQL instance must have **Automatic Backups** turned on to use CDC. Otherwise, binary logging will not be enabled at the MySQL level. Enabling automatic backups enables binary logging for the database instance. The backup retention period can be any value from one to 35 days. One day is enough for this walkthrough.

### Binary logging configuration

To use AWS DMS CDC, the following parameters must be set correctly in the parameter group attached to your database instances.

- `binlog_format` : "ROW"
- `binlog_row_image` : "Full"``
- `binlog_checksum` : "NONE"``

The default **binlog\_format** is "Mixed". AWS DMS requires the "ROW" format, and all columns before and after the imaging. We recommend that **binlog\_checksum** set to NONE.

### Binary logging retention hours

AWS DMS requires binary logs to be local to the Amazon RDS for MySQL database instance. To ensure that binary logs are available to AWS DMS, you should increase the length of time that the logs remain available in the database instance host. For example, to increase log retention to 24 hours, run the following command. 24 hours are enough for this walkthrough.



```
call mysql.rds_set_configuration('binlog retention hours', 24);
```

## VPC, Subnet and Network ACL configuration

In this walkthrough, the database instance and the replication instance are placed in the same VPC and the same subnet, so all you need to do is configure security groups, network ACLs, and route tables so that your Amazon RDS for MySQL database instance and AWS DMS replication instance can communicate within the same subnet. If you have source databases in a different subnet, VPC, or different location outside AWS, you need to configure your network to allow communication between your Amazon RDS for MySQL database instance and your AWS DMS replication instance.

### Inbound connection rule

To ensure that AWS DMS can access your database server, you need to make changes to the relevant security groups and network access control lists. AWS DMS only requires access to the MySQL database listener port (the default is 3306). The connection always starts from the AWS DMS replication instance to MySQL. Therefore, you add allowed connections from the replication instance to the ingress rule of the security group attached to the database instance. We recommend you add all subnet group ranges to the ingress rule, because replication instances are a managed service, and the IP address of a replication instance may change automatically.

You have now completed all necessary setup for your Amazon RDS for MySQL database instance. Next, create a replication instance.

## Step 1: Create a replication instance

To create an AWS DMS replication instance, do the following:

1. Sign in to the AWS Management Console, and open the [AWS DMS console](#).
2. If you are signed in as an AWS Identity and Access Management (IAM) user, you must have the appropriate permissions to access AWS DMS. For more information about the permissions required, see IAM permissions.
3. On the Welcome page, choose **Create replication instance** to start a database migration.
4. On the **Create replication instance** page, specify your replication instance information.

| For this parameter | Do this |
|--------------------|---------|
|--------------------|---------|

|                                |                                                                                                                                                                                                                      |
|--------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Name</b>                    | Enter <code>s3-datalake-migration-ri</code> . If you are using multiple replication servers or sharing an account, choose a name that helps you quickly differentiate between the different servers.                 |
| <b>Description</b>             | Enter <code>Migrate MySQL to [.shared]`S3 data lake`</code> .                                                                                                                                                        |
| <b>Instance class</b>          | Choose <code>dms.c5.9xlarge</code> . Each size and type of instance class has increasing CPU, memory, and I/O capacity.                                                                                              |
| <b>Engine version</b>          | Leave the default value, which is the latest stable version of the AWS DMS replication engine.                                                                                                                       |
| <b>Allocated storage (GiB)</b> | Choose <code>100</code> GiB.                                                                                                                                                                                         |
| <b>VPC</b>                     | Choose the virtual private cloud (VPC) in which your replication instance will launch. Select the same VPC in which your source is placed.                                                                           |
| <b>Multi AZ</b>                | In this scenario, choose <b>No</b> . If you choose <b>Yes</b> , AWS DMS creates a second replication server in a different Availability Zone for failover if there is a problem with the primary replication server. |
| <b>Publicly accessible</b>     | Choose <b>Yes</b> . If either your source or target database resides outside of the VPC in which your replication server resides, you must make your replication server policy publicly accessible.                  |

Once the creation of the replication instance starts, it usually becomes available in about ten minutes or more. The next endpoint can be created even when the replication instance is in the

creating status, but the connection test cannot be performed unless the replication instance is in the available status.

## Step 2: Create an AWS DMS source endpoint

To create a source endpoint, do the following:

1. Open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose **Endpoints**.
3. Choose **Create endpoint**.
4. On the **Create endpoint** page, enter the following information.

### Source endpoint 1:

#### Endpoint type

Choose **Source endpoint**, **Select RDS DB instance**, and choose the `dataLake-source-db1` RDS instance.

#### Endpoint identifier

Enter `mysql-dms-s3-source-1`

#### Source engine

Choose **MySQL**.

#### Access to endpoint database

Choose **Provide access information manually**.

#### Server name

Enter the Amazon RDS database server name.

#### Port

Enter **3306**.

#### Secure Socket Layer (SSL) mode

Choose **none**.

### User name

Enter **dms\_user**.

### Password

Enter the password that you created for the dms\_user user.

## Source endpoint 2:

### Endpoint type

Choose **Source endpoint, Select RDS DB instance**, and choose the dataLake-source-db2 RDS instance.

### Endpoint identifier

Enter **mysql-dms-s3-source-2**

### Source engine

Choose **MySQL**.

### Access to endpoint database

Choose **Provide access information manually**.

### Server name

Enter the [.shared] `RDS` database server name.

### Port

Enter **3306**.

### Secure Socket Layer (SSL) mode

Choose **none**.

**User name**

Enter **dms\_user**.

**Password**

Enter the password that you created for the `dms_user` user.

You can try testing the connection before you finish creating the endpoint. Test Connection attempts to connect from the replication instance to the source database and verify that the replication instance can connect to MySQL with the settings provided. If the connection test succeeds, go to the next step; otherwise, check if the values you set for the endpoint are correct. If correct, check if the network between the source and the replication instance is configured correctly.

### Step 3: Configure a target Amazon S3 bucket

To create the Amazon S3 bucket, do the following:

1. Open the Amazon S3 console at <https://s3.console.aws.amazon.com/s3/home>.
2. Choose **Create bucket**.
3. For **Bucket name**, enter *\*<your-bucket-name>\**. Note: The bucket name needs to be unique globally.
4. For **AWS Region**, choose the region that hosts your AWS DMS replication instance.
5. Leave the default values in the other fields and choose **Create bucket**.

To use Amazon S3 as an AWS Database Migration Service (AWS DMS) target endpoint, create an IAM role with write and delete access to the S3 bucket. Then add DMS (`dms.amazonaws.com`) as trusted entity in this IAM role. This is a minimum required assume role policy and policy document. For more information, see [Prerequisites for using Amazon S3 as a target](#).

Assume role policy:

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
```

```

        "Effect": "Allow",
        "Principal": {
            "Service": [
                "dms.amazonaws.com"
            ]
        },
        "Action": "sts:AssumeRole"
    }
]
}

```

### Policy:

```

{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Action": [
        "s3:PutObject",
        "s3:DeleteObject",
        "s3:PutObjectTagging"
      ],
      "Resource": [
        "arn:aws:s3:::mysql2s3walkthrough/*"
      ],
      "Effect": "Allow"
    },
    {
      "Action": "s3:ListBucket",
      "Resource": "arn:aws:s3:::mysql2s3walkthrough",
      "Effect": "Allow"
    }
  ]
}

```

To create a target endpoint, do the following:

1. Open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose **Endpoints**, and then choose **Create endpoint**.
3. On the **Create endpoint** page, enter the following information.

|                                |                                                                              |
|--------------------------------|------------------------------------------------------------------------------|
| <b>Endpoint type</b>           | Choose <b>Target endpoint</b> , and turn off <b>Select RDS DB instance</b> . |
| <b>Endpoint identifier</b>     | Enter <b>mysql-dms-s3-target</b> .                                           |
| <b>Target engine</b>           | Choose <b>Amazon S3</b> .                                                    |
| <b>Service access role ARN</b> | Enter the IAM role that can access your Amazon S3 data lake.                 |
| <b>Bucket name</b>             | Enter <b>&lt;your-bucket-name&gt;</b> .                                      |

Expand the **Endpoint settings** section, choose **Wizard**, and then choose **Add new setting** to add the settings as shown on the following image.

When using AWS DMS to migrate data to an Amazon Simple Storage Service (Amazon S3) data lake, you can change the default task behavior, such as file formats, partitioning, file sizing, etc. This leads to minimizing post-migration processing and helps downstream applications consume data efficiently. You can customize task behavior using endpoint settings and extra connection attributes (ECA). Most of the Amazon S3 endpoint settings and ECA settings overlap, except for a few parameters. In this walkthrough, we will configure Amazon S3 endpoint settings.

### Choose file format (dataFormat)

AWS DMS supports CSV and Parquet formats for outputting data to an S3 target. Each file format has its own benefits. Choose the right file format depending on your consumption pattern. Apache Parquet is an open-source file format that stores data in a columnar format, which is built to support efficient compression and encoding schemes providing storage space savings and performance benefits. CSV files are helpful when you plan to keep data in human readable format, or share or transfer Amazon S3 files into other downstream systems for further processing. In this scenario, we will use the CSV format.

### Date based partitioning (DatePartitionEnabled)

In addition to using optimized file formats like Parquet, another common approach for further optimization is to partition the data. AWS DMS supports date-based folder partitioning based on transaction commit dates. The data is stored in different folders based on a timestamp which has following benefits:

- Better management for your S3 objects.
- Limiting the size of each S3 folder.
- Optimizing data lake queries or other subsequent operations.

```
dms_sample/post_history/LOAD00000001.csv
dms_sample/post_history/LOAD00000002.csv
...
dms_sample/posts/LOAD00000001.csv
dms_sample/posts/LOAD00000002.csv
dms_sample/posts/LOAD00000003.csv
...
...
dms_sample/posts/2022/5/21/20220521-145815742.csv
dms_sample/posts/2022/5/21/20220521-145918391.csv
```

## Determine file size

By default, an AWS DMS task writes captured data to an Amazon S3 bucket either if the file size reaches 32 MB or if the previous file write was more than 60 seconds ago. These settings ensure that the data capture latency is low. However, this approach creates a large number of small files in the target Amazon S3 bucket. This value can be changed with `CdcMaxBatchInterval` in the S3 target endpoint settings.

However, we need to optimize this schema for cost and performance. When you use distributed processing frameworks such as Amazon Athena, AWS Glue or Amazon EMR, it is recommended to avoid having many small files (less than 64 MB). Small files tend to cause operational overhead in various distributed processing frameworks. Since we plan to use Amazon Athena to query data from our Amazon S3 bucket, we need to make sure our target file size is at least 64 MB.

In this scenario, we'll use the following endpoint settings: `MaxFileSize=64000`, `CdcMaxBatchInterval=3600` and `CdcMinFileSize=64000`. These settings ensure that AWS DMS does not write the file until its size reaches 64 MB or if the last file write was more than an hour ago.

## Serialize ongoing replication events

A common challenge when using Amazon S3 as a target involves identifying the ongoing replication event sequence when multiple records are updated at the same time on the source database. AWS DMS provides two options to help serialize such events for Amazon S3. You can



use the `TimeStampColumnName` endpoint setting or use transformation rules to include a LSN column. Here, we will discuss the first option. For more information about the second option, see [Step 6: Create an AWS DMS Task](#).

### Use the `TimeStampColumnName` endpoint setting

The `TimeStampColumnName` setting adds an additional `STRING` column to the target Parquet file created by AWS DMS. During ongoing replication, the column value represents the commit timestamp of the event in SQL Server. For the full load phase, the columns' values represent the timestamp of the data transfer to Amazon S3. The default format is `yyyy-MM-dd HH:mm:ss.SSSSSS`. This format provides a microsecond precision, but also depends on the source database transaction log timestamp precision.

### Include full load operation field

All files created during ongoing replication have the first column marked with I, U, or D. These symbols represent the DML operation on the source and stand for **Insert**, **Update**, or **Delete**. For full load files, you can add this column by configuring the following endpoint setting.

```
includeOpForFullLoad=true
```

This ensures that all full load files are marked with an I operation.

When you use this approach, new subscribers can consume the entire data set or prepare a fresh copy in case of any downstream processing issues.

AWS DMS outputs an extra column (Op) where each record has one of the DML flags (I: Insert, U: Update, or D: Delete) in addition to the existing columns in the source tables, indicating which operation generated the change at that time.

In the following example, a source table has a structure similar to the following:

| id | name  | age | year |
|----|-------|-----|------|
| 1  | Scott | 36  | 1986 |
| 2  | Mike  | 27  | 1995 |
| 3  | Bob   | 42  | 1980 |

For this example, we insert a record into this table such as the following:

```
INSERT INTO dms_example.users (id, name, age, birthday) VALUES (4, 'Kate', 23, 1999);
```

The generated record will look similar to the following:

```
I, 4, Kate, 23, 1999
```

To handle these changed data, you need to take the operation flag into consideration when querying the file output in the S3 bucket, or alternatively you can process those files using AWS Glue and store the output in another S3 bucket which can then be queried using Amazon Athena.

There are several possible methods depending on what software stack you want to achieve. The last section in this document, Next Steps, references specific examples.

In this scenario, we'll use the following settings:

### Endpoint 1:

```
{
  "ServiceAccessRoleArn": "arn:aws:iam::<ACCOUNT_ID>:role/mysql2s3-walkthrough-dms-s3-
target-access-role",
  "CsvRowDelimiter": "\\n",
  "CsvDelimiter": ",",
  "BucketName": "<S3_BUCKET_NAME>",
  "BucketFolder": "endpoint1",
  "CompressionType": "NONE",
  "DataFormat": "CSV",
  "EnableStatistics": true,
  "DatePartitionEnabled": true,
  "MaxFileSize": 64000,
  "CdcMaxBatchInterval": 3600,
  "CdcMinFileSize": 64000,
  "IncludeOpForFullLoad": true
}
```

### Endpoint 2:

```
{
  "ServiceAccessRoleArn": "arn:aws:iam::<ACCOUNT_ID>:role/mysql2s3-walkthrough-dms-s3-
target-access-role",
  "CsvRowDelimiter": "\\n",
```

```

"CsvDelimiter": ",",
"BucketName": "<S3_BUCKET_NAME>",
"BucketFolder": "endpoint2",
"CompressionType": "NONE",
"DataFormat": "CSV",
"EnableStatistics": true,
"DatePartitionEnabled": true,
"MaxFileSize": 64000,
"CdcMaxBatchInterval": 3600,
"CdcMinFileSize": 64000,
"IncludeOpForFullLoad": true
}

```

By using this configuration, data on two sharded database instances will be migrated to different bucket folders in the same bucket.

## Step 4: Create an AWS DMS Task

After you configure the replication instance and endpoints, the next step is creating the AWS DMS task. In this scenario, we will create a task that performs both full-load and CDC. To create a database migration task, do the following:

1. Open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Select **Database migration tasks**, and then choose **Create task**.
3. On the **Create database migration task** page, enter the following information.

### Replication task 1:

| For this parameter              | Do this                                                                        |
|---------------------------------|--------------------------------------------------------------------------------|
| <b>Task identifier</b>          | Enter <b>mysql-dms-s3-task-1</b>                                               |
| <b>Replication instance</b>     | Choose <b>datalake-migration-ri</b> (the value that you configured on Step 1). |
| <b>Source database endpoint</b> | Choose <b>mysql-dms-s3-source-1</b> (the value that you configured on Step 3). |
| <b>Target database endpoint</b> | Choose <b>mysql-dms-s3-target</b> (the value that you configured on Step 4).   |

|                                                     |                                                                     |
|-----------------------------------------------------|---------------------------------------------------------------------|
| <b>Migration type</b>                               | Choose <b>Migrate existing data and replicate ongoing changes</b> . |
| <b>Editing mode</b>                                 | Choose <b>Wizard</b> .                                              |
| <b>Custom CDC stop mode for source transactions</b> | Choose <b>Disable custom CDC stop mode</b> .                        |
| <b>Target table preparation mode</b>                | Choose <b>Drop and create</b>                                       |
| <b>Stop task after full load completes</b>          | Choose <b>Don't stop</b> .                                          |
| <b>Include LOB columns in replication</b>           | Choose <b>Limited LOB mode</b> .                                    |
| <b>Maximum LOB size (KB)</b>                        | Enter <b>1024</b>                                                   |
| <b>Enable validation</b>                            | Enter <b>1024</b>                                                   |
| <b>Enable validation</b>                            | Turn off because Amazon S3 does not support validation.             |
| <b>Enable CloudWatch logs</b>                       | Turn on.                                                            |

## Replication task 2:

|                                 |                                                                                |
|---------------------------------|--------------------------------------------------------------------------------|
| For this parameter              | Do this                                                                        |
| <b>Task identifier</b>          | Enter <b>mysql-dms-s3-task-2</b>                                               |
| <b>Replication instance</b>     | Choose <b>datalake-migration-ri</b> (the value that you configured on Step 1). |
| <b>Source database endpoint</b> | Choose <b>mysql-dms-s3-source-2</b> (the value that you configured on Step 3). |
| <b>Target database endpoint</b> | Choose <b>mysql-dms-s3-target</b> (the value that you configured on Step 4).   |

|                                                     |                                                                     |
|-----------------------------------------------------|---------------------------------------------------------------------|
| <b>Migration type</b>                               | Choose <b>Migrate existing data and replicate ongoing changes</b> . |
| <b>Editing mode</b>                                 | Choose <b>Wizard</b> .                                              |
| <b>Custom CDC stop mode for source transactions</b> | Choose <b>Disable custom CDC stop mode</b> .                        |
| <b>Target table preparation mode</b>                | Choose <b>Drop and create</b>                                       |
| <b>Stop task after full load completes</b>          | Choose <b>Don't stop</b> .                                          |
| <b>Include LOB columns in replication</b>           | Choose <b>Limited LOB mode</b> .                                    |
| <b>Maximum LOB size (KB)</b>                        | Enter <b>1024</b>                                                   |
| <b>Enable validation</b>                            | Enter <b>1024</b>                                                   |
| <b>Enable validation</b>                            | Turn off because Amazon S3 does not support validation.             |
| <b>Enable CloudWatch logs</b>                       | Turn on.                                                            |

Table mappings:

```
{
  "rules": [
    {
      "rule-type": "selection",
      "rule-id": 1,
      "rule-name": "1",
      "object-locator": {
        "schema-name": "dms_sample",
        "table-name": "%"
      },
      "rule-action": "include"
    },
    {
      "rule-type": "table-settings",
      "rule-id": 2,
      "rule-name": "2",
```

```
    "object-locator": {
      "schema-name": "dms_sample",
      "table-name": "post_history"
    },
    "parallel-load": {
      "type": "partitions-auto"
    }
  },
  {
    "rule-type": "table-settings",
    "rule-id": 3,
    "rule-name": "3",
    "object-locator": {
      "schema-name": "dms_sample",
      "table-name": "posts"
    },
    "parallel-load": {
      "type": "partitions-auto"
    }
  },
  {
    "rule-type": "table-settings",
    "rule-id": 4,
    "rule-name": "3",
    "object-locator": {
      "schema-name": "dms_sample",
      "table-name": "votes"
    },
    "parallel-load": {
      "type": "partitions-auto"
    }
  }
]
}
```

### Task settings:

```
{
  "TargetMetadata": {
    "SupportLobs": true,
    "LimitedSizeLobMode": true,
    "LobMaxSize": 1024,
  },
}
```

```
"FullLoadSettings": {
  "TargetTablePrepMode": "TRUNCATE_BEFORE_LOAD",
  "MaxFullLoadSubTasks": 49,
  "CommitRate": 50000
},
"Logging": {
  "EnableLogging": true
}
}
```

## Step 5: Run and monitor your AWS DMS Task

After you created your AWS Database Migration Service (AWS DMS) task, start your replication tasks. To start your AWS DMS task, do the following:

1. Open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Select **Database migration tasks**, and then choose **Create task**.
3. On the **Create database migration task** page, select your replication task.
4. Choose **Actions**, **“Restart / Resume”**.

## Step 6: Monitor your migration

### Task status and Table statistics

After you start the task, the full load operation starts loading tables. Your replication task status will be **“Running”** until full-load completes. After the AWS DMS task completes full load, the task status changes to the **Load complete, replication ongoing** phase. The following image shows the updated status of the task.

You can see the table load completion status in the **Table statistics** section and the corresponding target files in the Amazon S3 bucket. You can check the progress of replication on the **Table statistics** tab. AWS DMS first does full-load on each table. Meanwhile, the task status is **Running**, and at least one of the tables' Load states is **“Before Load”** or **“Full load”**. Tables that have been loaded are displayed as **“Table completed”**. When all tables have been fully loaded, the task status becomes **“Load completed, replication ongoing”**. The task continues to capture source changes and apply them to the target.

In this scenario, the full-load phase typically completes in about 20 minutes. If you don't use **partitions-auto** for table mapping, the same full-load phase takes about an hour. Parallel full load can significantly improve full load performance.

## Cloudwatch Metrics

The AWS DMS console shows CloudWatch statistics for each task. To see metrics, select the replication task and then select the **CloudWatch metrics** tab.

Task metrics are divided into statistics between the replication host and the source endpoint, and statistics between the replication host and the target endpoint. You can determine the total statistic for a task by adding two related statistics together. For example, you can determine the total latency, or replica lag, for a task by combining the **CDCLatencySource** and **CDCLatencyTarget** values.

**CDCLatencySource** is the gap, in seconds, between the last event captured from the source endpoint and current system time stamp of the AWS DMS instance. CDCLatencySource represents the latency between source and replication instance. High CDCLatencySource means the process of capturing changes from source is delayed. To identify latency in an ongoing replication, you can view this metric together with CDCLatencyTarget. If both CDCLatencySource and CDCLatencyTarget are high, investigate CDCLatencySource first.

**CDCLatencyTarget** is the gap, in seconds, between the first event timestamp waiting to commit on the target and the current timestamp of the AWS DMS instance. Target latency is the difference between the replication instance server time and the oldest unconfirmed event id forwarded to a target component. In other words, target latency is the timestamp difference between the replication instance and the oldest event applied but unconfirmed by the TRG endpoint. When CDCLatencyTarget is high, it indicates that the process of applying change events to the target is delayed.

These metrics are useful for knowing what state your tasks are in.

## Conclusion

In this walkthrough, we covered most prerequisites that help avoid configuration related errors. You can get started on your own migrations using the following documentation.

- [Getting started with Database Migration Service.](#)
- [Using a MySQL-compatible database as a source for AWS DMS](#)



- [Using Amazon S3 as a target](#)

If you observe issues when running your task, see [Troubleshooting migration tasks](#) and [Best practices](#) in the AWS DMS public documentation, or reach out to AWS Support for further assistance.

## Migrating an RDS PostgreSQL database to an S3 data lake

This walkthrough will help you understand the process of migrating data from Amazon Relational Database Service (Amazon RDS) for PostgreSQL database to Amazon Simple Storage Service (S3) using AWS Database Migration Service (AWS DMS).

In today's day and age, a data lake is a key component of an organization's data management strategy. Most organizations are seeing an increase in the amount of data they are collecting, and traditional data management strategies can be challenging to scale and operate. This results in siloed data with poor data quality, inconsistencies, and duplication.

As a part of this walkthrough, we'll build a data lake in Amazon S3 using data hosted in an Amazon RDS for PostgreSQL database. Amazon S3 is the largest and most performant cloud storage service. With Amazon S3, you can build a cost-effective, secure data lake with 99.999999999% (11 9s) of durability. Amazon S3 allows you to store and manage both structured and unstructured data at unlimited scale. For analytics, data lakes allow you to easily and cost-effectively create machine learning (ML)-based data visualization dashboards through services like Amazon QuickSight.

To illustrate the process, we'll migrate data from an example database using AWS DMS. AWS DMS is a managed service that lets you migrate between heterogeneous sources and targets (in our case, PostgreSQL and Amazon S3). Using AWS DMS, you can migrate your existing data and ensure that your source and target are synchronized through ongoing replication.

### Topics

- [Why AWS DMS?](#)
- [Use case](#)
- [Example data set](#)
- [Solution overview](#)
- [Prerequisites](#)

- [Step-by-step an Amazon RDS PostgreSQL database to an Amazon S3 data lake migration walkthrough](#)

## Why AWS DMS?

Data lakes typically require building, configuring, and maintaining multiple data ingestion pipelines from cloud and on-premises data stores. Traditionally, databases can be loaded once with data ingestion tools such as import, export, bulk copy, and so on. Ongoing changes are either not possible or are implemented by bookmarking the initial state. Setting up a data lake using these methods can present challenges ranging from increased load on the source database to overheads while carrying schema changes.

In contrast, AWS DMS extracts changes from the database transaction log generated by the database for recovery purposes. AWS DMS then takes these changes, converts them to the target format, and applies them to the target. This process provides near real-time replication to the target, reducing the complexity of replication monitoring.

## Use case

The following use case helps illustrate the challenge we're trying to solve.

Let's assume you run an insurance company. To improve customer service and to implement a delay detection mechanism, you need to collect and store your customers' claims history. If you can determine the relationship between the initial time it takes to initially register a claim, the time spent in the claim process, and the average number of claims processed per month, you can identify the delays in the claim process, and restore certain services to improve customer experience.

In this document, we'll walk through the migration of an PostgreSQL data warehouse hosted in an Amazon RDS PostgreSQL database to Amazon S3. We will show you the steps you can follow to migrate a large data warehouse dataset to an S3 data lake. We will cover various configurations and setups that you can do to achieve this goal to fulfill the critical business requirements mentioned below.


## Example data set

For this walkthrough we will use Insurance schema which includes 9 tables. The largest table is claim table which is a history table of all claims reported with 29 million rows. The total size

of source database is about 100 GB and has about 10+ years worth of claim history data. The remaining tables are mostly smaller dimension tables.

## Solution overview

The following diagram displays a high-level architecture of the solution, where we use AWS DMS to move data from PostgreSQL databases hosted on Amazon RDS to Amazon S3.



Nowadays, most organizations prefer to first create a data lake containing all the data, and then transform and move this data to their respective targets. Based on the use case, we will configure AWS DMS to replicate data from a single database instance containing multiple tables to an S3 bucket and folder.

To replicate data, you need to create and configure the following artifacts in AWS:

- **Replication Instance** — An AWS managed instance that hosts the AWS DMS engine. You control the type or size of the instance based on the workload you plan to migrate.
- **Source Endpoint** — An endpoint that provides connection details, data store type, and credentials to connect to a source database. For this use case, we will configure a source endpoints to point to a Amazon RDS for PostgreSQL database.
- **Target Endpoint** — AWS DMS supports several target systems including Amazon RDS, Amazon Aurora, Amazon Redshift, Amazon Kinesis Data Streams, Amazon S3, and more. For the use case, we will configure Amazon S3 as the target endpoint. In this case we will be using a single S3 bucket to hold the data from both the PostgreSQL sources.

- **Replication Task** — A task that runs on the replication instance and connects to endpoints to replicate data from the source database to the target database. In this case we will have a single migrating the insurance claim data from Amazon RDS PostgreSQL to S3.
- **Amazon Athena** — Amazon Athena is a managed service that makes it easier to run the interactive queries against large data sets by directly uploading them to Amazon S3 while it manages the infrastructure and data handling. With Athena, we just need to define the schema for our data and start querying with standard SQL.
- **Amazon QuickSight** — Amazon QuickSight powers data-driven organizations with unified business intelligence (BI) at hyper scale. With QuickSight, all users can meet varying analytic needs from the same source of truth through modern interactive dashboards, paginated reports, embedded analytics, and natural language queries.

## Prerequisites

The following prerequisites are required to complete this walkthrough:

- An AWS account with AWS Identity and Access Management (IAM) credentials that allows you to launch Amazon RDS and AWS Database Migration Service (AWS DMS) instances in your AWS Region. For information about IAM credentials, see [Create an IAM user](#).
- An understanding of the Amazon Virtual Private Cloud (Amazon VPC) service and security groups. For information about using Amazon VPC with Amazon RDS, see [Amazon Virtual Private Cloud \(VPCs\) and Amazon RDS](#). For information about Amazon RDS security groups, see [Controlling access with security groups](#).
- An understanding of the supported features and limitations of AWS DMS. For information about AWS DMS, see [Getting started with Database Migration Service](#).
- An understanding of how to work with PostgreSQL as a source and Amazon S3 data lake as a target. For information about working with PostgreSQL as a source, see [Using a PostgreSQL-compatible database as a source for AWS DMS](#). For information about working with Amazon S3 as a target, see [Using Amazon S3 as a target](#).
- An understanding of the supported data type conversion options for PostgreSQL and Amazon S3. For information about data types for PostgreSQL as a source, see [Source data types for PostgreSQL](#).
- An audit of your source PostgreSQL database. For each schema and all the objects under each schema, determine whether any of the objects are no longer being used. Deprecate these objects

on the source PostgreSQL database, because there's no need to migrate them if they aren't being used.

For more information about AWS DMS, see [Getting started with Database Migration Service](#).

## Step-by-step an Amazon RDS PostgreSQL database to an Amazon S3 data lake migration walkthrough

### Topics

- [Step 1: Create an AWS DMS replication instance](#)
- [Step 2: Configure a source Amazon RDS for PostgreSQL database](#)
- [Step 3: Create an AWS DMS source endpoint](#)
- [Step 4: Configure a target Amazon S3 bucket](#)
- [Step 5: Configure an AWS DMS target endpoint](#)
- [Step 6: Create an AWS DMS task](#)
- [Step 7: Run the AWS DMS tasks](#)
- [Conclusion](#)

### Step 1: Create an AWS DMS replication instance

To create an AWS Database Migration Service (AWS DMS) replication instance, see [Creating a replication instance](#). Usually, the full load phase is multi-threaded (depending on task configurations) and has a greater resource footprint than ongoing replication. Consequently, it's advisable to start with a larger instance class and then scale down once the tasks are in the ongoing replication phase. Moreover, if you intend to migrate your workload using multiple tasks, monitor your replication instance metrics and resize your instance accordingly.

For this use case, we will migrate a data set of the Insurance database, which is about 100 GB in size. Because we're performing a heterogeneous migration, we can start with a compute-optimized instance like c5.xlarge running the latest AWS DMS engine version. We can later scale up or down based on resource utilization during task execution.

To create an AWS DMS replication instance, do the following:

1. Sign in to the AWS Management Console, and open the [AWS DMS console](#).

2. If you are signed in as an AWS Identity and Access Management (IAM) user, you must have the appropriate permissions to access AWS DMS. For more information about the permissions required, see [IAM permissions](#).
3. On the Welcome page, choose **Create replication instance** to start a database migration.
4. On the **Create replication instance** page, specify your replication instance information.

| For this parameter             | Do the following                                                                                                                                                                                  |
|--------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Name</b>                    | Enter <code>datalake-migration-ri</code> . If you are using multiple replication servers or sharing an account, choose a name that helps you quickly differentiate between the different servers. |
| <b>Description</b>             | Enter <code>Migrate PostgreSQL to S3 data lake</code> .                                                                                                                                           |
| <b>Instance class</b>          | Choose <code>dms.c5.xlarge</code> . Each size and type of instance class has increasing CPU, memory, and I/O capacity.                                                                            |
| <b>Engine version</b>          | Leave the default value chosen, which is the latest stable version of the AWS DMS replication engine.                                                                                             |
| <b>Allocated storage (GiB)</b> | Choose <code>50</code> .                                                                                                                                                                          |
| <b>VPC</b>                     | Choose the virtual private cloud (VPC) in which your replication instance will launch. If possible, select the same VPC in which either your source or target database resides (or both).         |
| <b>Multi AZ</b>                | If you choose <b>Yes</b> , AWS DMS creates a second replication server in a different Availability Zone for failover if there is a problem with the primary replication server.                   |

**Publicly accessible**

If either your source or target database resides outside of the VPC in which your replication server resides, you must make your replication server policy publicly accessible.

## Step 2: Configure a source Amazon RDS for PostgreSQL database

One of the primary considerations when setting up AWS DMS replication is the load that it induces on the source database. During full load, AWS DMS tasks initiate two or three connections for each table that is configured for parallel load. Because AWS DMS settings and data volumes vary across tasks, workloads, and even across different runs of the same task, providing an estimate of resource utilization that applies for all use cases is difficult.

Ongoing replication is single-threaded, and it usually consumes fewer resources than full load. Providing estimates for change data capture (CDC) resource utilization has the same challenges described above.

For our source databases, we use an `m5.xlarge` Amazon RDS instance running PostgreSQL 13.4-R1. While the steps for Amazon RDS creation are out of scope for this walkthrough (for more information, see [Prerequisites](#)), make sure that your Amazon RDS instance has **Automatic Backups** turned on. If you plan to use CDC, you need to turn on logical replication to let DMS capture changes from Amazon RDS for PostgreSQL.

To enable logical replication (required for performing CDC) for an Amazon RDS for PostgreSQL database:

1. Use the AWS master user account for the PostgreSQL DB instance as the user account for the PostgreSQL source endpoint. The master user account has the required roles that allow it to set up CDC.
2. If you use an account other than the master user account, make sure to create several objects from the master account for the account that you use. For more information, see [Migrating an Amazon RDS for PostgreSQL database without using the master user account](#).
3. Set the `rds.logical_replication` parameter in your database parameter group to 1. This static parameter requires a reboot of the database instance to take effect. As part of applying this parameter, AWS DMS sets the `wal_level`, `max_wal_senders`, `max_replication_slots`, and `max_connections` parameters. These parameter changes can

- increase write ahead log (WAL) generation, so only set `rds.logical_replication` when you use logical replication slots.
4. The `wal_sender_timeout` parameter ends replication connections that are inactive longer than the specified number of milliseconds. The default is 60000 milliseconds (60 seconds). Setting the value to 0 (zero) disables the timeout mechanism, and is a valid setting for DMS.
  5. When setting `wal_sender_timeout` to a non-zero value, DMS requires a minimum of 10000 milliseconds (10 seconds), and fails if the value is between 0 and 10000. Keep the value less than 5 minutes to avoid causing a delay during a Multi-AZ failover of a DMS replication instance.
  6. Ensure the value of the `max_worker_processes` parameter in your Database Parameter Group is equal to or greater than the total combined values of `max_logical_replication_workers`, `autovacuum_max_workers`, and `max_parallel_workers`. A high number of background worker processes might impact application workloads on small instances. So, monitor performance of your database if you set `max_worker_processes` higher than the default value.

### Step 3: Create an AWS DMS source endpoint

After you configure the AWS Database Migration Service (AWS DMS) replication instance and the source Amazon RDS, ensure connectivity between both the components. To ensure that the replication instance can access the server and the port for the database, make changes to the relevant security groups and network access control lists. For more information about your network configuration, see [Setting up a network for a replication instance](#).

After you completed the network configurations, you can create a source endpoint. In this case, we create a source endpoint for Amazon RDS PostgreSQL.

To create a source endpoint, do the following:

1. Open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose **Endpoints**.
3. Choose **Create endpoint**.
4. On the **Create endpoint** page, enter the following information.

For this parameter

Do the following



|                                       |                                                                                                                                          |
|---------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Endpoint type</b>                  | Choose <b>Source endpoint</b> , turn on <b>Select Amazon RDS DB instance</b> , and choose <code>datalake-source-db-RDS</code> instance . |
| <b>Endpoint identifier</b>            | Enter <code>pg13rds-dms-s3-source</code>                                                                                                 |
| <b>Source engine</b>                  | Choose <b>PostgreSQL</b>                                                                                                                 |
| <b>Access to endpoint database</b>    | Choose <b>Provide access information manually</b> .                                                                                      |
| <b>Server name</b>                    | Enter the database server name on Amazon RDS.                                                                                            |
| <b>Port</b>                           | Enter 5432.                                                                                                                              |
| <b>Secure Socket Layer (SSL) mode</b> | Choose <b>none</b> .                                                                                                                     |
| <b>User name</b>                      | Enter <code>dms_user</code> .                                                                                                            |
| <b>Password</b>                       | Enter the password that you created for the <code>dms_user</code> user.                                                                  |

## Step 4: Configure a target Amazon S3 bucket

In this use case, we're migrating the Insurance schema to Amazon S3. To create the Amazon S3 bucket, do the following:

1. Open the Amazon S3 console at <https://s3.console.aws.amazon.com/s3/home>.
2. Choose **Create bucket**.
3. For **Bucket name**, enter `pg-dms-s3-target`.
4. For **AWS Region**, choose the region that hosts your AWS DMS replication instance.
5. Leave the default values in the other fields and choose **Create bucket**.

## Step 5: Configure an AWS DMS target endpoint

To use Amazon S3 as an AWS Database Migration Service (AWS DMS) target endpoint, create an IAM role with write and delete access to the S3 bucket. Then add DMS (`dms.amazonaws.com`) as

a trusted entity in this IAM role. For more information, see [Prerequisites for using Amazon S3 as a target](#).

To create a target endpoint, do the following:

1. Open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose **Endpoints**, and then choose **Create endpoint**.
3. On the **Create endpoint page**, enter the following information.

| For this parameter             | Do the following                                                                    |
|--------------------------------|-------------------------------------------------------------------------------------|
| <b>Endpoint type</b>           | Choose <b>Target endpoint</b> , and turn off <b>Select Amazon RDS DB instance</b> . |
| <b>Endpoint identifier</b>     | Enter <b>pg-dms-s3-target</b> .                                                     |
| <b>Target engine</b>           | Choose <b>Amazon S3</b> .                                                           |
| <b>Service access role ARN</b> | Enter the IAM role that can access your Amazon S3 data lake.                        |
| <b>Bucket name</b>             | Enter <b>&lt;your-name&gt;-datalake</b> .                                           |

Expand the **Endpoint settings** section, choose **Wizard**, and then choose **Add new setting** to add the settings as shown on the following image.

When using AWS DMS to migrate data to an Amazon Simple Storage Service (Amazon S3) data lake, you can change the default task behavior, such as file formats, partitioning, file sizing, and so on. This helps reduce post-migration processing so that consuming applications can access the data with lower latency. You can customize task behavior using endpoint settings and extra connection attributes (ECAs). Most of the Amazon S3 endpoint settings and ECA settings overlap, except for a few parameters. In this walkthrough, we will configure Amazon S3 endpoint settings.

## Endpoint settings

| Setting                                                                                                  | Value - A value is required                                                                 |                                                                                     |                         |
|----------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------|-------------------------|
|  CsvRowDelimiter        |  \n        |  | <button>Remove</button> |
|  CsvDelimiter           |  ,         |  | <button>Remove</button> |
|  CompressionType        |  NONE      |  | <button>Remove</button> |
|  EnableStatistics       |  true      |  | <button>Remove</button> |
|  DatePartitionEnabled   |  true      |  | <button>Remove</button> |
|  DatePartitionSequence  |  yyyyymmdd |  | <button>Remove</button> |
|  DatePartitionDelimiter |  none      |  | <button>Remove</button> |

## File Format and Data Partitioning

When using Amazon S3 as a target in an AWS DMS task, both full load and change data capture (CDC) data is written to comma-separated value (.csv) format by default. For more compact storage and faster query options, you also have the option to have the data written to Apache Parquet (.parquet) format. Each file format has its own benefits, CSV files are human-readable and when there is not too much data (less than 50 GB per database) being migrated CSV can be a good choice. Data in parquet files is stored in columnar format which is built to support efficient compression and encoding schemes providing storage space savings and performance benefits. In this walkthrough we will be using CSV as the file format for the Athena and Quicksight to consume.

AWS DMS writes data from a single source table into multiple files to the S3 target during full load and CDC as seen below. The size of these files can be modified by setting the extra connection attributes in the following link [https://docs.aws.amazon.com/dms/latest/userguide/CHAP\\_Target.S3.html#CHAP\\_Target.S3.Configuring](https://docs.aws.amazon.com/dms/latest/userguide/CHAP_Target.S3.html#CHAP_Target.S3.Configuring). This will make the processing of files easier for the application consuming this data as they will be in multiple smaller chunks.

```
schema_name/table_name1/LOAD000000001.csv
schema_name/table_name1/LOAD000000002.csv
```

```
...
schema_name/table_name2/LOAD000000001.csv
schema_name/table_name2/LOAD000000002.csv
schema_name/table_name2/LOAD000000003.csv
schema_name/table_name2/20220521-145815742.csv
schema_name/table_name2/20220521-145918391.csv
```

Additionally, to further optimize the consumption of data from the S3 bucket we can partition the data when loading it into the S3 bucket using AWS DMS. AWS DMS supports date based partitioning based on transactional commit dates for CDC and parallel load option for full load. Using both these options we can partition the data in the S3 bucket with a commit date for CDC and the partition columns date for full load as seen below.

```
schema_name/table_name1/20140912/LOAD000000001.csv
schema_name/table_name1/20140914/LOAD000000002.csv
...
...
schema_name/table_name2/20220615/20220615-203044023.csv
```

## Determine file size

By default, during ongoing replication AWS DMS tasks writes to Amazon S3 are triggered either if the file size reaches 32 KB or if the previous file write was more than 60 seconds ago. These settings ensure that the data capture latency is low. However, this approach creates numerous small files in the target Amazon S3 bucket.

Because we're migrating insurance data for an analytics use case, some latency is acceptable. However, we need to optimize this schema for cost and performance. When you use distributed processing frameworks such as Amazon Athena, it is recommended to avoid too many small files (less than 64 MB). Small files create management overhead for the driver node of the distributed processing framework.

Because we plan to use Amazon Athena to query data from our Amazon S3 bucket, we need to make sure our target file size is at least 64 MB.

Specify the following endpoint settings: `CdcMaxBatchInterval=3600` and `CdcMinFileSize=64000`. These settings ensure that AWS DMS writes the file until its size reaches 64 MB or if the last file write was more than an hour ago.

## Turn on S3 Partitioning

Partitioning in Amazon S3 structures your data by folders and subfolders that help efficiently query data. For example, if you receive insurance claim record data daily from different regions, and you query data for a specific region and find stats for a few months, then it is recommended to partition data by region, year, and month. In Amazon S3, the path for our use case looks as following depending on your setting:

```
s3://<claim-data-bucket-name>/<region>/<schemaname>/<tablename>/<year><month><day>

s3://insurance-policy-datalake
- s3://insurance-policy-datalake/US-WEST-DATA
- s3://insurance-policy-datalake/US-WEST-DATA/insurance
- s3://insurance-policy-datalake/US-WEST-DATA/insurance/claim/
- s3://insurance-policy-datalake/US-WEST-DATA/insurance/claim/20211123/
LOAD000000001.csv
- s3://insurance-policy-datalake/US-WEST-DATA/insurance/policy
- s3://insurance-policy-datalake/US-WEST-DATA/insurance/policy/LOAD000000001.csv
- s3://insurance-policy-datalake/US-WEST-DATA/insurance/policy/20211123/
- s3://insurance-policy-datalake/US-WEST-DATA/insurance/
policy/20211123/20211123-013830913.csv
- s3://insurance-policy-datalake/US-WEST-DATA/insurance/
policy/20211127/20211127-175902985.csv
```

In the above example we have used partitioning in both full load and CDC. Partitioning provides performance benefits because data scanning will be limited to the amount of data in the specific partition based on the filter condition in your queries. For our insurance claim data example, your queries might look as follows:

```
SELECT <column-list> FROM <Claim-table-name> WHERE <region> = <region-name> AND <year>
= <year-value>
```

If you use Amazon Athena to query data, partitioning helps reduce cost as Athena pricing is based on the amount of data that you scan when running queries.

To turn on partitioning for ongoing changes in the above format, use the following settings.

```
bucketFolder=US-WEST-DATA
DatePartitionedEnabled=true
DatePartitionSequence=YYYYMMDD
DatePartitionDelimiter=SLASH
```

## Other considerations

The preceding settings help optimize performance and cost. We also need to configure additional settings because:

- Our use case does not have a fixed end-date.
- We need to minimize issues arising from configurations or retroactive changes.
- We want to minimize recovery time in case of unforeseen issues.

## Serialize ongoing replication events

A common challenge when using Amazon S3 as a target involves identifying the ongoing replication event sequence when multiple records are updated at the same time on the source database. AWS DMS provides two options to help serialize such events for Amazon S3. You can use the `TimeStampColumnName` endpoint setting or use transformation rules to include LSN column. Here, we will discuss the first option. For more information about the second option, see [Using Amazon S3 as a target](#).

### Use the `TimeStampColumnName` endpoint setting

The `TimeStampColumnName` setting adds another `#STRING` column to the target Parquet file created by AWS DMS. During the ongoing replication, the column value represents the commit timestamp of the event in SQL Server. For the full load phase, the columns values represent the timestamp of the data transfer to S3. The default format is `yyyy-MM-dd HH:mm:ss.SSSSSS`. This format provides a microsecond precision but depends on the source database transaction log timestamp precision.

### Include full load operation field

All files created during the ongoing replication have the first column marked with I, U, or D. These symbols represent the DML operation on the source and stand for Insert, Update, or Delete operations. For full load files, you can add this column by configuring the endpoint setting.

```
includeOpForFullLoad=true
```

This ensures that all full load files are marked with an I operation. When you use this approach, new subscribers can consume the entire data set or prepare a fresh copy in case of any downstream processing issues.

## Step 6: Create an AWS DMS task

After you configure the replication instance and endpoints, you need to analyze your source database. A good understanding of the workload helps plan an effective migration approach and minimize configuration issues. Find some important considerations following and learn how they apply to our walkthrough.

### Size and number of records

The volume of migrated records affects the full load completion time. It is difficult to predict the full load time up front, but testing with a replica of a production instance should provide a baseline. Use this estimate to decide whether you should parallelize the full load by using multiple tasks or by using the parallel load option. The insurance schema includes 9 tables. The claim table is the largest table, containing about 800 million records. We can increase the number of tables loaded in parallel to 30 to accommodate the partitions in the table if the full load is slow. The default value for the number of tables loaded in parallel is 8.

### Transactions per second

While full load is affected by the number of records, the ongoing replication performance relies on the number of transactions on the source Amazon RDS. Performance issues during change data capture (CDC) generally stem from resource constraints on the source database, replication instance, target database, and network bandwidth or throughput. Knowing average and peak Transactions Per Second (TPS) on the source and recording CDC throughput and latency metrics help baseline AWS DMS performance and identify an optimal task configuration. For more information, see [Replication task metrics](#).

In this walkthrough, we will track the CDC latency and throughput values after the task moves into the ongoing replication phase to baseline AWS DMS performance.

### Unsupported data types

Identify the data types used in your tables and check that AWS DMS supports these data types. For more information, see [Source data types for PostgreSQL](#).

After running the initial load test, validate that AWS DMS converted the data as you expected. You can also initiate a pre-migration assessment to identify any unsupported data types in the migration scope. For more information, see [Specifying individual assessments](#).

## Task configuration

In this walkthrough, incremental changes to the source tables need to be migrated to the data lake. So, we will be using the Full Load + CDC option. For more information about the task creation steps and available configuration options, see [Creating a task](#). We will first focus on the following settings.

## Table mappings

Use selection rules to define the schemas and tables that the AWS DMS task will migrate. For more information, see [Selection rules and actions](#).

In this walkthrough, we are migrating all the tables (%) in the insurance schema. Another option is to include each table explicitly in the table mappings. However, that increases operational overhead by requiring repeated configurations. If we plan to add new tables to the source database in the future under the sales history schema, we should include all tables (%) in the table mapping.

## Note

Mapping rules are applied at the task level. You need to add a mapping rule to each task that replicates data to your data lake. For our use case we needed just one task.

## LOB settings

AWS DMS handles large binary object (LOB) columns differently compared to other data types. For more information, see [Migrating large binary objects \(LOBs\)](#).

A detailed explanation of LOB handling by AWS DMS is out of scope for this walkthrough. However, remember that increasing the LobMaxSize value increases the task's memory utilization. Because of that, it is recommended not to set LobMaxSize to a large value. For more information about LOB settings, see [Task Configuration](#).

The source data warehouse schema in this walkthrough does not have LOB data. However, in case there were any LOB columns to be migrated, we would have done further analysis on such columns. Because AWS DMS does not support Full LOB Mode for Amazon S3 endpoints, we need to identify a suitable LobMaxSize value.

## Parallel load

Though, we used a large instance class in previous run, overall improvement was not significant as the data volume is relatively large (29 million records in the claim table alone). To further optimize



the performance, we used [parallel-load ranges option](#). Below is the mapping rule used for that option. As seen below, 11 boundaries are defined to cover data from 2012 to 2023 in 11 ranges. With this option, full load finished in about 1 hour 34 minutes. As a result, we were able to reduce the time taken to complete full load to almost 60% as compared to initial load.

```
{
  "rules": [
    {
      "rule-type": "selection",
      "rule-id": "463842200",
      "rule-name": "463842200",
      "object-locator": {
        "schema-name": "insurance",
        "table-name": "claim"
      },
      "rule-action": "include",
      "filters": []
    },
    {
      "rule-type": "table-settings",
      "rule-id": "653647497",
      "rule-name": "653647497",
      "object-locator": {
        "schema-name": "insurance",
        "table-name": "claim"
      },
      "parallel-load": {
        "type": "ranges",
        "columns": [
          "claim_requested_timestamp"
        ],
        "boundaries": [
          [
            "2013-01-01 00:00:00"
          ],
          [
            "2014-01-01 00:00:00"
          ],
          [
            "2015-01-01 00:00:00"
          ],
          [
            "2016-01-01 00:00:00"
```

```

        ],
        [
            "2017-01-01 00:00:00"
        ],
        [
            "2018-01-01 00:00:00"
        ],
        [
            "2019-01-01 00:00:00"
        ],
        [
            "2020-01-01 00:00:00"
        ],
        [
            "2021-01-01 00:00:00"
        ],
        [
            "2022-01-01 00:00:00"
        ]
    ]
}
}
}
    ]
}

```

## Other task settings

Choose **Enable CloudWatch Logs** to upload the AWS DMS task execution log to Amazon CloudWatch. You can use these logs to troubleshoot issues because they include error and warning messages, start and end times of the run, configuration issues, and so on. Changes to the task logging setting, such as enabling debug or trace, can also be helpful to diagnose performance issues.

## Note

CloudWatch log usage is charged at standard rates. For more information, see [Amazon CloudWatch pricing](#).

For **Target table preparation mode**, choose one of the following options: Do nothing, Truncate, or Drop. Use Truncate in data pipelines where the downstream systems rely on a

fresh dump of clean data and do not rely on historical data. In this walkthrough, we choose **Do nothing** because we want to control the retention of files from previous runs.

For **Maximum number of tables to load in parallel**, enter the number of parallel threads that AWS DMS initiates during full load. You can increase this value to improve the full load performance and minimize the load time when you have numerous tables. Since we have several partitions that can be loaded in parallel, we used the maximum value of 49.

### Note

Increasing this parameter induces additional load on the source database, replication instance, and target database.

To create a database migration task, do the following:

1. Open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose **Database migration tasks**, and then choose **Create task**.
3. On the **Create database migration task** page, enter the following information.

| For This Parameter                                  | Do This                                                                              |
|-----------------------------------------------------|--------------------------------------------------------------------------------------|
| <b>Task identifier</b>                              | Enter <code>pg-dms-s3-task</code> .                                                  |
| <b>Replication instance</b>                         | Choose <code>datalake-migration-ri</code> (the value that you configured on Step 1). |
| <b>Source database endpoint</b>                     | Choose <code>pg-dms-s3-source</code> (the value that you configured on Step 3).      |
| <b>Target database endpoint</b>                     | Choose <code>pg-dms-s3-target</code> (the value that you configured on Step 4).      |
| <b>Migration type</b>                               | Choose <code>Migrate existing data and replicate ongoing changes</code> .            |
| <b>Editing mode</b>                                 | Choose <code>Wizard</code> .                                                         |
| <b>Custom CDC stop mode for source transactions</b> | Choose <code>Disable custom CDC stop mode</code> .                                   |

|                                                                                                          |                                                                         |
|----------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------|
| <b>Target table preparation mode</b>                                                                     | Choose Do nothing.                                                      |
| <b>Stop task after full load completes</b>                                                               | Choose Don't stop.                                                      |
| <b>Include LOB columns in replication</b>                                                                | Choose Limited LOB mode.                                                |
| <b>Maximum LOB size (KB)</b>                                                                             | Enter 32.                                                               |
| <b>Advanced task settings → Full load tuning settings → Maximum number of tables to load in parallel</b> | Enter 20.                                                               |
| <b>Enable validation</b>                                                                                 | Turn off because Amazon S3 does not support validation with CSV format. |
| <b>Enable CloudWatch logs</b>                                                                            | Turn on.                                                                |

Leave the default values in the other fields and choose **Create task**.

The task begins immediately. The **Database migration tasks** section shows you the status of the migration task.

## Step 7: Run the AWS DMS tasks

After you create your AWS Database Migration Service (AWS DMS) task, run the task a few times to identify the full load run time and ongoing replication performance. You can validate that initial configurations work as expected. You can do this by monitoring and documenting resource utilization on the source database, replication instance, and target database. These details make up the initial baseline and help determine if you need further optimizations.

After you start the task, the full load operation starts loading tables. You can see the table load completion status in the **Table Statistics** section and the corresponding target files in the Amazon S3 bucket.

In this scenario, we had one task migrating the insurance claim schema which was 100 GB in size. The `claim` table in the insurance claim schema was the largest among them which contained all the history data with respect to claims. In a regular run with no parallel-load enabled the task took 2hrs and 34 mins to complete this is because the `claim` table was being migrated as a whole.

The screenshot below shows table statistics with a r5.4xlarge replication instance with the parallel-load ranges option set. We were able to improve the performance of the task, and it completed in one hour and 32 minutes with the parallelism. In case you have a data set which is taking too long to migrate, using parallel-load and increasing the MaxFullLoadSubTasks setting could be a way to improve performance.

**Table statistics (8)**

↺

Export to CSV

Validate again

Reload table data

< 1 >

| <input type="checkbox"/> | Schema name ▾ | Table ▾      | Load state ▾    | Elapsed load time ▾ | Full load rows ▾ | Total rows ▾ | Validation state ▾ | Validation pending ▾ | Validation failed ▾ | V... |
|--------------------------|---------------|--------------|-----------------|---------------------|------------------|--------------|--------------------|----------------------|---------------------|------|
| <input type="checkbox"/> | insurance     | customer     | Table completed | 2 m 56 s            | 76,000,000       | 76,000,000   | Not enabled        | 0                    | 0                   | 0    |
| <input type="checkbox"/> | insurance     | vehicle      | Table completed | 5 m 56 s            | 198,205,129      | 198,205,129  | Not enabled        | 0                    | 0                   | 0    |
| <input type="checkbox"/> | insurance     | manufacturer | Table completed | 5 m 35 s            | 107,160,312      | 107,160,312  | Not enabled        | 0                    | 0                   | 0    |
| <input type="checkbox"/> | insurance     | claim        | Table completed | 1 h 32 m 43 s       | 28,967,200       | 28,967,200   | Not enabled        | 0                    | 0                   | 0    |
| <input type="checkbox"/> | insurance     | policy       | Table completed | 9 m 8 s             | 195,217,728      | 195,217,728  | Not enabled        | 0                    | 0                   | 0    |
| <input type="checkbox"/> | insurance     | mod_vehicle  | Table completed | 30 s                | 16,777,216       | 16,777,216   | Not enabled        | 0                    | 0                   | 0    |
| <input type="checkbox"/> | insurance     | occupation   | Table completed | 5 m 40 s            | 107,160,313      | 107,160,313  | Not enabled        | 0                    | 0                   | 0    |
| <input type="checkbox"/> | insurance     | model        | Table completed | 1 m 47 s            | 19,710,300       | 19,710,300   | Not enabled        | 0                    | 0                   | 0    |

We covered most prerequisites that help avoid configuration related errors. If you observe issues when running the task, see [Troubleshooting migration tasks in AWS Database Migration Service](#), [Best practices for AWS Database Migration Service](#), or reach out to AWS Support for further assistance.

Optionally, you could choose to validate the successful completion of the data migration by querying the S3 data using the Athena console. You can execute count queries or aggregation queries on key metric columns, and compare the results with the source database to validate the migration task.

After you complete the migration, validate that your data migrated successfully and delete the AWS DMS resources that you created.

## Conclusion

In this walkthrough, we carried out a step-by-step migration of an insurance claim history data warehouse from PostgreSQL to an AWS S3 data lake. The data lake is used by our example company for data visualization and analysis use cases. We achieved the crucial business requirements by using AWS DMS. Try out these steps to migrate your data to an S3 data lake and explore how you can centralize your data with a low-cost solution.


# Migrating SQL Server Databases to Amazon RDS for SQL Server

This walkthrough gets you started with homogeneous database migration from Microsoft SQL Server to Amazon Relational Database Service (Amazon RDS) for SQL Server. This guide provides a quick overview of the data migration process and provides suggestions on how to select the best option to use.

Customers looking to migrate self-managed SQL Server databases to Amazon RDS for SQL Server, can use one of the three main approaches.

- Use a native database migration method such as backup and restore.
- Use a managed service such as AWS DMS.
- Use a native tool for full load and a managed AWS DMS service for ongoing replication. We call this strategy the *hybrid approach*.

The following diagram shows the hybrid approach. Here, we use one of the three native tools for full load, and AWS DMS for ongoing replication.



The hybrid approach provides the simplicity of the native tools with additional built-in capabilities of AWS DMS. These include:

- Data validation
- Customizable source object selection rules
- Data filtering
- Renaming target tables or columns
- Data transformations
- Data partitioning

This document describes in detail the three full load migration methods. This guide helps you evaluate each method for your migration requirements. In the end, you can find a brief description of how to use AWS DMS for ongoing replication.

## Topics

- [Summary](#)
- [Full load SQL Server database migration](#)
- [Full load SQL Server database migration options performance comparison](#)
- [Migrate SQL Server database with AWS DMS ongoing replication](#)

## Summary

The following table helps understand how each migration approach fits to different use cases.

| SQL Server native tools           | Data transformation | Table filtering | Metadata rename | Migration of secondary objects | Data validation |
|-----------------------------------|---------------------|-----------------|-----------------|--------------------------------|-----------------|
| Backup and restore                | No                  | No              | No              | Yes                            | No              |
| Import and export wizard          | Yes                 | Yes             | Yes             | No                             | Yes             |
| SQL Server - Generate and Publish | No                  | Yes             | No              | No                             | No              |

| SQL Server native tools                               | Data transformation | Table filtering | Metadata rename | Migration of secondary objects | Data validation |
|-------------------------------------------------------|---------------------|-----------------|-----------------|--------------------------------|-----------------|
| Scripts<br>Wizard and bulk copy program utility (bcp) |                     |                 |                 |                                |                 |

You can see that the SQL Server backup and restore has the best performance among the three full load options. This is the preferred approach where the database size is less than 16 TiB and when you don't have transformation or filtering requirements. Backup and restore has the additional advantage of migrating your secondary database objects such as stored procedures, functions, and so on.

SQL Server Import and Export Wizard supports a wide range of features. Consider this approach as the next option to evaluate if you don't need to migrate secondary database objects such as views, stored procedures, triggers, and so on. Also, use this approach to overcome the backup and restore limitations. SQL Server Import and Export can also be used for smaller migrations where ease of use considerations override the minor performance gains provided by SQL Server Backup and Restore.

Using Generate and Publish Scripts Wizard and bulk copy program utility (bcp) is slower than SQL Server Import and Export Wizard. You can use this approach in some cases because in bcp you can parallelize the load. That said, data files created by bcp may be orders of magnitude larger than the original table size. Because of this, you might need a significant amount of storage space when you use bcp to migrate in parallel.

## Full load SQL Server database migration

The full load migration phase populates the target database with a copy of the source data. In each section, you can find detailed information about the full load method and their results to help you choose the one that fits your use case. For all three methods, we use the [dms\\_sample](#) database as an example. The dms\_sample database includes tables, views, indexes, stored procedures, and other database objects.

### Topics



- [SQL Server database backup and restore using Amazon S3](#)
- [SQL Server import and export wizard](#)
- [Generate and Publish Scripts wizard and Bulk Copy Program Utility](#)

## SQL Server database backup and restore using Amazon S3

Backup and restore is the easiest and usually the preferred method for the initial load of the target database. In this method, you create a full backup of your self-managed SQL Server database, transfer it to an Amazon S3 bucket, and restore it to your Amazon RDS for SQL Server instance. For more information, see [Importing and exporting SQL Server databases using native backup and restore](#) in the *Amazon RDS User Guide*.

The backup and restore method is suitable for the following use cases:

- Your database size is less than 64 TiB.
- You want to carry out a lift and shift migration with no changes or minimal changes to the database. For example, you want to migrate secondary database objects such as users, views, stored procedures, triggers, and so on in addition to your data.
- Network connectivity between your on-premises data center and AWS is often congested or has frequent disconnects. Backup and restore gives you the flexibility to transmit backup files during non-business hours.

The backup and restore method has the following limitations:

- Amazon RDS for SQL Server supports native restore of databases up to 64 TiB in size. For SQL Server Express Edition databases, Amazon RDS supports native restore of up to 10 GiB.
- On Multi-AZ database instances, you can only natively restore databases that are backed up in full recovery model.
- The Amazon S3 bucket where you store your data, has to be located in the same AWS Region as your target Amazon RDS for SQL Server database instance.
- Restoring backups from one time zone to a different time zone isn't recommended.
- You can't transform or filter data at a table-level when you use backup and restore.
- When you need to migrate a subset of tables, you can't use backup and restore.

## Migration steps

At a high level, the steps involved in backup and restore are the following:

- Perform a full backup of the source database.
- Copy the backup file to an Amazon S3 bucket.
- Restore the backup from the Amazon S3 bucket onto the target Amazon RDS for SQL Server database.

We use the [dms\\_sample](#) database in the following example.

### Perform full backup

First, perform a full back up of the source database. RDS for Microsoft SQL Server currently limits data files to 5 TiB. If the database backup size is less than 5 TiB, you can use the following command.

```
Use [dms_sample]
GO

BACKUP DATABASE [dms_sample] TO
DISK = 'C:\Backup\dms_sample.bak'
WITH NOFORMAT, NOINIT,
NAME = 'Full Backup of dms_sample', SKIP, NOREWIND, NOUNLOAD, STATS = 10
Go
```

If your database is larger than 5 TiB, split the backup files. Make sure that each file is less than 5 TiB in size. For example, the size of the `dms_sample` database is 15 TiB. This means that we use three backup files.

```
Use [dms_sample]
GO

BACKUP DATABASE [dms_sample] TO
DISK = 'C:\Backup\dms_sample1.bak',
DISK = 'C:\Backup\dms_sample2.bak',
DISK = 'C:\Backup\dms_sample3.bak'
WITH NOFORMAT, NOINIT,
NAME = 'Full Backup of dms_sample', SKIP, NOREWIND, NOUNLOAD, STATS = 10
Go
```

## Copy backup files to Amazon S3

Now, use the AWS CLI to upload the backup file to an Amazon S3 bucket.

```
aws s3 cp C:\Backup\dms_sample.bak s3://sampledatabaseuswest2/
```

For multiple backup files, use the folder path to copy the backup files to an Amazon S3 bucket.

```
aws s3 cp "C:\Backup" s3://sampledatabaseuswest2/ --recursive
```

Make sure that you define an AWS Identity and Access Management (IAM) role to access the option group. An option group can specify features, called options, that are available for a particular Amazon RDS DB instance. When you associate a DB instance with an option group, the specified options and option settings are enabled for that DB instance

When you create this IAM role, attach a trust relationship and a permissions policy. For more information, see [Manually creating an IAM role for native backup and restore](#).

We create the `sql-server-backup-restore` role, and then use it when we configure the target Amazon RDS database.

## Restore your backup to the target database

To restore your backup, do the following:

1. Create an option group for the target database.
  - a. In the Amazon RDS console, choose **Option groups**, and then choose **Create option group**.
  - b. For **Name**, enter **SQLServerrestore**.
  - c. For **Description**, enter **SQLServerrestore**.
  - d. For **Engine**, choose **sqlserver-se**.
  - e. For **Major engine version**, choose **14.00**.
  - f. Choose **Create**.
2. Add the `SQLSERVER_BACKUP_RESTORE` option and the `sql-server-backup-restore` role to this option group to access S3 bucket.
  - a. On the **Option groups** page, choose the option group that you created.
  - b. For **Options**, choose **Add option**. The **Add option** page opens.

- c. For **Option name**, choose **SQLSERVER\_BACKUP\_RESTORE**.
  - d. For **IAM role**, choose the `sql-server-backup-restore` role.
3. Modify your Amazon RDS for SQL Server DB instance and attach this option group.
  - a. In the Amazon RDS console, choose **Databases**, and then choose your target database.
  - b. Choose **Modify**. The **Modify DB instance** page opens.
  - c. In the **Additional configuration** section, choose **SQLServerrestore** for **Option group**.

Now, you can restore the backup file from Amazon S3 into the target Amazon RDS for SQL Server database. To restore your database, call the `rds_restore_database` stored procedure. For more information, see [Restoring a database](#).

```
exec msdb.dbo.rds_restore_database
@restore_db_name='DMS',
@s3_arn_to_restore_from='arn:aws:s3:::sampledatabaseuswest2/dms_sample.bak';
```

To restore multiple backup files, use the following command.

```
exec msdb.dbo.rds_restore_database
@restore_db_name='DMS',
@s3_arn_to_restore_from='arn:aws:s3:::sampledatabaseuswest2/dms_sample*';
```

The preceding statement returns the ID of the task. You can use the following command to check the status of the restore using this task ID.

```
exec msdb.dbo.rds_task_status
[@db_name='DMS'],
[@task_id=<ID_number>];
```

Finally, use the following SQL command to get the log sequence number (LSN) of the on-premises source database backup. Then use this LSN to set up the change data capture (CDC) task in AWS DMS.

```
Use [dms_sample]
GO

SELECT [Current LSN], [Begin Time], Description FROM fn_dblog(NULL, NULL) Where
[Transaction Name] = 'Backup:CommitDifferentialBase'
```

## SQL Server import and export wizard

Microsoft SQL Server Import and Export Wizard is a high-performance option for data migration. It uses the SQL Server Integration Services (SSIS) framework. For more information, see [Import and Export Data with the SQL Server Import and Export Wizard](#) and [SQL Server Integration Services](#).

The Import and Export Wizard is suitable for the following use cases:

- To achieve high migration performance.
- To transform data during the migration. You can use the wizard to create SSIS packages and modify them in Visual Studio with an SSIS extension to achieve this.
- To rename the target tables or schemas during the migration.
- To migrate only the tables and avoid the migration of the secondary database objects such as users, views, stored procedures, triggers, foreign keys or functions.

The migration performance is affected by resource constraints of the host where you run the wizard. During the migration, all data is funneled through this host.

### Migration steps

Use the following steps to migrate all the tables and views from the `dms_sample` database to your target database.

Disable all constraints on the target DB instance before to the migration. The Import and Export Wizard copies tables in a random order. This may lead to failures if you enforce referential integrity on the target.


```
EXEC sp_msforeachtable 'ALTER TABLE ? NOCHECK CONSTRAINT all'
```

Make sure that you capture the current log sequence number (LSN) from the source database before your start the full load. To capture the current LSN, use the following command.

```
SELECT max([Current LSN]) FROM fn_dblog(NULL, NULL)
```

Then you can use this LSN to set up the change data capture (CDC) task in AWS DMS.

Open the SQL Server Import and Export Wizard from the Windows Start menu. Connect to your source and target databases and select the source tables and views. The following image shows the SQL Server Import and Export Wizard application window.



Choose **Next**, then choose **Run immediately**, and then choose **Finish**. The SQL Server Import and Export Wizard starts the migration. You can monitor the progress of your migration using the Performing Operation screen. For more information, see [Performing Operation \(SQL Server Import and Export Wizard\)](#).

Make sure that you turn on constraints after you complete the migration.

```
EXEC sp_msforeachtable 'ALTER TABLE ? CHECK CONSTRAINT all'
```

For more information, see [Get started with this simple example of the Import and Export Wizard](#).

## Generate and Publish Scripts wizard and Bulk Copy Program Utility

You can use the SQL Server Generate and Publish Scripts wizard to create Transact-SQL scripts for objects in your database. Then you can run the Bulk Copy Program Utility (bcp) to copy data from your Microsoft SQL Server instance into data files. Also, you can use bcp to import data into a table from data files. For more information, see [How to: Generate a Script \(SQL Server Management Studio\)](#) and [bcp Utility](#).

This approach is suitable for the following use cases:

- You don't transform data during the migration.
- You don't rename tables or schemas during the migration.
- You use referential integrity on target tables. In this case, bcp automatically suspends RI constraints on target tables during data import.
- You can script and migrate all database objects using the SQL Server Generate and Publish Scripts wizard.

This approach has the following limitations:

- You need to create schemas on your target database before you can use migration scripts.
- This approach is slower than the Import and Export Wizard.
- Data transformations aren't supported.
- In bcp, the error messages are limited to 512 bytes. This can make troubleshooting complicated.
- You run the bcp command for each table. This increases the complexity for large migrations.


## Migration Steps

At a high level, the steps involved in this approach are the following:

- Use Microsoft Generate and Publish Scripts wizard to create Transact-SQL scripts from source database.
- Use the created Transact-SQL scripts to create database objects in Target database.
- Use SQL Server Bulk Copy Program Utility (bcp) to export data from the source database to data files. Then, use bcp to import data from the data files into the target database table.

The following example shows how to migrate the `dms_sample` database using Generate and Publish Scripts wizard and Bulk Copy Program Utility.

Generate a Transact-SQL script for the source database tables. You can save the script as single file or save in a new query window.



Next, create database objects on the target database using the script that you generated in the previous step.

Run the following command on the source database to capture the current log sequence number (LSN). Then use this LSN to set up the change data capture (CDC) task in AWS DMS.



```
SELECT max([Current LSN]) FROM fn_dblog(NULL, NULL)
```

Use the Windows command prompt to export the source tables to data files with the bcp utility.

```
bcp [database_name.] schema.table_name out "data_file" -c -t -S [server_name] -d  
[database_name] -U [login] -P [password]
```

Import the data files created in the previous step into the target database with the bcp utility.

```
bcp [database_name.] schema.table_name in "data_file" -S [server_name] -d  
[database_name] -c -t
```

You can create a .bat file with all the bcp scripts to avoid running script one by one. The following code example shows the contents of this .bat file.


```
bcp dbo.export1 out C:\BCP\export1.dat -c -t -S source-server-name -d dms_sample -U  
dms_user -P password  
bcp dbo.export2 out C:\BCP\export2.dat -c -t -S source-server-name -d dms_sample -U  
dms_user -P password  
bcp dbo.export3 out C:\BCP\export3.dat -c -t -S source-server-name -d dms_sample -U  
dms_user -P password  
bcp dbo.export1 in C:\BCP\export1.dat -c -t -S target-server-name -d dms_sample -U  
dms_user -P password  
bcp dbo.export2 in C:\BCP\export2.dat -c -t -S target-server-name -d dms_sample -U  
dms_user -P password  
bcp dbo.export3 in C:\BCP\export3.dat -c -t -S target-server-name -d dms_sample -U  
dms_user -P password
```

## Full load SQL Server database migration options performance comparison

To compare the full load migration performance for all three methods, we used a test environment. In this environment, we populated the `dms_sample` database with 410.90 GB of data. We used the same on-premise SQL Server source and RDS SQL Server target databases to load data three times. For these data loads, we used the following methods:

- Backup and restore.
- Import and export wizard.
- Generate and publish scripts wizard and bulk copy program utility (bcp).

The following image represents the performance comparison of the three migration methods. We expect similar performance trends for larger datasets.



The elapsed time shown in the diagram is the actual migration time. It doesn't include the time spent on implementing prerequisites.

For the backup and restore method, we spent 4.24 hours. This time includes:

- 1.66 hours to backup the database.
- 1.75 hours to copy the data from backup location to Amazon S3.
- 0.88 hours to restore the data from the S3 bucket to Amazon RDS for SQL Server.

For the import and export wizard, we spent 8.58 hours.

For the bcp method, we spent 199 hours. This time includes:

- 0.01 hours to generate scripts.
- 0.01 hours to run the generated script on Amazon RDS for SQL Server.
- 27.88 hours to run the bcp statements for unloading data from on-premise SQL Server.
- 171.1 hours to run the bcp statements for loading data into Amazon RDS for SQL Server.

## Migrate SQL Server database with AWS DMS ongoing replication

After you complete the full load, set up ongoing replication using AWS DMS to keep the source and target databases synchronized. To configure the ongoing replication task, open the AWS DMS console. On the **Create database migration task** page, follow these three steps.

- For **Migration type**, select **Replicate ongoing changes**.
- Under **CDC start mode for source transactions**, select **Specify a log sequence number**.
- Under **System change number**, enter the SQL Server log sequence number that you captured during the full load.

For more information, see [Continuous replication tasks](#).

## Migrating from Amazon RDS for Oracle to Amazon RDS for PostgreSQL and Aurora PostgreSQL

Amazon Relational Database Service (Amazon RDS) for PostgreSQL and Amazon Aurora PostgreSQL-Compatible Edition have evolved as a strong and cost-effective alternatives to Oracle without the need for a software license or a server to manage. The journey from Amazon RDS for Oracle to Amazon RDS for PostgreSQL and Aurora PostgreSQL has never been easier. This guide provides a quick overview of the process and considerations to be made when moving existing workloads to Amazon RDS for PostgreSQL or Aurora PostgreSQL and some of the tools that can assist in the process. It complements a large body of detailed online reference guidance on every aspects of a migration, and serves to provide a birds eye view of the process.

This document focuses on migrating custom applications where you control the source code. If you operate a packaged vendor application on Oracle, you must determine if the vendor supports the new platform.

### Topics

- [Can My Oracle Database Migrate?](#)
- [Migration Strategies](#)
- [The 12 Step Migration Process](#)
  - [Future State Architecture Design](#)
  - [Database Schema Conversion](#)

- [Application Conversion or Remediation](#)
- [Script/ETL/Report Conversion](#)
- [Integration with Third-Party Applications](#)
- [Data Migration Mechanism](#)
- [Testing and Bug Fixing](#)
- [Performance Tuning](#)
- [Setup, DevOps, Integration, Deployment, and Security](#)
- [Documentation and Knowledge Transfer](#)
- [Project Management and Version Control](#)
- [Post-Production Support](#)
- [Automation](#)
- [Platform Differences](#)

## Can My Oracle Database Migrate?

To quickly see if your workload qualifies as a migration candidate, please use the **DMA Connect Application and Database Questionnaire** to sort out migration obstacles specific to your application. Consider the following questions. The more you answer **No**, the easier the migration to Amazon RDS for PostgreSQL or Aurora PostgreSQL will be.

| Application Questions                                                                                           | Comments                                                                                                                  |
|-----------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------|
| Are there Oracle dependent parts of the application that you can't modify by yourself?                          | If you don't control all of the code it can be difficult to change the underlying database.                               |
| Is the application commercial off the shelf, and not available for PostgreSQL?                                  | Unless the commercial off-the shelf software (COTS) application also supports PostgreSQL, it will not be able to migrate. |
| Does the application use specific methods to connect to an Oracle database such as Oracle Call Interface (OCI)? | Refactoring OCI calls to ODBC is not impossible, but typically an involved process.                                       |
| Does the application use Oracle specific libraries?                                                             | It could be challenging finding PostgreSQL replacements for Oracle specific libraries.                                    |

| Database Questions                                                            | Comments                                                                                               |
|-------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------|
| Does the database use any third party packages?                               | It could be challenging finding PostgreSQL replacements for Oracle specific packages.                  |
| Does the database use any data cartridges?                                    | It could be challenging finding PostgreSQL replacements for Oracle specific cartridges.                |
| Does the application use Oracle Forms or Application Express (APEX)?          | Completely replacing Forms or APEX with a non-Oracle solution is substantial.                          |
| Does the database use SQLJ or .NET stored procedures?                         | You can refactor external stored procedure code for use with PostgreSQL, but it adds development work. |
| Does the database use Oracle Streams?                                         | Some refactoring is required to replace Oracle Streams with a PostgreSQL-compatible solution.          |
| Does the database use Oracle Multi Media?                                     | Some refactoring is required to replace Oracle Multi Media with a PostgreSQL-compatible solution.      |
| Does the database use Oracle Locator?                                         | Depending on feature use, such a solution may be refactored to work with PostGIS 3.1.                  |
| Does the database use Oracle Java Virtual Machine (JVM)?                      | Detaching a Java application from Oracle JVM can be involved development work.                         |
| Does the database use Oracle Machine Learning or formerly Advanced Analytics? | The solution will have to be refactored to use similar functionality on AWS.                           |

## Migration Strategies

The options for dealing with a legacy application have often been described as the 6 R's. For more information, see [6 Strategies for Migrating Applications to the Cloud](#).

- Re-host

- Re-platform
- Repurchase
- Refactor/Re-architect
- Retire
- Retain


This document describes the steps to migrate database instances running on Amazon RDS for Oracle to Aurora PostgreSQL or Amazon RDS for PostgreSQL. This also details out the steps to Re-platform and Refactor the application(s) running on these databases.

Re-platforming and re-architecting a database application ranges from modifying the code to work with a different cloud-native database to also adopting other cloud-native operations such as serverless application architectures like Kubernetes. This document deal with the changes necessary to migrate to a new database with pointers to other available documentation.

## The 12 Step Migration Process

You may have an Oracle database in Amazon RDS for both production or non-production purposes, and it may just be convenience and familiarity that steered you to Oracle even though there is a licensing cost to this choice. It is certainly easier to continue with the database you know than something new, but sometimes there are few remaining reasons do so.

Everyone's Oracle application is special, and nobody has the same setup and needs for the future. To provide a single framework for database migrations, this guide organizes the work in 12 steps. These steps cover what is in scope for most migrations. You can use these steps in sequence for multiple purposes and you shouldn't see them as a strictly linear process. You can consider these steps as an overall arch of a migration project where individual steps and activities can be overlapped or swapped to fit specific project conditions. The following image shows the 12 steps with an approximate share of effort in a typical project.

**Steps in migration****Effort breakdown**

Each step will be described at a high level in order to allow the reader to skip to relevant topics in the following chapters.

## 1. Future State Architecture Design

The understanding of the current design and its requirements together with those of the future state are addressed here with deployment diagrams and feature or component mappings for things that will change as a result of the migration. This step defines the scope and architectural view of the migration.

## 2. Database Schema Conversion

Because we are migrating a database application from Oracle to Amazon RDS or Aurora, the database schema needs to change. Subtle differences in functionality and syntax need to accommodate the new platform and comprehensive tooling exists to automate this step. In this step we include replacements for any Oracle specific database feature which works differently on PostgreSQL.

## 3. Application Conversion or Remediation

The Oracle application may be implemented in any programming language like Java or C#, and often abstracts from the nature of the underlying database through an object relational model (ORM). But it is also common to have some reliance on the database syntax directly in the application code, and this step covers the necessary changes to the application code to work

with Amazon RDS for PostgreSQL. In this step we also include operating system dependencies like direct file access which may need to change on the new platform.

#### **4. Script/ETL/Report Conversion**

An Oracle application may move data in and out sideways in addition to the application for the purpose of reporting or data import/export. This can happen by executing a stored procedure or through external scripting and SQL\*Loader. Such scripts and PL/SQL statements and the operational framework will need to be modified to work with PostgreSQL.

#### **5. Integration with Third-Party Applications**

Few applications are islands and often connect to other applications and monitoring. These dependencies may be affected by the move to the PostgreSQL database platform. The monitoring of the database may need to use native AWS tools or the third-party applications use Oracle specific means of communication. These dependencies may already support PostgreSQL or suitable replacements will need to be found.

#### **6. Data Migration Mechanism**

As we move from one database platform to another the data needs to move as well. This will happen several times through the migration, first for testing purposes and later for production cutover. If there are multiple customers of the database application they may need to be migrated at different times once the application has been migrated.

#### **7. Testing and Bug Fixing**

Migration touches all the stored procedures and functions and may affect substantial parts of the application code. For this reason good testing is required both at the unit and system functional level.

#### **8. Performance Tuning**

Due to database platform differences and syntax, certain constructs or combinations of data objects may perform differently on the new platform. The performance tuning part of the migration resolves any bottlenecks that might have been created.

#### **9. Setup, DevOps, Integration, Deployment, and Security**

How the application is put together and deployed may be impacted by the migration, and many customers take the opportunity to embrace infrastructure as code for the first time in the context of a migration. In this step we also focus on the impact to application security. In this step we also address cutover planning.



## 10 Documentation and Knowledge Transfer

In order to support the application going forward it may be necessary to document the changes that happened to the application and the operational environment. Maintenance of the application will have been impacted by the change of database and certain application behavior may have changed. This is especially important if the migration is done by a different team from those maintaining the application.

## 11 Project Management and Version Control

A migration certainly involves people with different skills and often entirely different teams, and maybe an outside party. A successful project needs to be well planned and coordinated to execute on a predictable schedule. Version control is a crucial foundation for a migration since database code may not be managed in the same way as application code.

## 12 Post-Production Support

After the application is live, the migration team may need to stay around for a while to address any emerging problems on the new platform that were not caught by testing.

# Automation

This document references the freely available AWS Schema Conversion Tool (AWS SCT) for code conversion and the AWS Database Migration Service for data migration. For more information, see [Installing, verifying, and updating Schema Conversion Tool](#) and [Database Migration Service](#).

## Oracle application future state architecture design

When you migrate an Oracle application to use a different database like PostgreSQL you must capture the architecture of the existing application to ensure that all considerations are covered, we call that the current state architecture. The current state architecture describes the part of the application that matters to the migration from an architectural point of view. The same is true for the future state architecture which takes the new database platform into account. We don't need to describe everything, but some things, like external dependencies, are very important, and help us determine what work to do.

You may already have some favorite drawing tools for architecture diagrams such as Lucidchart, Visio or the freely available [Diagrams.net](#) which are all great choices as they supports AWS infrastructure symbols along with many others to describe the current and future environment. But the tool is less important than what is captured in the diagrams.

The architecture diagrams also serve the important role of defining the context of what is inside and outside the scope of work as a team collaborates on the task.

## Current State Architecture

There may be existing documentation on the database application which should be examined for currency and relevance. Let us review what is important for the migration work before we decide if more documentation is needed.


A **network diagram** is useful because it typically connects servers to each other, and servers to databases. It may also show the division into multiple availability or disaster recovery zones. This is useful because it shows potential server and network dependencies that must be addressed in the new architecture. A network diagram may also highlight important security considerations like multiple networks and internet connectivity.

A **component diagram** is useful if the application is comprised of multiple parts using different technologies which each may present the migration with their own challenges to address.

A **class diagram** is useful if it shows a specific persistence layer or a query factory where the migration can focus while leaving the rest of the application untouched.

A **data flow diagram** is useful because it directly shows parts of the value chain of information flowing inside and outside the application highlighting what additional code may needs to be changed.

The following image shows a simple network diagram that can help easily communicate current architecture.



## Future State Architecture


The future state architecture envisions the application using the new database, and potentially other services in the environment. It's a new version of the current state architecture diagrams with

certain parts replaced with the new components. This document will focus mainly on replacing Amazon RDS for Oracle with Amazon RDS for PostgreSQL or Aurora PostgreSQL.

## Transition Architecture

Depending on how involved your migration is, you may need a transition architecture by which we mean, infrastructure that is there only for migration purposes. Examples of transition architecture includes AWS DMS servers and other mediating or transformation platforms. Such infrastructure has to be provisioned, secured and removed after the migration to avoid additional vulnerability and cost.

The following image shows a transition architecture diagram.



For more information, see [AWS Well-Architected Framework](#).

## Oracle database schema conversion

Relational databases contain a tabular structure for data using basic data types and procedural code in the form of triggers, functions and stored procedures. Oracle uses the PL/SQL dialect which is different from the PL/pgSQL dialect of PostgreSQL and while some table definitions and queries may look the same, others will require modification. Doing so manually would be a substantial task, but fortunately there are freely available AWS tools to automate this job (See Automation section in the Introduction).

The AWS Schema Conversion Tool (AWS SCT) is capable of connecting to Oracle and reading all PL/SQL code directly from the source Oracle database and converting it to PostgreSQL PL/pgSQL.

AWS SCT will retrieve the DDL for tables, views, triggers, stored procedures and functions from the database, parse it and generate the equivalent PostgreSQL code.

Based on experience, AWS SCT fully converts 90+% of the database code which leaves less than 10% for the database expert to improve.

## Process

At a high level, the database conversion process works like this:

- Download and install AWS SCT (Linux or Windows).
- Download and install Oracle database drivers (you probably have those already).
- Download and install the PostgreSQL database drivers for Amazon RDS or Aurora PostgreSQL.
- Run AWS SCT and create a migration assessment report. For more information, see [Creating migration assessment reports](#).
- Run AWS SCT and automatically convert the database code. For more information, see [Converting database schemas](#).
- Fix any warnings and error in the database code conversion.

AWS SCT operates with default assumptions about mappings between Oracle and PostgreSQL which may or may not be optimal for your particular application due to the data you have in the database. Certain data type mappings may need to be changed to ensure good performance. As an example, a NUMBER datatype in Oracle is an extremely versatile container which without further qualification may be too expensive for the application. In this case you would look at the type of data contained in the NUMERIC column and its requirements for precision and scale, and then determine the best match for that in PostgreSQL with the appropriate precision and scale.

Once AWS SCT has automatically converted the DDL code, the developer needs to investigate any warnings and errors which need manual remediation. Warnings and Errors can happen for many reasons. AWS SCT does not have 100% coverage of all syntactical situations, and code inside the database can be corrupted or encrypted preventing AWS SCT from reading it. In these situations, the output DDL code is marked up with comments about the problem AWS SCT had with conversion, and ask the developer for help.

## Exceptions

There are exceptions to the automatic code conversion by AWS SCT like SQLJ, .NET Stored Procedures, Spatial data, RDF Graphs. But in each case there are good candidate replacement features like Lambda functions, PostGIS and Neptune.

## Interactive and Batch Modes

AWS SCT offers both an interactive GUI and a command line interface (CLI) which are useful in different situations. The user interface is good in a more interactive situations where the user needs to explore the schema and perhaps select only part of it for conversion. The CLI is good for automation in situations where DDL code might be coming from a different source such as reports. For more information, see [Script/ETL/Report Conversion](#).

## Schema Drift

If the original database schema changes during the timeframe of migration, this can be detected in AWS SCT which can compare the old and the new database schema and highlight the object that need to be updated. If an object was converted 100% or with few manual changes, that object can be converted again and remediated.

For more information, see [AWS Schema Conversion Tool User Guide](#), [Oracle Database 19c To Amazon Aurora with PostgreSQL Compatibility \(12.4\) Migration Playbook](#), and [AWS Schema Conversion Tool CLI and Interactive Mode Reference](#).

## Oracle application conversion or remediation

The Oracle application may be written in a variety of languages like C++, C# and Java, each with their own patterns for calling Oracle. A common case is the use of an object relational model (ORM) between the application code and the database which reduce the amount of PL/SQL that needs to be changed. Examples include Entity Framework and Hibernate which are also supported on PostgreSQL.

Oracle uses the PL/SQL dialect which is different from the PL/pgSQL dialect of PostgreSQL and while some table definitions and queries may look the same, others will require modification. Doing so manually would be a substantial task, but the freely available AWS Schema Conversion Tool (AWS SCT).

AWS SCT is capable of identifying and replacing embedded PL/SQL in the application code with the equivalent PostgreSQL code. For more information, see [Automation](#).

In addition to using AWS SCT, you must also examine the source code for possible issues like:

- Specific ORM or other data access framework and versions or in use and confirm its compatibility with the target engine.
- Modify database connection as appropriate for the new engine.
- Modify any table/entity mapping configuration or code as appropriate for the converted schema.
- Identify and refactor any vendor-specific driver functionality in use in the code.

## Process

At a high level, the application conversion process works like this:

1. Perform the database conversion. This is necessary because the PL/SQL conversion needs to know the schema of the database. For more information, see [Database Schema Conversion](#).
2. Run AWS SCT and automatically convert the application code. For more information, see [Converting application SQL](#).
3. Fix any warnings and errors in the application code conversion.

## Exceptions

There are exceptions to the automated application code conversion process. If the application uses the native Oracle Call Interface (OCI). In this case the developer must refactor the code to use ODBC or JDBC.

## Database migration script/ETL/report conversion

ETL is an acronym that stands for Extract, Transform and Load. The ETL process plays a central role in data integration strategies. ETL allows businesses to gather data from multiple sources and consolidate it into a single, centralized location. ETL also makes it possible for different types of data to work together.

ELT is similar to ETL. However, the primary difference between them is that the data transformation processes occur after the Raw data from the source have been extracted and loaded into a staging area. The transformation of the data may occur in the destination database or in the middle tier or via serverless tools that might reduce the cost of the data processing.

Transforming the data is a critical process that may provide significant value to the data. It's also the stage where the data could be cleansed, standardized, deduplicated, verified, sorted, shared, and much more.

The role of ELT or ETL in database migration projects is critical for any successful migration.

For the remainder of this document, ETL will also refer to ELT patterns.

ETL can be implemented in the database itself, in external scripts or in third-party tools such as Informatica, Talend, and so on. If the ETL is done using Oracle stored procedure, the freely available AWS Schema Conversion Tool (AWS SCT) is capable of converting the ETL code to AWS Glue. For more information, see [Automation](#).

## Process for Conversion to AWS Glue

If Python/Glue is a desired future state architecture for ETL code, and the ETL is implemented in the database, the conversion process works like this:

1. Perform the database conversion. This is necessary because the PL/SQL conversion needs to know the schema of the database. For more information, see [Database Schema Conversion](#).
2. Run AWS SCT, select the code involved in ETL and automatically convert the ETL code to AWS Glue. For more information, see [Converting ETL processes](#).
3. Fix any warnings and errors in the ETL code conversion.

## Process for Conversion of Stored Procedures

If ETL or report process is implemented in the database, then the database conversion takes care of converting the code, and only the method to call the stored procedures need to change.

## Process for Conversion of Scripts, Reports, and Third-Party ETL

If the ETL or Report code is available in scripts or hosted in third-party tools and those tools will be used in the future, then a custom process will have to be implemented:

1. Perform the database conversion. This is necessary because the PL/SQL conversion needs to know the schema of the database. For more information, see [Database Schema Conversion](#).
2. Extract PL/SQL statements from the third-party ETL or reporting tool into flat files, unless already available.

3. Write YAML configuration files for AWS SCT CLI to convert external files.
4. Run AWS SCT CLI on the external scripts using the YAML configuration files. For more information, see [AWS Schema Conversion Tool CLI and Interactive Mode Reference](#).
5. Fix any warnings and errors in the ETL or report code conversion.
6. Insert the converted PL/pgSQL code back into the third-party ETL or reporting tool, unless they stay as flat files.

## Oracle application migration and integration with third-party applications

Few applications are islands and your Oracle application is likely to integrate with other applications that are not themselves going to be migrated. Examples include ETL, reporting, and monitoring applications for alerts and logs. For more information, see [Script/ETL/Report Conversion](#).

If these third-party applications connect directly to the Oracle database, they are going to be affected by the migration. If they are packaged applications, the vendor may offer support for Amazon RDS and Aurora PostgreSQL and if they are custom, you may need to modify them to work with the migrated application. There are a wealth of resources on the partner network which complement any solution from AWS.

AWS native tools such as [Amazon Simple Notification Service](#), [Amazon RDS Performance Insights](#), [Amazon CloudWatch](#), and [Amazon Relational Database Service](#) are already integrated with the Amazon RDS and Aurora PostgreSQL database platform and are recommended for a full picture of the ongoing performance.

For more information, see [Engage with Amazon Web Services Partners](#).

## Amazon RDS for Oracle data migration mechanism

For testing purposes and for production cutover, data needs to be migrated from the old Amazon RDS for Oracle instance to the new Amazon RDS or Aurora PostgreSQL instance. Such a data migration requires knowledge of data type mapping and possibly incremental loading, depending on the size of the data and migration window.

For this purpose AWS Database Migration Service (AWS DMS) can be used to connect source and target databases to replicate the contents of the data in the most optimal way.



## Process

1. Create a replication server.
2. Create source and target endpoints that have connection information about your data stores.
3. Create one or more migration tasks to migrate data between the source and target data stores.


After you configured AWS DMS, you can perform the following operations:

- A full data migration from Oracle to PostgreSQL.
- An ongoing replication from Oracle to PostgreSQL.

Depending on the type of data in the database, you may need to optimize AWS DMS for handling certain data types like LOBS which you can read more about in the product guidance.

## Reverse Migration

Normally you just fall back to the old system if a migration fails during smoke testing, and in most cases you may decide to fix forward after cutover, in which case you fix any unforeseen bugs in the migrated system. But in some cases you may decide to have the option of migrating production data back from the new system to the original system after having been in production for a time. In those cases, a reverse data migration mechanism must be configured.



For more information, see [What is Database Migration Service?](#) and [Migrating Oracle databases with near-zero downtime.](#)

## Oracle database migration testing and bug fixing

Testing can be manual or automated. We recommend that you use an automated framework for testing. During migration, you will need to run the test multiple times, so having an automated testing framework helps speed up the bug fixing and optimization cycles.

### Unit Testing

Unit testing at the data level after migration can range from comparing every last bit of data in source and target by comparing extracted CSV files, but more realistically, custom aggregation queries should be constructed to incorporate large amounts of the migrated data and compare the results.

Unit tests validate individual units of your code, independent from any other components. Unit tests check the smallest unit of functionality and should have few reasons to fail.

Database objects need to be validated after migrating the DDL of an Oracle database to PostgreSQL. Database objects includes packages, tables, views, sequences, triggers, primary and foreign keys, indexes, constraints.

A typical way to perform unit testing on the converted database is to script out calls to stored procedures and functions and compare the returned data with external tools such as standard Linux/Unix tooling of diff.

The application needs to be validated with new and existing test case scenarios based on documented changes on database objects such as field names, types, minimum and maximum values, length, mandatory fields, field level validations etc.

### Functional Testing

Functional testing of the application is done by exercising user stories and comparing the results on the source and target system. This is typically a manual process, but third-party tools do exist to make automated regression tests of the UI (e.g. Selenium).

Functional testing of the database is largely done through the application, but there may be additional direct database use cases that can only be done directly on the database such as ETL for imports and extracts. In these cases, the data can be compared automatically before and after using standard Linux/Unix tooling like diff on extracted CSV files for example.

Functional testing of reports involve visual inspection to see that all fields are correctly displayed and comparison of the semantic values between the old and the new reports.

## Load Testing

In order to stress the migrated system and test its performance you may perform load testing which is typically done on a system that is scaled the same as production and requires a means of simulating load on the system. It is sometimes limited to running specific well-known expensive operations rather than user traffic.

## Standard Operating Procedures

Standard Operating Procedures (SOP) may be affected by a migration of database application. Database management procedures change when going to PostgreSQL and some procedures may be unnecessary when going to the highly managed Aurora PostgreSQL.

In any case, all existing operational procedures need to be tested and their language updated to reflect the new environment. For more information, see [Managing Amazon Aurora PostgreSQL](#).

## Monitoring

Monitoring of the database will be affected by the migration and some metrics may change which could affect how SLA is monitored. The way operational staff go from detecting a problem to diving into the underlying details may be affected. For more information, see [Monitor Amazon RDS for PostgreSQL and Amazon Aurora for PostgreSQL database log errors and set up notifications using Amazon CloudWatch](#).

## Cutover

Cutover procedures are the planned event where everything goes the way you want, but it still needs to be tested.

## Fallback

Fallback is when you have both old and new systems in sync with the new one operating as primary and you decide to switch back to the original which is still in sync.

## Rolling back the Oracle database migration to PostgreSQL

Rollback is usually the scenario when you don't have an ongoing replication mechanism to keep old and new in sync, so in the event of a no-go decision during the cutover, you abandon the new system and go back to the original.

## Migrate Back

In some rare cases, you may decide to include the option of migrating production from the new system back to the old system after having the cutover. If you include this scenario, it must be tested.

For more information, see [Testing Amazon Aurora PostgreSQL by using fault injection queries](#), [Automate benchmark tests for Amazon Aurora PostgreSQL](#), [Validating database objects after migration](#), and [Validate database objects after migrating from Oracle to Amazon Aurora PostgreSQL](#).

## Oracle database migration performance tuning

Any migration is likely to slightly change the performance of individual queries in the application and in stored procedures and functions. Depending on the context, those small differences may not matter in reality. But it is a good idea to deliberately compare the performance of operations that are known to be slow in the original system because any difference in performance is likely to be greater. Such testing is usually confined to specific long running ETL jobs and reports. Other performance issues may show up during functional or load testing and will be addressed as bugs.

## Oracle database migration to PostgreSQL setup, DevOps, integration, deployment, and security

Deployment to production is the culmination of the migration activity and is a high stakes effort which requires good planning and benefits from well tested automation.

With DevOps, you can create a process that helps easily deploy and update your virtual architecture in a scalable and repeatable way. This reduces the risk of human error. Furthermore, DevOps allow us to deploy much faster than humans which may be a factor in large deployments.

## Wave Planning

For any application or cluster of applications there is an important question of sequencing because every cutover window, for example, a weekend can only accommodate so much work. This means that larger portfolio may need to be migrated in multiple waves, and this makes wave planning necessary.

Wave planning considers that some parts of the application will move while other stay behind under different network, connectivity and security conditions. Different parts of the application

may also be under different ownership, so wave planning becomes the place where all stakeholders coordinate their efforts. Wave planning is a matter of minimizing risk during the overall migration.

## Infrastructure Automation

Infrastructure automation is a code layer that wraps API calls to a cloud provider with commands to provision infrastructure. Such code typically in YAML scripts are easily learned with any coding experience and are scalable and powerful. This layer will allow you to spin up one or one hundred web nodes nearly simultaneously. This layer is not designed to configure files on a server, install software on a server, or run commands on a server - That comes in the next section.

[Terraform](#) represents a cross cloud incarnation of this idea. The downside of Terraform is that its cross-cloud and open source nature makes it slower to adopt new features and provide detailed provisioning, often months after a new feature or configuration is released.

[CloudFormation](#) is a native AWS language in JSON or YAML format. Use AWS CloudFormation to write infrastructure code more specifically to specific features because it doesn't have to work with other clouds.

## Configuration Management

Configuration management systems manages configuration of software and state of files on a server or group of servers. These systems however are capable of much more than that, they also allow functionality such as installing software, running local commands, starting services and more in the same scalable way.

[Ansible](#) is a lightweight and easily installed tool that is configured using YAML files. It doesn't require a local daemon to be installed on the instances that Ansible is managing. The way that ansible does all this is through open source functions that essentially wrap cli commands that are run on remote hosts over Secure Shell Protocol (SSH). This allows for a plethora of functionality from database manipulation to package installation through simple changes in pre-written functions at the YAML level. Beyond a large library of open source function one can easily write custom functions (in python) or simply use the cli function to run any cli command through ansible remotely and in a scalable fashion. Some environments could be prevented from using Ansible due to limited or highly restricted SSH access to resources due to security protocols and standards.

[Puppet](#) works in a primary and secondary system that communicates over https (443) and is configured using its own language called [puppet](#). It's often found in enterprise level deployments, configuration management platforms based on a locally deployed daemon called a node. Puppet

differs from other similar platforms like Chef is its methodology in regards to how a resource acquires a desired state. Puppet takes a declarative approach, which is to say it defines the end state it requires, but makes not design on how it is achieved. Due to its fairly high level of technical investment in regards to its programming language, puppet is generally not recommended for smaller deployments as the investment.

[Chef](#) has a lot in common with puppet like a similar primary and secondary model, they both communicate over https, and are configured using a programming language. Where they differ is in terms of how they handles state. Chef takes an imperative approach, which is to say you as the end user have nearly full control on how a resource acquires a desired state which it achieved through Ruby as its configuration system. This type of deployment provides more flexibility as well as being easy to adopt if you are already using Ruby.

## Code Repository

A code repository offers safe storage of code and a change capture log which facilitates parallel development of a codebase by many developers simultaneously with the use of code branches, and integration with CI/CD pipelines. Other files than application source code might be stored in a Git repository such as infrastructure as code (IaC), database reports; recorded state changes or even short logs. Importantly you should never store credentials in the code repository. Credentials should be handled in other ways discussed below to avoid sensitive files being deleted or scrubbed, remnants left behind of the original state. The two most widely used code repositories are [GitHub](#) and [GitLab](#) with similar features and functionality.

## Secrets Management

A vault is for credentials. There are several options when it comes to Vault, many of which are baked into a lot of the technologies already described. Ansible has ansible vault, Jenkins has a functionality for storing and parametrize credentials which can be used at a smaller scale as a vault. However most of these baked in vaults don't generally have the effectiveness and capabilities of a dedicated vault.

[HashiCorp Vault](#) is a dedicated vault that differs from a secrets manager that comes packaged with another product is its varied capabilities around secret management. HashiCorp is an industry leader in this regards with capabilities such as dynamic secrets that can be generated on the fly for database or application credentials, data encryption, leasing and renewal which always credentials to be expired and rotated as well as the ability to revoke credentials remotely. In general, if an enterprise requires a wide range of credentials stored across a range a technologies, Vault is generally a good option to start.


## Orchestration

Orchestration is the glue that binds DevOps together. Many of the described technologies can be executed manually or from a scheduled script, however building on this idea of removing one of the largest points of failures such as humans from the actual deployment and management of infrastructure we can eliminate many of the pitfalls that can arise in the process of executing deployment scripts. Orchestration allows you to create a repeatable timeline, with logic gates, to deploy your infrastructure in exactly the order with the configurations chosen. This also allows deployments themselves to be tested before a change or deployment to production. Each of the configuration management platforms discussed above generally have an orchestration platform, they are generally focused on scheduling jobs within their particular vertical. For example, AWX for Ansible is mostly limited to scheduling ansible jobs.

[Jenkins](#) is managed through a GUI running on a primary node that is deployed on a resource within the company. There is also functionality to allow secondary nodes deployed on micro services for larger scale.

The process is arranged in a series of [Groovy](#) files called Jenkinsfile that dictate what action will be taken during each step in the pipeline. These jobs then can be scheduled jobs that periodically run and kick any job, from a ansible jobs that runs periodically to prevent configuration drift, to reporting jobs that execute a series of database calls. Generally Jenkins can connect any modern codebase and technology together.

An example pipeline you might run on Jenkins: use git to pull configuration scripts for the pipeline, then use a Terraform job from the cloned codebase to deploy a server, use Ansible to install and configure a database, then push a status file with information on the pipeline run back to git all while using Vault to manage the secrets for both access right for Jenkins and configuration of users on the database.



## Oracle database migration to PostgreSQL documentation and knowledge transfer

As a result of the migration, the operation of the database and the future development of the application and database will have been affected due to infrastructure and technology changes. If the migration is done by a separate team, it is vital that these changes be documented.

There may be a need for additional PostgreSQL training to operate the database and develop for PostgreSQL going forward.

## Oracle database migration to PostgreSQL project management and version control

Experience tells us that the steps take different amounts of efforts across a typical project.

When planning a migration project, tooling like AWS SCT can provide an important data point in the form of how much manual work needs to be done to fix database and application code that was not automatically converted. Using the above rules of thumb shares of the overall project, an initial plan can be created.

## Oracle database migration to PostgreSQL post-production support

After production cutover, there are a few possibilities for what can happen. You will have either decided to fix forward as the old system is being decommissioned, or you have decided to way a certain amount of time, a bake-in time, with production on the new system, during which a decision to abandon the new system can be made. Abandoning the new system has the following flavors:

- Roll back, all new data is lost.
- Roll back and reapply all new transactions.
- Roll back and migrate new production data back.
- Maintain a live replication back to the old system until bake-in period is over.

During this time, defects are tracked and triaged for possibly triggering the rollback or being fixed forward. Help desk will have been trained in the new system differences and will be able to detect if an end user inquiry just requires training or it may be a defect.



Beyond acceptance migration criteria, the application may have well defined KPIs defined already which can be observed when in production on the new system and compared to historical KPIs.

For more information, see [How to Migrate Your Oracle Database to PostgreSQL](#).

## Oracle and PostgreSQL platform differences

This section discusses some of the differences between Oracle and PostgreSQL to illustrate opportunities and challenges with migrating an Oracle application. This overview is by no means an exhaustive, however these are common challenges you may encounter when administering PostgreSQL after a background with Oracle.

### Range and List Partitions

Along with possible performance difference, architecturally partitions on Oracle and PostgreSQL act quite differently. On Oracle you can define a Range, for example each month or year being a partition, or a list, where every occurrence of say the letter "N" or "Y" in a char field is partitioned at the table definition level and PostgreSQL handles these operations differently. PostgreSQL operates with a "parent" table that holds no data and a "child" table that defines the partitions themselves and holds the data. The parent table is created first, the child tables is then defined with corresponding constraints to create the partition. You must supply a trigger to insert into the parent table and have the data be routed to the correct partition. For more information, see [Strategy for Migrating Partitioned Tables from Oracle to Amazon RDS for PostgreSQL and Amazon Aurora with PostgreSQL Compatibility](#).

### Data Types

Some PostgreSQL data types are much easier to work with than their corresponding Oracle types. For example, the Text type can store up to 1 GB of text and can be handled in SQL just like the char and varchar fields. They don't require special large object functions like character large objects (CLOBs) do.

However, there are some important differences to note. You can use the Numeric field in PostgreSQL to map any Number data types. But when you use it for joins (such as for a foreign key), it is less performant than using an int or bigint data type. This is a typical area where custom data type mapping should be considered.

The PostgreSQL Timestamp with time zone field is slightly different from and corresponds to the Oracle Timestamp with local time zone. These small differences can cause either performance issues or subtle application bugs that require thorough testing.

For more information, see [Migration tips for developers converting Oracle and SQL Server code to PostgreSQL](#).

## Transaction Control and Exception Handling

PostgreSQL Multiversion Concurrency Control (MVCC) is very different from Oracle rollback segments, even though they both provide ACID transactions. PostgreSQL creates a snapshot state taken at the start of the transaction, and essentially copies the data to a temporary page while the transaction is in flight. This can affect both the queries and the application, as well as hardware considerations. Because open transactions require temporary space to hold their snapshot state during the transaction, a workload that requires many open transactions in addition to possible deadlocks, needs to consider the location and parameters surrounding temporary table space.

Unlike Oracle, PostgreSQL uses auto-commit for transactions by default. However, there are two options to support explicit transactions, which are similar to the default behavior in Oracle (non-auto-commit). You can use the `START TRANSACTION` or `BEGIN TRANSACTION` statements and then `COMMIT` or `ROLLBACK`; or you can simply set `AUTOCOMMIT` to `OFF` at the session or system level.

PostgreSQL does not allow transaction control inside of PL/pgSQL like commit or roll back inside a stored procedure. The caller must perform the transaction management. If your existing PL/SQL contains explicit commit and rollback code it must be modified. When a run-time exception has occurred in a transaction, you must roll back that transaction before you can execute any new statement on the connection. Finally, exception handling in PL/pgSQL, using a `BEGIN...EXCEPTION...END` block to let your code catch any errors that occur. This block automatically creates a savepoint before the block, and rolls back to that savepoint when an exception occurs. You can then determine what logic to execute based on whether there was an error. Exception blocks are expensive however, due to the created savepoint. If you don't need to catch an error, or if you are planning to simply raise the error back to the calling application, don't use the exception block at all. Let the original error flow up to the application.

For more information, see [Oracle Database 19c To Amazon Aurora with PostgreSQL Compatibility \(12.4\) Migration Playbook](#).


## Migrating from SAP ASE to Amazon Aurora MySQL

Following, you can find a high-level outline and a step-by-step walkthrough that show the migration process of an on-premises SAP ASE database to Amazon Aurora MySQL-Compatible

Edition using AWS Database Migration Service (AWS DMS). Amazon Aurora is a highly available and managed relational database service with automatic scaling and high-performance features. The combination of MySQL compatibility with Aurora enterprise database capabilities provides an ideal target for commercial database migrations.

This walkthrough covers all steps in the migration from initial analysis of the source database to final cutover of applications to the target database.

The following diagram shows the basic architecture for the migration.



We use the **pubs2** database for SAP ASE as the example database in the rest of this document.

## Topics

- [Prerequisites for migrating from SAP AWS to Amazon Aurora MySQL](#)
- [Preparation and assessment for migrating from SAP ASE to Amazon Aurora MySQL](#)
- [SAP ASE to Amazon Aurora MySQL database code conversion and data loading](#)

- [Best practices for migrating from SAP ASE to Amazon Aurora MySQL](#)

## Prerequisites for migrating from SAP AWS to Amazon Aurora MySQL

The following prerequisites are required to complete this walkthrough:

- Familiarity with Amazon Relational Database Service (Amazon RDS), the applicable database technologies, and SQL.
- Understand the supported features and limitations of AWS Database Migration Service (AWS DMS). For more information, see [What Is Database Migration Service?](#).
- Accomplish the prerequisites required for using an SAP ASE database as a source for AWS DMS. For more information, see [Prerequisites for using an SAP ASE database as a source](#).
- Understand the limitations on using SAP ASE as a source and MySQL as a target for AWS DMS. For more information, see [Limitations on using SAP ASE as a source](#) and [Limitations on using a MySQL-compatible database as a target](#).
- Accomplish the prerequisites required for using a MySQL-compatible database as a target for AWS DMS. For more information, see [Using a MySQL-compatible database as a target](#).
- Set up the network for AWS DMS replication. This includes configuring VPC, private subnets, availability zone, and adding connections on the source firewall if it exists. For more information, see [Setting up a network for a replication instance](#).
- Download and install AWS Schema Conversion Tool (AWS SCT) with the required SAP ASE and MySQL JDBC drivers. For more information, see [Installing, verifying, and updating the Schema Conversion Tool](#).
- Know the recommendations on the most efficient way to use AWS DMS. For more information, see [Best practices](#).

## Preparation and assessment for migrating from SAP ASE to Amazon Aurora MySQL

Preparation and assessment of your source database is the initial phase. Before you start moving data, you should monitor and analyze the source database schema for the data lifecycle. To provide the best migration solution, you should have a good understanding on the workload, data access patterns, and data dependencies. Make sure to consider the following items:

- Character set.

- Largest table size.
- Largest LOB size.
- Integration with other databases or OS.

## Determine the Character Set

To find out the default character set and sort order for your SAP ASE database, run the following query:

```
exec sp_default_charset
```

If your application uses a different character set, you can find it from your session by checking the global variable @@client\_csname or @@client\_csid.

If you have a non-default character set that you want to migrate, use an extra connect attribute to specify the character set for the source database. For example, if your default character set is UTF8, specify charset=utf8 as an extra connect attribute to correctly migrate data.

## Determine the Largest Table Size

Explore your largest and busiest tables to find out their size and rate of change. This gives you an accurate estimate of where time will be spent when you do the initial data migration using the AWS Database Migration Service (AWS DMS) full load feature.

You can parallelize the load on the table level with one task to save time by using parallel load feature. For more information, see [Using parallel load for selected tables, views, and collections](#).

In SAP ASE, you can run only one replication thread for each database. Because of that, you can start one AWS DMS task at one time for each database. You can't run multiple tasks, which is common when you migrate from other database engines. For more information, see [Limitations on using SAP ASE as a source](#).

For SAP ASE version 15 and later, query sysobjects to list the top 10 in row count and space used. Use the following query:

```
select top 10 convert(varchar(30),o.name) AS table_name,  
    row_count(db_id(), o.id) AS row_count,  
    data_pages(db_id(), o.id, 0) AS pages,
```

```
data_pages(db_id(), o.id, 0) * (@@maxpagesize/1024) AS kbs
from sysobjects o
where type = 'U'
order by kbs DESC, table_name ASC
```

The output of this query is shown following.

```
table_name | row_count | pages | kbs |
salesdetail | 116 | 2 | 8 |
```

## Determine the Largest LOB Size

Large objects (LOBs) typically take the longest to migrate, unlike data types such as number and character, because of time spent encoding, storing, decoding, and retrieving them. You should identify tables with the TEXT, UNITEXT, and IMAGE data types, because AWS DMS converts these objects to LOB.

AWS DMS recommends to identify the size of LOB columns and choose limited or full mode, and max LOB size in the LOB settings appropriately. You can use the following dynamic SQL to generate a query for each table.

```
select 'select max(datalength('+ c.name +'))/1024 KB_SIZE from dbo.' + o.name + ';'
from sysobjects o,
     syscolumns c
where o.type = 'U' and
     o.id = c.id and
     c.type in (34,35,174);
```

The output of this query is shown following.

```
select max(datalength(pic))/1024 KB_SIZE from dbo.au_pix;
select max(datalength(copy))/1024 KB_SIZE from dbo.blurbs;
```

After you run the preceding queries, compare the results and choose the top one.

For example, we found the largest LOB column in our SAP ASE database is 51 KB. We used this number as input in our task settings with limited LOB mode.

The speed of the full load is improved with the limited LOB mode compared to the full LOB mode. For performance reasons, AWS DMS recommends to use limited LOB mode and increase the

maximum LOB size to cover the actual size you find from your query. For more information, see [How can I improve the speed of a migration task that has LOB data?](#).

## Document Integrations with Other Databases or Applications

Replace remote objects or interfaces in your code with other AWS services or equivalent external services.

For example, in the SAP ASE database, you may send out email using the `xp_sendmail` procedure. Because Amazon Aurora MySQL doesn't provide native support for sending emails, redesign the process. For example, use an AWS Lambda function to send email from the database. For more information, see [AWS Lambda](#) and [Sending notifications from Amazon Aurora MySQL](#).

### Note

For database links from the source database to a remote server in SAP ASE, update data using foreign data wrappers (FDW).

## SAP ASE to Amazon Aurora MySQL database code conversion and data loading

This section covers two major database migration tasks: code conversion and data load. You can use AWS Schema Conversion Tool (AWS SCT) to convert database schema objects such as tables, views, procedures, functions, and so on. Then you can use AWS Database Migration Service (AWS DMS) to load data.

### Database Schema Conversion

To convert your database schema and code objects from SAP ASE to Amazon Aurora MySQL, follow these steps.

1. Download and install AWS Schema Conversion Tool (AWS SCT) with the required SAP ASW and MySQL JDBC drivers. For more information, see [https://docs.aws.amazon.com/SchemaConversionTool/latest/userguide/CHAP\\_Installing.html](https://docs.aws.amazon.com/SchemaConversionTool/latest/userguide/CHAP_Installing.html).
2. Create a new AWS SCT project, add your source and target databases, and add a mapping rule. For more information, see [Creating a new project](#), [Adding database servers](#), and [Creating mapping rules](#).

3. Convert your database schema. For more information, see [Converting database schemas](#).
4. Save the converted SQL scripts. For more information, see [Saving and applying your converted schema](#).
5. Run these scripts against your target MySQL database. First, create tables with primary keys only. Then add the foreign keys and secondary indexes after you complete the full load.

## Migrate an SAP ASE Database to Amazon Aurora MySQL Using AWS DMS

This section covers the steps that you follow to migrate an SAP ASE database to Amazon Aurora MySQL using AWS DMS.

AWS DMS creates the schema in the target if the schema doesn't exist. However, AWS DMS only creates the tables with primary keys. AWS DMS doesn't create foreign keys or secondary indexes. Even the default values may be missing.

The best practice is to create the schema objects using the scripts that AWS SCT generated in the prior step, then start AWS DMS to load table data.

### Create a Replication Instance

To start data migration, create an AWS DMS replication instance. For performance reasons, AWS DMS recommends creating the replication instance in the same AWS Region as your target Amazon Aurora database. For more information, see [Creating a replication instance](#).

### Create a Source Endpoint

Create a source endpoint for SAP ASE and test the connection using the preceding replication instance.

- On the AWS DMS console, choose **Endpoints**.
- Choose **Create endpoint**.
- For **Endpoint type**, select **Source endpoint**.
- Enter your desired endpoint configuration.



## Endpoint configuration

### Endpoint identifier [Info](#)

A label for the endpoint to help you identify it.

Source-SAP-ASE-Database

### Descriptive Amazon Resource Name (ARN) - *optional*

A friendly name to override the default DMS ARN. You cannot modify it after creation.

Friendly-ARN-name

### Source engine

The type of database engine this endpoint is connected to.

SAP Sybase ASE

### Access to endpoint database

- ☐ AWS Secrets Manager
- ☒ Provide access information manually

### Server name

Server name

### Port

The port the database runs on for this endpoint.

5000

### Secure Socket Layer (SSL) mode

The type of Secure Socket Layer enforcement

none

You can use your own on-premises name server and a hostname instead of the IP address. For more information, see [Using your own on-premises name server](#).

- Select the endpoint that you created, and choose **Test connection** from the **Actions** drop-down menu.

To use Transport Layer Security (TLS) for an SAP ASE database version 15.7 and higher, use the Adaptive Server Enterprise 16.03.06 extra connection attribute (ECA) provider. Use the following example:

```
provider=Adaptive Server Enterprise 16.03.06;
```

Make sure that you open the database port to the IP range of your replication instance before you test the connection. If the firewall is open but you still experience a connection issue, please contact AWS support.

## Create a Target Endpoint

Create a target endpoint for your Amazon Aurora MySQL database.

- On the AWS DMS console, choose **Endpoints**.
- Choose **Create endpoint**.
- For **Endpoint type**, select **Target endpoint**.
- Enter your desired endpoint configuration.

### Endpoint configuration

#### Endpoint identifier [Info](#)

A label for the endpoint to help you identify it.

Target-Aurora-MySQL-Database

#### Descriptive Amazon Resource Name (ARN) - *optional*

A friendly name to override the default DMS ARN. You cannot modify it after creation.

Friendly-ARN-name

#### Target engine

The type of database engine this endpoint is connected to.

Amazon Aurora MySQL

#### Access to endpoint database

- ☐ AWS Secrets Manager
- ☒ Provide access information manually

#### Server name

aurora-mysql.cluster-exampleclustername.eu-west-1.rds.amazonaws.com

#### Port

The port the database runs on for this endpoint.

3306

#### Secure Socket Layer (SSL) mode

The type of Secure Socket Layer enforcement

none

#### User name [Info](#)

mysqluser

#### Password [Info](#)

••••••••

- Test the connection using the preceding replication instance.

To establish the connection, make sure that you edit the security group for your Amazon Aurora DB instance. Also, open the 3306 port on your MySQL database to the private IP or IP range of the replication instance.

- On the Amazon Relational Database Service (Amazon RDS) console, choose your Amazon Aurora MySQL DB instance.
- On the **Connectivity & security** tab, locate your security group name under **Security**.
- Choose the security group link. A new security group interface page opens.
- Choose **Inbound rules**.
- Choose **Edit inbound rules**.
- Add the IP range of the replication instance.

## Create a Migration Task

Create a migration task using the source and target endpoints that you created on the preceding step. For more information, see [Creating a task](#). After you create your task, AWS DMS sets its status to **Ready**. When you start or resume the task, AWS DMS changes the status to **Starting** or **Running**.

To monitor the process, choose **Task Monitoring**, **Table Statistics**, **Logs**. For more information, see [Monitoring Database Migration Service metrics](#) and [How can I enable monitoring for an database migration task?](#)

## Cutover Procedures

When the AWS DMS task finishes the full load and applies cached changes, the task moves to the change data capture (CDC) stage. At this point, you can perform the cutover to Amazon Aurora. You run SQL queries to validate data and use AWS services to set up backup and monitor jobs.

To perform the cutover, do the following:

- Analyze the database queries in Amazon Aurora MySQL and test the performance of critical queries.
- Shut down all the application servers and stop all the client connections to SAP ASE. Close any user sessions if necessary.
- Verify that the target data has been synced with the source database.

- Stop the AWS DMS task.
- Create the foreign keys and secondary indexes in Amazon Aurora MySQL if you didn't create them before the CDC stage started.
- Validate tables, views, procedures, functions, and triggers within your schema.
- Switch the application servers, clients, and jobs to the Amazon Aurora MySQL database.
- Create the CloudWatch alarms based on your desired DB metrics. For more information, see [Key Metrics for Amazon Aurora](#) and [Monitoring an Amazon Aurora DB Cluster](#).
- Add a reader node to an existing Amazon Aurora MySQL cluster. By default, Amazon Aurora replicates data across three Availability Zones in one Region at the storage level. This architecture is fault tolerant by design. For enhanced availability, add a reader node for a production database to automate failover in case of instance failure. Modify the database cluster to enable failover in case of instance failure. For more information, see [High Availability for Amazon Aurora](#).

## Troubleshooting

For more information about troubleshooting issues with AWS DMS, see [Troubleshooting migration tasks in Database Migration Service](#).

For more information about troubleshooting issues specific to using AWS DMS with SAP ASE databases, see [Troubleshooting issues with SAP ASE](#).

For more information about troubleshooting issues specific to using AWS DMS with Amazon Aurora MySQL databases, see [Troubleshooting issues with MySQL](#) and [Troubleshooting issues with Amazon Aurora MySQL](#).

## Best practices for migrating from SAP ASE to Amazon Aurora MySQL

- When the AWS Database Migration Service (AWS DMS) task completes the full load, AWS DMS stops this task. You can take this opportunity to add or enable your foreign keys or constraints and triggers in your target. If your migration type is **Migrate existing data and replicate ongoing changes**, resume the task to pick up the cached changes.
- When you have tasks running, you can monitor the on-premises source host, your replication instance, and your target Amazon Aurora database. Make sure that you create alarms and get notified on key metrics, such as central processing unit (CPU) utilization, freeable memory, and IOPS.

- Before you start the AWS DMS migration, make sure that you disable foreign keys and triggers on the target database. Additionally, make sure that your user has privileges on AWS DMS and Amazon Relational Database Service (Amazon RDS).
- AWS DMS recommends to periodically monitor the exception tables using the following query:

```
select STATEMENT from admin."awsdms_apply_exceptions" where TASK_NAME in ('TASK  
NAME')
```

- Monitoring the AWS DMS CloudWatch log for errors and warnings. In case of any problem with migration, AWS DMS records a corresponding warning or error message in the log.
- Set up monitoring of source and target database latency to understand the replication lag.
- Use the AWS DMS data validation feature to detect data issues.
- Test the steps designed for migration to understand any unforeseen issues.

This walkthrough covers proven end-to-end steps for migrating an SAP ASE database to Amazon Aurora MySQL using AWS DMS. The walkthrough also includes basic instructions that show how to perform a similar migration.

# Migrating Databases to the Amazon Web Services Cloud Using the Database Migration Service

You can use several AWS tools and services to migrate data from an external database to AWS. Depending on the type of database migration you are doing, you may find that the native migration tools for your database engine are also effective.

AWS Database Migration Service (AWS DMS) helps you migrate databases to AWS efficiently and securely. The source database can remain fully operational during the migration, minimizing downtime to applications that rely on the database. AWS DMS can migrate your Oracle data to the most widely used commercial and open-source databases on AWS.

AWS DMS migrates data, tables, and primary keys to the target database. All other database elements are not migrated. If you are migrating an Oracle database to Amazon Aurora MySQL-Compatible Edition, for example, you would want to use the AWS Schema Conversion Tool in conjunction with AWS DMS.

The AWS Schema Conversion Tool (AWS SCT) makes heterogeneous database migrations easy by automatically converting the source database schema and a majority of the custom code, including views, stored procedures, and functions, to a format compatible with the target database. Any code that cannot be automatically converted is clearly marked so that it can be manually converted. You can use this tool to convert your source Oracle databases to an Amazon Aurora MySQL, MySQL, or PostgreSQL target database on either Amazon RDS or EC2.

It is important to understand that AWS DMS and AWS SCT are two different tools and serve different needs.

- AWS DMS takes a minimalist approach and creates only those objects required to efficiently migrate the data for example tables with primary key. Therefore, we will use DMS to load the tables with data without any foreign keys or constraints. (We can also use the SCT to generate the table scripts and create it on the target before performing the load via DMS).
- We use AWS SCT:
  - To identify the issues, limitations and actions for the schema conversion
  - To generate the target schema scripts including foreign key and constraints
  - To convert code such as procedures and views from source to target and apply it on target

The size and type of Oracle database migration you want to do greatly determines the tools you should use. For example, a heterogeneous migration, where you are migrating from an Oracle database to a different database engine on AWS, is best accomplished using AWS DMS. A homogeneous migration, where you are migrating from an Oracle database to an Oracle database on AWS, is best accomplished using native Oracle tools.

## Topics

- [Migrating an On-Premises Oracle Database to Amazon Aurora MySQL](#)
- [Migrating an Amazon RDS for Oracle Database to Amazon Aurora MySQL](#)
- [Migrating a SQL Server Database to Amazon Aurora MySQL](#)
- [Migrating a SQL Server AlwaysOn Database on Primary Replica to Amazon Aurora PostgreSQL](#)
- [Migrating an Amazon RDS for Oracle Database to an Amazon S3 Data Lake](#)
- [Migrating an Amazon RDS for SQL Server Database to an Amazon S3 Data Lake](#)
- [Migrating an Oracle Database to PostgreSQL](#)
- [Migrating Oracle databases to Amazon Aurora MySQL with DMS Schema Conversion](#)
- [Migrating Oracle databases to Amazon RDS for PostgreSQL with DMS Schema Conversion](#)
- [Migrating SQL Server databases to Amazon Aurora PostgreSQL with DMS Schema Conversion](#)
- [Migrating SQL Server databases to Amazon RDS for MySQL with DMS Schema Conversion](#)
- [Migrating an Amazon RDS for Oracle Database to Amazon Redshift](#)
- [Migrating a BigQuery Project to Amazon Redshift](#)
- [Migrating a MySQL-Compatible Database to Amazon Aurora MySQL](#)
- [Migrating a MariaDB Database to Amazon RDS for MySQL or Amazon Aurora MySQL](#)
- [Migrating from MongoDB to Amazon DocumentDB](#)

## Migrating an On-Premises Oracle Database to Amazon Aurora MySQL

Following, you can find a high-level outline and also a complete step-by-step walkthrough that both show the process for migrating an on-premises Oracle database (the source endpoint) to an Amazon Aurora MySQL-Compatible Edition (the target endpoint) using AWS Database Migration Service (AWS DMS) and the AWS Schema Conversion Tool (AWS SCT).

AWS DMS migrates your data from your Oracle source into your Aurora MySQL target. AWS DMS also captures data manipulation language (DML) and data definition language (DDL) changes that happen on your source database and apply these changes to your target database. This way, AWS DMS helps keep your source and target databases in synch with each other. To facilitate the data migration, DMS creates tables and primary key indexes on the target database if necessary.

However, AWS DMS doesn't migrate your secondary indexes, sequences, default values, stored procedures, triggers, synonyms, views and other schema objects not specifically related to data migration. To migrate these objects to your Aurora MySQL target, use the AWS Schema Conversion Tool.

We highly recommend that you follow along using the Amazon sample database. To find a tutorial that uses the sample database and instructions on how to get a copy of the sample database, see [Working with the Sample Database for Migration](#).

If you've used AWS DMS before or you prefer clicking a mouse to reading, you probably want to work with the high-level outline. If you need the details and want a more measured approach (or run into questions), you probably want the step-by-step guide.

### **Topic: Migration from On-Premises Oracle to Aurora MySQL or Amazon RDS for MySQL**

**Time:**

**Cost:**

**Source Database:** Oracle

**Target Database:** Amazon Aurora MySQL/MySQL

**Restrictions:**

**Oracle Edition:** Enterprise, Standard, Express and Personal

**Oracle Version:** 10g (10.2 and later), 11g, 12c or higher

**MySQL or Related Database Version:** 5.5, 5.6, 5.7, MariaDB, Amazon Aurora MySQL



## Costs

For this walkthrough, you provision AWS Database Migration Service (AWS DMS) resources. You can use a t2.large replication instance with 50 GB of storage to keep your replication logs. Also, you provision an Amazon Aurora MySQL DB instance. You can use a db.r3.large Aurora MySQL DB instance with 10 GB of storage. Provisioning these resources will incur charges to your user by the hour.

To estimate what it will cost to run this walkthrough on AWS, you can use the AWS Pricing Calculator. For more information, see <https://calculator.aws/> and [Database Migration Service pricing](#).

To avoid additional charges, delete all resources after you complete the walkthrough.

## Migration from Oracle to Aurora MySQL using AWS DMS high-level outline

To migrate your data from Oracle to Aurora MySQL using AWS DMS, you take the following steps. If you've used AWS DMS before or prefer clicking a mouse to reading, the following summary should help you kick-start your migration. To get the details about migration or if you run into questions, see the step-by-step guide.

### Working with the Sample Database for Migration

We recommend working through the following outline and guide by using the sample Oracle database provided by Amazon. This database mimics a simple sporting event ticketing system. The scripts to generate the sample database can be found at <https://github.com/aws-samples/aws-database-migration-samples/tree/master/oracle/sampleldb/v1>.

To build the sample database, go to the `oracle/sampleldb/v1` folder and follow the instructions in the README.md file.

The sample creates approximately 8-10 GB of data. The sample database also includes a *ticketManagment* package, which you can use to generate some transactions. To generate transactions, log into SQL\*Plus or SQL Developer and run the following as `dms_sample`:

```
SQL>call generateTicketActivity(1000,0.01);
```

The first parameter is the transaction delay in seconds, the second is the number of transactions to generate. The preceding procedure simply sells tickets to people. You'll see updates to the tables: `sporting_event_ticket`, and `ticket_purchase_history`.

Once you've sold some tickets, you can transfer them using the command following:

```
SQL>call generateTransferActivity(100,0.1);
```

The first parameter is the transaction delay in seconds, the second is the number of transactions to generate. This procedure also updates `sporting_event_ticket` and `ticket_purchase_history`.

## Step 1: Prepare Your Oracle Source Database

To use AWS DMS to migrate data from an Oracle source database requires some preparation and we also recommend a few additional steps as best practices.

- **AWS DMS user** — It's a good practice to create a separate user for the specific purpose of migrating your data. This user should have the minimal set of privileges required to migrate your data. You can find specific details regarding those privileges later. If you are simply interested in testing AWS DMS on a non-production database, any DBA user will be sufficient.
- **Supplemental logging** — To capture changes, you must enable supplemental logging in order to use DMS. To enable supplemental logging at the database level issue the following command.

```
ALTER DATABASE ADD SUPPLEMENTAL LOG DATA
```

Additionally, AWS DMS requires for each table being migrated, you set at least key-level supplemental logging. AWS DMS automatically adds this supplemental logging for you if you include the following extra connection parameter for your source connection.

```
addSupplementalLogging=Y
```

- **Source database** – To migrate your data, the AWS DMS replication server needs access to your source database. Make sure that your firewall rules give the AWS DMS replication server ingress.

## Step 2: Launch and Prepare Your Aurora MySQL Target Database

Following are some things to consider when launching your Aurora MySQL instance:

- For best results, we recommend that you locate your Aurora MySQL instance and your replication instance in the same VPC and, if possible, the same Availability Zone.
- We recommend that you create a separate user with minimal privileges for migrating your data. The AWS DMS user needs the following privileges on all databases to which data is being migrated.

```
ALTER, CREATE, DROP, INDEX, INSERT, UPDATE, DELETE, SELECT
```

Additionally, AWS DMS needs complete access to the `awsdms_control` database. This database holds information required by AWS DMS specific to the migration. To provide access, run the following command.

```
ALL PRIVILEGES ON awsdms_control.* TO 'dms_user'
```

### Step 3: Launch a Replication Instance

The AWS DMS service connects to your source and target databases from a replication instance. Here are some things to consider when launching your replication instance:

- For best results, we recommend that you locate your replication instance in the same VPC and Availability Zone as your target database, in this case Aurora MySQL.
- If either your source or target database is outside of the VPC where you launch your replication server, the replication server must be publicly accessible.
- AWS DMS can consume a fair bit of memory and CPU. However, it's easy enough to scale up if necessary. If you anticipate running several tasks on a single replication server or if your migration involves a large number of tables, consider using one of the larger instances.
- The default storage is usually enough for most migrations.

### Step 4: Create a Source Endpoint

For AWS DMS to access your Oracle source database you'll need to create a source endpoint. The source endpoint defines all the information required for AWS DMS to connect to your source database from the replication server. Following are some requirements for the source endpoint.

- Your source endpoint needs to be accessible from the replication server. To allow this, you will likely need to modify your firewall rules to whitelist the replication server. You can find the IP address of your replication server in the AWS DMS Management Console.
- For AWS DMS to capture changes, Oracle requires supplemental logging be enabled. If you want AWS DMS to enable supplemental logging for you, add the following to the extra connection attributes for your Oracle source endpoint.

```
addSupplementalLogging=Y
```

## Step 5: Create a Target Endpoint

For AWS DMS to access your Aurora MySQL target database you'll need to create a target endpoint. The target endpoint defines all the information required for DMS to connect to your Aurora MySQL database.

- Your target endpoint needs to be accessible from the replication server. You might need to modify your security groups to make the target endpoint accessible.
- If you've pre-created the database on your target, it's a good idea to disable foreign key checks during the full load. To do so, add the following to your extra connection attributes.

```
initstmt=SET FOREIGN_KEY_CHECKS=0
```

## Step 6: Create and Run a Migration Task

A migration task tells AWS DMS where and how you want your data migrated. When you create a migration task, consider setting migration parameters as shown following.

**Endpoints and replication server** — Choose the endpoints and replication server created before.

**Migration type** — In most cases you'll want to choose **migrate existing data and replication ongoing changes**. With this option, AWS DMS loads your source data while capturing changes to that data. When the data is fully loaded, AWS DMS applies any outstanding changes and keeps the source and target databases in sync until the task is stopped.

**Target table preparation mode** — If you're having AWS DMS create your tables, **choose drop tables on target**. If you're using some other method to create your target tables such as the AWS Schema Conversion Tool, choose **truncate**.

**LOB parameters** — If you're just trying AWS DMS, choose **include LOB columns in replication, Limited LOB mode**, and set your **max LOB size to 16** (which is 16k.) For more information regarding LOBs, read the details in the step-by-step guide.

**Enable logging** — To help with debugging migration issues, always enable logging.

**Table mappings** — When migrating from Oracle to Aurora MySQL, we recommend that you convert your schema, table, and column names to lowercase. To do so, create a custom table mapping. The following example migrates the schema DMS\_SAMPLE and converts schema, table and column names to lower case.

```
{
  "rules": [
    {
      "rule-type": "selection",
      "rule-id": "1",
      "rule-name": "1",
      "object-locator": {
        "schema-name": "DMS_SAMPLE",
        "table-name": "%"
      },
      "rule-action": "include"
    },
    {
      "rule-type": "transformation",
      "rule-id": "6",
      "rule-name": "6",
      "rule-action": "convert-lowercase",
      "rule-target": "schema",
      "object-locator": {
        "schema-name": "%"
      }
    },
    {
      "rule-type": "transformation",
      "rule-id": "7",
      "rule-name": "7",
      "rule-action": "convert-lowercase",
      "rule-target": "table",
      "object-locator": {
        "schema-name": "%",
        "table-name": "%"
      }
    }
  ]
}
```

```
    },
    {
      "rule-type": "transformation",
      "rule-id": "8",
      "rule-name": "8",
      "rule-action": "convert-lowercase",
      "rule-target": "column",
      "object-locator": {
        "schema-name": "%",
        "table-name": "%",
        "column-name": "%"
      }
    }
  ]
}
```

## Step-by-step Oracle to Aurora MySQL using AWS DMS migration walkthrough

Following, you can find step-by-step instructions for migrating an Oracle database from an on-premises environment to Amazon Aurora MySQL. These instructions assume that you have already done the setting up steps for using AWS DMS located at [Setting up for Database Migration Service](#).

### Topics

- [Step 1: Configure Your Oracle Source Database](#)
- [Step 2: Configure Your Aurora Target Database](#)
- [Step 3: Create a Replication Instance](#)
- [Step 4: Create Your Oracle Source Endpoint](#)
- [Step 5: Create Your Aurora MySQL Target Endpoint](#)
- [Step 6: Create a Migration Task](#)
- [Step 7: Monitor Your Migration Task](#)
- [Troubleshooting](#)

## Step 1: Configure Your Oracle Source Database

To use Oracle as a source for AWS Database Migration Service (AWS DMS), you must first ensure that ARCHIVELOG MODE is on to provide information to LogMiner. AWS DMS uses LogMiner to read information from the archive logs so that AWS DMS can capture changes.

For AWS DMS to read this information, make sure the archive logs are retained on the database server as long as AWS DMS requires them. If you configure your task to begin capturing changes immediately, you should only need to retain archive logs for a little longer than the duration of the longest running transaction. Retaining archive logs for 24 hours is usually sufficient. If you configure your task to begin from a point in time in the past, archive logs need to be available from that time forward. For more specific instructions for enabling ARCHIVELOG MODE and ensuring log retention for your on-premises Oracle database see the [Oracle documentation](#).

To capture change data, AWS DMS requires supplemental logging to be enabled on your source database for AWS DMS. Minimal supplemental logging must be enabled at the database level. AWS DMS also requires that identification key logging be enabled. This option causes the database to place all columns of a row's primary key in the redo log file whenever a row containing a primary key is updated (even if no value in the primary key has changed). You can set this option at the database or table level.

To configure your Oracle source database, do the following:

### 1. Enable database-level supplemental logging

Run the following command to enable supplemental logging at the database level, which AWS DMS requires:

```
ALTER DATABASE ADD SUPPLEMENTAL LOG DATA;
```

### 2. Enable identification key supplemental logging

Use the following command to enable identification key supplemental logging at the database level. AWS DMS requires supplemental key logging at the database level unless you allow AWS DMS to automatically add supplemental logging as needed or enable key-level supplemental logging at the table level:

```
ALTER DATABASE ADD SUPPLEMENTAL LOG DATA (PRIMARY KEY) COLUMNS;
```

### 3. (Optional) Enable key level supplemental logging at the table level

Your source database incurs a small bit of overhead when key level supplemental logging is enabled. Therefore, if you are migrating only a subset of your tables, you might want to enable key level supplemental logging at the table level. To enable key level supplemental logging at the table level, use the following command.

```
alter table table_name add supplemental log data (PRIMARY KEY) columns;
```

If a table does not have a primary key you have two options:

- You can add supplemental logging to all columns involved in the first unique index on the table (sorted by index name.)
- You can add supplemental logging on all columns of the table.

To add supplemental logging on a subset of columns in a table, that is those involved in a unique index, run the following command.

```
ALTER TABLE table_name ADD SUPPLEMENTAL LOG GROUP example_log_group (ID,NAME)  
ALWAYS;
```

To add supplemental logging for all columns of a table, run the following command.

```
alter table table_name add supplemental log data (ALL) columns;
```

### 4. Create or configure a database account to be used by AWS DMS

We recommend that you use a user with the minimal privileges required by AWS DMS for your AWS DMS connection. AWS DMS requires the following privileges.

```
CREATE SESSION  
SELECT ANY TRANSACTION  
SELECT on V_$ARCHIVED_LOG  
SELECT on V_$LOG  
SELECT on V_$LOGFILE  
SELECT on V_$DATABASE  
SELECT on V_$THREAD  
SELECT on V_$PARAMETER  
SELECT on V_$NLS_PARAMETERS
```



```
SELECT on V_$TIMEZONE_NAMES
SELECT on V_$TRANSACTION
SELECT on ALL_INDEXES
SELECT on ALL_OBJECTS
SELECT on ALL_TABLES
SELECT on ALL_USERS
SELECT on ALL_CATALOG
SELECT on ALL_CONSTRAINTS
SELECT on ALL_CONS_COLUMNS
SELECT on ALL_TAB_COLS
SELECT on ALL_IND_COLUMNS
SELECT on ALL_LOG_GROUPS
SELECT on SYS.DBA_REGISTRY
SELECT on SYS.OBJ$
SELECT on DBA_TABLESPACES
SELECT on ALL_TAB_PARTITIONS
SELECT on ALL_ENCRYPTED_COLUMNS
* SELECT on all tables migrated
```

If you want to capture and apply changes (CDC) you also need the following privileges.

```
EXECUTE on DBMS_LOGMNR
SELECT on V_$LOGMNR_LOGS
SELECT on V_$LOGMNR_CONTENTS
LOGMINING /* For Oracle 12c and higher. */
* ALTER for any table being replicated (if you want to add supplemental logging)
```

For Oracle versions before 11.2.0.3, you need the following privileges. If views are exposed, you need the following privileges.

```
SELECT on DBA_OBJECTS /* versions before 11.2.0.3 */
SELECT on ALL_VIEWS (required if views are exposed)
```

## Step 2: Configure Your Aurora Target Database

As with your source database, it's a good idea to restrict access of the user you're connecting with. You can also create a temporary user that you can remove after the migration.

```
CREATE USER 'dms_user'@'%' IDENTIFIED BY 'dms_user';
GRANT ALTER, CREATE, DROP, INDEX, INSERT, UPDATE, DELETE,
SELECT ON <target database(s)>.* TO 'dms_user'@'%';
```

AWS DMS uses some control tables on the target in the database `awsdms_control`. The following command ensures that your `dms_user` has the necessary access to the `awsdms_control` database:

```
GRANT ALL PRIVILEGES ON awsdms_control.* TO 'dms_user'@'%';  
flush privileges;
```

### Step 3: Create a Replication Instance

An AWS DMS replication instance performs the actual data migration between source and target. The replication instance also caches the changes during the migration. How much CPU and memory capacity a replication instance has influences the overall time required for the migration. Use the following procedure to set the parameters for a replication instance.

To create an AWS DMS replication instance, do the following:

1. Sign in to the [AWS Management Console](https://console.aws.amazon.com/dms/v2/), and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/> and choose **Replication instances**. If you are signed in as an AWS Identity and Access Management (IAM) user, you must have the appropriate permissions to access AWS DMS. For more information on the permissions required, see [IAM Permissions](#).
2. Choose **Create replication instance**.
3. On the **Create replication instance** page, specify your replication instance information as shown following.

| For This Parameter    | Do This                                                                                                                                                               |
|-----------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Name</b>           | If you plan to launch multiple replication instances or share a user, choose a name that helps you quickly differentiate between the different replication instances. |
| <b>Description</b>    | A good description gives others an idea of what the replication instance is being used for and can prevent accidents.                                                 |
| <b>Instance class</b> | AWS DMS can use a fair bit of memory and CPU. If you have a large database (many tables) or use a number of LOB data types, setting up a larger instance is probably  |

| For This Parameter | Do This                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
|--------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                    | <p>better. As described following, you might be able to boost your throughput by running multiple tasks. Multiple tasks consume more resources and require a larger instance. Keep an eye on CPU and memory consumption as you run your tests. If you find you are using the full capacity of the CPU or swap space, you can easily scale up.</p>                                                                                                                                                                                                                                                     |
| <b>VPC</b>         | <p>Here you can choose the VPC where your replication instance will be launched. We recommend that, if possible, you select the same VPC where either your source or target database is (or both). AWS DMS needs to access your source and target database from within this VPC. If either or both of your database endpoints are outside of this VPC, modify your firewall rules to allow AWS DMS access.</p>                                                                                                                                                                                        |
| <b>Multi-AZ</b>    | <p>If you choose Multi-AZ, AWS DMS launches a primary and secondary replication instance in separate Availability Zones. In the case of a catastrophic disk failure, the primary replication instance automatically fails over to the secondary, preventing an interruption in service. In most situations, if you are performing a migration, you won't need Multi-AZ. If your initial data load takes a long time and you need to keep the source and target databases in sync for a significant portion of time, you might consider running your migration server in a Multi-AZ configuration.</p> |

| For This Parameter         | Do This                                                                                                                                                                  |
|----------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Publicly accessible</b> | If either your source or your target database are outside of the VPC where your replication instance is, you need to make your replication instance publicly accessible. |

4. In the **Advanced** section, set the **Allocated storage (GB)** parameter, and then choose **Next**.

| For This Option               | Do This                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
|-------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Allocated storage (GB)</b> | <p>Storage is primarily consumed by log files and cached transactions. For caches transactions, storage is used only when the cached transactions need to be written to disk. Therefore, AWS DMS doesn't use a significant amount of storage. Some exceptions include the following:</p> <ul style="list-style-type: none"> <li>* <i>Very large tables that incur a significant transaction load.</i> Loading a large table can take some time, so cached transactions are more likely to be written to disk during a large table load.</li> <li>* <i>Tasks that are configured to pause prior to loading cached transactions.</i> In this case, all transactions are cached until the full load completes for all tables. With this configuration, a fair amount of storage might be consumed by cached transactions.</li> <li>* <i>Tasks configured with tables being loaded into Amazon Redshift.</i> However, this configuration isn't an issue when Aurora MySQL is the target.</li> </ul> |

| For This Option                 | Do This                                                                                                                                                                                                                                |
|---------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                                 | In most cases, the default allocation of storage is sufficient. However, it's always a good idea to pay attention to storage related metrics and scale up your storage if you find you are consuming more than the default allocation. |
| <b>Replication Subnet Group</b> | If you run in a Multi-AZ configuration, you need at least two subnet groups.                                                                                                                                                           |
| <b>Availability Zone</b>        | If possible, locate your primary replication server in the same Availability Zone as your target database.                                                                                                                             |
| <b>VPC Security group(s)</b>    | With security groups you can control ingress and egress to your VPC. With AWS DMS you can associate one or more security groups with the VPC where your replication server launches.                                                   |
| <b>KMS key</b>                  | With AWS DMS, all data is encrypted at rest using a KMS encryption key. By default, AWS DMS creates a new encryption key for your replication server. However, you can use an existing key if desired.                                 |

## Step 4: Create Your Oracle Source Endpoint

While your replication instance is being created, you can specify the Oracle source endpoint using the [AWS Management Console](#). However, you can only test connectivity after the replication instance has been created, because the replication instance is used to test the connection.

To specify source or target database endpoints, do the following:

1. In the AWS DMS console, choose **Endpoints** on the navigation pane.
2. Choose **Create endpoint**. The **Create database endpoint** page appears, as shown following.

## Create database endpoint

A database endpoint is used by the replication server to connect to a database. The database specified in the endpoint can be on-premise, on RDS, in EC2 or in the cloud. Details should be specified in the form below. It is recommended that you test your endpoint connections here to avoid errors during processing.

Endpoint type\* ☒ Source ☐ Target ⓘ

Endpoint identifier\*  ⓘ

Source engine\*  ⓘ

Server name\*

Port\*  ⓘ

SSL mode\*  ⓘ

User name\*

Password\*

► Advanced

- Specify your connection information for the source Oracle database. The following table describes the source settings.

| For This Parameter  | Do This                                                                                                             |
|---------------------|---------------------------------------------------------------------------------------------------------------------|
| Endpoint type       | Choose <b>Source</b> .                                                                                              |
| Endpoint Identifier | Enter an identifier for your Oracle endpoint. The identifier for your endpoint must be unique within an AWS Region. |
| Source Engine       | Choose <b>oracle</b> .                                                                                              |
| Server name         | Enter an IP address that AWS DMS can use to connect to your database from the replication server.                   |

| For This Parameter | Do This                                                                                                                                                                                                             |
|--------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Port</b>        | Enter the port which your database is listening for connections (the Oracle default is 1521).                                                                                                                       |
| <b>SSL mode</b>    | Choose a Secure Sockets Layer (SSL) mode if you want to enable connection encryption for this endpoint. Depending on the mode you select, you might need to provide certificate and server certificate information. |
| <b>Username</b>    | Enter the user name. We recommend that you create a user specific to your migration.                                                                                                                                |
| <b>Password</b>    | Provide the password for the user name preceding.                                                                                                                                                                   |

4. Choose the **Advanced** tab to set values for extra connection strings and the encryption key.

| For This Option                    | Do This                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
|------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Extra connection attributes</b> | <p>Here you can add values for extra attributes that control the behavior of your endpoint. A few of the most relevant attributes are listed here. For the full list, see the documentation. Separate multiple entries from each other by using a semi-colon (;).</p> <p>* <b>addSupplementalLogging:</b> AWS DMS will automatically add supplemental logging if you enable this option (addSupplementalLogging=Y).</p> <p>* <b>useLogminerReader:</b> By default AWS DMS uses Oracle LogMiner to capture change data from the logs. AWS DMS can also parse the logs using its proprietary technology. If you</p> |

| For This Option | Do This                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
|-----------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                 | <p>use Oracle 12c and need to capture changes to tables that include LOBS, set this to No (useLogminerReader=N).</p> <p>* <b>numberDataTypeScale:</b> Oracle supports a NUMBER data type that has no precision or scale. By default, NUMBER is converted to a number with a precision of 38 and scale of 10, number(38,10). Valid values are 0—38 or -1 for FLOAT.</p> <p>* <b>archivedLogDestId:</b> This option specifies the destination of the archived redo logs. The value should be the same as the DEST_ID number in the \$archived_log table. When working with multiple log destinations (DEST_ID), we recommend that you specify a location identifier for archived redo logs. Doing so improves performance by ensuring that the correct logs are accessed from the outset. The default value for this option is 0.</p> |
| KMS key         | Choose the encryption key to use to encrypt replication storage and connection information. If you choose <b>(Default) aws/dms</b> , the default AWS KMS key associated with your user and region is used.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |

Before you save your endpoint, you can test it. To do so, select a VPC and replication instance from which to perform the test. As part of the test AWS DMS refreshes the list of schemas associated with the endpoint. (The schemas are presented as source options when creating a task using this source endpoint.)



## Step 5: Create Your Aurora MySQL Target Endpoint

Next, you can provide information for the target Amazon Aurora MySQL database by specifying the target endpoint settings.

To specify a target database endpoint, do the following:

1. In the AWS DMS console, choose **Endpoints** on the navigation pane.
2. Choose **Create endpoint**. The **Create endpoint** appears, as shown following.
3. Specify your connection information for the target Aurora MySQL database. The following table describes the target settings.

| For This Parameter             | Do This                                                                                                                                                                                       |
|--------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Endpoint type                  | Choose <b>Target endpoint</b> .                                                                                                                                                               |
| Endpoint identifier            | Enter an identifier for your Aurora MySQL endpoint. The identifier for your endpoint must be unique within an AWS Region.                                                                     |
| Target engine                  | Choose <b>Amazon Aurora MySQL</b> .                                                                                                                                                           |
| Access to endpoint database    | Choose <b>Provide access information manually</b> .                                                                                                                                           |
| Server name                    | Enter the writer endpoint for your Aurora MySQL instance. The writer endpoint is the primary instance.                                                                                        |
| Port                           | Enter the port assigned to the instance.                                                                                                                                                      |
| Secure Socket Layer (SSL) mode | Choose an SSL mode if you want to enable connection encryption for this endpoint. Depending on the mode you select, you might need to provide certificate and server certificate information. |

| For This Parameter | Do This                                                                                                                           |
|--------------------|-----------------------------------------------------------------------------------------------------------------------------------|
| User name          | Enter the user name for the user you are using for the migration. We recommend that you create a user specific to your migration. |
| Password           | Provide the password for the user name preceding.                                                                                 |

4. Define additional specific settings for your endpoints using wizard or editor in **Endpoint settings**.
5. Choose the encryption key to use to encrypt replication storage and connection information in **KMS key**. If you choose **(Default) aws/dms**, the default AWS Key Management Service (AWS KMS) key associated with your user and region is used.
6. Add tags to organize your DMS resources in **Tags**. You can use tags to manage your IAM roles and policies, and track your DMS costs.

Prior to saving your endpoint, you have an opportunity to test it in **Test endpoint connection (optional)**. To do so you'll need to choose a VPC and replication instance from which to perform the test.

## Step 6: Create a Migration Task

When you create a migration task you tell AWS DMS exactly how you want your data migrated. Within a task you define which tables you'd like migrated, where you'd like them migrated, and how you'd like them migrated. If you're planning to use the change capture and apply capability of AWS DMS it's important to know transactions are maintained within a single task. In other words, you should migrate all tables that participate in a single transaction together in the same task.

Using an AWS DMS task, you can specify what schema to migrate and one of the following types of migration:

- Migrate existing data
- Migrate existing data and replicate ongoing changes
- Replicate data changes only

This walkthrough migrates existing data only.

To create a migration task, do the following:

1. On the navigation pane, choose **Tasks**.
2. Choose **Create Task**.
3. On the **Create Task** page, specify the task options. The following table describes the settings.

| For This Option             | Do This                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
|-----------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Task name</b>            | It's always a good idea to give your task a descriptive name that helps organization.                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| <b>Task description</b>     | Enter a description for the task.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
| <b>Source endpoint</b>      | Select your source endpoint.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| <b>Target endpoint</b>      | Select your target endpoint.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| <b>Replication instance</b> | Select a replication instance on which to run the task. Remember, your source and target endpoints must be accessible from this instance.                                                                                                                                                                                                                                                                                                                                                                                                     |
| <b>Migration type</b>       | <p>You can use three different migration types with AWS DMS.</p> <p><i>1. Migrate existing data</i></p> <p>If you select this option, AWS DMS migrates only your existing data. Changes to your source data aren't captured and applied to your target. If you can afford taking an outage for the duration of the full load, migrating with this option is simple and straight forward. This method is also good to use when creating test copies of your database.</p> <p><i>2. Migrate existing data and replicate ongoing changes</i></p> |

| For This Option | Do This                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
|-----------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                 | <p>With this option, AWS DMS captures changes while migrating your existing data. AWS DMS continues to capture and apply changes even after the bulk data has been loaded. Eventually the source and target databases will be in sync, allowing for a minimal downtime migration. To do this, take the following steps:</p> <ul style="list-style-type: none"><li>* Shut the application down</li><li>* Let the final change flow through to the target</li><li>* Perform any administrative tasks such as enabling foreign keys and triggers</li><li>* Start the application pointing to the new target database</li></ul> <p>Note that AWS DMS loads the bulk data table-by-table, &lt;n&gt; tables at a time. As the full load progresses, AWS DMS begins applying cached changes to the target tables as soon as possible. During the bulk load, referential integrity is violated, therefore existing foreign keys must be disabled for the full load. Once the full load is complete, your target database has integrity and changes are applied as transactions.</p> <p><i>3. Replicate data changes only</i></p> <p>In some cases you might choose to load bulk data using a different method. This approach</p> |

| For This Option             | Do This                                                                                                                                                      |
|-----------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                             | generally only applies to homogeneous migrations.                                                                                                            |
| <b>Start task on create</b> | In most situations having the task start immediately is fine. Sometimes you might want to delay the start of a task, for instance, to change logging levels. |

4. Next, set the Advanced settings as shown following.

| For This Option                      | Do This                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                 |
|--------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Target table preparation mode</b> | <p>AWS DMS allows you to specify how you would like your target tables prepared prior to loading.</p> <p><b>Do nothing</b> - When you select this option, AWS DMS does nothing to prepare your tables. Your table structure remains as is and any existing data is left in the table. You can use this method to consolidate data from multiple systems.</p> <p><b>Drop tables on target</b> - Typically you use this option when you want AWS DMS to create your target table for you. When you select this option, AWS DMS drops and recreates the tables to migrate before migration.</p> <p><b>Truncate</b> - Select this option if you want to pre-create some or all of the tables on your target system, maybe with the AWS Schema Conversion Tool. When you select this option, AWS DMS truncates a target table prior to loading it. If the target table doesn't exist, AWS DMS creates the table for you.</p> |

| For This Option                           | Do This                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                           |
|-------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Include LOB columns in replication</b> | <p>Large objects, (LOBs) can sometimes be difficult to migrate between systems. AWS DMS offers a number of options to help with the tuning of LOB columns. To see which and when datatypes are considered LOBS by AWS DMS, see the AWS DMS documentation.</p> <p><b>Don't include LOB columns</b> - When you migrate data from one database to another, you might take the opportunity to rethink how your LOBs are stored, especially for heterogeneous migrations. If you want to do so, there's no need to migrate the LOB data.</p> <p><b>Full LOB mode</b> - In <b>full LOB mode</b> AWS DMS migrates all LOBs from source to target regardless of size. In this configuration, AWS DMS has no information about the maximum size of LOBs to expect. Thus, LOBs are migrated one at a time, piece by piece. Full LOB mode can be quite slow.</p> <p><b>Limited LOB mode</b> - In <b>limited LOB mode</b>, you set a maximum size LOB that AWS DMS should accept. Doing so allows AWS DMS to pre-allocate memory and load the LOB data in bulk. LOBs that exceed the maximum LOB size are truncated and a warning is issued to the log file. In <b>limited LOB mode</b> you get significant performance gains over <b>full LOB mode</b>. We recommend that you use <b>limited LOB mode</b> whenever possible.</p> <p>Note that with Oracle, LOBs are treated as VARCHAR data types whenever possible. This approach means AWS DMS fetches</p> |

| For This Option              | Do This                                                                                                                                                                                                                                                                                                                                                        |
|------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                              | them from the database in bulk, which is significantly faster than other methods. The maximum size of a VARCHAR in Oracle is 64K, therefore a limited LOB size of less than 64K is optimal when Oracle is your source database.                                                                                                                                |
| <b>Max LOB size (K)</b>      | When a task is configured to run in <b>limited LOB mode</b> , this option determines the maximum size LOB that AWS DMS accepts. Any LOBs that are larger than this value will be truncated to this value.                                                                                                                                                      |
| <b>LOB chunk size (K)</b>    | When a task is configured to use <b>full LOB mode</b> , AWS DMS retrieves LOBs in pieces. This option determines the size of each piece. When setting this option, pay particular attention to the maximum packet size allowed by your network configuration. If the LOB chunk size exceeds your maximum allowed packet size, you might see disconnect errors. |
| <b>Custom CDC start time</b> | This parameter pertains to tasks configured to replicate data changes only. It tells AWS DMS where to start looking for changes in the change stream.                                                                                                                                                                                                          |
| <b>Enable logging</b>        | Always enable logging.                                                                                                                                                                                                                                                                                                                                         |

## 5. Set additional parameters.

| For This Option                                 | Do This                                                                                  |
|-------------------------------------------------|------------------------------------------------------------------------------------------|
| <b>Create control table(s) in target schema</b> | AWS DMS requires some control tables in the target database. By default those tables are |

| For This Option                                     | Do This                                                                                                                                                                                         |
|-----------------------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                                                     | created in the same database as your data. This parameter allows you to tell AWS DMS to puts those artifacts somewhere else.                                                                    |
| <b>Maximum number of tables to load in parallel</b> | AWS DMS performs a table-by-table load of your data. This parameter allows you to control how many tables AWS DMS will load in parallel. The default is 8, which is optimal in most situations. |

## 6. Specify any table mapping settings.

Table mappings tell AWS DMS which tables a task should migrate from source to target. Table mappings are expressed in JSON, though some settings can be made using the [AWS Management Console](#). Table mappings can also include transformations such as changing table names from the upper case to lower case.

AWS DMS generates default table mappings for each (non-system) schema in the source database. In most cases you'll want to customize your table mapping. To customize your table mapping select the custom radio button. For details on creating table mappings see the AWS DMS documentation. The following table mapping does these things:

- It includes the DMS\_SAMPLE schema in the migration.
- It excludes the tables NFL\_DATA, MLB\_DATA, NAME\_DATE, and STADIUM\_DATA.
- It converts the schema, table, and column names to lower case.

```
{
  "rules": [
    {
      "rule-type": "selection",
      "rule-id": "1",
      "rule-name": "1",
      "object-locator": {
        "schema-name": "DMS_SAMPLE",
        "table-name": "%"
      },
      "rule-action": "include"
    },
  ],
}
```



```
{
  "rule-type": "selection",
  "rule-id": "2",
  "rule-name": "2",
  "object-locator": {
    "schema-name": "DMS_SAMPLE",
    "table-name": "MLB_DATA"
  },
  "rule-action": "exclude"
},
{
  "rule-type": "selection",
  "rule-id": "3",
  "rule-name": "3",
  "object-locator": {
    "schema-name": "DMS_SAMPLE",
    "table-name": "NAME_DATA"
  },
  "rule-action": "exclude"
},
{
  "rule-type": "selection",
  "rule-id": "4",
  "rule-name": "4",
  "object-locator": {
    "schema-name": "DMS_SAMPLE",
    "table-name": "NFL_DATA"
  },
  "rule-action": "exclude"
},
{
  "rule-type": "selection",
  "rule-id": "5",
  "rule-name": "5",
  "object-locator": {
    "schema-name": "DMS_SAMPLE",
    "table-name": "NFL_STADIUM_DATA"
  },
  "rule-action": "exclude"
},
{
  "rule-type": "transformation",
  "rule-id": "6",
```

```

    "rule-name": "6",
    "rule-action": "convert-lowercase",
    "rule-target": "schema",
    "object-locator": {
      "schema-name": "%"
    }
  },
  {
    "rule-type": "transformation",
    "rule-id": "7",
    "rule-name": "7",
    "rule-action": "convert-lowercase",
    "rule-target": "table",
    "object-locator": {
      "schema-name": "%",
      "table-name": "%"
    }
  },
  {
    "rule-type": "transformation",
    "rule-id": "8",
    "rule-name": "8",
    "rule-action": "convert-lowercase",
    "rule-target": "column",
    "object-locator": {
      "schema-name": "%",
      "table-name": "%",
      "column-name": "%"
    }
  }
]
}

```

## Step 7: Monitor Your Migration Task

Three sections in the console provide visibility into what your migration task is doing:

- Task monitoring — The **Task Monitoring** tab provides insight into your full load throughput and also your change capture and apply latencies.
- Table statistics — The **Table Statistics** tab provides detailed information on the number of rows processed, type and number of transactions processed, and also information on DDL operations.

- **Logs** — From the **Logs** tab you can view your task's log file, (assuming you turned logging on.) If for some reason your task fails, search this file for errors. Additionally, you can look in the file for any warnings. Any data truncation in your task appears as a warning in the log file. If you need to, you can increase the logging level by using the AWS Command Line Interface (AWS CLI).

## Troubleshooting

The two most common areas people have issues with when working with Oracle as a source and Aurora MySQL as a target are: supplemental logging and case sensitivity.

- **Supplemental logging** — With Oracle, in order to replication change data supplemental logging must be enabled. However, if you enable supplemental logging at the database level, it sometimes still need to enable it when creating new tables. The best remedy for this is to allow DMS to enable supplemental logging for you using the extra connection attribute:

```
addSupplementalLogging=Y
```

- **Case sensitivity** — Oracle is case-insensitive (unless you use quotes around your object names). However, text appears in the upper case. Thus, AWS DMS defaults to naming your target objects in the upper case. In most cases, you'll want to use transformations to change schema, table and column names to lower case.

For more tips, see the AWS DMS troubleshooting section in the [Troubleshooting migration tasks in Database Migration Service](#).

To troubleshoot issues specific to Oracle, see the Oracle troubleshooting section: [Troubleshooting issues with Oracle](#).

To troubleshoot Aurora MySQL issues, see the MySQL troubleshooting section: [Troubleshooting issues with MySQL](#).

## Migrating an Amazon RDS for Oracle Database to Amazon Aurora MySQL

This walkthrough gets you started with heterogeneous database migration from Amazon RDS for Oracle to Amazon Aurora MySQL-Compatible Edition using AWS Database Migration Service (AWS DMS) and the AWS Schema Conversion Tool (AWS SCT). This is an introductory exercise so does

not cover all scenarios but will provide you with a good understanding of the steps involved in executing such a migration.

It is important to understand that AWS DMS and AWS SCT are two different tools and serve different needs. They don't interact with each other in the migration process. At a high level, the steps involved in this migration are:

1. Using AWS SCT to:

- Run the conversion report for Oracle to Amazon Aurora MySQL to identify the issues, limitations, and actions required for the schema conversion.
  - Generate the schema scripts and apply them on the target before performing the data load via AWS DMS. AWS SCT will perform the necessary code conversion for objects like procedures and views.
2. Identify and implement solutions to the issues reported by AWS SCT. For example, an object type like Oracle Sequence that is not supported in the Amazon Aurora MySQL can be handled using the `auto_increment` option to populate surrogate keys or develop logic for sequences at the application layer.
3. Disable foreign keys or any other constraints which may impact the AWS DMS data load.
4. AWS DMS loads the data from source to target using the Full Load approach. Although AWS DMS is capable of creating objects in the target as part of the load, it follows a minimalistic approach to efficiently migrate the data so it doesn't copy the entire schema structure from source to target.
5. Perform post-migration activities such as creating additional indexes, enabling foreign keys, and making the necessary changes in the application to point to the new database.

This walkthrough uses a custom AWS CloudFormation template to create an Amazon RDS DB instances for Oracle and Amazon Aurora MySQL. It then uses a SQL command script to install a sample schema and data onto the Amazon RDS Oracle DB instance that you then migrate to Amazon Aurora MySQL.

This walkthrough takes approximately two hours to complete. The estimated cost to complete it, using AWS resources, is about \$5.00. Be sure to follow the instructions to delete resources at the end of this walkthrough to avoid additional charges.

## Topics

- [Costs](#)

- [Prerequisites for migrating from Amazon RDS for Oracle to Amazon Aurora MySQL](#)
- [Migration architecture for migrating from Amazon RDS for Oracle database to Amazon Aurora MySQL-Compatible Edition](#)
- [Step-by-step Amazon Relational Database Service to Amazon Aurora MySQL-Compatible Edition migration walkthrough](#)
- [AWS DMS migration from Amazon RDS for Oracle next steps](#)

## Costs

For this walkthrough, you provision Amazon Relational Database Service (Amazon RDS) resources by using AWS CloudFormation and also AWS Database Migration Service (AWS DMS) resources. Provisioning these resources will incur charges to your user by the hour. The AWS Schema Conversion Tool incurs no cost; it is provided as a part of AWS DMS.

Although you'll need only a minimum of resources for this walkthrough, some of these resources are not eligible for AWS Free Tier. At the end of this walkthrough, you'll find a section in which you delete the resources to avoid additional charges. Delete the resources as soon as you complete the walkthrough. For more information, see [chap-rdsoracle2aurora.steps.deleteresources](#).

To estimate what it will cost to run this walkthrough on AWS, you can use the AWS Pricing Calculator. For more information, see <https://calculator.aws/> and [Database Migration Service pricing](#).

The following table shows AWS DMS and Amazon RDS resources that you use for this walkthrough. You can specify these values in the AWS Pricing Calculator.

| AWS service                                                                           | Instance Type | Storage and I/O                         |
|---------------------------------------------------------------------------------------|---------------|-----------------------------------------|
| Amazon RDS for Oracle DB instance, License Included (Standard Edition Two), Single AZ | db.m3.medium  | Single AZ, 10 GB storage, GP2           |
| Amazon Aurora MySQL DB instance                                                       | db.r3.large   | Single AZ, 10 GB storage, 1 million I/O |

| AWS service                  | Instance Type | Storage and I/O                                        |
|------------------------------|---------------|--------------------------------------------------------|
| AWS DMS replication instance | t2.small      | 50 GB of storage for keeping replication logs included |

## Prerequisites for migrating from Amazon RDS for Oracle to Amazon Aurora MySQL

The following prerequisites are also required to complete this walkthrough:

- Familiarity with Amazon RDS, the applicable database technologies, and SQL.
- The custom scripts that include creating the tables to be migrated and SQL queries for confirming the migration, as listed following:
  - `Oracle-HR-Schema-Build.sql` — SQL statements to build the **HR** schema.
  - `Oracle_Aurora_For_DMSDemo.template` — an AWS CloudFormation template.

These scripts are available at the following link: [dms-sbs-RDSOracle2Aurora.zip](#).

Each step in the walkthrough also contains a link to download the file involved or includes the exact query in the step.

- A user with AWS Identity and Access Management (IAM) credentials that allow you to launch Amazon Relational Database Service (Amazon RDS) and AWS Database Migration Service (AWS DMS) instances in your AWS Region. For information about IAM credentials, see [Setting up for Amazon RDS](#).
- Basic knowledge of the Amazon Virtual Private Cloud (Amazon VPC) service and of security groups. For information about using Amazon VPC with Amazon RDS, see [Amazon VPC VPCs and Amazon RDS](#). For information about Amazon RDS security groups, see [Amazon RDS Security Groups](#).
- An understanding of the supported features and limitations of AWS DMS. For information about AWS DMS, see <https://docs.aws.amazon.com/dms/latest/userguide/Welcome.html>.
- Knowledge of the supported data type conversion options for Oracle and Amazon Aurora MySQL. For information about data types for Oracle as a source, see [Using an Oracle database as a source](#). For information about data types for Amazon Aurora MySQL as a target, see [Using a MySQL-Compatible database as a target](#).

For more information about AWS DMS, see [Getting started with Database Migration Service](#).

## Migration architecture for migrating from Amazon RDS for Oracle database to Amazon Aurora MySQL-Compatible Edition

This walkthrough uses AWS CloudFormation to create a simple network topology for database migration that includes the source database, the replication instance, and the target database in the same VPC. For more information about AWS CloudFormation, see the [AWS CloudFormation documentation](#).


We will provision the AWS resources that are required for this AWS Database Migration Service (AWS DMS) walkthrough through AWS CloudFormation. These resources include a VPC and Amazon Relational Database Service (Amazon RDS) instances for Oracle and Amazon Aurora MySQL-Compatible Edition. We provision through AWS CloudFormation because it simplifies the process, so we can concentrate on tasks related to data migration. When you create a stack from the AWS CloudFormation template, it provisions the following resources:

- A VPC with CIDR (10.0.0.0/24) with two public subnets in your region, DBSubnet1 at the address 10.0.0.0/26 in Availability Zone 1 (AZ 1) and DBSubnet2 at the address 10.0.0.64/26, in AZ 2.
- A DB subnet group that includes DBSubnet1 and DBSubnet2.
- Oracle RDS Standard Edition Two with these deployment options:
  - License Included
  - Single-AZ setup
  - db.m3.medium or equivalent instance class
  - Port 1521
  - Default option and parameter groups
- Amazon Aurora MySQL DB instance with these deployment options:
  - No replicas
  - db.r3.large or equivalent instance class
  - Port 3306
  - Default option and parameter groups
- A security group with ingress access from your computer or 0.0.0.0/0 (access from anywhere) based on the input parameter

We have designed the AWS CloudFormation template to require few inputs from the user. It provisions the necessary AWS resources with minimum recommended configurations. However, if you want to change some of the configurations and parameters, such as the VPC CIDR block and Amazon RDS instance types, feel free to update the template.

We will use the [AWS Management Console](#) to provision the AWS DMS resources, such as the replication instance, endpoints, and tasks. You will install client tools such as SQL Workbench/J and the AWS Schema Conversion Tool (AWS SCT) on your local computer to connect to the Amazon RDS instances.

Following is an illustration of the migration architecture for this walkthrough.



## Step-by-step Amazon Relational Database Service to Amazon Aurora MySQL-Compatible Edition migration walkthrough

In the following sections, you can find step-by-step instructions for migrating an Amazon Relational Database Service (Amazon RDS) for Oracle database to Amazon Aurora MySQL-Compatible Edition. These steps assume that you have already prepared your source database as described in preceding sections.

### Topics

- [Step 1: Launch the RDS Instances in a VPC by Using the AWS CloudFormation Template](#)



- [Step 2: Install the SQL Tools and AWS Schema Conversion Tool on Your Local Computer](#)
- [Step 3: Test Connectivity to the Oracle DB Instance and Create the Sample Schema](#)
- [Step 4: Test the Connectivity to the Aurora MySQL DB Instance](#)
- [Step 5: Use the AWS Schema Conversion Tool to Convert the Oracle Schema to Aurora MySQL](#)
- [Step 6: Validate the Schema Conversion](#)
- [Step 7: Create an AWS DMS Replication Instance](#)
- [Step 8: Create AWS DMS Source and Target Endpoints](#)
- [Step 9: Create and Run Your AWS DMS Migration Task](#)
- [Step 10: Verify That Your Data Migration Completed Successfully](#)
- [Step 11: Delete Walkthrough Resources](#)

## Step 1: Launch the RDS Instances in a VPC by Using the AWS CloudFormation Template

Before you begin, you'll need to download an AWS CloudFormation template. Follow these instructions:

1. Download the following archive to your computer: [dms-sbs-RDSOracle2Aurora.zip](#).
2. Extract the AWS CloudFormation template (Oracle\_Aurora\_For\_DMSDemo.template) from the archive.
3. Copy and paste the Oracle\_Aurora\_For\_DMSDemo.template file into your current directory.

Now you need to provision the necessary AWS resources for this walkthrough. Do the following:

1. Sign in to the AWS Management Console and open the AWS CloudFormation console at <https://console.aws.amazon.com/cloudformation>.
2. Choose **Create stack** and then choose **With new resources (standard)**.
3. On the **Specify template** section of the **Create stack** page, choose **Upload a template file**.
4. Click **Choose file**, and then choose the Oracle\_Aurora\_For\_DMSDemo.template file that you extracted from the dms-sbs-RDSOracle2Aurora.zip archive.
5. Choose **Next**. On the **Specify Details** page, provide parameter values as shown following.

| For This Parameter      | Do This                                                                                                                                                                                                                                                                                                                                                                    |
|-------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Stack Name</b>       | Enter DMSdemo.                                                                                                                                                                                                                                                                                                                                                             |
| <b>OracleDBName</b>     | Provide a unique name for your database. The name should begin with a letter. The default is ORCL.                                                                                                                                                                                                                                                                         |
| <b>OracleDBUsername</b> | Specify the admin (DBA) user for managing the Oracle instance. The default is oraadmin.                                                                                                                                                                                                                                                                                    |
| <b>OracleDBPassword</b> | Provide the password for the admin user. The default is oraadmin123 .                                                                                                                                                                                                                                                                                                      |
| <b>AuroraDBUsername</b> | Specify the admin (DBA) user for managing the Aurora MySQL instance. The default is auradmin .                                                                                                                                                                                                                                                                             |
| <b>AuroraDBPassword</b> | Provide the password for the admin user. The default is auradmin123 .                                                                                                                                                                                                                                                                                                      |
| <b>ClientIP</b>         | Specify the IP address in CIDR (x.x.x.x/32) format for your local computer. You can get your IP address from <a href="https://whatsmyip.org">whatsmyip.org</a> . Your RDS instances' security group will allow ingress to this IP address. The default is access from anywhere (0.0.0.0/0), which is not recommended; you should use your IP address for this walkthrough. |

# Specify stack details

## Stack name

Stack name

Stack name can include letters (A-Z and a-z), numbers (0-9), and dashes (-).

## Parameters

Parameters are defined in your template and allow you to input custom values when you create or update a stack.

### Source Oracle Database Configuration

OracleDBName

Enter Oracle Database name

OracleDBUsername

Enter database Admin username for Oracle

OracleDBPassword

Enter password for Oracle Admin user

### Target Aurora Database Configuration

AuroraDBUsername

Enter database Admin username for Aurora

AuroraDBPassword

Enter password for Aurora Admin user

### Enter IP address for DB Security group Configuration

ClientIP

The IP address range that can be used to connect to the RDS instances from your local machine. It must be a valid IP CIDR range of the form x.x.x.x/x. Pls get your address using [checkip.amazonaws.com](http://checkip.amazonaws.com) or [whatsmyip.org](http://whatsmyip.org)

6. Choose **Next**. On the **Configure stack options** page, shown following, choose **Next**.

## Configure stack options

### Tags

You can specify tags (key-value pairs) to apply to resources in your stack. You can add up to 50 unique tags for each stack. [Learn more](#) 

RemoveAdd tag

### Permissions

Choose an IAM role to explicitly define how CloudFormation can create, modify, or delete resources in the stack. If you don't choose a role, CloudFormation uses permissions based on your user credentials. [Learn more](#) 

#### IAM role - optional

Choose the IAM role for CloudFormation to use for all operations performed on the stack.

Remove

## Stack failure options

### Behavior on provisioning failure

Specify the roll back behavior for a stack failure. [Learn more](#) 

- ☒ **Roll back all stack resources**  
Roll back the stack to the last known stable state.
- ☐ **Preserve successfully provisioned resources**  
Preserves the state of successfully provisioned resources, while rolling back failed resources to the last known stable state. Resources without a last known stable state will be deleted upon the next stack operation.

## Advanced options

You can set additional options for your stack, like notification options and a stack policy. [Learn more](#) 

### ► Stack policy

Defines the resources that you want to protect from unintentional updates during a stack update.

### ► Rollback configuration

Specify alarms for CloudFormation to monitor when creating and updating the stack. If the operation breaches an alarm threshold, CloudFormation rolls it back. [Learn more](#) 

### ► Notification options

### ► Stack creation options

7. On the **Review** page, review the details, and if they are correct, scroll down and choose **Create stack**. You can get the estimated cost of running this AWS CloudFormation template by choosing **Estimate cost** at the **Template** section on top of the page.

## Review DMSdemo

Step 1: Specify template
Edit

### Template

Template URL

[https://s3-external-1.amazonaws.com/cf-templates-grtw7rybb40t-us-east-1/20212514ci-Oracle\\_Aurora\\_For\\_DMSDemo.template](https://s3-external-1.amazonaws.com/cf-templates-grtw7rybb40t-us-east-1/20212514ci-Oracle_Aurora_For_DMSDemo.template)

Stack description

This CloudFormation sample template OracletoAurora\_DMS creates Oracle and Aurora instances in a VPC which can be used to test the migration using AWS DMS service. You will be billed for the AWS resources used if you create a stack from this template

[Estimate cost](#)

Step 2: Specify stack details
Edit

### Parameters (6)

| Key              | Value           |
|------------------|-----------------|
| AuroraDBPassword | *****           |
| AuroraDBUsername | auradmin        |
| ClientIP         | 54.239.6.187/32 |
| OracleDBName     | ORCL            |
| OracleDBPassword | *****           |
| OracleDBUsername | oraadmin        |

8. AWS can take about 20 minutes or more to create the stack with Amazon RDS for Oracle and Amazon Aurora MySQL instances.

The screenshot shows the AWS CloudFormation console. At the top, there are buttons for 'Create Stack', 'Actions', and 'Design template'. Below these, there's a filter set to 'Active' and a search bar. The main table lists stacks with columns: Stack Name, Created Time, Status, and Description. The 'DMSdemo' stack is highlighted with a red circle. Below the table, the 'Events' tab is selected, showing a log of events for the 'DMSdemo' stack.

| Stack Name                          | Created Time                 | Status             | Description                               |
|-------------------------------------|------------------------------|--------------------|-------------------------------------------|
| <b>DMSdemo</b>                      | 2016-08-18 14:57:51 UTC-0700 | CREATE_IN_PROGRESS | This CloudFormation sample template C     |
| Basic CodeDeploy                    | 2016-08-10 13:56:20 UTC-0700 | CREATE_COMPLETE    | Basic CodeDeploy                          |
| CloudFormation sample template      | 2016-08-09 13:51:47 UTC-0700 | CREATE_COMPLETE    | Create instances ready for CodeDeploy     |
| CloudFormation sample template      | 2016-08-04 15:06:43 UTC-0700 | ROLLBACK_COMPLETE  | Create instances ready for CodeDeploy     |
| CloudFormation sample template      | 2016-08-01 10:56:18 UTC-0700 | CREATE_COMPLETE    | Create instances ready for CodeDeploy     |
| Build a VPC                         | 2016-07-11 10:07:14 UTC-0700 | CREATE_COMPLETE    | Builds a VPC, a subnet, a VPC with public |
| Attach an IAM user to an IAM group  | 2016-07-01 17:24:20 UTC-0700 | CREATE_COMPLETE    | Attaches an IAM user to an IAM group      |
| Create an IAM user                  | 2016-07-01 17:18:56 UTC-0700 | CREATE_COMPLETE    | Creates an IAM user                       |
| Assign an IAM permissions to an IAM | 2016-07-01 17:10:18 UTC-0700 | CREATE_COMPLETE    | Assigns an IAM permissions to an IAM      |

| 2016-08-18        | Status             | Type                       | Logical ID | Status reason  |
|-------------------|--------------------|----------------------------|------------|----------------|
| 14:57:51 UTC-0700 | CREATE_IN_PROGRESS | AWS::CloudFormation::Stack | DMSdemo    | User Initiated |

9. After the stack is created, choose **Stack**, select the DMSdemo stack, and then choose **Outputs**. Record the JDBC connection strings, **OracleJDBCConnectionString** and **AuroraJDBCConnectionString**, for use later in this walkthrough to connect to the Oracle and Aurora MySQL DB instances.

CloudFormation &gt; Stack List &gt; Stack Detail: DMSdemo

## DMSdemo

Other Actions ▾

Update Stack

Stack name: DMSdemo

Stack ID: arn:aws:cloudformation:us-west-2:100137374868:stack/DMSdemo/23aa30e0-6628-11e6-bd48-50d5ca0184d2

Status: **CREATE\_COMPLETE**

Status reason:

Description:

## ▼ Outputs

| Key                        | Value                                                                                                   | Description                                |
|----------------------------|---------------------------------------------------------------------------------------------------------|--------------------------------------------|
| OracleJDBCConnectionString | jdbc:oracle:thin:@do1xaskn0mfp18y.cq1w4tcs0mg7.us-west-2.rds.amazonaws.com:1521:ORCL                    | JDBC connection string for Oracle database |
| Regionname                 | us-west-2                                                                                               |                                            |
| StackName                  | DMSdemo                                                                                                 |                                            |
| AuroraJDBCConnectionString | jdbc:mysql://dmsdemo-auroracluster-1u1oynqny35jwv.cluster-cq1w4tcs0mg7.us-west-2.rds.amazonaws.com:3306 | JDBC connection string for Aurora database |

**Note**

Oracle 12c SE Two License version 12.1.0.2.v4 is available in all regions. However, Amazon Aurora MySQL is not available in all regions. Amazon Aurora MySQL is currently available in US East (N. Virginia), US West (Oregon), EU (Ireland), Asia Pacific (Tokyo), Asia Pacific (Mumbai), Asia Pacific (Sydney), and Asia Pacific (Seoul). If you try to create a stack in a region where Aurora MySQL is not available, creation fails with the error `Invalid DB Engine for AuroraCluster`.

## Step 2: Install the SQL Tools and AWS Schema Conversion Tool on Your Local Computer

Next, you need to install a SQL client and the AWS Schema Conversion Tool (AWS SCT) on your local computer.




This walkthrough assumes you will use the SQL Workbench/J client to connect to the RDS instances for migration validation. A few other software tools you might want to consider are the following:


- [JACK DB](#), an online web interface to work with RDS databases (Oracle and Aurora MySQL) over JDBC
- [DBVisualizer](#)
- [Oracle SQL Developer](#)

To install the SQL client software, do the following:

1. Download SQL Workbench/J from [the SQL Workbench/J website](#), and then install it on your local computer. This SQL client is free, open-source, and DBMS-independent.
2. Download the JDBC driver for your Oracle database release. For more information, go to <https://www.oracle.com/jdbc>.
3. Download the MySQL JDBC driver (`0.jar` file). For more information, go to <https://dev.mysql.com/downloads/connector/j/>.
4. Using SQL Workbench/J, configure JDBC drivers for Oracle and Aurora MySQL to set up connectivity, as described following.
  - a. In SQL Workbench/J, choose **File**, then choose **Manage Drivers**.
  - b. From the list of drivers, choose **Oracle**.
  - c. Choose the Open icon, then choose the `0.jar` file for the Oracle JDBC driver that you downloaded in the previous step. Choose **OK**.




- d. From the list of drivers, choose MySQL.
- e. Choose the Open icon, then choose the MySQL JDBC driver that you downloaded in the previous step. Choose **OK**.



To install the AWS Schema Conversion Tool and the required JDBC drivers, do the following:

1. Download the AWS Schema Conversion Tool from [Installing, verifying, and updating the Schema Conversion Tool](#).
2. Launch the AWS Schema Conversion Tool.
3. In the AWS Schema Conversion Tool, choose **Global settings** from **Settings**.
4. In **Global settings**, choose **Driver**, and then choose **Browse** for **Oracle driver path**. Locate the JDBC Oracle driver and choose **OK**. Next, choose **Browse** for **MySQL driver path**. Locate the JDBC MySQL driver and choose **OK**. Choose **OK** to close the dialog box.



### Step 3: Test Connectivity to the Oracle DB Instance and Create the Sample Schema

After the AWS CloudFormation stack has been created, test the connection to the Oracle DB instance by using SQL Workbench/J and then create the **HR** sample schema.


To test the connection to your Oracle DB instance and create the sample schema, do the following:

1. In SQL Workbench/J, choose **File**, then choose **Connect window**. Create a new connection profile using the following information as shown following

| For This Parameter      | Do This                                                 |
|-------------------------|---------------------------------------------------------|
| <b>New profile name</b> | Enter <code>RDSOracleConnection</code> .                |
| <b>Driver</b>           | Choose <code>Oracle (oracle.jdbc.OracleDriver)</code> . |

| For This Parameter | Do This                                                                                                                                    |
|--------------------|--------------------------------------------------------------------------------------------------------------------------------------------|
| <b>URL</b>         | Use the <b>OracleJDBCConnectionString</b> value you recorded when you examined the output details of the DMSdemo stack in a previous step. |
| <b>Username</b>    | Enter oraadmin.                                                                                                                            |
| <b>Password</b>    | Provide the password for the admin user that you assigned when creating the Oracle DB instance using the AWS CloudFormation template.      |

2. To test the connection, choose **Test**. Choose **OK** to close the dialog box, then choose **OK** to create the connection profile.



**Note**

If your connection is unsuccessful, ensure that the IP address you assigned when creating the AWS CloudFormation template is the one you are attempting to connect from. This is the most common issue when trying to connect to an instance.

3. Create the HR schema you will use for migration using a custom SQL script (Oracle-HR-Schema-Build.sql). To obtain this script, do the following:
  - a. Download the following archive to your computer: [dms-sbs-RDSOracle2Aurora.zip](#).
  - b. Extract the SQL script(Oracle-HR-Schema-Build.sql) from the archive.
  - c. Copy and paste the Oracle-HR-Schema-Build.sql file into your current directory.
4. Open the provided SQL script in a text editor. Copy the entire script.
5. In SQL Workbench/J, paste the SQL script in the Default.wksp window showing **Statement 1**.
6. Choose **SQL**, then choose **Execute All**.

When you run the script, you will get an error message indicating that user **HR** does not exist. You can ignore this error and run the script. The script drops the user before creating it, which generates the error.

7. Verify the object types and count in **HR** Schema were created successfully by running the following SQL query.

```
Select OBJECT_TYPE, COUNT(*) from dba_OBJECTS where owner='HR'
GROUP BY OBJECT_TYPE;
```

The results of this query should be similar to the following:

| OBJECT_TYPE | COUNT(*) |
|-------------|----------|
| INDEX       | 8        |
| PROCEDURE   | 2        |
| SEQUENCE    | 3        |
| TABLE       | 7        |
| VIEW        | 1        |

8. Verify the number of constraints in the **HR** schema by running the following SQL query:

```
Select CONSTRAINT_TYPE,COUNT(*) from dba_constraints where owner='HR'
AND (CONSTRAINT_TYPE IN ('P','R')OR SEARCH_CONDITION_VC NOT LIKE '%NOT NULL%')
```

```
GROUP BY CONSTRAINT_TYPE;
```

The results of this query should be similar to the following:

| CONSTRAINT_TYPE | COUNT(*) |
|-----------------|----------|
| R               | 10       |
| P               | 7        |
| C               | 1        |

9. Analyze the **HR** schema by running the following:

```
BEGIN
    dbms_stats.gather_schema_stats('HR');
END;
/
```

10. Verify the total number of tables and number of rows for each table by running the following SQL query:

```
SELECT table_name, num_rows from dba_tables where owner='HR' order by 1;
```

The results of this query should be similar to the following:

| TABLE_NAME  | NUM_ROWS |
|-------------|----------|
| COUNTRIES   | 25       |
| DEPARTMENTS | 27       |
| EMPLOYEES   | 107      |
| JOBS        | 19       |
| JOB_HISTORY | 10       |
| LOCATIONS   | 23       |
| REGIONS     | 4        |

11. Verify the relationships of the tables. Check the departments with employees greater than 10 by running the following SQL query:

```
Select b.department_name,count(*) from HR.Employees a,HR.departments b where
    a.department_id=b.department_id
group by b.department_name having count(*) > 10
order by 1;
```

The results of this query should be similar to the following:

|                 |           |
|-----------------|-----------|
| DEPARTMENT_NAME | COUNT( *) |
| Sales           | 34        |
| Shipping        | 45        |


### Step 4: Test the Connectivity to the Aurora MySQL DB Instance

Next, test your connection to your Aurora MySQL DB instance.

1. In SQL Workbench/J, choose **File**, then choose **Connect window**. Choose the Create a new connection profile icon. using the following information: Connect to the Aurora MySQL DB instance in SQL Workbench/J by using the information as shown following:

| For This Parameter | Do This                                                                                                                                     |
|--------------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| New profile name   | Enter RDSAuroraConnection .                                                                                                                 |
| Driver             | Choose MySQL (com.mysql.jdbc.Driver) .                                                                                                      |
| URL                | Use the <b>AuroraJDBCConnectionString</b> value you recorded when you examined the output details of the DMSdemo stack in a previous step.  |
| Username           | Enter auradmin.                                                                                                                             |
| Password           | Provide the password for the admin user that you assigned when creating the Aurora MySQL DB instance using the AWS CloudFormation template. |


2. Test the connection by choosing **Test**. Choose **OK** to close the dialog box, then choose OK to create the connection profile.

**Note**

If your connection is unsuccessful, ensure that the IP address you assigned when creating the AWS CloudFormation template is the one you are attempting to connect from. This is the most common issue when trying to connect to an instance.

3. Log on to the Aurora MySQL instance by using the master admin credentials.
4. Verify your connectivity to the Aurora MySQL DB instance by running a sample SQL command, such as `SHOW DATABASES;`.





## Step 5: Use the AWS Schema Conversion Tool to Convert the Oracle Schema to Aurora MySQL

Before you migrate data to Aurora MySQL, you convert the Oracle schema to an Aurora MySQL schema. [This video covers all the steps of this process.](#)

To convert an Oracle schema to an Aurora MySQL schema using AWS Schema Conversion Tool (AWS SCT), do the following:

1. Launch AWS SCT. In AWS SCT, choose **File**, then choose **New Project**. Create a new project named DMSDemoProject, specify the **Location** of the project folder, and then choose **OK**.
2. Choose **Add source** to add a source Oracle database to your project, then choose **Oracle**, and choose **Next**.
3. Enter the following information, and then choose **Test Connection**.

| For This Parameter     | Do This                                                                                |
|------------------------|----------------------------------------------------------------------------------------|
| <b>Connection name</b> | Enter Amazon RDS for Oracle. AWS SCT displays this name in the tree in the left panel. |
| <b>Type</b>            | Choose <b>SID</b> .                                                                    |
| <b>Server name</b>     | Use the <b>OracleJDBCConnectionString</b> value you used to connect to the Oracle      |

| For This Parameter | Do This                                                                                                                                                                                                                                                                                                                                                                                           |
|--------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                    | DB instance, but remove the JDBC prefix information. For example, a sample connection string you use with SQL Workbench/J might be "jdbc:oracle:thin:@do1xa4grferti8y.cqiw4tcs0mg7.us-west-2.rds.amazonaws.com:1521:ORCL". For AWS SCT <b>Server name</b> , you remove "jdbc:oracle:thin:@//" and ":1521" to use just the server name: "do1xa4grferti8y.cqiw4tcs0mg7.us-west-2.rds.amazonaws.com" |
| <b>Server port</b> | Enter 1521.                                                                                                                                                                                                                                                                                                                                                                                       |
| <b>Oracle SID</b>  | Enter ORCL.                                                                                                                                                                                                                                                                                                                                                                                       |
| <b>User name</b>   | Enter oraadmin.                                                                                                                                                                                                                                                                                                                                                                                   |
| <b>Password</b>    | Enter the password for the admin user that you assigned when creating the Oracle DB instance using the AWS CloudFormation template.                                                                                                                                                                                                                                                               |

**Add source**

**CONNECTION** SSL

Specify parameters for new connections to Oracle

Connection name: Amazon RDS for Oracle

Type: SID

Server name: do1xa4grferti8y.cqiw4tcs0mg7.us-west-2.rds.amazonaws.com

Server port: 1521

Oracle SID: ORCL

User name: oraadmin

Password: ••••••••

☐ Use SSL

☐ Store password

Test connection Previous Connect

4. Choose **OK** to close the alert box, then choose **Connect** to close the dialog box and to connect to the Oracle DB instance.
5. Choose **Add target** to add a target Amazon Aurora MySQL database to your project, then choose **Amazon Aurora (MySQL compatible)**, and choose **Next**.
6. Enter the following information and then choose **Test Connection**.

| For This Parameter     | Do This                                                                        |
|------------------------|--------------------------------------------------------------------------------|
| <b>Connection name</b> | Enter Aurora MySQL. AWS SCT displays this name in the tree in the right panel. |

| For This Parameter | Do This                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                          |
|--------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Server name</b> | Use the <b>AuroraJDBCConnectionString</b> value you used to connect to the Aurora MySQL DB instance, but remove the JDBC prefix information and the port suffix. For example, a sample connection string you use with SQL Workbench/J might be "jdbc:mysql://dmsdemo-auroracluster-1u1ogdfg35v.cluster-cqiw4tcs0mg7.us-west-2.rds.amazonaws.com:3306". For AWS SCT <b>Server name</b> , you remove "jdbc:mysql://" and ":3306" to use just the server name: "dmsdemo-auroracluster-1u1ogdfg35v.cluster-cqiw4tcs0mg7.us-west-2.rds.amazonaws.com" |
| <b>Server port</b> | Enter 3306.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                      |
| <b>User name</b>   | Enter auradmin.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| <b>Password</b>    | Enter the password for the admin user that you assigned when creating the Oracle DB instance using the AWS CloudFormation template.                                                                                                                                                                                                                                                                                                                                                                                                              |

- Choose **OK** to close the alert box, then choose **Connect** to connect to the Amazon Aurora MySQL DB instance.
- In the tree in the left panel, select only the **HR** schema. In the tree in the right panel, select your target Aurora MySQL database. Choose **Create mapping**.

▼ Server mappings

| <input type="checkbox"/> | Source path                                | Target path                       |  |
|--------------------------|--------------------------------------------|-----------------------------------|--|
| <input type="checkbox"/> | //Servers/Amazon RDS for Oracle/Schemas/HR | //Servers/Aurora MySQL/Schemas/hr |  |
| <input type="checkbox"/> |                                            |                                   |  |
| <input type="checkbox"/> |                                            |                                   |  |
| <input type="checkbox"/> |                                            |                                   |  |
| <input type="checkbox"/> |                                            |                                   |  |
| <input type="checkbox"/> |                                            |                                   |  |
| <input type="checkbox"/> |                                            |                                   |  |
| <input type="checkbox"/> |                                            |                                   |  |
| <input type="checkbox"/> |                                            |                                   |  |

Create mapping
 Delete selected mappings
 Main view

9. Choose **Main view**. In the tree in the left panel, right-click the **HR** schema and choose **Create report**.

10. Check the report and the action items it suggests. The report discusses the type of objects that can be converted by using AWS SCT, along with potential migration issues and actions to resolve these issues. For this walkthrough, you should see something like the following:


#### Database objects with conversion actions for Amazon Aurora (MySQL compatible)

Of the total 27 database storage object(s) and 51 database code object(s) in the source database, we identified 17 (63%) database storage object(s) and 40 (78%) database code object(s) that can be converted to Amazon Aurora (MySQL compatible) automatically or with minimal changes. We found 1 encrypted object(s).

10 (37%) database storage object(s) require 11 complex user action(s) to complete the conversion.

11 (22%) database code object(s) require 11 complex user action(s) to complete the conversion.

**Figure: Conversion statistics for database storage objects**



You can optionally save the report as .csv or .pdf format for later analysis.

- 11 Choose **Action Items**, and review any recommendations that you see.
- 12 In the tree in the left panel, right-click the **HR** schema and then choose **Convert schema**.
- 13 Choose **Yes** for the confirmation message. AWS SCT then converts your schema to the target database format.
- 14 In the tree in the right panel, choose the converted **hr** schema, and then choose **Apply to database** to apply the schema scripts to the target Aurora MySQL instance.
- 15 Choose the **hr** schema, and then choose **Refresh from Database** to refresh from the target database.

The database schema has now been converted and imported from source to target.

## Step 6: Validate the Schema Conversion

To validate the schema conversion, you compare the objects found in the Oracle and Aurora MySQL databases using SQL Workbench/J.

1. In SQL Workbench/J, choose **File**, then choose **Connect window**. Choose the **RDS Aurora Connection** you created in an earlier step. Click **OK**.
2. Run the following script to verify the number of object types and count in the **HR** schema in the target Aurora MySQL database. These values should match the number of objects in the source Oracle database:

```
SELECT a.OBJECT_TYPE, COUNT(*)
FROM
(
  SELECT OBJECT_TYPE
    ,OBJECT_SCHEMA
    ,OBJECT_NAME
  FROM (
    SELECT 'TABLE' AS OBJECT_TYPE
    ,TABLE_NAME AS OBJECT_NAME
    ,TABLE_SCHEMA AS OBJECT_SCHEMA
  FROM information_schema.TABLES
  where  TABLE_TYPE='BASE TABLE'
  UNION
  SELECT 'VIEW' AS OBJECT_TYPE
    ,TABLE_NAME AS OBJECT_NAME
```

```

,TABLE_SCHEMA AS OBJECT_SCHEMA
FROM information_schema.VIEWS
UNION

SELECT 'INDEX' AS OBJECT_TYPE
,CONCAT (
CONSTRAINT_TYPE
,' : '
,CONSTRAINT_NAME
,' : '
,TABLE_NAME
) AS OBJECT_NAME
,TABLE_SCHEMA AS OBJECT_SCHEMA
FROM information_schema.TABLE_CONSTRAINTS
where constraint_type='PRIMARY KEY'
UNION
SELECT ROUTINE_TYPE AS OBJECT_TYPE
,ROUTINE_NAME AS OBJECT_NAME
,ROUTINE_SCHEMA AS OBJECT_SCHEMA
FROM information_schema.ROUTINES
UNION
SELECT 'TRIGGER' AS OBJECT_TYPE
,CONCAT (
TRIGGER_NAME
,' : '
,EVENT_OBJECT_SCHEMA
,' : '
,EVENT_OBJECT_TABLE
) AS OBJECT_NAME
,TRIGGER_SCHEMA AS OBJECT_SCHEMA
FROM information_schema.triggers
) R
WHERE R.OBJECT_SCHEMA = 'HR'
order by 1) a
GROUP BY a.OBJECT_TYPE;

```

The output from this query should be similar to the following:

| OBJECT_TYPE | COUNT(*) |
|-------------|----------|
| INDEX       | 7        |
| PROCEDURE   | 2        |
| TABLE       | 7        |
| TRIGGER     | 10       |

VIEW

1

Next, run the following query to get table constraints information:

```
SELECT CONSTRAINT_TYPE, COUNT(*)
FROM information_schema.TABLE_CONSTRAINTS where constraint_schema='HR'
GROUP BY CONSTRAINT_TYPE;
```

The output from this query should be similar to the following:

| CONSTRAINT_TYPE | COUNT(*) |
|-----------------|----------|
| FOREIGN KEY     | 10       |
| PRIMARY KEY     | 7        |
| UNIQUE          | 7        |

## Step 7: Create an AWS DMS Replication Instance

After we validate the schema structure between source and target databases, as described preceding, we proceed to the core part of this walkthrough, which is the data migration. The following illustration shows a high-level view of the migration process.



A DMS replication instance performs the actual data migration between source and target. The replication instance also caches the transaction logs during the migration. How much CPU and memory capacity a replication instance has influences the overall time required for the migration.

To create an AWS DMS replication instance, do the following:

1. Sign in to the AWS Management Console, select [AWS Database Migration Service](#) (AWS DMS) and choose **Create replication instance**. If you are signed in as an AWS Identity and Access Management (IAM) user, you must have the appropriate permissions to access AWS DMS. For more information about the permissions required, see [IAM Permissions](#).



2. On the **Create replication instance** page, specify your replication instance information as shown following.

| For This Parameter                            | Do This                                                                                                     |
|-----------------------------------------------|-------------------------------------------------------------------------------------------------------------|
| <b>Name</b>                                   | Enter <code>DMSdemo-repserver</code> .                                                                      |
| <b>Descriptive Amazon Resource Name (ARN)</b> | Skip this optional field.                                                                                   |
| <b>Description</b>                            | Enter a brief description, such as <code>DMS demo replication server</code> .                               |
| <b>Instance class</b>                         | Choose <b>dms.t3.medium</b> . This instance class is large enough to migrate a small set of tables.         |
| <b>Engine version</b>                         | Choose <b>3.4.5</b> . This is the latest AWS DMS version, which includes all new features and enhancements. |
| <b>Allocated storage (GiB)</b>                | Choose <b>50</b> . This storage space is enough for your migration project.                                 |
| <b>VPC</b>                                    | Choose <code>DMSDemoVPC</code> , which is the VPC that was created by the AWS CloudFormation stack.         |
| <b>Multi-AZ</b>                               | Choose <code>Dev or test workload (Single-AZ)</code> .                                                      |
| <b>Publicly accessible</b>                    | Leave this item selected.                                                                                   |

3. For the **Advanced**, **Maintenance**, and **Tags** sections, leave the default settings as they are, and choose **Create**.

## Step 8: Create AWS DMS Source and Target Endpoints

While your replication instance is being created, you can specify the source and target database endpoints using the [AWS Management Console](#). However, you can only test connectivity after the replication instance has been created, because the replication instance is used in the connection.

1. Specify your connection information for the source Oracle database and the target Amazon Aurora MySQL database. The following table describes the source settings.

| For This Parameter  | Do This                                                                                                                                                                     |
|---------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Endpoint Identifier | Enter <code>Orasource</code> (the Amazon RDS for Oracle endpoint).                                                                                                          |
| Source Engine       | Choose <b>oracle</b> .                                                                                                                                                      |
| Server name         | Provide the Oracle DB instance name. This is the <b>Server name</b> you used for AWS SCT, such as <code>"do1xa4grferti8y.cqiw4tcs0mg7.us-west-2.rds.amazonaws.com"</code> . |
| Port                | Enter 1521.                                                                                                                                                                 |
| SSL mode            | Choose <b>None</b> .                                                                                                                                                        |
| Username            | Enter <code>oraadmin</code> .                                                                                                                                               |
| Password            | Provide the password for the Oracle DB instance.                                                                                                                            |
| SID                 | Provide the Oracle database name.                                                                                                                                           |

The following table describes the target settings.

| For This Parameter  | Do This                                                          |
|---------------------|------------------------------------------------------------------|
| Endpoint Identifier | Enter <code>Aurtarget</code> (the Amazon Aurora MySQL endpoint). |

| For This Parameter | Do This                                                                                                                                                                                         |
|--------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Target Engine      | Choose <b>aurora</b> .                                                                                                                                                                          |
| Servename          | Provide the Aurora MySQL DB instance name. This is the <b>Server name</b> you used for AWS SCT, such as "dmsdemo-auroracluster-1u1oyqny35jwv.cluster-cqiw4tcs0mg7.us-west-2.rds.amazonaws.com". |
| Port               | Enter 3306.                                                                                                                                                                                     |
| SSL mode           | Choose <b>None</b> .                                                                                                                                                                            |
| Username           | Enter auradmin.                                                                                                                                                                                 |
| Password           | Provide the password for the Aurora MySQL DB instance.                                                                                                                                          |

The completed page should look like the following:

## Connect source and target database endpoints

✓ Replication instance created successfully.

Your database endpoint can be on-premise, in EC2, RDS or in the cloud. Define the connection details below. It is recommended that you test your endpoint connections here to avoid errors later.

### Source database connection details

Endpoint identifier\*  ⓘ

Source engine\*  ⓘ

Server name\*

Port\*  ⓘ

SSL mode\*  ⓘ

User name\*

Password\*

SID\*

▶ Advanced

### Target database connection details

Endpoint identifier\*  ⓘ

Target engine\*  ⓘ

Server name\*

Port\*  ⓘ

SSL mode\*  ⓘ

User name\*

Password\*

▶ Advanced

Run test

- In order to disable foreign key checks during the initial data load, you must add the following commands to the target Aurora MySQL DB instance. In the **Advanced** section, shown following, type the following commands for **Extra connection attributes**: `initstmt=SET FOREIGN_KEY_CHECKS=0; autocommit=1`

The first command disables foreign key checks during a load, and the second command commits the transactions that DMS executes.

**SID\***

▶ **Advanced**

**Run test**

✓ Connection tested successfully

▼ **Advanced**

**Extra connection attributes**

**KMS master key**  ⓘ

**Description** Default master key that protects my DMS replication instance volumes when no other key is defined

**Account** 100137374868

**Key ARN** arn:aws:kms:us-west-2:100137374868:key/ef0d4b5e-a9ad-74830e851659

**Run test**

✓ Connection tested successfully

3. Choose **Next**.

## Step 9: Create and Run Your AWS DMS Migration Task

Using a AWS DMS task, you can specify what schema to migrate and the type of migration. You can migrate existing data, migrate existing data and replicate ongoing changes, or replicate data changes only. This walkthrough migrates existing data only.

1. On the **Create Task** page, specify the task options. The following table describes the settings.

| For This Parameter      | Do This                              |
|-------------------------|--------------------------------------|
| <b>Task name</b>        | Enter <code>migratehrschema</code> . |
| <b>Task description</b> | Enter a description for the task.    |

| For This Parameter   | Do This                                                                                             |
|----------------------|-----------------------------------------------------------------------------------------------------|
| Source endpoint      | Shows <code>orasource</code> (the Amazon RDS for Oracle endpoint).                                  |
| Target endpoint      | Shows <code>aurtarget</code> (the Amazon Aurora MySQL endpoint).                                    |
| Replication instance | Shows <code>DMSdemo-repserver</code> (the AWS DMS replication instance created in an earlier step). |
| Migration type       | Choose <b>Migrate existing data</b> .                                                               |
| Start task on create | Select this option.                                                                                 |

The page should look like the following:

## Create task

A task can contain one or more table mappings which define what data is moved from the source to the target. If a table does not exist on the target, it can be created automatically.

|                                                                                        |                                                    |  |
|----------------------------------------------------------------------------------------|----------------------------------------------------|--|
| Task name*                                                                             | <input type="text" value="migrateschema"/>         |  |
| Task description*                                                                      | <input type="text" value="migrates HR schema"/>    |  |
| Source endpoint                                                                        | orasource                                          |  |
| Target endpoint                                                                        | aurtarget                                          |  |
| Replication instance                                                                   | dmsdemo-repserver                                  |  |
| Migration type*                                                                        | <input type="text" value="Migrate existing data"/> |  |
| Start task on create                                                                   | <input checked="" type="checkbox"/>                |  |
| ▶ Task Settings                                                                        |                                                    |  |
| ▶ Table mappings                                                                       |                                                    |  |
| <div><a href="#">Cancel</a> <a href="#">Previous</a> <a href="#">Create task</a></div> |                                                    |  |

- Under **Task Settings**, choose **Do nothing** for **Target table preparation mode**, because you have already created the tables through Schema Migration Tool. Because this migration doesn't contain any LOBs, you can leave the LOB settings at their defaults.

Optionally, you can select **Enable logging**. If you enable logging, you will incur additional Amazon CloudWatch charges for the creation of CloudWatch logs. For this walkthrough, logs are not necessary.

▼ Task Settings

Target table preparation mode\* ☒ Do nothing ⓘ  
☐ Drop tables on target  
☐ Truncate

Include LOB columns in replication\* ☐ Don't include LOB columns ⓘ  
☐ Full LOB mode  
☒ Limited LOB mode

Max LOB size (kb)\*  ⓘ

Enable logging ☐

[Advanced Settings](#)

► Table mappings

[Cancel](#) [Previous](#) [Create task](#)

3. Leave the Advanced settings at their default values.

4. Choose **Table mappings**, choose **Default** for **Mapping method**, and then choose HR for **Schema to migrate**.

The completed section should look like the following.

Enable logging ☐

[Advanced Settings](#)

▼ Table mappings

Mapping method\* ☒ Default ☐ Custom ⓘ

Schema to migrate  ⓘ  
DMS will create the schema on the target if it does not already exist


[Show JSON](#)

[Cancel](#) [Previous](#) [Create task](#)

5. Choose **Create task**. The task will begin immediately.



The Tasks section shows you the status of the migration task.



You can monitor your task if you choose **Enable logging** when you set up your task. You can then view the CloudWatch metrics by doing the following:

1. On the navigation pane, choose **Tasks**.
2. Choose your migration task (migratehrschema).
3. Choose the **Task monitoring** tab, and monitor the task in progress on that tab.

## Step 10: Verify That Your Data Migration Completed Successfully

When the migration task completes, you can compare your task results with the expected results.

1. On the navigation pane, choose **Tasks**.
2. Choose your migration task (migratehrschema).
3. Choose the **Table statistics** tab, shown following.


The screenshot displays the AWS DMS console interface for a task named 'migrateschema'. At the top, there are buttons for 'Create task', 'Start/Resume', 'Stop', and 'Delete', along with refresh, settings, and help icons. A filter bar is present above a table with columns: ID, Status, Source, Target, and Type. The 'migrateschema' task is selected, showing a status of 'Load complete', source 'orasource', target 'aurtarget', and type 'Full Load'.

Below the task summary, the 'Table statistics' tab is active. It shows a refresh button and another filter bar. A table lists statistics for various tables in the 'migrateschema' schema. All tables are in a 'Table completed' state with 0 inserts, deletes, or updates. The 'Full Load Rows' column shows the total rows migrated for each table.

| Table       | State           | Inserts | Deletes | Updates | DDLs | Full Load Rows |
|-------------|-----------------|---------|---------|---------|------|----------------|
| COUNTRIES   | Table completed | 0       | 0       | 0       | 0    | 25             |
| DEPARTMENTS | Table completed | 0       | 0       | 0       | 0    | 27             |
| EMPLOYEES   | Table completed | 0       | 0       | 0       | 0    | 107            |
| JOBS        | Table completed | 0       | 0       | 0       | 0    | 19             |
| JOB_HISTORY | Table completed | 0       | 0       | 0       | 0    | 10             |
| LOCATIONS   | Table completed | 0       | 0       | 0       | 0    | 23             |
| REGIONS     | Table completed | 0       | 0       | 0       | 0    | 4              |

4. Connect to the Amazon Aurora MySQL instance by using SQL Workbench/J, and then check if the database tables were successfully migrated from Oracle to Aurora MySQL by running the SQL script shown following.

```
SELECT TABLE_NAME, TABLE_ROWS
FROM INFORMATION_SCHEMA.TABLES
WHERE TABLE_SCHEMA = 'HR' and TABLE_TYPE='BASE TABLE' order by 1;
```



- Run the following query to check the relationship in tables; this query checks the departments with employees greater than 10.

```
SELECT B.DEPARTMENT_NAME, COUNT(*)
FROM HR.EMPLOYEES A, HR.DEPARTMENTS B
WHERE A.DEPARTMENT_ID=B.DEPARTMENT_ID
GROUP BY B.DEPARTMENT_NAME HAVING COUNT(*) > 10
ORDER BY 1;
```

The output from this query should be similar to the following.

```
department_name count(*)
Sales            34
Shipping         45
```

Now you have successfully completed a database migration from an Amazon RDS for Oracle database instance to Amazon Aurora MySQL.

## Step 11: Delete Walkthrough Resources

After you have completed this walkthrough, perform the following steps to avoid being charged further for AWS resources used in the walkthrough. It's necessary that you do the steps in order, because some resources cannot be deleted if they have a dependency upon another resource.

1. On the navigation pane, choose **Tasks**, choose your migration task (migratehrschema), and then choose **Delete**.
2. On the navigation pane, choose **Endpoints**, choose the Oracle source endpoint (orasource), and then choose **Delete**.
3. Choose the Amazon Aurora MySQL target endpoint (auratarget), and then choose **Delete**.
4. On the navigation pane, choose **Replication instances**, choose the replication instance (DMSdemo-repserver), and then choose **Delete**.

Next, you must delete your AWS CloudFormation stack, DMSdemo.

1. Sign in to the AWS Management Console and open the [AWS CloudFormation console](#).

Note that if you are signed in as an AWS Identity and Access Management (IAM) user, you must have the appropriate permissions to access AWS CloudFormation.

2. Choose your AWS CloudFormation stack, DMSdemo.
3. For **Actions**, choose **Delete stack**.

The status of the stack changes to DELETE\_IN\_PROGRESS while AWS CloudFormation AWS CloudFormation cleans up the resources associated with the DMSdemo stack. When AWS CloudFormation is finished cleaning up resources, it removes the stack from the list.

## AWS DMS migration from Amazon RDS for Oracle next steps

You can explore several other features of AWS DMS that were not included in this walkthrough, including the following:

- The AWS DMS change data capture (CDC) feature, for ongoing replication of data.
- Transformation actions that let you specify and apply transformations to the selected schema or table as part of the migration process.

For more information, see [Getting started with Database Migration Service](#).

## Migrating a SQL Server Database to Amazon Aurora MySQL

Using this walkthrough, you can learn how to migrate a Microsoft SQL Server database to an Amazon Aurora MySQL-Compatible Edition database using the AWS Schema Conversion Tool (AWS

SCT) and AWS Database Migration Service (AWS DMS). AWS DMS migrates your data from your SQL Server source into your Aurora MySQL target.

AWS DMS doesn't migrate your secondary indexes, sequences, default values, stored procedures, triggers, synonyms, views, and other schema objects that aren't specifically related to data migration. To migrate these objects to your Aurora MySQL target, use AWS SCT.

To estimate what it will cost to run this walkthrough on AWS, you can use the AWS Pricing Calculator. For more information, see <https://calculator.aws/>.

To avoid additional charges, delete all resources after you complete the walkthrough.

## Topics

- [Prerequisites for Migrating from a SQL Server database to Amazon Aurora MySQL](#)
- [Step-by-step SQL Server database to Amazon Aurora MySQL migration walkthrough](#)
- [SQL Server database migration to Amazon Aurora MySQL troubleshooting](#)

## Prerequisites for Migrating from a SQL Server database to Amazon Aurora MySQL

The following prerequisites are required to complete this walkthrough:

- An understanding of Amazon Relational Database Service (Amazon RDS), the applicable database technologies, and SQL.
- Create a user with AWS Identity and Access Management (IAM) credentials that allows you to launch Amazon RDS and AWS Database Migration Service (AWS DMS) instances in your AWS Region. For information about IAM credentials, see [Setting up for Amazon RDS](#).
- An understanding of the Amazon Virtual Private Cloud (Amazon VPC) service and security groups. For information about using Amazon VPC with Amazon RDS, see [Amazon Virtual Private Cloud \(VPCs\) and Amazon RDS](#). For information about Amazon RDS security groups, see [Amazon RDS Security Groups](#).
- An understanding of the supported features and limitations of AWS DMS. For information about AWS DMS, see <https://docs.aws.amazon.com/dms/latest/userguide/Welcome.html>.
- An understanding of how to work with Microsoft SQL Server as a source and Amazon Aurora MySQL as a target. For information about working with SQL Server as a source, see [Using a SQL Server Database as a Source](#). Aurora MySQL is a MySQL-compatible database. For information

about working with Aurora MySQL as a target, see [Using a MySQL-Compatible database as a target](#).

- An understanding of the supported data type conversion options for SQL Server and Aurora MySQL. For information about data types for SQL Server as a source, see [Source Data Types for Microsoft SQL Server](#). For information about data types for Aurora MySQL; as a target, see [Target Data Types for MySQL](#).
- Size your target Aurora MySQL database host. DBAs should be aware of the load profile of the current source SQL Server database host. Consider CPU, memory, and IOPS. With Amazon RDS, you can size up the target database host, or reduce it, after the migration. If this is the first time that you're migrating to Aurora MySQL, we recommended that you have extra capacity to account for performance issues and tuning opportunities.
- Audit your source SQL Server database. For each schema and all the objects under each schema, determine whether any of the objects are no longer being used. Deprecate these objects on the source SQL Server database, because there's no need to migrate them if they aren't being used.
- Decide between these migration options: migrate existing data only or migrate existing data and replicate ongoing changes.
  - If you migrate existing data only, the migration is a one-time data transfer from a SQL Server source database to the Aurora MySQL target database. If the source database remains open to changes during the migration, these changes must be applied to the target database after the migration is complete.

 **Note**

If the SQL Server database is an Amazon RDS database, replication is not supported, and you must use the option to migrate existing data only.

- If you migrate existing data and replicate ongoing changes, one option is to replicate the source database changes. Replication keeps the source and target databases in sync with each other during the migration process and can reduce database downtime. With this option, you complete an initial sync operation and then configure MS-REPLICATION. This option requires the Standard, Enterprise, or Developer SQL Server edition. You enable MS-REPLICATION for each SQL Server instance that you want to use as a database source.
- If you want to migrate existing data and replicate ongoing changes, another option is change data capture (CDC) instead of replication. This option allows AWS DMS to perform ongoing migration of data. In the case of CDC, AWS DMS uses the CDC tables to enable ongoing

database migration. This option requires the Standard, Enterprise or Developer edition of SQL Server.

For more information about AWS DMS, see [Getting started with Database Migration Service](#).

## Step-by-step SQL Server database to Amazon Aurora MySQL migration walkthrough

The following steps provide instructions for migrating a Microsoft SQL Server database to an Amazon Aurora MySQL database. These steps assume that you have already prepared your source database as described in [Prerequisites](#).

### Topics


- [Step 1: Install the SQL Drivers and AWS Schema Conversion Tool on Your Local Computer](#)
- [Step 2: Configure Your Microsoft SQL Server Source Database](#)
- [Step 3: Configure Your Aurora MySQL Target Database](#)
- [Step 4: Use AWS SCT to Convert the SQL Server Schema to Aurora MySQL](#)
- [Step 5: Create an AWS DMS Replication Instance](#)
- [Step 6: Create AWS DMS Source and Target Endpoints](#)
- [Step 7: Create and Run Your AWS DMS Migration Task](#)
- [Step 8: Cut Over to Aurora MySQL](#)

### Step 1: Install the SQL Drivers and AWS Schema Conversion Tool on Your Local Computer

First, install the SQL drivers and the AWS Schema Conversion Tool (AWS SCT) on your local computer. Do the following:

1. Download the JDBC driver for Microsoft SQL Server [mssql-jdbc-7.2.2.jre11.jar](#).
2. Download the [JDBC driver for Aurora MySQL](#). Amazon Aurora MySQL uses the MySQL driver.
3. Install AWS SCT and the required JDBC drivers.
  - a. See [Installing, verifying, and updating the Schema Conversion Tool](#), and choose the appropriate link to download AWS SCT.
  - b. Start AWS SCT, and choose **Settings, Global settings**.

- c. In **Global settings**, choose **Drivers**, and then choose **Browse** for **Microsoft SQL Server driver path**. Locate the JDBC driver for SQL Server, and choose **OK**.
- d. Choose **Browse** for **MySQL driver path**. Locate the JDBC driver you downloaded for Aurora MySQL, and choose **OK**.



- e. Choose **OK** to close the **Global settings** dialog box.

## Step 2: Configure Your Microsoft SQL Server Source Database

After installing the SQL drivers and AWS Schema Conversion Tool, you can configure your Microsoft SQL Server source database using one of several options, depending on how you plan to migrate your data.

When configuring your source database, you can choose to migrate existing data only, migrate existing data and replicate ongoing changes, or migrate existing data and use change data capture (CDC) to replicate ongoing changes. For more information about these options, see [Prerequisites](#).

### Migrating existing data only

No configuration steps are necessary for the SQL Server database. You can move on to [Step 3: Configure Your Aurora MySQL Target Database](#).



**Note**

If the SQL Server database is an Amazon RDS database, replication is not supported, and you must use the option for migrating existing data only.

## Migrating existing data and replicating ongoing changes

To configure MS-REPLICATION, complete the following steps:

1. In Microsoft SQL Server Management Studio, open the context (right-click) menu for the **Replication** folder, and then choose **Configure Distribution**.
2. In the **Distributor** step, choose **db\_name will act as its own distributor**. SQL Server creates a distribution database and log.

For more information, see [Microsoft documentation](#).

When the configuration is complete, your server is enabled for replication. Either a distribution database is in place, or you have configured your server to use a remote distribution database.

**Note**

Replication requires a primary key for all tables that are being replicated. If your tables don't have primary keys defined, consider using CDC instead.

## Migrating existing data and using change data capture (CDC) to replicate ongoing changes

To configure MS-CDC, complete the following steps:

1. Connect to SQL Server with a login that has SYSADMIN role membership.
2. For each database containing data that is being migrated, run the following command within the database context:

```
use [DBname]
EXEC sys.sp_cdc_enable_db
```

3. For each table that you want to configure for ongoing migration, run the following command:

```
EXEC sys.sp_cdc_enable_table @source_schema = N'schema_name', @source_name =  
N'table_name', @role_name = NULL;
```

For more information, see [Microsoft documentation](#).

#### Note

- If you are migrating databases that participate in an Always On Availability Group, it is best practice to use replication for migration. To use this option, publishing must be enabled, and a distribution database must be configured for each node of the Always On Availability Group. Additionally, ensure you are using the name of the availability group listener for the database rather than the name of the server currently hosting the availability group database for the target server name. These requirements apply to each instance of SQL Server in the cluster and must not be configured using the availability group listener.
- If your database isn't supported for MS-REPLICATION or MS-CDC (for example, if you are running the Workgroup Edition of SQL Server), some changes can still be captured, such as INSERT and DELETE statements, but other DML statements such as UPDATE and TRUNCATE TABLE will not be captured. Therefore, a migration with continuing data replication is not recommended in this configuration, and a static one time migration (or repeated one time full migrations) should be considered instead.

For more information about using MS-REPLICATION and MS-CDC, see [Configuring a Microsoft SQL Server Database as a Replication Source](#).

## Step 3: Configure Your Aurora MySQL Target Database

AWS DMS migrates the data from the SQL Server source into an Amazon Aurora MySQL target. In this step, you configure the Aurora MySQL target database.

1. Create the AWS DMS user to connect to your target database, and grant Superuser or the necessary individual privileges (or for Amazon RDS, use the master username).

Alternatively, you can grant the privileges to an existing user.

```
CREATE USER 'aurora_dms_user' IDENTIFIED BY 'password';

GRANT ALTER, CREATE, CREATE TEMPORARY TABLES, DROP, INDEX, INSERT, UPDATE, DELETE,
SELECT ON target_database.* TO 'aurora_dms_user';
```

2. AWS DMS uses control tables on the target in the database `awsdms_control`. Use the following command to ensure that the user has the necessary access to the `awsdms_control` database:

```
GRANT ALL PRIVILEGES ON awsdms_control.* TO 'aurora_dms_user';
FLUSH PRIVILEGES;
```

## Step 4: Use AWS SCT to Convert the SQL Server Schema to Aurora MySQL

Before you migrate data to Amazon Aurora MySQL, convert the Microsoft SQL Server schema to an Aurora MySQL schema using the AWS Schema Conversion Tool (AWS SCT). [This video covers all the steps of this process.](#)

To convert a SQL Server schema to an Aurora MySQL schema, do the following:

1. Launch AWS SCT. In AWS SCT, choose **File**, then choose **New Project**. Create a new project named AWS Schema Conversion Tool SQL Server to Aurora MySQL, specify the **Location** of the project folder, and then choose **OK**.
2. Choose **Add source** to add a source Microsoft SQL Server database to your project, then choose **Microsoft SQL Server**, and choose **Next**.
3. Enter the following information, and then choose **Test connection**.

| Parameter       | Description                                                                           |
|-----------------|---------------------------------------------------------------------------------------|
| Connection name | Enter Microsoft SQL Server. AWS SCT displays this name in the tree in the left panel. |
| Server name     | Enter the server name.                                                                |
| Server port     | Enter the SQL Server port number. The default is 1433.                                |

| Parameter     | Description                                  |
|---------------|----------------------------------------------|
| Instance name | Enter the SQL Server database instance name. |
| User name     | Enter the SQL Server admin user name.        |
| Password      | Enter the password for the admin user.       |

**Add source**

**CONNECTION** SSL

Specify parameters for new connections to Microsoft SQL Server

Connection name: Microsoft SQL Server

Server name: ec2-52-17-21-76.eu-west-1.compute.amazonaws.com

Server port: 1433

Instance name: GOLD\_TEST\_SS\_PG

Authentication: SQL Server Authentication

User name: min\_privs

Password: [masked]

☐ Use SSL

☐ Store password


Test connection Previous Connect

- Choose **OK** to close the alert box. Then choose **Connect** to close the dialog box and connect to the Microsoft SQL Server database instance. AWS SCT displays the structure of the Microsoft SQL Server database instance in the left panel.
- Choose **Add target** to add a target Amazon Aurora MySQL database to your project, then choose **Amazon Aurora (MySQL compatible)**, and choose **Next**.

6. Enter the following information and then choose **Test Connection**.

| Parameter              | Description                                                                    |
|------------------------|--------------------------------------------------------------------------------|
| <b>Connection name</b> | Enter Aurora MySQL. AWS SCT displays this name in the tree in the right panel. |
| <b>Server name</b>     | Enter the server name.                                                         |
| <b>Server port</b>     | Enter the SQL Server port number. The default is 3306.                         |
| <b>User name</b>       | Enter the Aurora MySQL admin user name.                                        |
| <b>Password</b>        | Enter the password for the admin user.                                         |

- Choose **OK** to close the alert box. Then choose **Connect** to close the dialog box and connect to the Aurora MySQL database instance.
- In the tree in the left panel, select the schema to migrate. In the tree in the right panel, select your target Aurora MySQL database. Choose **Create mapping**.
- Choose **Main view**. In the tree in the left panel, right-click the HR schema and choose **Create report**.



10 Open the context (right-click) menu for the schema to migrate, and then choose **Convert schema**.

11. Choose **Yes** for the confirmation message. AWS SCT analyzes the schema, creates a database migration assessment report, and converts your schema to the target database format.
12. Choose **Assessment Report View** from the menu to check the database migration assessment report. The report breaks down by each object type and by how much manual change is needed to convert it successfully.

Generally, packages, procedures, and functions are more likely to have some issues to resolve because they contain the most custom Transact-SQL code. AWS SCT also provides hints about how to fix these objects.

13. Choose the **Action Items** tab.

The **Action Items** tab shows each issue for each object that requires attention.

For each conversion issue, you can complete one of the following actions:

- Modify the objects on the source SQL Server database so that AWS SCT can convert the objects to the target Aurora MySQL database.
  - i. Modify the objects on the source SQL Server database.
  - ii. Repeat the previous steps to convert the schema and check the assessment report.
  - iii. If necessary, repeat this process until there are no conversion issues.
  - iv. Choose **Main View** from the menu. Open the context (right-click) menu for the target Aurora MySQL schema, and choose **Apply to database** to apply the schema changes to the Aurora MySQL database, and confirm that you want to apply the schema changes.
- Instead of modifying the source schema, modify scripts that AWS SCT generates before applying the scripts on the target Aurora MySQL database.
  - i. Choose **Main View** from the menu. Open the context (right-click) menu for the target Aurora MySQL schema name, and choose **Save as SQL**. Next, choose a name and destination for the script.
  - ii. In the script, modify the objects to correct conversion issues.

You can also exclude foreign key constraints, triggers, and secondary indexes from the script because they can cause problems during the migration. After the migration is complete, you can create these objects on the Aurora MySQL database.

- iii. Run the script on the target Aurora MySQL database.

For more information, see [Converting Database Schema to Amazon RDS](#).

14. (Optional) Use AWS SCT to create migration rules.

- a. Choose **Mapping view** and then choose **New migration rule**.
- b. Create additional migration transformation rules that are required based on the action items.
- c. Save the migration rules.
- d. Choose **Export script for DMS** to export a JSON format of all the transformations that the AWS DMS task will use. Choose **Save**.

## Step 5: Create an AWS DMS Replication Instance

After validating the schema structure between source and target databases, continue with the core part of this walkthrough, which is the data migration. The following illustration shows a high-level view of the migration process.



An AWS DMS replication instance performs the actual data migration between source and target. The replication instance also caches the transaction logs during the migration. The amount of CPU and memory capacity a replication instance has influences the overall time that is required for the migration.

For information about best practices for using AWS DMS, see [AWS Database Migration Service Best Practices](#).

To create an AWS DMS replication instance, do the following:

1. Sign in to the AWS Management Console, and open the [AWS DMS console](#).
2. In the console, choose **Create migration**. If you are signed in as an AWS Identity and Access Management (IAM) user, you must have the appropriate permissions to access AWS DMS. For more information about the permissions required, see [IAM Permissions](#).
3. On the Welcome page, choose **Next** to start a database migration.
4. On the **Create replication instance** page, specify your replication instance information.

| Parameter                  | Description                                                                                                                                                                                                                                    |
|----------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Name</b>                | Select a name for your replication instance. If you are using multiple replication servers or sharing a user, choose a name that helps you quickly differentiate between the different servers.                                                |
| <b>Description</b>         | Enter a brief description.                                                                                                                                                                                                                     |
| <b>Instance class</b>      | Select the type of replication server to create. Each size and type of instance class has increasing CPU, memory, and I/O capacity. Generally, t2 instances are for lower load tasks, and the c4 instances are for higher load and more tasks. |
| <b>VPC</b>                 | Choose the virtual private cloud (VPC) in which your replication instance will launch. If possible, select the same VPC in which either your source or target database resides (or both).                                                      |
| <b>Multi-AZ</b>            | If you choose <b>Yes</b> , AWS DMS creates a second replication server in a different Availability Zone for failover if there is a problem with the primary replication server.                                                                |
| <b>Publicly accessible</b> | If either your source or target database resides outside of the VPC in which your replication server resides, you must make your replication server policy publicly accessible.                                                                |

5. For the **Advanced** section, specify the following information.



| Parameter                       | Description                                                                                                                                                                                                                                                                                         |
|---------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Allocated storage (GB)</b>   | Amount of storage on the replication server for the AWS DMS task logs, including historical tasks logs. AWS DMS also uses disk storage to cache certain data while it replicates it from the source database to the target. Additionally, more storage generally enables better IOPS on the server. |
| <b>Replication Subnet Group</b> | If you are running in a Multi-AZ configuration, you need at least two subnet groups.                                                                                                                                                                                                                |
| <b>Availability zone</b>        | Generally, performance is better if you locate your primary replication server in the same Availability Zone as your target database.                                                                                                                                                               |
| <b>VPC Security Group(s)</b>    | Security groups enable you to control ingress and egress to your VPC. AWS DMS lets you associate one or more security groups with the VPC in which your replication server is launched.                                                                                                             |
| <b>KMS key</b>                  | With AWS DMS, all data is encrypted at rest using a KMS encryption key. By default, AWS DMS creates a new encryption key for your replication server. However, you might choose to use an existing key.                                                                                             |

For information about the KMS key, see [Setting an Encryption Key and Specifying KMS Permissions](#).

6. Click **Next**.

## Step 6: Create AWS DMS Source and Target Endpoints

While your replication instance is being created, you can specify the source and target database endpoints using the [AWS Management Console](#). However, you can test connectivity only after the replication instance has been created, because the replication instance is used in the connection.

1. In the AWS DMS console, specify your connection information for the source SQL Server database and the target Aurora MySQL database. The following table describes the source settings.

| Parameter           | Description                                                                        |
|---------------------|------------------------------------------------------------------------------------|
| Endpoint Identifier | Enter a name, such as <code>SQLServerSource</code> .                               |
| Source Engine       | Choose <b>sqlserver</b> .                                                          |
| Server name         | Provide the SQL Server DB instance server name.                                    |
| Port                | Enter the port number of the database. The default for SQL Server is 1433.         |
| SSL mode            | Choose an SSL mode if you want to enable encryption for your connection's traffic. |
| User name           | Enter the name of the user you want to use to connect to the source database.      |
| Password            | Provide the password for the user.                                                 |
| Database name       | Provide the SQL Server database name.                                              |

The following table describes the advanced source settings.

| Parameter                   | Description                                                                         |
|-----------------------------|-------------------------------------------------------------------------------------|
| Extra connection attributes | Extra parameters that you can set in an endpoint to add functionality or change the |

| Parameter      | Description                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
|----------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                | <p>behavior of AWS DMS. A few of the most relevant attributes are listed here. Use a semicolon (;) to separate multiple entries.</p> <ul style="list-style-type: none"> <li>• <code>safeguardpolicy</code> - Changes the behavior of SQL Server by opening transactions to prevent the transaction log from being truncated while AWS DMS is reading the log. Valid values are <code>EXCLUSIVE_AUTOMATIC_TRUNCATION</code> or <code>RELY_ON_SQL_SERVER_REPLICATION_AGENT</code> (default).</li> <li>• <code>useBCPFullLoad</code> - Directs AWS DMS to use BCP (bulk copy) for data loading. Valid values are Y or N. When the target table contains an identity column that does not exist in the source table, you must disable the use of BCP for loading the table by setting the parameter to N.</li> <li>• <code>BCPPacketSize</code> - If BCP is enabled for data loads, then enter the maximum packet size used by BCP. Valid values are 1 – 100000 (default 16384).</li> <li>• <code>controlTablesFileGroup</code> - Specifies the file group to use for the control tables that the AWS DMS process creates in the database.</li> </ul> |
| <b>KMS key</b> | Enter the KMS key if you choose to encrypt your replication instance's storage.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   |

The following table describes the target settings.

| Parameter           | Description                                                                        |
|---------------------|------------------------------------------------------------------------------------|
| Endpoint Identifier | Enter a name, such as <code>AuroraTarget</code> .                                  |
| Target Engine       | Choose <b>aurora</b> .                                                             |
| Server name         | Provide the Aurora MySQL DB server name for the primary instance.                  |
| Port                | Enter the port number of the database. The default for Aurora MySQL is 3306.       |
| SSL mode            | Choose <b>None</b> .                                                               |
| User name           | Enter the name of the user that you want to use to connect to the target database. |
| Password            | Provide the password for the user.                                                 |

The following table describes the advanced target settings.

| Parameter                          | Description                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
|------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Extra connection attributes</b> | <p>Extra parameters that you can set in an endpoint to add functionality or change the behavior of AWS DMS. A few of the most relevant attributes are listed here. Use a semicolon to separate multiple entries.</p> <ul style="list-style-type: none"><li>• <code>targetDbType</code> - By default, AWS DMS creates a different database for each schema that is being migrated. If you want to combine several schemas into a single database, set this option to <code>targetDbType=SPECIFIC_DATABASE</code> .</li><li>• <code>initstmt</code> - Use this option to invoke the MySQL <code>initstmt</code> connection parameter and accept anything MySQL <code>initstmt</code> accepts. For an Aurora MySQL target, it's often useful to disable foreign key checks by setting this option to <code>initstmt=SET FOREIGN_KEY_CHECKS=0</code> .</li></ul> |
| <b>KMS key</b>                     | Enter the KMS key if you choose to encrypt your replication instance's storage.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |

The following is an example of the completed page.

### Connect source and target database endpoints

 Replication instance created successfully.

Your database endpoint can be on-premise, in EC2, RDS or in the cloud. Define the connection details below. It is recommended that you test your endpoint connections here to avoid errors later.

| Source database connection details                                | Target database connection details                               |
|-------------------------------------------------------------------|------------------------------------------------------------------|
| Endpoint identifier* <input type="text" value="SQLServSource"/> ⓘ | Endpoint identifier* <input type="text" value="AuroraTarget"/> ⓘ |
| Source engine* <input type="text" value="sqlserver"/> ⓘ           | Target engine* <input type="text" value="aurora"/> ⓘ             |
| Server name* <input type="text"/>                                 | Server name* <input type="text"/>                                |
| Port* <input type="text" value="1433"/> ⓘ                         | Port* <input type="text" value="3306"/> ⓘ                        |
| SSL mode* <input type="text" value="none"/> ⓘ                     | SSL mode* <input type="text" value="none"/> ⓘ                    |
| User name* <input type="text" value="sqlservadmin"/> ⓘ            | User name* <input type="text" value="auroraadmin"/> ⓘ            |
| Password* <input type="password" value="*****"/> ⓘ                | Password* <input type="password" value="*****"/> ⓘ               |
| Database name* <input type="text" value="sqlserv"/>               | <input type="button" value="Advanced"/>                          |
| <input type="button" value="Advanced"/>                           | <input type="button" value="Run test"/>                          |
| <input type="button" value="Run test"/>                           |                                                                  |

For information about extra connection attributes, see [Using Extra Connection Attributes](#).

2. After the endpoints and replication instance are created, test the endpoint connections by choosing **Run test** for the source and target endpoints.
3. Drop foreign key constraints and triggers on the target database.

During the full load process, AWS DMS does not load tables in any particular order, so it might load the child table data before parent table data. As a result, foreign key constraints might be violated if they are enabled. Also, if triggers are present on the target database, they might change data loaded by AWS DMS in unexpected ways.

```
ALTER TABLE 'table_name' DROP FOREIGN KEY 'fk_name';

DROP TRIGGER 'trigger_name';
```

4. If you dropped foreign key constraints and triggers on the target database, generate a script that enables the foreign key constraints and triggers.

Later, when you want to add them to your migrated database, you can just run this script.

5. (Optional) Drop secondary indexes on the target database.

Secondary indexes (as with all indexes) can slow down the full load of data into tables because they must be maintained and updated during the loading process. Dropping them can improve the performance of your full load process. If you drop the indexes, you must to add them back later, after the full load is complete.

```
ALTER TABLE 'table_name' DROP INDEX 'index_name';
```

6. Choose **Next**.

## Step 7: Create and Run Your AWS DMS Migration Task

Using an AWS DMS task, you can specify what schema to migrate and the type of migration. You can migrate existing data, migrate existing data and replicate ongoing changes, or replicate data changes only.

1. In the AWS DMS console, on the **Create task** page, specify the task options. The following table describes the settings.

| Parameter                   | Description                                                                                                                              |
|-----------------------------|------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Task name</b>            | Enter a name for the migration task.                                                                                                     |
| <b>Task description</b>     | Enter a description for the task.                                                                                                        |
| <b>Source endpoint</b>      | Shows the SQL Server source endpoint.<br><br>If you have more than one endpoint for the user, choose the correct endpoint from the list. |
| <b>Target endpoint</b>      | Shows the Aurora MySQL target endpoint.                                                                                                  |
| <b>Replication instance</b> | Shows the AWS DMS replication instance.                                                                                                  |
| <b>Migration type</b>       | Choose an option.                                                                                                                        |

| Parameter                   | Description                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |
|-----------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                             | <ul style="list-style-type: none"><li>• <b>Migrate existing data</b> - AWS DMS migrates only your existing data. Changes to your source data aren't captured and applied to your target. If you can afford to take an outage for the duration of the full load, then this is the simplest option. You can also use this option to create test copies of your database. If the source SQL Server database is an Amazon RDS database, you must choose this option.</li><li>• <b>Migrate existing data and replicate ongoing changes</b> - AWS DMS captures changes while migrating your existing data. AWS DMS continues to capture and apply changes even after the bulk data has been loaded. Eventually the source and target databases are in sync, allowing for a minimal downtime.</li><li>• <b>Replicate data changes only</b> - Bulk load data using a different method. This approach generally applies only to homogeneous migrations.</li></ul> |
| <b>Start task on create</b> | In most situations, you should choose this option. Sometimes, you might want to delay the start of a task, for example, if you want to change logging levels.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            |

The page should look similar to the following:



### Create task

A task can contain one or more table mappings which define what data is moved from the source to the target. If a table does not exist on the target, it can be created automatically.

Task name\*

SQLServer2Aurora

Task description\*

Migrate schema from SQL Server to Aurora

Source endpoint

sqlservsource

Target endpoint

auroratarget

Replication instance

sqlserver2aurora

Migration type\*

Migrate existing data

Start task on create

☒

2. Under **Task settings**, specify the settings. The following table describes the settings.

| Parameter                     | Description                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
|-------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Target table preparation mode | <div>Choose an option.</div> <ul style="list-style-type: none"><li>• <b>Do nothing</b> - AWS DMS does nothing to prepare your tables. Your table structure remains the same, and any existing data remains in the table. You can use this method to consolidate data from multiple systems.</li><li>• <b>Drop tables on target</b> - AWS DMS creates your target tables for you. AWS DMS drops and re-creates the tables to migrate before migration. AWS DMS creates the table and a primary key only for heterogeneous migrations.</li><li>• <b>Truncate</b> - AWS DMS truncates a target table before loading it. If the target table doesn't exist, then AWS DMS creates it.</li></ul> |

| Parameter                                 | Description                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    |
|-------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                                           | <div> <b>Important</b><br/>If the AWS Schema Conversion Tool already created the tables on the target, choose <b>Do nothing</b> or <b>Truncate</b>.</div>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |
| <b>Include LOB columns in replication</b> | <p>Choose an option.</p> <ul style="list-style-type: none"><li>• <b>Don't include LOB columns</b> - Do not migrate LOB data.</li><li>• <b>Full LOB mode</b> - AWS DMS migrates all LOBs (large objects) from the source to the target regardless of size. In this configuration, AWS DMS has no information about the maximum size of LOBs to expect. Thus, LOBs are migrated one at a time, piece by piece. Full LOB mode can be relatively slow.</li><li>• <b>Limited LOB mode</b> - You set a maximum size LOB that AWS DMS accepts. This option enables AWS DMS to pre-allocate memory and load the LOB data in bulk. LOBs that exceed the maximum LOB size are truncated, and a warning is issued to the log file. In limited LOB mode, you get significant performance gains over full LOB mode. We recommend that you use limited LOB mode whenever possible.</li></ul> |

| Parameter                | Description                                                                                                                                                                       |
|--------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Max LOB size (kb)</b> | When <b>Limited LOB mode</b> is selected, this option determines the maximum LOB size that AWS DMS accepts. Any LOBs that are larger than this value are truncated to this value. |
| <b>Enable logging</b>    | It's best to select <b>Enable logging</b> . If you enable logging, you can see any errors or warnings that the task encounters, and you can troubleshoot those issues.            |

3. Leave the Advanced settings at their default values.
4. If you created and exported mapping rules with AWS SCT in the last step in [Step 4: Convert the SQL Server Schema to Aurora MySQL](#), choose **Table mappings**, and select the **JSON** tab. Then select **Enable JSON editing**, and enter the table mappings you saved.

If you did not create mapping rules, then proceed to the next step.

5. Choose **Create task**. The task starts immediately.

The **Tasks** section shows you the status of the migration task.

| <div> Create task Modify Start/Resume Stop Delete </div> <div> Filter: <input type="text"/> </div> |          |               |              |           |            |              |               |
|----------------------------------------------------------------------------------------------------|----------|---------------|--------------|-----------|------------|--------------|---------------|
| ID                                                                                                 | Status   | Source        | Target       | Type      | Complete % | Elapsed time | Tables loaded |
| sqlserver2aurora                                                                                   | Creating | sqlservsource | auroratarget | Full Load | 0          |              | 0             |

If you chose **Enable logging** during setup, you can monitor your task. You can then view the Amazon CloudWatch metrics.

1. On the navigation pane, choose **Tasks**.
2. Choose your migration task.
3. Choose the **Task monitoring** tab, and monitor the task in progress on that tab.

When the full load is complete and cached changes are applied, the task stops on its own.

4. On the target Aurora MySQL database, if you disabled foreign key constraints and triggers, enable them using the script that you saved previously.

5. On the target Aurora MySQL database, re-create the secondary indexes if you removed them previously.
6. If you chose to use AWS DMS to replicate changes, in the AWS DMS console, start the AWS DMS task by choosing **Start/Resume** for the task.

Important replication instance metrics to monitor include the following:

- CPU
- FreeableMemory
- DiskQueueDepth
- CDCLatencySource
- CDCLatencyTarget

The AWS DMS task keeps the target Aurora MySQL database up to date with source database changes. AWS DMS keeps all the tables in the task up to date until it's time to implement the application migration. The latency is zero, or close to zero, when the target has caught up to the source.

For more information, see [Monitoring DMS tasks](#).

## Step 8: Cut Over to Aurora MySQL

To move connections from your Microsoft SQL Server database to your Amazon Aurora MySQL database, do the following:

1. End all SQL Server database dependencies and activities, such as running scripts and client connections. Ensure that the SQL Server Agent service is stopped.

The following query should return no results other than your connection:

```
SELECT session_id, login_name from sys.dm_exec_sessions where session_id > 50;
```

2. Kill any remaining sessions (other than your own).

```
KILL session_id;
```

3. Shut down the SQL Server service.
4. Let the AWS DMS task apply the final changes from the SQL Server database on the Amazon Aurora MySQL database.

5. In the AWS DMS console, stop the AWS DMS task by choosing **Stop** for the task, and then confirming that you want to stop the task.

## SQL Server database migration to Amazon Aurora MySQL troubleshooting

When you work with Microsoft SQL Server as a source database and Amazon Aurora MySQL as a target database, the two most common problem areas are SQL Server change data capture (CDC) and foreign keys.

- **MS-CDC:** If you are using MS-CDC with SQL Server for the migration, errors that are related to permissions or errors during change data capture are common. These types of errors usually result when one of the prerequisites was not met. For example, the most common overlooked prerequisite is a full database backup.
- **Foreign keys:** During the full load process, AWS DMS does not load tables in any particular order, so it might load the child table data before parent table data. As a result, foreign key constraints might be violated if they are enabled. You should disable foreign keys on the Aurora MySQL target database. You can enable the foreign keys on the target after the migration is complete.

For more tips, see the AWS DMS troubleshooting section in the [Troubleshooting migration tasks](#).

To troubleshoot issues specific to SQL Server, see the SQL Server troubleshooting section:

- [Troubleshooting Microsoft SQL Server Specific Issues](#)

To troubleshoot Aurora MySQL issues, see the Aurora MySQL troubleshooting section and the MySQL troubleshooting section:

- [Troubleshooting Amazon Aurora MySQL Specific Issues](#)
- [Troubleshooting MySQL Specific Issues](#)

# Migrating a SQL Server AlwaysOn Database on Primary Replica to Amazon Aurora PostgreSQL

In this walkthrough we will cover the process of migrating a database from SQL Server AlwaysOn Primary Replica to Amazon Aurora PostgreSQL using AWS Database Migration Service (AWS DMS). We will highlight common migration issues, and methods to overcome them. We will also guide you through the process of our automatic SQL scripts to arrange the tables, prepare the JSON table mappings, and explore methods of distributing tables across multiple DMS tasks for optimal efficiency.

## Why Amazon Aurora PostgreSQL?

Most organizations use online transaction process (OLTP) database with mixed workloads running on SQL Server AlwaysOn platform. Because of the advanced capabilities and cost-effectiveness of open-source databases, many corporations are moving away from legacy, on-premise SQL Server AlwaysOn environments running high-profile workloads to robust, cloud-based, highly scalable, and resilient solutions.

Organizations prefer to migrate their data to PostgreSQL because it's an open-source database solution which offers advanced RDBMS capabilities without commercial licensing costs. PostgreSQL is also backed by community base support which isn't dependent on any specific vendor. Running critical workloads within a robust, secure, and redundant cloud base infrastructure also brings resiliency benefits without high cap-ex costs of maintaining multiple data centers. For more information, see [Working with Amazon Aurora PostgreSQL](#). For the latest features and key benefits, see [Amazon Aurora PostgreSQL](#).

## Common database migration challenges

Following are some common migration problems that could potentially drain project resources and derail data migration project timelines.

- Underestimating the complexity of the table structure - A typical application user may not be aware of the specific table fields that hold all the data elements. A manual data migration process often results in incomplete, inaccurate, and outdated information being transferred to the target endpoint.
- Lack of integrated workflow processes - A database migration typically involves disparate teams using various tools to interact with the data. When you use spreadsheets or other manual

methods to document data specifications, human errors can easily occur, resulting in wasted time, and resources or incomplete migration.

- Inability to validate data transformation specifications - With ongoing data changes on the source endpoint, it's quite difficult to manually validate all migrated data. Sample-based data validation often results in missed discrepancies which may have negative repercussions post-migration.


## Why AWS DMS?

AWS DMS is a managed-service which provides an out-of-box migration solution that helps you mitigate the aforementioned migration challenges. AWS DMS offers the following key benefits. For complete list of feature benefits, see [AWS Database Migration Service Features](#).

- Cost-effectiveness – you pay only for the compute and log storage resources used during your migration.
- Ongoing data replication capability that allows you to complete your application cutover without disrupting your DevOps and business processes. You can also enable data validation on a task to ensure that no outdated information is being transferred to the target endpoint.
- Automatically analyze the source table schema and structure, and retrieve the specific table fields that hold all the data elements, and reduce the need for you to manually gather the data specifications.
- Ability to integrate with other AWS services such as CloudWatch. With CloudWatch, you can create custom alarms that watch DMS metrics and send notifications when a threshold limit is reached. For more information, see [Monitoring AWS DMS tasks](#).


## Migration overview

The following image shows a high-level architecture of the AWS DMS replication workflow. AWS DMS migration consists of an EC2 replication instance which hosts the DMS software. The replication instance handles the execution of one or more DMS tasks. Each task replicates a specific set of table data from the SQL Server AlwaysOn primary replica source to the Amazon Aurora PostgreSQL target endpoint. For more information, see [Working with an AWS DMS replication instance](#). For a complete migration playbook, see [Microsoft SQL Server to Amazon Aurora PostgreSQL Migration Playbook](#).



In the rest of this document, we'll migrate a sample financial institution database from SQL Server AlwaysOn to Aurora PostgreSQL. The database includes tables containing large object (LOB) data types with either a primary key or unique key. Tables with these characteristics pose different migration challenges which will also be discussed. The entity relationship diagram of the sample database is shown below.





## Prerequisites for migrating SQL Server AlwaysOn databases on primary replica to Amazon Aurora PostgreSQL

The following prerequisites are required to complete this migration:

- An AWS account with AWS Identity and Access Management (IAM) credentials that allow you to launch Amazon RDS and AWS Database Migration Service (AWS DMS) instances in your AWS Region. For information about IAM credentials, see [Create an IAM User](#).
- A general understanding of Amazon Virtual Private Cloud (Amazon VPC), DNS, and security groups concepts. For information about using Amazon VPC with Amazon RDS, see [Amazon Virtual Private Cloud \(VPCs\) and Amazon RDS](#). For information about Amazon RDS security groups, see [Amazon RDS Security Groups](#). For information about network setup to support AWS DMS replication instances, see [Setting up network for replication instance](#). For information about AWS DMS using Route53 for endpoint name resolution, see [Using Amazon Route 53 Resolver with AWS DMS](#).
- A general understanding of using Microsoft SQL Server as a source and Amazon Aurora PostgreSQL as a target endpoint in an AWS DMS] based migration. For information about working with SQL Server as a source, see [Using a SQL Server Database as a Source](#). Aurora PostgreSQL is a PostgreSQL compatible database. For information about working with Aurora PostgreSQL as a target, see [Using a PostgreSQL database as a Target](#).
- An understanding of your source table structure, backup policy, and resource constraints.
- Convert your SQL Server database schema to PostgreSQL using a optional tool such as [Schema Conversion Tool \(AWS SCT\)](#)

## Step-by-step SQL Server AlwaysOn databases on primary replica to Amazon Aurora PostgreSQL migration walkthrough

The following standard steps assume you have prepared your source and target endpoint as described in the above prerequisites. We also assume that you have converted your SQL Server database schema to PostgreSQL using the [Schema Conversion Tool \(AWS SCT\)](#) first, and that you've created all database objects on the target database.

### Step 1: Configure SQL Server database for Replication or Change Data Capture

In this walkthrough we create a “migrate existing data and replicate ongoing changes” DMS migration task. This type of DMS task will perform an initial copy of all existing data from source to the target, and then it will transition into replicating ongoing as changes occurring on the source endpoint. The migration mode also provides flexibility of reloading the target tables when needed. For more information, see [Prerequisites for using ongoing replication \(CDC\) from a SQL Server source](#).

Following is the list of tables in our sample database. We use the AWS DMS Best Practice Support Scripts for SQL Server to gather details about the tables. Table level info will be useful while designing the migration approach and selecting DMS task settings in later steps.

|    | TABLERNAME                           | HasLOB | HasUniqueKey | HasPrimaryKey | HasUniqueConstraint | ROWS      | TotalSpaceMB |
|----|--------------------------------------|--------|--------------|---------------|---------------------|-----------|--------------|
| 1  | tbl_Client                           | No     | Yes          | Yes           | No                  | 32767     | 6.63         |
| 2  | tbl_ClientAccount                    | No     | Yes          | Yes           | No                  | 997210    | 19.38        |
| 3  | tbl_ClientAccountStatusType          | No     | Yes          | Yes           | No                  | 6         | 0.07         |
| 4  | tbl_ClientAccountType                | No     | Yes          | Yes           | No                  | 10        | 0.07         |
| 5  | tbl_ClientClientAccount              | No     | Yes          | Yes           | No                  | 62383434  | 1271.29      |
| 6  | tbl_ClientOverDraftLog               | Yes    | Yes          | Yes           | No                  | 1012604   | 130.07       |
| 7  | tbl_ClientTransactionLog             | No     | Yes          | Yes           | No                  | 311918075 | 16094.73     |
| 8  | tbl_ClientTransactionType            | No     | Yes          | Yes           | No                  | 5         | 0.07         |
| 9  | tbl_Employee                         | No     | Yes          | Yes           | No                  | 997210    | 43.57        |
| 10 | tbl_FailedClientTransactionErrorType | No     | Yes          | Yes           | No                  | 5         | 0.07         |
| 11 | tbl_FailedClientTransactionLog       | Yes    | Yes          | Yes           | No                  | 312607694 | 13324.98     |
| 12 | tbl_LoginClientAccount               | No     | No           | No            | No                  | 65483     | 4.07         |
| 13 | tbl_LoginErrorLog                    | Yes    | Yes          | Yes           | No                  | 312607695 | 17895.48     |
| 14 | tbl_SavingsClientInterestRates       | No     | Yes          | Yes           | No                  | 8         | 0.07         |
| 15 | tbl_UserLogins                       | No     | Yes          | Yes           | No                  | 32767     | 1.63         |
| 16 | tbl_UserSecurityAnswers              | No     | Yes          | Yes           | No                  | 32767     | 0.95         |
| 17 | tbl_UserSecurityQuestions            | No     | Yes          | Yes           | No                  | 10        | 0.07         |

LOB in SQL Server is a data type designed to store large amounts of data. Data replication performance could be impacted when LOB data types are replicated and DMS task settings may need to be adjusted accordingly. Those task settings are discussed later in this document. For more information, see [LOB support for source database in an AWS DMS task](#).

Next, we follow the AWS DMS SQL Server source endpoint public documentation to configure the distribution database on each SQL Server AlwaysOn replica. AWS DMS ongoing change replication supports either the Microsoft SQL Server Replication (MS-Replication) or Microsoft Change Data Capture (MS-CDC) feature. For more information concerning setup of distribution database and enabling MS-CDC to support AWS DMS replication task, see [Setting up ongoing replication using the sysadmin role with self-managed SQL Server](#). For more information about using MS-REPLICATION and MS-CDC, see [Configuring a Microsoft SQL Server Database as a Replication Source](#).

## Step 2: Create an AWS DMS replication instance

To create an AWS DMS replication instance, follow the steps below:

1. Sign in to the AWS Management Console, and open the [AWS DMS console](#).

2. In the console, choose **Create replication instance**. If you are signed in as an AWS Identity and Access Management (IAM) user, you must have the appropriate permissions to access AWS DMS. For more information about the permissions required, see [IAM Permissions](#).
3. On the **Create replication instance** page, specify your replication instance information. For this walkthrough, both endpoints reside in the same AWS region. We configure our AWS replication instance using the same endpoint VPC. For more information concerning how to select the best instance class to support your data migration, see [Choosing the right AWS DMS replication instance class for your migration](#).

| Parameter             | Value                            | Explanation                                                                                                                                                                                                  |
|-----------------------|----------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Name</b>           | <b>replication-test</b>          | Helps quickly differentiate between the different servers.                                                                                                                                                   |
| <b>Description</b>    | DMS replication instance         | Helps identify the purpose of the server.                                                                                                                                                                    |
| <b>Instance class</b> | dms.c5.xlarge                    | c5.xlarge EC2 class provides 2 vCPU, 4 GB RAM, and up to 10 Gbps base network bandwidth. It also provides better performance over a general purpose t3.medium EC2 class with up to 5 Gbps network bandwidth. |
| <b>VPC</b>            | vpc-08xxxxxxxxxxxxxe             | Using the same VPC as the SQL server endpoint since the instance is hosted in the AWS network.                                                                                                               |
| <b>Multi-AZ</b>       | Dev or test workload (Single-AZ) | Testing DMS replication workload. If it's production data, you will choose <b>Yes</b> to create a standby replication server to support Multi-AZ,                                                            |

|                            |    |                                                                                        |
|----------------------------|----|----------------------------------------------------------------------------------------|
|                            |    | and to support high availability.                                                      |
| <b>Publicly accessible</b> | No | Publicly accessible is not needed since source and endpoint reside in the AWS network. |

4. For the **Advanced** section, specify the following information. For more information, see [Working with an AWS DMS replication instance](#). For information about the KMS key, see [Setting an Encryption Key and Specifying KMS Permissions](#).

| Parameter                       | Value                       | Explanation                                                                                                         |
|---------------------------------|-----------------------------|---------------------------------------------------------------------------------------------------------------------|
| <b>Allocated storage</b>        | 80 GB                       | Allocated storage size based on 1.5x of the migrating database size                                                 |
| <b>Replication Subnet Group</b> | default-vpc-08xxxxxxxxxxxxe | Using single replication subnet group with single AZ                                                                |
| <b>Availability zone (AZ)</b>   | us-east-2c                  | Using same availability zone as SQL server source endpoint which reside in AWS network                              |
| <b>VPC Security Group(s)</b>    | xxx-sec-group               | Using the same VPC security group as source and target database since both endpoints reside in the same AWS region. |

5. Click **Next**.

## Step 3: Create an AWS DMS source endpoint for SQL server

You can specify the source or target database endpoints using the <https://console.aws.amazon.com/AWS Management Console>].

1. In the AWS DMS console, specify your connection information for the source SQL Server database. The following table describes the source settings used in this walkthrough.

| Parameter                  | Value                                           | Explanation                                                                                                                            |
|----------------------------|-------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------|
| <b>Endpoint Identifier</b> | <b>sqlserver-source</b>                         | Helps quickly differentiate between the different endpoints.                                                                           |
| <b>Source Engine</b>       | <b>sqlserver</b>                                | Define the endpoint engine                                                                                                             |
| <b>Server name</b>         | listener-xxxxxx.us-east-2.compute.amazonaws.com | Setting the endpoint to use SQL Server AlwaysOn Listener's fully qualify domain name                                                   |
| <b>Port</b>                | 1433                                            | Setting the endpoint to use default port 1433 on the SQL Server AlwaysOn Listener's fully qualify domain                               |
| <b>SSL mode</b>            | none                                            | Setting SSL mode to none since the replication instance will reside in the same AWS VPC and region as the source and target endpoints. |
| <b>User name</b>           | dmsuser                                         | Setting the endpoint to use the dmsuser to connect to the SQL Server endpoint.                                                         |

|                      |                 |                                                                                        |
|----------------------|-----------------|----------------------------------------------------------------------------------------|
| <b>Password</b>      | strong password | Setting the endpoint to use the password to connect to the SQL Server endpoint.        |
| <b>Database name</b> | BankDatabase    | Setting the endpoint to use the BankDatabase after successful login to the SQL Server. |

Note that using SQL Server with dynamic ports may result in frequent DMS task failures as every time the SQL Server service is restarted, several settings will need to be adjusted to work with the new port number. For more information about AWS DMS source endpoint settings such as providing support for SQL Server AlwaysOn read-only replica, see [Using extra connection attributes for SQL Server as source endpoint when working with a secondary availability group replica](#). For source endpoint limitation, see [Limitations on using SQL Server as a source for AWS DMS](#).

2. Test the endpoint connection by choosing **Run test** for the source endpoints.

## Step 4: Configure and verify Aurora PostgreSQL database DMS user account

In this step, you need to configure and verify that the DMS user account has required permissions on the Aurora PostgreSQL target database. Our walkthrough assumes that your target database objects were also precreated using the AWS SCT tool.

1. Create the AWS DMS user with one of the following permissions on your Aurora PostgreSQL database if it does not exist. Your PostgreSQL target endpoint requires minimum user permissions to run an AWS DMS migration. For more information, see [Security requirements when using PostgreSQL database as target for AWS Database Migration Service](#).
2. Verify whether DMS user account has privileges by executing the following tests:
  - a. Log into the Aurora PostgreSQL database.
  - b. Create a new sample test table. Example: `create table sample_test (num int, description varchar(100));`
  - c. Alter the new sample\_test table. Example: `alter table sample_test add primary key (num);`

- d. Insert a new record into the sample\_test table. Example: `insert into sample_test (num, description) values (1, 'test');`
  - e. Delete the new record from the sample\_test table. Example: `delete sample_test where num = 1;`
  - f. Truncate all records on the sample\_test table. Example: `truncate sample_test;`
  - g. Drop the sample\_test table. Example: `drop table sample_test;`
3. Script out table constraints and triggers to a post-deployment script file which will be used to recreate those objects after the migration is complete.
  4. Drop table constraints (FKs, Check constraints, defaults) and triggers.

## Step 5: Configure an AWS DMS target endpoint for Aurora PostgreSQL

You can create the endpoint using the <https://console.aws.amazon.com/AWS Management Console>].

1. In the AWS DMS console, specify your connection information for the target Aurora PostgreSQL database.
2. Check "Select RDS DB instance" and then choose Aurora PostgreSQL instance if it's created under the current user and region. The following table describes the target settings.

| Parameter                  | Value                                                              | Explanation                                                        |
|----------------------------|--------------------------------------------------------------------|--------------------------------------------------------------------|
| <b>Endpoint Identifier</b> | <b>postgres-target</b>                                             | Helps quickly differentiate between the different endpoints.       |
| <b>Target Engine</b>       | Amazon Aurora PostgreSQL                                           | Define the endpoint engine                                         |
| <b>Server name</b>         | postgresql-source-instance.<br>xxxxxxx.us-east-2.rds.amazonaws.com | Setting the endpoint to use PostgreSQL fully qualified domain name |
| <b>Port</b>                | 5432                                                               | Setting the endpoint to use default port 5432 on the               |



|                      |                 |                                                                                                                                    |
|----------------------|-----------------|------------------------------------------------------------------------------------------------------------------------------------|
|                      |                 | PostgreSQL fully qualify domain                                                                                                    |
| <b>SSL mode</b>      | none            | Setting SSL mode to none since replication instance will be resided in the same AWS VPC and region as source and target endpoints. |
| <b>User name</b>     | dmsuser         | Set the endpoint to use the dmsuser to connect to the PostgreSQL endpoint.                                                         |
| <b>Password</b>      | strong password | Set the endpoint to use the password to connect to the PostgreSQL endpoint.                                                        |
| <b>Database name</b> | BankDatabase    | Set the endpoint to use the BankDatabase after successfully logging in to the Aurora PostgreSQL database.                          |

For the purpose of this walkthrough, we'll use the following advanced target endpoint settings. For information about other AWS DMS endpoint settings for a PostgreSQL target, see [Endpoint Settings for PostgreSQL as target endpoint](#).

| Parameter         | Value              | Explanation                                                                                                                              |
|-------------------|--------------------|------------------------------------------------------------------------------------------------------------------------------------------|
| Endpoint settings | executeTimeout=300 | Setting executionTimeout to 5 minutes to support a longer execution duration of DMS' replication query due to a slow SQL server endpoint |

3. Test the endpoint connections by choosing **Run test** for the target endpoints.

## Step 6: Create an AWS DMS migration task(s)

Depending on your workload pattern as well as your business requirements, distributing tables across multiple DMS tasks can help improve migration performance, and provide ease of troubleshooting. Isolating tables with certain characteristics in a separate DMS task may be beneficial to the overall migration. For example, tables with LOBs or without primary or unique key may present unique challenges, and may require different DMS task settings. Following, we describe a method of grouping the tables into 3 categories depending on their characteristics. Organizing tables according to the following criteria will help achieve optimal performance during the migration. For our walkthrough, we prepare an automatic SQL query script to help you determine the table arrangement. The script also dynamically generates the table mappings in JSON format so you can simply paste the output into the DMS task setting. For more information, see [Working with diagnostic support scripts in AWS DMS](#).

Following, are the three main categories:

1. Special tables – which require special handling due to their characteristics
2. Large tables – which require special handling due to their size
3. General tables – all other tables which do not meet the previous criteria

### Special tables

Special tables are tables that do not follow OLTP workload practices. Data loading typically involves truncating the table first, and then reloading all the data from scratch. Due to the nature of how data are loaded to the special table, it is not a good candidate for the “migrate existing data and replicate ongoing changes” DMS task. DMS task will replicate the large bulk workloads in this case using a row-by-row replication mode. Instead, we recommend you to change the replication task type to “migrate existing data”, incorporate the tables into the same DMS task that perform the truncate, and reload data operation. For example, the tbl\_LoginErrorLog table is a standalone table. The ETL process can truncate the data in tbl\_LoginErrorLog table first, and then reload the entire table.

### Large tables

Large tables are tables that contain partitioned, wide table columns, or a large amount of data that can require a specialized management process to maintain. DMS tasks containing large tables generally do not perform well when using the default, single threaded full load per table. For partitioned tables, you can leverage parallel load by configuring the auto-partition option in

the table mapping. For non-partitioned tables, you can accomplish parallel load by utilizing the ranges-partitioning option which allows range boundaries to be specified manually.

For our walkthrough, we use the automatic SQL query script to dynamically generate the task mappings output. By examining the table overview report generated earlier, we notice that the `tbl_ClientOverDraftLog` table meets our large table criteria. Next, we examine the `tbl_ClientOverDraftLog` table structure, and notice the table has a sequential integer data type on the `ClientTransactionID` column. We input the schema and table name in the automatic SQL query script to generate a JSON format output that will include the range boundaries for a given large table.

The following image shows the output generated by the automatic SQL query script for the `tbl_ClientOverDraftLog` table. You may copy the JSON section of the output to your DMS task mappings. You can also modify the task mapping JSON output to include transformation rules such as renaming the target table name. You can then repeat this process for the other tables that you consider to be large tables.

 Messages

```
--- DMS Task Mapping for dbo.tbl_ClientOverDraftLog with Boundaries ---
{"rules":[{"rule-type":"selection",

"rule-id":"1",

"rule-name":"1",

"object-locator":{"schema-name":"dbo",

"table-name":"tbl_ClientOverDraftLog"},

"rule-action":"include"},

{"rule-type":"table-settings",

"rule-id":"2",

"rule-name":"2",

"object-locator":{"schema-name":"dbo",

"table-name":"tbl_ClientOverDraftLog"},

"parallel-load":{"Types":"ranges",

"columns":["ClientAccountID"],

"boundaries":[["16972"],

["35106"],

["68999"],

["102888"],

["136794"],

["170727"]]}]}

Completion time: 2022-12-22T20:19:15.9322939+00:00
```

## General tables

General tables are table that do not meet either the special or large table categories. The automatic SQL query script will automatically arrange the general tables in different tasks. It also avoids accidentally misplacing the same table in multiple tasks.

### Determine table placement

The following shows a sample assignment of tables to different groups, generated by the query mentioned above. The sample query groups the tables based on row count.

### Note

In this walkthrough we assume that there are table dependencies on the PostgreSQL target, such as foreign keys, constraints, and triggers. Because the data will be migrated using multiple DMS tasks, there is a possibility of temporary lapses in referential integrity during CDC migration. Because of that we will remove those types of constraints from the target database for the duration of the migration. We will recreate them after the migration is complete, as at that point, the data will be consistent. For our sample scenario, tables with over 1 million records are considered to be large tables. The row size helps reduce the amount of task creation generation by the script. It also avoids over-utilizing the SQL Server source endpoint when the instance does not have adequate resources to support the concurrent retrieval.

For example, the `dbo.tbl_ClientOverDraftLog` table has over 1 million records, and belongs to its own DMS Full Load + CDC task called `dbo_tbl_ClientOverDraftLog table only`. Tables in GroupNum 8 contain primary keys, and are not particularly large which means that they do not pose a challenge from a migration perspective, so they are all placed in the same group. Isolating certain types of tables into their own tasks will allow you to enable additional tasks or table settings to help improve replication performance.

A AWS DMS task loads tables in alphabetical order based on the table name. If specific table loading order is required, you should include the load-order setting, and place the dependent tables together in the same task. If your table contains dependencies to tables in another database, you can't place both tables in the same DMS task. For tables with cross database dependencies, you must load these tables using their own tasks, and then remove referential integrity constraints on the target database for the duration of the migration. For more information see [Table and collection settings rules and operations](#).

|    | tableName                                | totalRows | TotalSpaceMB | DMS_Task_Groups                                                                   | GroupNum |
|----|------------------------------------------|-----------|--------------|-----------------------------------------------------------------------------------|----------|
| 1  | dbo.tbl_ClientClientAccount              | 62383434  | 3813.87      | Dedicate a Full Load + CDC task for dbo_tbl_ClientClientAccount table only        | 1        |
| 2  | dbo.tbl_ClientOverDraftLog               | 1012604   | 520.28       | Dedicate a Full Load + CDC task for dbo_tbl_ClientOverDraftLog table only         | 2        |
| 3  | dbo.tbl_ClientTransactionLog             | 311918075 | 144852.54    | Dedicate a Full Load + CDC task for dbo_tbl_ClientTransactionLog table only       | 3        |
| 4  | dbo.tbl_Employee                         | 997210    | 217.85       | Dedicate a Full Load + CDC task for dbo_tbl_Employee table only                   | 4        |
| 5  | dbo.tbl_FailedClientTransactionLog       | 312607694 | 53299.91     | Dedicate a Full Load + CDC task for dbo_tbl_FailedClientTransactionLog table only | 5        |
| 6  | dbo.tbl_LoginErrorLog                    | 312607695 | 53686.43     | Dedicate a Full Load + CDC task for dbo_tbl_LoginErrorLog table only              | 6        |
| 7  | dbo.tbl_LoginClientAccount               | 65483     | 8.14         | Include this table in nonUnique dbo Full Load + CDC task                          | 7        |
| 8  | dbo.tbl_FailedClientTransactionErrorType | 5         | 0.14         | Include this table in primaryKey dbo Full Load + CDC task                         | 8        |
| 9  | dbo.tbl_SavingsClientInterestRates       | 8         | 0.21         | Include this table in primaryKey dbo Full Load + CDC task                         | 8        |
| 10 | dbo.tbl_UserLogins                       | 32767     | 4.90         | Include this table in primaryKey dbo Full Load + CDC task                         | 8        |
| 11 | dbo.tbl_UserSecurityAnswers              | 32767     | 2.84         | Include this table in primaryKey dbo Full Load + CDC task                         | 8        |
| 12 | dbo.tbl_UserSecurityQuestions            | 10        | 0.14         | Include this table in primaryKey dbo Full Load + CDC task                         | 8        |
| 13 | dbo.tbl_ClientTransactionType            | 5         | 0.28         | Include this table in primaryKey dbo Full Load + CDC task                         | 8        |
| 14 | dbo.tbl_Client                           | 32767     | 106.13       | Include this table in primaryKey dbo Full Load + CDC task                         | 8        |
| 15 | dbo.tbl_ClientAccount                    | 997210    | 96.91        | Include this table in primaryKey dbo Full Load + CDC task                         | 8        |
| 16 | dbo.tbl_ClientAccountStatusType          | 6         | 0.14         | Include this table in primaryKey dbo Full Load + CDC task                         | 8        |
| 17 | dbo.tbl_ClientAccountType                | 10        | 0.14         | Include this table in primaryKey dbo Full Load + CDC task                         | 8        |

To simplify the process of creating the table mappings JSON, you can use an automatic SQL query script to generate the JSON format output for the specific group of tables. The following shows a sample script output for groupNum 8. You can simply copy the entire JSON output into the DMS task mapping rules.

```

Results
DMS Task Mapping
--- GroupNum: 8 ---

{
  "rules": [
    {
      "rule-type": "selection", "rule-id": "1", "rule-name": "1", "object-locator": { "schema-name": "dbo", "table-name": "tbl_Client" }, "rule-action": "include" },
    {
      "rule-type": "selection", "rule-id": "2", "rule-name": "2", "object-locator": { "schema-name": "dbo", "table-name": "tbl_ClientAccount" }, "rule-action": "include" },
    {
      "rule-type": "selection", "rule-id": "3", "rule-name": "3", "object-locator": { "schema-name": "dbo", "table-name": "tbl_ClientAccountStatusType" }, "rule-action": "include" },
    {
      "rule-type": "selection", "rule-id": "4", "rule-name": "4", "object-locator": { "schema-name": "dbo", "table-name": "tbl_ClientAccountType" }, "rule-action": "include" },
    {
      "rule-type": "selection", "rule-id": "5", "rule-name": "5", "object-locator": { "schema-name": "dbo", "table-name": "tbl_ClientTransactionType" }, "rule-action": "include" },
    {
      "rule-type": "selection", "rule-id": "6", "rule-name": "6", "object-locator": { "schema-name": "dbo", "table-name": "tbl_FailedClientTransactionErrorType" }, "rule-action": "include" },
    {
      "rule-type": "selection", "rule-id": "7", "rule-name": "7", "object-locator": { "schema-name": "dbo", "table-name": "tbl_SavingsClientInterestRates" }, "rule-action": "include" },
    {
      "rule-type": "selection", "rule-id": "8", "rule-name": "8", "object-locator": { "schema-name": "dbo", "table-name": "tbl_UserLogins" }, "rule-action": "include" },
    {
      "rule-type": "selection", "rule-id": "9", "rule-name": "9", "object-locator": { "schema-name": "dbo", "table-name": "tbl_UserSecurityAnswers" }, "rule-action": "include" },
    {
      "rule-type": "selection", "rule-id": "10", "rule-name": "10", "object-locator": { "schema-name": "dbo", "table-name": "tbl_UserSecurityQuestions" }, "rule-action": "include" }
  ]
}

```



Create Replication Tasks

Follow the steps below for each DMS task you need to create. For our walkthrough, we'll be creating total of 8 Full Load + CDC tasks based on script output from the previous steps.

1. In the AWS DMS console, on the **Create task** page, specify the task options as show below.

| Parameter        | Value                                               | Explanation                                                        |
|------------------|-----------------------------------------------------|--------------------------------------------------------------------|
| Task name        | sql-2-postgres-dynamic-task                         | Helps quickly differentiate between the different tasks.           |
| Task description | migrating data from SQL server to Aurora PostgreSQL | Helps identify the purpose of the task.                            |
| Source endpoint  | sqlserver-source                                    | Task will be using this SQL server source endpoint                 |
| Target endpoint  | postgres-target                                     | Task will be replicating to this Aurora PostgreSQL target endpoint |

|                             |                                                     |                                                                                                           |
|-----------------------------|-----------------------------------------------------|-----------------------------------------------------------------------------------------------------------|
| <b>Replication instance</b> | replication-instance                                | Task will be handled by this DMS replication instance                                                     |
| <b>Migration type</b>       | Migrate existing data and replicate ongoing changes | Task will be performing full load first and then replicate ongoing data changes when full load completes. |
| <b>Start task on create</b> | Enable                                              | Start task after task creation                                                                            |

2. Under Task settings, specify the settings as shown below. For the purpose of this walkthrough, we set `TargetTablePrepMode` to `TRUNCATE_BEFORE_LOAD`, that way the target schema created using SCT will not be dropped and recreated by DMS. Limited LOB mode is enabled to properly handle the LOB data type contained in the tables. For more information, see [Specifying task settings for AWS Database Migration Service Tasks](#).

| Parameter                                 | Value       | Explanation                                                                                                                                                                                                                                                                                                                                 |
|-------------------------------------------|-------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Target table preparation mode</b>      | Truncate    | For our sample walk through, the data on the target Aurora PostgreSQL database can be dropped therefore we set the target preparation mode to truncate. For production replication, you might want to set this setting to "Do Nothing" and manually execute either the drop or the truncate before running the DMS task from the beginning. |
| <b>Include LOB columns in replication</b> | Limited LOB | Following best practices, set the task to use Limited LOB mode to provide better replication performance.                                                                                                                                                                                                                                   |



|                          |        |                                                                                                                                                                                           |
|--------------------------|--------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Max LOB size (kb)</b> | 32     | For our walkthrough, we are setting limited LOB mode to 32 kilobytes to reduce the LOB size, and provide better replication performance.                                                  |
| <b>Enable logging</b>    | Enable | Enable logging to provide insight on the task errors or warnings.                                                                                                                         |
| <b>Batch Apply</b>       | TRUE   | For our walkthrough, foreign keys constraints and triggers are dropped on the target endpoint. We enable Batch Apply to quickly apply the transactions to the Aurora PostgreSQL database. |

3. Leave the Advanced settings at their default values.
4. If you created your table mappings' JSON using one of the queries mentioned above, choose **Table mappings**, and select the **JSON** tab. Then select **Enable JSON editing**, and paste the table mappings you generated using the scripts. If you did not create mapping rules using the scripts, in the **Selection rules** section, specify the settings as shown in the table below. For information about AWS SCT table mapping rules, see [AWS SCT Mapping](#).

| For this parameter         | Do this                | Explanation                                                                                   |
|----------------------------|------------------------|-----------------------------------------------------------------------------------------------|
| <b>Schema name is</b>      | Choose Enter a schema. | Schema related object                                                                         |
| <b>Schema name is like</b> | Type %.                | % is a wildcard character in this section, and means all schema names on the source endpoint. |
| <b>Table name is like</b>  | Type %.                | % is a wildcard character in this, section and means all                                      |

|        |                 |                                                                            |
|--------|-----------------|----------------------------------------------------------------------------|
|        |                 | tables that are owned by the same schema name on the source endpoint.      |
| Action | Choose Include. | Include all tables, because we specified % for schema name and table name. |

5. Choose **Create task**. The task starts immediately.

If you enabled the **Start Task On Create** option, the task will start automatically after its creation.

Step 7: Verify AWS DMS replication task status

1. In the AWS DMS console, choose **Database Migration Tasks** page.
2. Choose the newly created Database Migration Task from step 6.
3. Click Actions and then choose Restart the full load task from the beginning if task did not start immediately after task creation in step 6.

The **Tasks** section shows you the status of the migration task.

The following shows the table statistic output when starting the DMS Full Load + CDC task from the beginning. Notice that the DMS table statistics show replicating tables that were included in the table mappings JSON.

Summary

Status

Load complete, replication ongoing

Type

Full load, ongoing replication

Source

sqlserver-source

Target

postgres-target

Overview details

Table statistics

CloudWatch metrics

Mapping rules

Premigration assessments

Tags

Table statistics (11)

Total rows include loaded source table rows from Inserts, Deletes, Updates, DDLs, and Full load rows.

Find schema

Schema name

Table

Load state

Elapsed load time

Inserts

Deletes

Updates

DDLs

Applied inserts

Applied deletes

Applied updates

Applied DDLs

Full load rows

Total rows

Validation state

Validation pending

Validation failed

Validation suspended

Validation started

dbo

tbl\_UserLogins

Table completed

2 s

0

0

0

0

0

0

0

0

32,767

32,767

Not enabled

0

0

0

dbo

tbl\_UserSecurityAnswers

Table completed

2 s

0

0

0

0

0

0

0

0

32,767

32,767

Not enabled

0

0

0

dbo

tbl\_UserSecurityQuestions

Table completed

< 1 s

0

0

0

0

0

0

0

0

10

10

Not enabled

0

0

0

dbo

tbl\_FailedClientTransactionErrorType

Table completed

< 1 s

0

0

0

0

0

0

0

0

5

5

Not enabled

0

0

0

dbo

tbl\_ClientAccountStatusType

Table completed

< 1 s

0

0

0

0

0

0

0

0

6

6

Not enabled

0

0

0

dbo

tbl\_Client

Table completed

9 s

0

0

0

0

0

0

0

0

32,767

32,767

Not enabled

0

0

0

dbo

tbl\_ClientTransactionLog

Table completed

6 s

0

0

0

0

0

0

0

0

328,320

328,320

Not enabled

0

0

0

dbo

tbl\_ClientAccountType

Table completed

< 1 s

0

0

0

0

0

0

0

0

10

10

Not enabled

0

0

0

dbo

tbl\_ClientClientAccount

Table completed

4 s

0

0

0

0

0

0

0

0

65,483

65,483

Not enabled

0

0

0

dbo

tbl\_ClientTransactionType

Table completed

< 1 s

0

0

0

0

0

0

0

0

5

5

Not enabled

0

0

0

dbo

tbl\_SavingsClientInterestRates

Table completed

< 1 s

0

0

0

0

0

0

0

0

8

8

Not enabled

0

0

0


Export to CSV

If you chose **Enable logging** during setup, you can monitor your task in Amazon CloudWatch.

1. On the navigation pane, choose **Tasks**.
2. Choose your migration task.
3. Choose the **Task monitoring** tab, and monitor the task in progress on that tab.
4. When the full load is complete and cached changes are applied, the task stops on its own.
5. If you chose to use AWS DMS to replicate changes, in the AWS DMS console, start the AWS DMS task by choosing **Start/Resume** for the task.
6. Important replication instance metrics to monitor include the following: ..CPU ..FreeableMemory ..DiskQueueDepth ..CDCLatencySource ..CDCLatencyTarget

For more information about monitoring AWS DMS task status and metrics, see [Monitoring AWS DMS tasks](#). For more information about viewing DMS error log events, see [Task Logs](#).

You can use the MS SQL Server activity monitor included with SQL Server Management Studio to examine the resource utilization on the source database instance as shown below. You will notice that DMS creates multiple sessions on the source instance. When DMS is configured for parallel load, you will see a separate session for each partition. Configuring additional ranges for parallel load may increase performance of the full load phase of the migration as long as none of the resources are constrained on the endpoint. The number of sessions utilized by DMS will be reduced once the full load phase is complete. You may need to reduce the number of concurrently loaded tables if you observe high resource utilization on the source or target.



## Cleanup

The purpose of these steps is to make sure you have redirected your database application traffic to your new PostgreSQL database after successful migration. Also, to ensure that there are no other users connecting to the old MS SQL database.

- To confirm that there aren't clients connecting to the MS SQL database, the following query should return no results other than your individual connection:

```
--- SQL Server releases after 2020 might require adjusting query column names due to  
version deprecating functionality.
```

```
SELECT session_id, login_name,  
       db_name(database_id) database_name,  
       program_name, HOST_NAME  
FROM sys.dm_exec_sessions where session_id > 50;
```

- If the AWS DMS task is still replicating ongoing changes, allow some time after stopping the application(s) connecting to the old database. After ensuring that DMS is no longer replicating any changes by examining CloudWatch metrics, stop the AWS DMS task.
- Then, re-run the script from point 1, and if there are any remaining sessions (other than your own), consider checking the database application settings to make sure it was redirected to the new PostgreSQL database. If your application is still using the SQL Server database, you may need to resume the DMS task to replicate all data. Data might have been changed by the application after the initial task stop. Correct your application connection settings again, and then repeat the cleanup process.
- Once full load and CDC ongoing replication completes, and you are done using DMS tasks, you can run the scripts (for example, `constraint-putback`) saved from AWS SCT to apply the foreign keys, constraints, and triggers on the Aurora PostgreSQL database.
- Run a last full database backup on the old SQL Server database.
- Shut down SQL Server services if the SQL server can be decommissioned.

## Conclusion

AWS DMS is a robust data migration service which can greatly simplify the process of migrating your database. In this walkthrough, we've shown you how to determine if you need to create one or more DMS tasks to support your data migration to Aurora PostgreSQL. We shared with you our automatic SQL queries to help you arrange your tables, prepare tasks for parallel load, and enable CDC on tables with or without primary keys. We recommend that you review our public documentation for additional information of DMS features and enhancements which can further improve your overall database migration process.

# Migrating an Amazon RDS for Oracle Database to an Amazon S3 Data Lake

This walkthrough helps you understand the process of migrating data from Amazon Relational Database Service (Amazon RDS) for Oracle to Amazon Simple Storage Service (Amazon S3) using AWS Database Migration Service (AWS DMS).

Most organizations use Online Transaction Processing (OLTP) database engines to host their transactional workloads. These engines are optimized for high-transaction volumes such as an online order processing application. However, these engines typically perform poorly for analytical applications, such as business intelligence or building predictive models using machine learning. For these use cases, a popular solution is to build a data lake for analysis.

In this document, we build a data lake in Amazon S3 using data hosted in an RDS for Oracle database. Amazon S3 is the largest and most performant cloud storage service. With Amazon S3, you can build a cost-effective, secure data lake. Amazon S3 provides 99.999999999% (11 9s) of data durability and makes it possible to store and manage both structured and unstructured data at unlimited scale.

To illustrate the process, we use AWS DMS to migrate data from an example database. AWS DMS is a managed service that helps migrate between heterogeneous sources and targets. In our case, we migrate an Oracle database to Amazon S3. AWS DMS support not only the migration of your existing data, but also ensures that the source and target are synchronized for ongoing transactions.

## Topics

- [Why use AWS DMS?](#)
- [Example data set](#)
- [Solution overview](#)
- [Prerequisites for migrating an RDS for Oracle database to an Amazon S3 data lake](#)
- [Step-by-step Amazon RDS for Oracle database to Amazon S3 data lake migration walkthrough](#)
- [Step-by-step Amazon RDS for Oracle database to Amazon S3 data lake migration conclusion](#)


## Why use AWS DMS?

You can use a SQL-level mechanism to source ongoing changes. This approach impacts your source database performance or requires that you implement additional logic. For example, you can use SQL filters on last updated timestamps or add triggers to capture DML changes. In contrast, AWS DMS mines changes from the database transaction logs, which are generated by the database for recovery purposes. AWS DMS then takes those changes, converts them into the target format, and applies them to the target. This process minimizes overhead on the source and provides near-real time replication to the target.

In the rest of this document, we guide you through the steps that you take to migrate the example Oracle database into Amazon S3. In the next sections, we describe the characteristics of the database. Then we build the replication resources in AWS DMS that we use to migrate the database, paying close attention to matching the AWS DMS configuration with our particular use case.

## Example data set

For this walkthrough, we use the [Sales History Oracle sample data set](#). The Sales History schema includes 8 tables. The largest table is the sales table which is a fact table with 5 billion rows. The total size of this source database is about 200 GB and has about 20 years worth of sales history data in 96 partitions. The remaining tables are mostly smaller dimension tables. The following diagram shows the data model for our sample use case.



Note that the costs and sales tables don't have primary keys. However, the sales table is partitioned on a date column. This date column is important to sequence the latest version of a sales record for analysis purposes.

The company loads data into its data warehouse regularly to gather statistics for these reports. The company also runs reports on different distribution channels through which its sales are delivered. When the company runs special promotions on its products, it analyzes the impact of the promotions on sales. It also analyzes sales by different geographical regions.


The company in our use case does high volume of business, so it runs business statistics reports and uses machine learning algorithms to aid in decision-making. Most of this analysis is time-sensitive, and they analyze past data trends to get insights on business operations.

The company's data scientists want to explore the data to decide which data to use for model training. Once this data discovery phase is complete, the data will be used to build predictive models using machine learning algorithms. Once the data is migrated to S3, it is used for training machine learning (ML) models using AWS ML managed-services. These models will be used for demand product forecasting and inventory replenishment.

The business also requires that the initial transfer of data from Oracle to Amazon S3 must complete within an 8 hour window.

## Solution overview

The following diagram shows the architecture of a migration from RDS for Oracle to Amazon S3 using AWS DMS.



The Amazon RDS for Oracle database contains the example sales history data set. AWS Database Migration Service (AWS DMS) contains several components used to host the replication engine. Amazon S3 provides storage for the data lake tables and downstream applications for machine learning and analytics consume the data lake information.

To run this walkthrough, create the following resources in AWS DMS.

- **Replication instance** — An AWS managed instance that hosts the AWS DMS engine. You control the type or size of the instance based on the workload you plan to migrate.
- **Source endpoint** — An endpoint that provides connection details, data store type, and credentials to connect to a source database. For this use case, we configure the source endpoint to point to the Amazon RDS for Oracle database.
- **Target endpoint** — AWS DMS supports several target systems including Amazon RDS, Amazon Aurora, Amazon Redshift, Amazon Kinesis Data Streams, Amazon S3, and so on. For this use case, we configure Amazon S3 as the target endpoint.
- **Replication task** — A task that runs on the replication instance and connects to endpoints to replicate data from the source database to the target database.



In the rest of this document, we show how to configure each of these components to migrate the sales history data set. We start with the prerequisites to complete this walkthrough, and then continue with the step-by-step migration procedure and conclusion.

## Prerequisites for migrating an RDS for Oracle database to an Amazon S3 data lake

The following prerequisites are required to complete this walkthrough:

- Understand Amazon Relational Database Service (Amazon RDS), the applicable database technologies, and SQL.
- Create a user with AWS Identity and Access Management (IAM) credentials that allows you to launch Amazon RDS and AWS Database Migration Service (AWS DMS) instances in your AWS Region. For information about IAM credentials, see [Create an administrative user](#).
- Understand the Amazon Virtual Private Cloud (Amazon VPC) service and security groups. For information about using Amazon VPC with Amazon RDS, see [Amazon VPC VPCs and Amazon RDS](#). For information about Amazon RDS security groups, see [Controlling access with security groups](#).
- Understand the supported features and limitations of AWS DMS. For information about AWS DMS, see [What is Database Migration Service?](#).
- Understand how to work with Oracle as a source and Amazon S3 data lake as a target. For information about working with Oracle as a source, see [Using an Oracle database as a source](#). For information about working with Amazon S3 as a target, see [Using Amazon S3 as a target](#).
- Understand the supported data type conversion options for Oracle and Amazon S3. For information about data types for Oracle as a source, see [Source data types for Oracle](#). For information about data types for Amazon S3 as a target (Parquet only), see [Target data types for S3 Parquet](#).
- Audit your source Oracle database. For each schema and all the objects under each schema, determine whether any of the objects are no longer being used. Deprecate these objects on the source Oracle database, because there's no need to migrate them if they aren't being used.

For more information about AWS DMS, see [Getting started with Database Migration Service](#).

To estimate what it will cost to run this walkthrough on AWS, you can use the AWS Pricing Calculator. For more information, see <https://calculator.aws/>.

To avoid additional charges, delete all resources after you complete the walkthrough.

## Step-by-step Amazon RDS for Oracle database to Amazon S3 data lake migration walkthrough

The following steps provide instructions for migrating an Amazon RDS for Oracle database to an Amazon S3 data lake. These steps assume that you have already prepared your source database as described in previously.

### Topics

- [Step 1: Create an AWS DMS Replication Instance](#)
- [Step 2: Configure a Source Amazon RDS for Oracle Database](#)
- [Step 3: Create an AWS DMS Source Endpoint](#)
- [Step 4: Create a Target Amazon S3 Bucket](#)
- [Step 5: Configure an AWS DMS Target Endpoint](#)
- [Step 6: Create an AWS DMS Task](#)
- [Step 7: Run the AWS DMS Task](#)

### Step 1: Create an AWS DMS Replication Instance

An AWS DMS replication instance hosts the software migrates data between the source and target. The replication instance also caches the transaction logs during the migration. The CPU and memory capacity of the replication instance influences the overall time needed for the migration. Make sure that you consider the specifics of your particular use case when you determine the size of your replication instance. A full load task consumes a lot of memory if it is run multithreaded. For more information, see [Choosing the right replication instance for your migration](#).

For our use case, we have a limited time window of 8 hours to complete the full load, and the sales table that includes 197 GB of data. Our goal is to fit into the 8 hour window. Therefore, we scale the replication instance to accommodate these requirements.

Each type of instance class has different CPU, memory, and I/O capacity. Sizing the replication instance should be based on factors like data volume, transaction frequency, large objects (LOBs) within storage of the data migration, and so on. We initially chose a DMS t3.medium instance running the latest AWS DMS engine version. This instance completed the migration in 18 hours. We then upgraded to a DMS c5.12xlarge instance. This instance size, combined with the proper task configuration, brought the full load time to under 8 hours.

We also upgraded the storage of the replication instance to 200 GB, and as a result, 600 IOPS were available for our replication instance. By default, DMS allocates 50 GB of storage to a replication instance. This may not be sufficient for use cases where more tasks are running on same replication instance or when running tasks with parallel load for large tables. With 600 IOPS, we saved several minutes of migration time. For more information about storage volume performance and burst I/O credits, see [General Purpose SSD \(gp2\) volumes](#).

Because we replicate production data in this walkthrough, we use the Multi-AZ deployment option for our replication instance for high availability. Also, we didn't make this replication instance publicly accessible for additional security.

For information about best practices for using AWS DMS, see [Database Migration Service Best Practices](#).

### To create an AWS DMS replication instance

1. Sign in to the AWS Management Console, and open the [AWS DMS console](#).
2. Choose **Replication instances**, then choose **Create replication instance**.
3. On the **Create replication instance** page, enter the following information.

| Parameter             | Action                                                                                                                                                                                                     |
|-----------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Name</b>           | Enter <code>oracle-s3-migration-replication-instance</code> . If you use multiple replication servers or sharing a user, choose a name that helps you quickly differentiate between the different servers. |
| <b>Description</b>    | Enter <code>Replication instance that supports Oracle to S3 data lake migration</code> . You can change the description to fit your use case.                                                              |
| <b>Instance class</b> | Choose <code>dms.c5.12xlarge</code> .                                                                                                                                                                      |
| <b>VPC</b>            | Choose the virtual private cloud (VPC) where AWS DMS launches your replication instance. If possible, select the same VPC in which                                                                         |

| Parameter                  | Action                                                   |
|----------------------------|----------------------------------------------------------|
|                            | either your source or target database resides (or both). |
| <b>Multi AZ</b>            | Choose <b>Yes</b> .                                      |
| <b>Publicly accessible</b> | Turn off this option.                                    |

#### 4. Choose **Create**.

## Step 2: Configure a Source Amazon RDS for Oracle Database

In this step, we configure the source Oracle database. Make sure that AWS DMS can access the database transaction logs and capture the data changes. Also, make sure that you set permissions for AWS DMS to access tables and database catalogs.

AWS DMS can help users to migrate historical data from Oracle source database and also replicate the ongoing changes to a centralized data lake. To use Oracle as a source for AWS DMS, you must turn on the `ARCHIVELOG MODE`.

Make sure that your database server retains the archive logs as long as AWS DMS requires them. If you configure your task to begin capturing changes immediately, then you should only need to retain archive logs for a little longer than the duration of the longest running transaction. Retaining archive logs for 24 hours is usually sufficient. If you configure your task to begin from a point in time in the past, then archive logs must be available from that time forward. For more information about turning on the `ARCHIVELOG MODE` and ensuring log retention for your Oracle database, see the [Oracle documentation](#).

To capture change data, AWS DMS requires that you turn on supplemental logging on your source database. Minimal supplemental logging must be turned on at the database level. AWS DMS also requires that you turn on identification key logging. This option causes the database to place all columns of a row's primary key in the redo log file whenever you update a row that contains a primary key. This occurs even if there is a change of value in any of the columns other than the primary key columns. You can set this option at the database or table level.

- Create or configure a database user AWS DMS. We recommend that you use a user with the minimal privileges required by AWS DMS for your connection. AWS DMS requires the following privileges.

```
CREATE SESSION
SELECT ANY TRANSACTION
SELECT on V_$ARCHIVED_LOG
SELECT on V_$LOG
SELECT on V_$LOGFILE
SELECT on V_$DATABASE
SELECT on V_$THREAD
SELECT on V_$PARAMETER
SELECT on V_$NLS_PARAMETERS
SELECT on V_$TIMEZONE_NAMES
SELECT on V_$TRANSACTION
SELECT on ALL_INDEXES
SELECT on ALL_OBJECTS
SELECT on ALL_TABLES
SELECT on ALL_USERS
SELECT on ALL_CATALOG
SELECT on ALL_CONSTRAINTS
SELECT on ALL_CONS_COLUMNS
SELECT on ALL_TAB_COLS
SELECT on ALL_IND_COLUMNS
SELECT on ALL_LOG_GROUPS
SELECT on SYS.DBA_REGISTRY
SELECT on SYS.OBJ$
SELECT on DBA_TABLESPACES
SELECT on ALL_TAB_PARTITIONS
SELECT on ALL_ENCRYPTED_COLUMNS
SELECT on <<all tables migrated>>
```

- For tables with primary key, turn on supplemental logging at key level.

```
ALTER TABLE table_name ADD SUPPLEMENTAL LOG DATA (PRIMARY KEY) COLUMNS;
```

- Data warehouse databases can have fact tables without a primary key. In our sample database, the SALES table doesn't have a primary key. This table is part of a full load and CDC task, however it does not have primary key or unique indexes. So we can add supplemental logging on all columns of the table.

```
ALTER TABLE SH.SALES ADD SUPPLEMENTAL LOG DATA (ALL) COLUMNS;
```

For more information, see [Working with an Oracle database as a source](#).

## Step 3: Create an AWS DMS Source Endpoint

In this step, we configure a source endpoint. AWS DMS uses this endpoint to connect to the source database to read data as well as changes to the data via transaction logs. You can use Extra Connection Attributes for the source endpoint to configure how AWS DMS captures changes to the data.

After you configure the AWS Database Migration Service (AWS DMS) replication instance and the source RDS for Oracle instance, ensure connectivity between these two instances. To ensure that the replication instance can access the server and the port for the database, make changes to the relevant security groups and network access control lists. For more information about your network configuration, see [Setting up a network for a replication instance](#).

AWS DMS can stream the changes to the data from REDO logs using either Logminer or Binary reader protocols. You can choose this protocol when you create your source endpoint. For detailed comparison on which mode to pickup for CDC replication, see [Using Oracle LogMiner or Binary Reader for CDC](#).

Logminer option is easier to set up. However, since our source Oracle database workload involves ETL jobs that result in high volume of transactions, we choose Binary Reader since it offers better performance for ongoing replication.

After you completed the network configurations, you can create a source endpoint.

### To create a source endpoint

1. Sign in to the AWS Management Console, and open the [AWS DMS console](#).
2. Choose **Endpoints**, then choose **Create endpoint**.
3. On the **Create endpoint** page, enter the following information.

| Parameter           | Action                                                                                                                                               |
|---------------------|------------------------------------------------------------------------------------------------------------------------------------------------------|
| Endpoint type       | Choose <b>Source endpoint</b> , turn on <b>Select RDS DB instance</b> , and choose an RDS for Oracle instance that you created for this walkthrough. |
| Endpoint identifier | Enter <code>datalake-source-db</code> .                                                                                                              |

| Parameter                                       | Action                                                                                                                                                           |
|-------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Source engine                                   | Choose <b>Oracle</b> .                                                                                                                                           |
| Access to endpoint database                     | Choose <b>Provide access information manually</b> . Alternatively, you can choose to provide a secret from AWS Secrets Manager that includes connection details. |
| Server name                                     | Enter the database server name on Amazon RDS.                                                                                                                    |
| Port                                            | Enter 1521.                                                                                                                                                      |
| Secure Socket Layer (SSL) mode                  | Choose <b>none</b> .                                                                                                                                             |
| User name                                       | Enter the name of the user that you created for your RDS for Oracle database.                                                                                    |
| Password                                        | Enter the password that you created for your Oracle DB user.                                                                                                     |
| SID/Service name                                | Enter SH.                                                                                                                                                        |
| Endpoint settings - Extra connection attributes | Enter <code>useLogminerReader=N;useBfile=Y; .</code>                                                                                                             |

4. Choose **Create endpoint**.

## Step 4: Create a Target Amazon S3 Bucket

Before you create the target endpoint, you create an Amazon S3 bucket for the target data lake.

To create the Amazon S3 bucket, do the following:

1. Sign in to the AWS Management Console, and open the [Amazon S3 console](#).
2. Choose **Create bucket**.
3. For **Bucket name**, enter `s3-dataLake`.
4. For **AWS Region**, choose the region that hosts your AWS DMS replication instance.

5. Keep the default values in the other fields and choose **Create bucket**.

You can also plan to optimize the storage cost from Amazon S3 using [Intelligent-Tiering](#) and [Lifecycle policies](#) when storing huge volume of data.

Now, you have the Amazon S3 bucket for your data lake. Next, you can create a target endpoint for this bucket.

## Step 5: Configure an AWS DMS Target Endpoint

In this section, we walk through the configuration for setting up target data lake AWS DMS endpoint. You will also select appropriate options to store files in data lake.

To use Amazon S3 as an AWS Database Migration Service (AWS DMS) target endpoint, create an IAM role with write and delete access to the AWS DMS bucket. Then add `dms.amazonaws.com` as a trusted entity in this IAM role. For more information, see [Prerequisites for using Amazon S3 as a target](#).

When you use AWS DMS to migrate data to an Amazon Simple Storage Service (Amazon S3) data lake, you can change the default task behavior, such as file formats, partitioning, file sizing, and so on. This leads to minimizing post-migration processing and helps downstream applications consume data efficiently. You can customize task behavior using endpoint settings and extra connection attributes (ECA). Most of the AWS DMS endpoint settings and ECA settings overlap, except for a few parameters. In this section of walkthrough, we configure AWS DMS endpoint settings.

### Choose File Format

For this walkthrough, we use the Parquet file format to help the data scientists consume data for data exploration and data discovery activities. Apache Parquet is a columnar format, which is built to support efficient compression and encoding schemes providing storage space savings and performance benefits.

Specify the following endpoint settings.

```
DataFormat=parquet
ParquetVersion=PARQUET_2_0
```



## Determine File Size

By default, during ongoing replication AWS DMS task write calls to Amazon S3 are triggered either if the file size reaches 32 KB or if the previous file write was more than 60 seconds ago. These settings ensure that the data capture latency is less than a minute. However, this approach creates numerous small files in target Amazon S3 bucket.

Because we migrate our source Sales History database schema for a machine learning use case, some latency is acceptable. However, we need to optimize this schema for cost and performance.

During the data discovery phase performed by the data scientists, it is helpful to have large files for efficient analysis using the tools of their choice. We recommend that you set the size of the target file to at least 64 MB. Specify the following endpoint settings: `CdcMaxBatchInterval=3600` and `CdcMinFileSize=64000`. These settings ensure that AWS DMS writes the file until its size reaches 64 MB or if the last file write was more than an hour ago.

### Note

Parquet files created by AWS DMS are usually smaller than the specified `CdcMinFileSize` setting because Parquet data compression ratio varies depending on the source data set. The size of CSV files created by AWS DMS is equal to the value specified in `CdcMinFileSize`.

## Turn on S3 Partitioning

Partitioning in Amazon S3 structures your data by folders and subfolders that help efficiently query data. For example, if you receive sales record data daily from different regions, and you query data for a specific region and find stats for a few months, then you can partition data by {Product/source/region}, year, and month.

The following example shows the path in Amazon S3 for our use case.

```
s3://<sales-analytics-bucket-name>/<Project/Source/Region>/<schemaname>/<tablename>/  
<year>-<month>-<day>
```

```
s3://s3-datalake
```

- s3://s3-datalake/oracledb
- s3://s3-datalake/oracledb/Sales
  - s3://s3-datalake/oracledb/Sales/Products/
    - s3://s3-datalake/oracledb/Sales/Products/LOAD000000001.parquet

```
- s3://s3-datalake/Oracledb/Sales/Customer
- s3://s3-datalake/Oracledb/Sales/Customer/LOAD00000001.parquet
- s3://s3-datalake/Oracledb/Sales/Sales/Products/2022-10-23/
- s3://s3-datalake/Oracledb/Sales/Sales/
Products/2022-10-23/20221023-013830913.parquet
- s3://s3-datalake/Oracledb/Sales/Sales/
Products/2022-10-24/20221024-175902985.parquet
```

Partitioning provides performance benefits because data scanning will be limited to the amount of data in the specific partition based on the filter condition in your queries. For our sales data example, data scientists' queries might look as follows:

```
SELECT <column-list> FROM <sales-hist-table-name> WHERE <region> = <region-name> AND
<date> = <date-to-query>
```

When performing data exploration, the data scientists can consume the incremental load using partitions. Partitioning the data helps read only latest data from the Amazon S3 bucket. In this case, you explore the latest data and use it for training the models to determine latest sales trends.

The following code example shows how to turn on partitioning for ongoing changes.

```
bucketFolder=Oracledb
DatePartitionedEnabled=true
DatePartitionSequence=YYYYMMDD
DatePartitionDelimiter=DASH
```

#### Note

The date partition delimiter is chosen as DASH because it creates prefixes in the format YYYY-MM-DD rather than YYYY/MM/DD format. The advantage of using DASH is that it makes the 3 console view better with the files from each date (YYYY-MM-DD) being a single folder rather than having different folders for Year, month, and date. This will also let users query for a particular date in a simpler manner.

## Serialize Ongoing Replication Events

A common challenge when using Amazon S3 as a target involves identifying the ongoing replication event sequence when multiple records are updated at the same time on the source database.

AWS DMS provides two options to help serialize such events for Amazon S3. You can use the `TimeStampColumnName` endpoint setting or use transformation rules to include the Oracle System Change Number (SCN) column. The `TimeStampColumnName` setting adds another `STRING` column to the target file created by AWS DMS. During the ongoing replication, the column value represents the commit timestamp of the event in the Oracle database. For the full load phase, the column values represent the timestamp of data transfer to Amazon S3. The second option adds another column to include Oracle SCN. You can use this field when the source database might have transactions that are occurring within a microsecond or if the source database doesn't offer microsecond level precision.

Because the sales history table doesn't have a primary key column, we add the `Timestamp` column according to the option to add **`TimeStampColumnName`** which will serve as a unique identifier during data exploration and model training phases of machine learning. We chose the option of timestamp over Oracle SCN because partitioning the data by timestamp will help data scientists for data exploration based on various criteria such as seasonal demand or product promotions.

This setting is done at a task level. Make sure that you repeat it for each task separately that migrates data from the Oracle database endpoint.

For more information about this option, see [the section called "Step 6: Create an AWS DMS Task"](#).

### To create a target endpoint

1. Sign in to the AWS Management Console, and open the [AWS DMS console](#).
2. Choose **Endpoints**, then choose **Create endpoint**.
3. On the **Create endpoint** page, enter the following information.

| Parameter               | Action                                                                       |
|-------------------------|------------------------------------------------------------------------------|
| Endpoint type           | Choose <b>Target endpoint</b> , and turn off <b>Select RDS DB instance</b> . |
| Endpoint identifier     | Enter <code>oracle-datalake-target</code> .                                  |
| Target engine           | Choose <b>Amazon S3</b> .                                                    |
| Service access role ARN | Enter the IAM role that can access your Amazon S3 data lake.                 |

| Parameter     | Action               |
|---------------|----------------------|
| Bucket name   | Enter s3-data-lake . |
| Bucket folder | Enter Oracledb.      |

4. Expand the **Endpoint settings** section, choose **Wizard**, and then choose **Add new setting** to add the following information.

| Parameter              | Action         |
|------------------------|----------------|
| CdcMinFileSize         | 64000          |
| CdcMaxBatchInterval    | 3600           |
| CdcPath                | Oracledb       |
| DataFormat             | parquet        |
| DatePartitionDelimiter | DASH           |
| DatePartitionEnabled   | TRUE           |
| DatePartitionSequence  | YYYYMMDD       |
| ParquetVersion         | PARQUET_2_0    |
| TimestampColumnName    | sourcetscolumn |

5. Choose **Create endpoint**.

## Step 6: Create an AWS DMS Task

Before you create the replication task, it is important to understand the workload on the source database and usage pattern of the tables being replicated. This helps plan an effective migration approach and minimize any configuration or workload related issues. In this section, we first review the important considerations and then learn how to configure our walkthrough DMS task accordingly by applying table mappings and task settings.

## Considerations Before Creating an AWS DMS Task

### Size and number of records

The volume of migrated records affects the full load completion time. It is difficult to predict the full load time upfront, but testing with a replica of a production instance should provide a baseline. Use this estimate to decide whether you should parallelize full load by using multiple tasks or by using the parallel load option.

To speed up the full load of large tables such as sales table in our use case, we can increase the number of tables and partitions loaded in parallel up to 49. The default value for the number of tables and partitions loaded in parallel is eight. For more information about parallel load task settings, see [Full-load task settings](#).

The `MaxFullLoadSubTasks` parameter controls number of tables or partitions loaded in parallel during full load.

### Transactions per second

While full load is affected by the number of records, the ongoing replication performance relies on the number of transactions on the source Oracle database. Performance issues during change data capture (CDC) generally stem from resource constraints on the source database, replication instance, target database, and network bandwidth or throughput. Knowing average and peak TPS on the source and recording CDC throughput and latency metrics helps baseline AWS DMS performance and identify an optimal task configuration. For more information, see [Replication task metrics](#).

In this walkthrough, the source database is a data warehouse where transaction volume is not always high because the data is loaded on a periodic basis from the Online Transaction Processing (OLTP) layer. Also, we run a heterogeneous data migration using the AWS DMS parallel load to migrate large tables with improved performance. For more information, see [Using parallel load for selected tables, views, and collections](#).

This approach requires a replication instance with higher compute capacity if the data volume is huge. We chose the compute intensive c5 class replication instance to speed up the process.

If you are not sure about your data volumes or performance expectations from the migration task, start with general t3 class instances, and then migrate to c5 class instances for compute intensive tasks or r5 class instances for memory intensive tasks. You should monitor the task metrics continuously and choose the appropriate instance class that best suits your needs.

## Unsupported data types

Identify data types used in tables and check that AWS DMS supports these data types. For more information, see [Source data types for Oracle](#).

Validate that the target Amazon S3 has the corresponding data types. For more information, see [Target data types for S3 Parquet](#).

After you run the initial load test, validate that AWS DMS converted data as you expected. You can also initiate a pre-migration assessment to identify any unsupported data types in the migration scope. For more information, see [Specifying individual assessments](#).

## Source Database Workload

Running AWS DMS replication tasks for large tables can add to the workload on the source database especially during the full load phase when AWS DMS reads whole tables from source database without any filters to restrict rows. When you use filters in AWS DMS task table mapping, confirm that appropriate indexes exist on the source tables and indexes are actually being used during full load. Regularly monitor the source database to identify any workload related issues. For more information, see [Using table mapping to specify task settings](#).

### Note

The previous list isn't complete. For more information, see [Best practices](#).

Combining the considerations from the previous list, we start with a single task that migrates all eight tables in parallel. Based on the full load run time and resource utilization metrics on the source Oracle database instance and replication instance, we used AWS DMS parallel load option to further improve full load performance.

## Task Configuration

In this walkthrough, we migrate the incremental changes to the fact tables to the data lake. To do so, we use the Full Load + CDC option. For more information about the AWS DMS task creation steps and available configuration options, see [Creating a task](#).

We will first focus on the following settings.

## Table mappings

Use selection rules to define the schemas and tables that the AWS DMS task migrates. For more information, see [Selection rules and actions](#).

In this walkthrough, we are migrating all the tables (%) in the sales history SH schema. Another option is to include each table explicitly in the table mappings. However, that increases operational overhead by requiring repeated configurations. If we plan to add new tables to source database in future under the sales history schema, we should include all tables (%) in table mapping.

#### Note

Mapping rules are applied at the task level. Make sure that you add a mapping rule to each task that replicates data to your data lake. For our use case we need a single task.

## LOB settings

AWS DMS handles large binary objects (LOBs) columns differently compared to other data types. For more information, see [Migrating large binary objects \(LOBs\)](#).

A detailed explanation of LOB handling by AWS DMS is out of scope for this walkthrough. However, remember that increasing the LOB Max Size increases the task's memory utilization. Because of that, we recommended that you don't set LOB Max Size to a large value. For more information about LOB settings, see [Task Configuration](#).

The source data warehouse schema in this walkthrough doesn't include LOB data. When you migrate LOB columns, make sure that you perform analysis on these columns. Because AWS DMS doesn't support Full LOB mode for Amazon S3 endpoints, we need to identify a suitable LOB Max Size.

## Parallel load

Though, we used significantly large instance class in previous run, overall improvement wasn't significant because the data volume is relatively large. The sales fact table includes 5 billion records. To further optimize the performance, we used parallel-load ranges option. For more information, see [Using parallel load for selected tables, views, and collections](#).

The following code example shows the mapping rule that we used. As you can see, we defined 16 boundaries to cover data from 1998 to 2026 in 16 ranges. With this option, full load finished in about 6.5 hours. As a result, we reduced the time taken to complete full load to almost one third as compared to initial load.

```
{
  "rules": [
    {
      "rule-type": "selection",
      "rule-id": "653647496",
      "rule-name": "653647496",
      "object-locator": {
        "schema-name": "SH",
        "table-name": "SALES"
      },
      "rule-action": "include",
      "filters": []
    },
    {
      "rule-type": "table-settings",
      "rule-id": "653647497",
      "rule-name": "653647497",
      "object-locator": {
        "schema-name": "SH",
        "table-name": "SALES"
      },
      "parallel-load": {
        "type": "ranges",
        "columns": [
          "TIME_ID"
        ],
        "boundaries": [
          [
            "1998-01-01 00:00:00"
          ],
          [
            "2000-01-01 00:00:00"
          ],
          [
            "2002-01-01 00:00:00"
          ],
          [
            "2004-01-01 00:00:00"
          ],
          [
            "2006-01-01 00:00:00"
          ],
          [
```



```

        "2008-01-01 00:00:00"
    ],
    [
        "2010-01-01 00:00:00"
    ],
    [
        "2012-01-01 00:00:00"
    ],
    [
        "2014-01-01 00:00:00"
    ],
    [
        "2016-01-01 00:00:00"
    ],
    [
        "2018-01-01 00:00:00"
    ],
    [
        "2020-01-01 00:00:00"
    ],
    [
        "2022-01-01 00:00:00"
    ],
    [
        "2024-01-01 00:00:00"
    ],
    [
        "2026-01-01 00:00:00"
    ]
]
}
}
.
.
.
    ]
}

```

You can also use the `partitions-auto` option instead of `ranges` option because the `SALES` table is already partitioned. In our testing, we found that with the `ranges` option, full load finishes faster. So, we chose `ranges` option.

## Other task settings

Choose **Enable CloudWatch Logs** to upload the AWS DMS task run log to Amazon CloudWatch. You can use these logs to troubleshoot issues because they include error and warning messages, start and end times of the run, configuration issues, and so on. To diagnose performance issues, you can use changes to the task logging setting, such as enabling debug or trace.

 **Note**

CloudWatch log usage is charged at standard rates. For more information, see [Amazon CloudWatch pricing](#).

For **Target table preparation mode**, choose one of the following options: **Do nothing**, **Truncate**, and **Drop**. Use **Truncate** in data pipelines where the downstream systems rely on a fresh dump of clean data and do not rely on historical data. In this walkthrough, we choose **Do nothing** because we want to control the retention of files from previous runs.

For **Maximum number of tables to load in parallel**, enter the number of parallel threads that AWS DMS initiates during the full load. You can increase this value to improve the full load performance and minimize the load time when you have numerous tables. Because we have several partitions that AWS DMS can load in parallel, we used the maximum value of 49.

 **Note**

Increasing this parameter induces additional load on the source database, replication instance, and target database.

## To create a database migration task

1. Sign in to the AWS Management Console, and open the [AWS DMS console](#).
2. Choose **Database migration tasks**, then choose **Create task**.
3. On the **Create database migration task** page, enter the following information.

| Parameter       | Action                         |
|-----------------|--------------------------------|
| Task identifier | Enter Oracle-to-S3-data-lake . |

| Parameter                                                                                       | Action                                                                                        |
|-------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------|
| Replication instance                                                                            | Choose <b>oracle-s3-migration-replication-instance</b> . You configured this value in Step 1. |
| Source database endpoint                                                                        | Choose <b>datalake-source-db</b> . You configured this value in Step 3.                       |
| Target database endpoint                                                                        | Choose <b>oracle-datalake-target</b> . You configured this value in Step 5.                   |
| Migration type                                                                                  | Choose <b>Migrate existing data and replicate ongoing changes</b> .                           |
| Editing mode                                                                                    | Choose <b>Wizard</b> .                                                                        |
| Custom CDC stop mode for source transactions                                                    | Choose <b>Disable custom CDC stop mode</b> .                                                  |
| Target table preparation mode                                                                   | Choose <b>Do nothing</b> .                                                                    |
| Stop task after full load completes                                                             | Choose <b>Don't stop</b> .                                                                    |
| Include LOB columns in replication                                                              | Choose <b>Limited LOB mode</b> .                                                              |
| Maximum LOB size (KB)                                                                           | Enter 32.                                                                                     |
| Advanced task settings, Full load tuning settings, Maximum number of tables to load in parallel | Enter 49.                                                                                     |
| Enable validation                                                                               | Turn off because Amazon S3 doesn't support validation.                                        |
| Enable CloudWatch logs                                                                          | Turn on.                                                                                      |

4. Keep the default values for other parameters, and choose **Create task**.

AWS DMS runs the task immediately. The **Database migration tasks** section displays the status of the migration task.

## Step 7: Run the AWS DMS Task

After you create your AWS Database Migration Service (AWS DMS) task, do a test run to identify the full load run time and ongoing replication performance. You can validate that initial configurations work as expected. You can do this by monitoring and documenting resource utilization on the source database, replication instance, and target database. These details make up the initial baseline and help determine if you need further optimization.

After you started the task, the full load operation starts loading tables. You can see the table load completion status in the **Table Statistics** section and the corresponding target files in the Amazon S3 bucket.

The following image shows table statistics with c5.12xlarge replication instance with parallel-load ranges option. The full load completed in 6.5 hours. This means that we achieved our goal of completing full load in less than 8 hours.

Summary

Status

Load complete, replication ongoing

Type

Full load, ongoing replication

Source

dms-walkthrough-ora19a-src

Target

dms-walkthrough-s3-tgt

Overview details

Table statistics

CloudWatch metrics

Mapping rules

Premigration assessments

Tags

Table statistics (11)

Total rows include loaded source table rows from Inserts, Deletes, Updates, DDLs, and Full load rows.

Find schema

Export to CSV

| <input type="checkbox"/> | Schema name ▾ | Table ▾                    | Load state ▾    | Elapsed load time ▾ | Inserts ▾ | Deletes ▾ | Updates ▾ | DDLs ▾ | Applied inserts ▾ | Applied deletes ▾ | Applied updates ▾ | Applied DDLs ▾ | Full load rows ▾ |
|--------------------------|---------------|----------------------------|-----------------|---------------------|-----------|-----------|-----------|--------|-------------------|-------------------|-------------------|----------------|------------------|
| <input type="checkbox"/> | SH            | SALE_STATS                 | Table completed | 2 s                 | 0         | 0         | 0         | 0      | 0                 | 0                 | 0                 | 0              | 3                |
| <input type="checkbox"/> | SH            | QUARTERLY_REGIONAL_SALES   | Table completed | 1 s                 | 0         | 0         | 0         | 0      | 0                 | 0                 | 0                 | 0              | 33               |
| <input type="checkbox"/> | SH            | PRODUCTS                   | Table completed | 1 s                 | 0         | 0         | 0         | 0      | 0                 | 0                 | 0                 | 0              | 72               |
| <input type="checkbox"/> | SH            | PROMOTIONS                 | Table completed | 1 s                 | 0         | 0         | 0         | 0      | 0                 | 0                 | 0                 | 0              | 503              |
| <input type="checkbox"/> | SH            | CHANNELS                   | Table completed | 3 s                 | 0         | 0         | 0         | 0      | 0                 | 0                 | 0                 | 0              | 5                |
| <input type="checkbox"/> | SH            | COSTS                      | Table completed | 4 s                 | 0         | 0         | 0         | 0      | 0                 | 0                 | 0                 | 0              | 82,112           |
| <input type="checkbox"/> | SH            | SALES                      | Table completed | 6 h 39 m 31 s       | 0         | 0         | 0         | 0      | 0                 | 0                 | 0                 | 0              | 5,013,024,433    |
| <input type="checkbox"/> | SH            | COUNTRIES                  | Table completed | 3 s                 | 0         | 0         | 0         | 0      | 0                 | 0                 | 0                 | 0              | 23               |
| <input type="checkbox"/> | SH            | TIMES                      | Table completed | 5 s                 | 0         | 0         | 0         | 0      | 0                 | 0                 | 0                 | 0              | 10,227           |
| <input type="checkbox"/> | SH            | SUPPLEMENTARY_DEMOGRAPHICS | Table completed | 3 s                 | 0         | 0         | 0         | 0      | 0                 | 0                 | 0                 | 0              | 4,500            |
| <input type="checkbox"/> | SH            | CUSTOMERS                  | Table completed | 3 s                 | 0         | 0         | 0         | 0      | 0                 | 0                 | 0                 | 0              | 55,500           |

We also monitored the CloudWatch metrics such as compute, memory, network to identify the resource usage of AWS DMS instances. You have to identify the resource constraint and scale-up to the AWS DMS instance class that serves your workloads better. You could also scale-down the AWS DMS instance to a t3 or r5 instance class based on the transaction volume for your ongoing replication task.

Because we turned on the parallel-load option, the I/O load on the replication instance is expected to increase. We described in [the section called “Step 1: Create an AWS DMS Replication Instance”](#)

that you should monitor the **Write IOPS** and **Read IOPS** metrics in CloudWatch to make sure that the total IOPS (write + read IOPS) doesn't exceed the total IOPS available for your replication instance. If it does, make sure that you allocate more storage to scale for better I/O performance. For more information, see [Monitoring replication tasks using Amazon CloudWatch](#).

We covered most prerequisites that help avoid errors related to configuration. If you observe issues when running the task, then see [Troubleshooting migration tasks in Database Migration Service](#), [Best practices for Database Migration Service](#), or reach out to AWS Support for further assistance.

Optionally, you could choose to validate the successful completion of the data migration by querying the Amazon S3 data through Athena console. You can run count or aggregation queries on key metric columns and compare with the source database to validate the migration task. AWS DMS also provides data validation features to verify successful migration of the data. For more information, see [Data validation](#).

After you completed the migration, validate that your data migrated successfully and delete the cloud resources that you created.

## Step-by-step Amazon RDS for Oracle database to Amazon S3 data lake migration conclusion

In this walkthrough, we covered all steps that you need to take to migrate a sales history data warehouse from Oracle to an Amazon S3 data lake. Our example company can use this data lake for machine learning and analysis use cases. We achieved the crucial business requirements by using AWS DMS. Try out these steps to migrate your data to an Amazon S3 data lake and explore how you can centralize your data with a low-cost solution. To learn more about AWS DMS service, see [Database Migration Service Documentation](#).

## Migrating an Amazon RDS for SQL Server Database to an Amazon S3 Data Lake

This walkthrough gets you started with the process of migrating from an Amazon Relational Database Service (Amazon RDS) for Microsoft SQL Server to Amazon Simple Storage Service (Amazon S3) cloud data lake using AWS Database Migration Service (AWS DMS).

For most organizations, data is distributed across multiple systems and data stores to support varied business needs. On-premises data stores struggle to scale performance as data sizes and

formats grow exponentially for analytics and reporting purposes. These limitations in data storage and management limit efficient and comprehensive analytics.

Amazon S3 based data lakes provide reliable and scalable storage, where you can store structured, semi-structured and unstructured datasets for varying analytics needs. You can integrate Amazon S3 based data lakes with distributed processing frameworks such as Apache Spark, Apache Hive, and Presto to decouple compute and storage, so that both can scale independently.

## Topics

- [Why Amazon S3?](#)
- [Why AWS DMS?](#)
- [Solution overview](#)
- [Prerequisites for migrating from an Amazon RDS for SQL Server database to an Amazon S3 data lake](#)
- [Step-by-step Amazon RDS for SQL Server database to an Amazon S3 data lake migration walkthrough](#)

## Why Amazon S3?

Amazon S3 is an object storage service for structured, semi-structured, and unstructured data that offers industry-leading scalability, data availability, security, and performance. With a data lake built on Amazon S3, you can use native AWS services, optimize costs, organize data, and configure fine-tuned access controls to meet specific business, organizational, and compliance requirements.

Amazon S3 is designed for 99.999999999% (11 9s) of data durability. The service automatically creates and stores copies of all uploaded S3 objects across multiple systems. This means your data is available when needed and protected against failures, errors, and threats.

Amazon S3 is secure by design, scalable on demand, and durable against the failure of an entire AWS Availability Zone. You can use AWS native services and integrate with third-party service providers to run applications on your data lake.

## Why AWS DMS?

Data lakes typically require building, configuring, and maintaining multiple data ingestion pipelines from cloud and on-premises data stores.

Traditionally, databases can be loaded once with data ingestion tools such as import, export, bulk copy, and so on. Ongoing changes are either not possible or are implemented by bookmarking the initial state. Setting up a data lake using these methods can present challenges ranging from increased load on the source database to overheads while carrying schema changes.

AWS DMS supports a one-time load and near-real-time ongoing replication making the data migration seamless, while supporting multiple source and target database platforms. One of the common use cases is the need to derive insights on data stored in several sources. For example, you may need to identify monthly sales for a specific year on sales data stored on different database instances.


As a part of this walkthrough, we will configure AWS DMS to move data from an Amazon RDS for SQL Server database instance to Amazon S3 for a sales analytics use case.

### Note


This introductory exercise doesn't cover all use cases of migrating to Amazon S3 but provides an overview of the migration process using AWS DMS. This example covers commonly faced problems and describes best practices to follow when migrating to an Amazon S3 data lake.

## Solution overview

The following diagram displays a high-level architecture of the solution, where we use AWS DMS to move data from Microsoft SQL Server databases hosted on Amazon Relational Database Service (Amazon RDS) to Amazon Simple Storage Service (Amazon S3).



The following diagram shows the structure of the Amazon S3 bucket from the preceding diagram.



To replicate data, you need to create and configure the following artifacts in AWS DMS:

- **Replication Instance** — An AWS managed instance that hosts the AWS DMS engine. You control the type or size of the instance based on the workload you plan to migrate.
- **Source Endpoint** — An endpoint that provides connection details, data store type, and credentials to connect to a source database. For this use case, we will configure the source endpoint to point to the Amazon RDS for SQL Server database.
- **Target Endpoint** — AWS DMS supports several target systems including Amazon RDS, Amazon Aurora, Amazon Redshift, Amazon Kinesis Data Streams, Amazon S3, and more. For the use case, we will configure Amazon S3 as the target endpoint.
- **Replication Task** — A task that runs on the replication instance and connects to endpoints to replicate data from the source database to the target database

For this walkthrough, we will use the AdventureWorks sample database on an Amazon RDS for SQL Server instance as the base data for the walkthrough. The AdventureWorks database holds sales, marketing, and order data. We will use AWS DMS to move sales data from the source database to Amazon S3 object store, which can be used as a data lake for downstream analytics needs.

### Note

You can refer to [the section called “Migrating a SQL Server Always On Database to Amazon Web Services”](#) for details on migrating from a Microsoft SQL Server Always On database instance.

We will create an AWS DMS task, which will perform a one-time full load to migrate a point in time snapshot and will then stream incremental data to the target Amazon S3 bucket. This way, sales data in the S3 bucket will be kept in sync with the source database.



## Prerequisites for migrating from an Amazon RDS for SQL Server database to an Amazon S3 data lake

The following prerequisites are required to complete this walkthrough:

- Understand Amazon Relational Database Service (Amazon RDS), the applicable database technologies, and SQL.
- Create a user with AWS Identity and Access Management (IAM) credentials that allows you to launch Amazon RDS and AWS Database Migration Service (AWS DMS) instances in your AWS Region. For information about IAM credentials, see [Create an IAM user](#).
- Understand the Amazon Virtual Private Cloud (Amazon VPC) service and security groups. For information about using Amazon VPC with Amazon RDS, see [Amazon Virtual Private Cloud \(VPCs\) and Amazon RDS](#). For information about Amazon RDS security groups, see [Controlling access with security groups](#).
- Understand the supported features and limitations of AWS DMS. For information about AWS DMS, see [What is Database Migration Service](#).
- Understand how to work with Microsoft SQL Server as a source and Amazon S3 data lake as a target. For information about working with SQL Server as a source, see [Using a Microsoft SQL Server database as a source](#). For information about working with Amazon S3 as a target, see [Using Amazon S3 as a target](#).
- Understand the supported data type conversion options for SQL Server and Amazon S3. For information about data types for SQL Server as a source, see [Source data types for SQL Server](#). For information about data types for Amazon S3 as a target (Parquet only), see [Target data types for S3 Parquet](#).
- Audit your source SQL Server database. For each schema and all the objects under each schema, determine whether any of the objects are no longer being used. Deprecate these objects on the source SQL Server database, because there's no need to migrate them if they aren't being used.

For more information about AWS DMS, see [Getting started with Database Migration Service](#).

To estimate what it will cost to run this walkthrough on AWS, you can use the AWS Pricing Calculator. For more information, see <https://calculator.aws/>.

To avoid additional charges, delete all resources after you complete the walkthrough.

# Step-by-step Amazon RDS for SQL Server database to an Amazon S3 data lake migration walkthrough

The following steps provide instructions for migrating an Amazon RDS for SQL Server database to an Amazon S3 data lake. These steps assume that you have already prepared your source database as described in [the section called “Prerequisites for migrating from an Amazon RDS for SQL Server database to an Amazon S3 data lake”](#).

## Topics

- [Step 1: Create an AWS DMS Replication Instance](#)
- [Step 2: Configure a Source Amazon RDS for SQL Server Database](#)
- [Step 3: Create an AWS DMS Source Endpoint](#)
- [Step 4: Configure a Target Amazon S3 Bucket](#)
- [Step 5: Configure an AWS DMS Target Endpoint](#)
- [Step 6: Create an AWS DMS Task](#)
- [Step 7: Run the AWS DMS Task](#)

## Step 1: Create an AWS DMS Replication Instance

To create an AWS Database Migration Service (AWS DMS) replication instance, see [Creating a replication instance](#). Usually, the full load phase is multi-threaded (depending on task configurations) and has a greater resource footprint than ongoing replication. Consequently, it's advisable to start with a larger instance class and then scale down once the task is in the ongoing replication phase. Moreover, if you intend to migrate your workload using multiple tasks, monitor your replication instance metrics and re-size your instance accordingly.

For this use case, we will migrate a subset (the Sales schema) of the AdventureWorks database, which is over 3 GB in size. Because we perform a heterogenous migration without many LOB columns, we can start with a compute optimized instance like c5.xlarge running the latest AWS DMS engine version. We can later scale up or down based on resource utilization during task execution.

### Note

Scaling replication instance during full load and ongoing replication phases is usually based on CloudWatch metrics such as CPU, memory, I/O, and so on. Choosing the appropriate

replication instance class and size depends on several factors such as number of tasks, table size, DML activity, size of transactions, Large Objects (LOB), and so on. This is out of scope for this walkthrough. To learn more about these topics, see [Choosing replication instance types](#) and [Sizing a replication instance](#).

To create an AWS DMS replication instance, do the following:

1. Sign in to the AWS Management Console, and open the [AWS DMS console](#).
2. If you are signed in as an AWS Identity and Access Management (IAM) user, you must have the appropriate permissions to access AWS DMS. For more information about the permissions required, see [IAM permissions](#).
3. On the Welcome page, choose **Create replication instance** to start a database migration.
4. On the **Create replication instance** page, specify your replication instance information.

| For This Parameter             | Do This                                                                                                                                                                                       |
|--------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Name</b>                    | Enter <code>datalake-migration-ri</code> . If you are using multiple replication servers or sharing a user, choose a name that helps you quickly differentiate between the different servers. |
| <b>Description</b>             | Enter <code>Migrate SQL Server to Amazon S3 data lake</code> .                                                                                                                                |
| <b>Instance class</b>          | Choose <code>dms.c5.xlarge</code> . Each size and type of instance class has increasing CPU, memory, and I/O capacity.                                                                        |
| <b>Engine version</b>          | Leave the default value, which is the latest stable version of the AWS DMS replication engine.                                                                                                |
| <b>Allocated storage (GiB)</b> | Choose 50.                                                                                                                                                                                    |
| <b>VPC</b>                     | Choose the virtual private cloud (VPC) in which your replication instance will launch. If                                                                                                     |

| For This Parameter         | Do This                                                                                                                                                                         |
|----------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                            | possible, select the same VPC in which either your source or target database resides (or both).                                                                                 |
| <b>Multi AZ</b>            | If you choose <b>Yes</b> , AWS DMS creates a second replication server in a different Availability Zone for failover if there is a problem with the primary replication server. |
| <b>Publicly accessible</b> | If either your source or target database resides outside of the VPC in which your replication server resides, you must make your replication server policy publicly accessible. |

## 5. Choose **Create**.

## Step 2: Configure a Source Amazon RDS for SQL Server Database

One of the primary considerations when setting up AWS DMS replication is the load that it induces on the source database. During full load, AWS DMS tasks initiate two or three connections for each table that is configured for parallel load. Because AWS DMS settings and data volumes vary across tasks, workloads, and even across different runs of the same task, providing an estimate of resource utilization that applies for all use cases is difficult.

Ongoing replication is single-threaded and it usually consumes less resources than full load. Providing estimates for change data capture (CDC) resource utilization has the same challenges described before.

That said, you can estimate the expected increase in load on your source Amazon RDS instance, by running test AWS DMS tasks on replicas of your source Amazon RDS for SQL Server instance and monitoring the CPU, memory, IO and throughput metrics.

For our source database, we use an `m5.xlarge` Amazon RDS instance running Microsoft SQL Server 2019. While the steps for Amazon RDS for SQL Server creation are out of scope for this walkthrough (for more information, see [the section called “Prerequisites for migrating from an Amazon RDS for SQL Server database to an Amazon S3 data lake”](#)), make sure that your Amazon

RDS instance has **Automatic Backups** turned on so that the recovery model for the database is set to **FULL**. This is a pre-requisite for ongoing replication with AWS DMS. You can turn on these settings when you create or modify an existing Amazon RDS instance.

The following image displays the database settings required for ongoing replication with AWS DMS.

**Backup**

☒ **Enable automated backups**  
Creates a point-in-time snapshot of your database

**Backup retention period** [Info](#)  
Choose the number of days that RDS should retain automatic backups for this instance.

7 days ▼

**Backup window** [Info](#)  
Select the period for which you want automated backups of the database to be created by Amazon RDS.

☒ **Select window**

☐ No preference

**Start time** : 09 ▼ : 00 ▼ UTC

**Duration** : 0.5 ▼ hours

☒ **Copy tags to snapshots**

To perform the full load phase, AWS DMS requires read privileges to the tables in scope for migration. For more information about required permissions, see [Permissions for full load only tasks](#).

Connect to the Amazon RDS for SQL Server instance and run the following queries. Use a login with master user privileges for both full load and CDC.

```
USE AdventureWorks;
CREATE LOGIN dms_user WITH PASSWORD = 'password'
CREATE USER dms_user FOR LOGIN dms_user
ALTER ROLE [db_datareader] ADD MEMBER dms_user
ALTER ROLE [db_owner] ADD MEMBER dms_user
GRANT VIEW DATABASE STATE to dms_user

USE master;
GRANT VIEW SERVER STATE TO dms_user
```

**Note**

Here, we create a new user to perform the migration. You can skip this step if you plan to use existing logins and users that have the required privileges.

Turn on MS-CDC for your Amazon RDS for SQL Server database instance at the database level.

```
exec msdb.dbo.rds_cdc_enable_db 'AdventureWorks'
```

Because we migrate all tables in the Sales schema of the AdventureWorks database, we need to identify the total number of tables.

```
SELECT TABLE_SCHEMA, TABLE_NAME, TABLE_TYPE
FROM information_schema.tables
WHERE TABLE_SCHEMA = 'Sales'
ORDER BY TABLE_NAME
```

Then we need to divide tables in the following groups:

- Tables with a primary key.
- Tables with a unique index without primary key.
- Tables without a primary key and unique index.

We use the information\_schema to identify tables that have a primary key or a unique index without a primary key.

```
SELECT a.TABLE_SCHEMA, a.TABLE_NAME, a.CONSTRAINT_TYPE, CONSTRAINT_NAME
FROM information_schema.table_constraints a
JOIN information_schema.tables b ON a.TABLE_SCHEMA = b.TABLE_SCHEMA
AND a.TABLE_NAME = b.TABLE_NAME
WHERE b.TABLE_TYPE = 'BASE TABLE'
AND a.TABLE_SCHEMA = 'Sales'
AND a.CONSTRAINT_TYPE in ('UNIQUE', 'PRIMARY KEY')
ORDER BY a.TABLE_SCHEMA, a.TABLE_NAME
```

The query results show that the task has 19 tables and all of them have primary keys. For all these tables, run the following query to turn on MS-CDC at the table level.

```
exec sys.sp_cdc_enable_table
@source_schema = N'Sales',
@source_name = N'table_name',
@role_name = NULL,
@supports_net_changes = 1
```

Now, set the retention period for changes to be available on the source using the following commands. Set the `pollinginterval` value to 86399 seconds to increase the retention of changes on the Amazon RDS for SQL Server instance.

```
EXEC sys.sp_cdc_change_job @job_type = 'capture', @pollinginterval = 86399
exec sys.sp_cdc_stop_job @job_type = 'capture'
exec sys.sp_cdc_start_job @job_type = 'capture'
exec sys.sp_cdc_help_jobs
```

Set the polling interval on your secondary database to 86399 seconds too. For most use cases these settings should be enough. For databases that have a large number of transactions, you need to make additional configuration changes to make sure that the transaction log has optimal retention. For more information, see [Optional settings when using Amazon RDS for SQL Server as a source](#).

For more information about ongoing replication, see [Setting up ongoing replication on a Cloud SQL Server DB instance](#).

#### Note

AWS DMS does not support replicating ongoing changes from views. For more information, see [Selection rules and actions](#).

In this walkthrough, we focus on migrating the tables and do not include views in the migration scope. You should also look at estimating the number of records in the tables you are going to migrate as this is a useful consideration while configuring AWS DMS tasks.

## Step 3: Create an AWS DMS Source Endpoint

After you configured the AWS Database Migration Service (AWS DMS) replication instance and the source Amazon RDS for SQL Server instance, ensure connectivity between these two instances. To ensure that the replication instance can access the server and the port for the database, make

changes to the relevant security groups and network access control lists. For more information about your network configuration, see [Setting up a network for a replication instance](#).

After you completed the network configurations, you can create a source endpoint.

To create a source endpoint, do the following:

1. Open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose **Endpoints**.
3. Choose **Create endpoint**.
4. On the **Create endpoint** page, enter the following information.

| For This Parameter                    | Do This                                                                                                                          |
|---------------------------------------|----------------------------------------------------------------------------------------------------------------------------------|
| <b>Endpoint type</b>                  | Choose <b>Source endpoint</b> , turn on <b>Select RDS DB instance</b> , and choose <code>datalake-source-db</code> RDS instance. |
| <b>Endpoint identifier</b>            | Enter <code>datalake-source-db</code> .                                                                                          |
| <b>Source engine</b>                  | Choose <b>Microsoft SQL Server</b> .                                                                                             |
| <b>Access to endpoint database</b>    | Choose <b>Provide access information manually</b> .                                                                              |
| <b>Server name</b>                    | Enter the database server name on Amazon RDS.                                                                                    |
| <b>Port</b>                           | Enter <b>1433</b> .                                                                                                              |
| <b>Secure Socket Layer (SSL) mode</b> | Choose <b>none</b> .                                                                                                             |
| <b>User name</b>                      | Enter <code>dms_user</code> .                                                                                                    |
| <b>Password</b>                       | Enter the password that you created for the <code>dms_user</code> user.                                                          |
| <b>Database name</b>                  | Enter <b>AdventureWorks</b> .                                                                                                    |

5. Choose **Create endpoint**.



**Note**

To migrate a Microsoft SQL Server Always On database, you need to use different configurations. For more information, see [the section called “Migrating a SQL Server Always On Database to Amazon Web Services”](#).

## Step 4: Configure a Target Amazon S3 Bucket

You can integrate Amazon S3 with other AWS and third-party services to take advantage of the following:

- Data analysis using Amazon Athena query engine. This service helps reduce cost as you do not pay for dedicated resources and instead pay based on the amount data being scanned.
- Perform extract, transform, and load (ETL) operations using distributed processing frameworks such as Spark with Amazon EMR or AWS Glue.
- Implement machine learning use cases, because Amazon S3 can store granular time series data spanning years in raw form, in conjunction with Amazon SageMaker.

Because in this use case we migrate the Sales schema to Amazon S3, we need to account for future use cases of the migrated data before we set up Amazon S3 bucket and AWS DMS endpoints.

To create the Amazon S3 bucket, do the following:

1. Open the Amazon S3 console at <https://s3.console.aws.amazon.com/s3/home>.
2. Choose **Create bucket**.
3. For **Bucket name**, enter **adventure-works-datalake**.
4. For **AWS Region**, choose the region that hosts your AWS DMS replication instance.
5. Leave the default values in the other fields and choose **Create bucket**.

## Step 5: Configure an AWS DMS Target Endpoint

To use Amazon S3 as an AWS Database Migration Service (AWS DMS) target endpoint, create an IAM role with write and delete access to the S3 bucket. Then add DMS (dms.amazonaws.com) as

*trusted entity* in this IAM role. For more information, see [Prerequisites for using Amazon S3 as a target](#).

When using AWS DMS to migrate data to an Amazon Simple Storage Service (Amazon S3) data lake, you can change the default task behavior, such as file formats, partitioning, file sizing, and so on. This leads to minimizing post-migration processing and helps downstream applications consume data efficiently. You can customize task behavior using endpoint settings and extra connection attributes (ECA). Most of the Amazon S3 endpoint settings and ECA settings overlap, except for a few parameters. In this walkthrough, we will configure Amazon S3 endpoint settings.

## Choose File Format

AWS DMS supports data replication through comma-separated values (CSV) or Apache Parquet file formats. Each file format has its own benefits. Choose the right file format depending on your consumption pattern.

Apache Parquet is a columnar format, which is built to support efficient compression and encoding schemes providing storage space savings and performance benefits. With Parquet, you can specify compression schemes for each column to improve query performance when using `avg()`, `max()`, or other column level aggregation operations. That is why Parquet is popular for data lake and analytics use cases.

CSV files are helpful when you plan to keep data in human readable format, share or transfer Amazon S3 files into other downstream systems for further processing.

For this walkthrough, we will use the Parquet file format. Specify the following endpoint settings.

```
DataFormat=parquet
ParquetVersion=PARQUET_2_0
```

## Determine File Size

By default, during ongoing replication AWS DMS tasks writes to Amazon S3 are triggered either if the file size reaches 32 KB or if the previous file write was more than 60 seconds ago. These settings ensure that the data capture latency is less than a minute. However, this approach creates a large number of small files in target Amazon S3 bucket.

Because we migrate our source `Sales` database schema for an analytics use case, some latency is acceptable. However, we need to optimize this schema for cost and performance. When you

use distributed processing frameworks such as Amazon Athena, AWS Glue or Amazon EMR, it is recommended to avoid too many small files (less than 64 MB). Small files create management overhead for the driver node of the distributed processing framework.

Because we plan to use Amazon Athena to query data from our Amazon S3 bucket, we need to make sure our target file size is at least 64 MB. Specify the following endpoint settings: `CdcMaxBatchInterval=3600` and `CdcMinFileSize=64000`. These settings ensure that AWS DMS writes the file until its size reaches 64 MB or if the last file write was more than an hour ago.

### Note

Parquet files created by AWS DMS are usually smaller than the specified `CdcMinFileSize` setting because Parquet data compression ratio varies depending on the source data set. The size of CSV files created by AWS DMS is equal to the value specified in `CdcMinFileSize`.

## Turn on S3 Partitioning

Partitioning in Amazon S3 structures your data by folders and subfolders that help efficiently query data. For example, if you receive sales record data daily from different regions and you query data for a specific region and find stats for a few months, then it is recommended to partition data by region, year, and month. In Amazon S3, the path for our use case looks as following:

```
s3://<sales-data-bucket-name>/<region>/<schemaname>/<tablename>/<year>/<month>/<day>

s3://adventure-works-datalake
- s3://adventure-works-datalake/US-WEST-DATA
- s3://adventure-works-datalake/US-WEST-DATA/Sales
  - s3://adventure-works-datalake/US-WEST-DATA/Sales/CreditCard/
    - s3://adventure-works-datalake/US-WEST-DATA/Sales/CreditCard/
      LOAD00000001.parquet
  - s3://adventure-works-datalake/US-WEST-DATA/Sales/SalesPerson
    - s3://adventure-works-datalake/US-WEST-DATA/Sales/SalesPerson/
      LOAD00000001.parquet
      - s3://adventure-works-datalake/US-WEST-DATA/Sales/SalesPerson/2021/11/23/
        - s3://adventure-works-datalake/US-WEST-DATA/Sales/
          SalesPerson/2021/11/23/20211123-013830913.parquet
        - s3://adventure-works-datalake/US-WEST-DATA/Sales/
          SalesPerson/2021/11/27/20211127-175902985.parquet
```

Partitioning provides performance benefits because data scanning will be limited to the amount of data in the specific partition based on the filter condition in your queries. For our sales data example, your queries might look as follows:

```
SELECT <column-list> FROM <sales-table-name> WHERE <region> = <region-name> AND <year>
= <year-value>
```

If you use Amazon Athena to query data, partitioning helps reduce cost as Athena pricing is based on the amount of data that you scan when running queries.

To turn on partitioning for ongoing changes in the preceding format, use the following queries.

```
bucketFolder=US-WEST-DATA
DatePartitionedEnabled=true
DatePartitionSequence=YYYYMMDD
DatePartitionDelimiter=SLASH
```

## Other Considerations

The preceding settings help optimize performance and cost. We also need to configure additional settings because:

- Our use case does not have a fixed end-date.
- We need to minimize issues arising from misconfigurations or retroactive changes.
- We want to minimize recovery time in case of unforeseen issues.

## Serialize ongoing replication events

A common challenge when using Amazon S3 as a target involves identifying the ongoing replication event sequence when multiple records are updated at the same time on the source database.

AWS DMS provides two options to help serialize such events for Amazon S3. You can use the `TimeStampColumnName` endpoint setting or use transformation rules to include LSN column. Here, we will discuss the first option. For more information about the second option, see [the section called “Step 6: Create an AWS DMS Task”](#).

## Use the `TimeStampColumnName` endpoint setting

The `TimeStampColumnName` setting adds an additional `STRING` column to the target Parquet file created by AWS DMS. During the ongoing replication, the column value represents the commit timestamp of the event in SQL Server. For the full load phase, the columns values represent the timestamp of data transfer to Amazon S3.

The default format is `yyyy-MM-dd HH:mm:ss.SSSSSS`. This format provides a microsecond precision but depends on the source database transaction log timestamp precision. The following image shows the seven microseconds difference between two operations in the `sourceRecordTime` field.

```
{
  "Op": "U",
  "sourceRecordTime": "2021-11-29 02:43:02.990000",
  "BusinessEntityID": 274,
  "Bonus": 1,
  "CommissionPct": 0.1,
  "SalesYTD": 559697.5639,
  "SalesLastYear": 0,
  "rowguid": "48754992-9EE0-4C0E-8C94-9451604E3E02",
  "ModifiedDate": "2010-12-28T00:00:00.000Z",
  "transact-id": "0000006c:00018542:0003"
}
{
  "Op": "U",
  "sourceRecordTime": "2021-11-29 02:43:02.997000",
  "BusinessEntityID": 275,
  "TerritoryID": 2,
  "Bonus": 1,
  "CommissionPct": 0.1,
  "SalesYTD": 3763178.1787,
  "SalesLastYear": 1750406.4785,
  "rowguid": "1E0A7274-3064-4F58-88EE-4C6586C87169",
  "ModifiedDate": "2011-05-24T00:00:00.000Z",
  "transact-id": "0000006c:00018544:0002"
}
```

#### Note

Because `TimeStampColumnName` is an endpoint setting, all tasks that use this endpoint, will include this column for all tables.

## Include full load operation field

All files created during the ongoing replication, have the first column marked with I, U, or D. These symbols represent the DML operation on the source and stand for **Insert**, **Update**, or **Delete** operations.

For full load files, you can add this column by configuring the endpoint setting.

```
includeOpForFullLoad=true
```

This ensures that all full load files are marked with an I operation.

When you use this approach, new subscribers can consume the entire data set or prepare a fresh copy in case of any downstream processing issues.

## Create a Target Endpoint

After you completed all settings configurations, you can create a target endpoint.

To create a target endpoint, do the following:

1. Open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose **Endpoints**, and then choose **Create endpoint**.
3. On the **Create endpoint** page, enter the following information.

| For This Parameter      | Do This                                                                      |
|-------------------------|------------------------------------------------------------------------------|
| Endpoint type           | Choose <b>Target endpoint</b> , and turn off <b>Select RDS DB instance</b> . |
| Endpoint identifier     | Enter <b>adventure-works-datalake-target</b> .                               |
| Target engine           | Choose <b>Amazon S3</b> .                                                    |
| Service access role ARN | Enter the IAM role that can access your Amazon S3 data lake.                 |
| Bucket name             | Enter <b>adventure-works-datalake</b> .                                      |
| Bucket folder           | Enter <b>US-WEST-DATA</b> .                                                  |

4. Expand the **Endpoint settings** section, choose **Wizard**, and then choose **Add new setting** to add the settings as shown on the following image.

| Endpoint settings                                     |                                                 |                         |
|-------------------------------------------------------|-------------------------------------------------|-------------------------|
| Setting                                               | Value - A value is required                     |                         |
| <input type="text" value="DataFormat"/> X             | <input type="text" value="parquet"/> X          | <button>Remove</button> |
| <input type="text" value="ParquetVersion"/> X         | <input type="text" value="PARQUET_2_0"/> X      | <button>Remove</button> |
| <input type="text" value="CdcMinFileSize"/> X         | <input type="text" value="64000"/> X            | <button>Remove</button> |
| <input type="text" value="CdcMaxBatchInterval"/> X    | <input type="text" value="3600"/> X             | <button>Remove</button> |
| <input type="text" value="DatePartitionEnabled"/> X   | <input type="text" value="true"/> X             | <button>Remove</button> |
| <input type="text" value="DatePartitionSequence"/> X  | <input type="text" value="YYYYMMDD"/> X         | <button>Remove</button> |
| <input type="text" value="DatePartitionDelimiter"/> X | <input type="text" value="SLASH"/> X            | <button>Remove</button> |
| <input type="text" value="IncludeOpForFullLoad"/> X   | <input type="text" value="true"/> X             | <button>Remove</button> |
| <input type="text" value="TimestampColumnName"/> X    | <input type="text" value="sourceRecordTime"/> X | <button>Remove</button> |

5. Choose **Create endpoint**.

## Step 6: Create an AWS DMS Task

After you configured the replication instance and endpoints, you need to analyze your source database. A good understanding of the workload helps plan an effective migration approach and minimize configuration issues. Find some of the important considerations following and learn how they apply to our walkthrough.

### Size and number of records

The volume of migrated records affects the full load completion time. It is difficult to predict the full load time upfront, but testing with a replica of a production instance should provide a baseline. Use this estimate to decide whether you should parallelize full load by using multiple tasks or by using the parallel load option.

The Sales schema includes 19 tables. The CreditCard table is the largest table containing 100,000 records. We can increase the number of tables loaded in parallel to 19 if the full load is slow. The default value for the number of tables loaded in parallel is eight.

## Transactions per second

While full load is affected by the number of records, the ongoing replication performance relies on the number of transactions on the source Amazon RDS for SQL Server database. Performance issues during change data capture (CDC) generally stem from resource constraints on the source database, replication instance, target database, and network bandwidth or throughput. Knowing average and peak TPS on the source and recording CDC throughput and latency metrics help baseline (AWS DMS) performance and identify an optimal task configuration. For more information, see [Replication task metrics](#).

In this walkthrough, we will track the CDC latency and throughput values after the task moves into the ongoing replication phase to baseline AWS DMS performance.

## LOB columns

AWS DMS handles large binary objects (LOBs) columns differently compared to other data types. For more information, see [Migrating large binary objects \(LOBs\)](#).

Because AWS DMS does not support **Full LOB mode** for Amazon S3 endpoints, we need to identify a suitable **LOB Max Size**.

A detailed explanation of LOB handling by AWS DMS is out of scope for this walkthrough. However, remember that increasing the **LOB Max Size** increases the tasks memory utilization. Because of that, it is recommended not to set **LOB Max Size** to a large value.

For more information about LOB settings, see [the section called "Task Configuration"](#).

## Unsupported data types

Identify data types used in tables and check that AWS DMS supports these data types. For more information, see [Source data types for SQL Server](#).

Validate that the target Amazon S3 has the corresponding data types. For more information, see [Target data types for S3 Parquet](#).

After running the initial load test, validate that AWS DMS converted data as you expected. You can also initiate a pre-migration assessment to identify any unsupported data types in the migration scope. For more information, see [Specifying individual assessments](#).



**Note**

The preceding list is not complete. For more information, see [Best practices](#).

Combining the considerations from the preceding list, we start with a single task that migrates all 19 tables. Based on the full load run time and resource utilization metrics on the source SQL Server database instance and replication instance, we can evaluate if we should parallelize the load further to improve performance.

## Task Configuration

In an AWS DMS task, you can specify the schema or table to migrate, the type of migration, and the configurations for the migration. You can choose one of the following options for your task.

- **Full Load only** — migrate existing data.
- **Full Load + CDC** — migrate existing data and replicate ongoing changes.
- **CDC only** — replicate ongoing changes.

For more information about the task creation steps and available configuration options, see [Creating a task](#).

In this walkthrough, we will focus on the following settings.

## Table mappings

Use selection rules to define the schemas and tables that the AWS DMS task will migrate. For more information, see [Selection rules and actions](#).

Because we need to identify monthly sales for a specific year, one possible approach can restrict the migration to `SalesOrder1%` tables in the `Sales` schema and keep adding new tables to the task when additional reporting is required. This approach saves cost and minimizes the load, but increases operational overhead by requiring repeated configurations, performance baselining, and so on. For the walkthrough, we will migrate all tables (%) from the `Sales` schema.

## Using transformation rules to include LSN column

In the previous section we discussed using the `TimestampColumnName` endpoint setting to serialize ongoing replication events. For more information about using the

TimestampColumnName endpoint setting, see [the section called “Serialize ongoing replication events”](#).

Because the source database transaction log precision is limited to milliseconds, multiple events can have the same timestamp. To address this issue, you can use task level transformation rules to include source table headers to the Amazon S3 target files as described in the task creation section.

Source table headers add an additional column that contains the log sequence number (LSN) value of the operation from the source SQL Server database instance. You can use this information in our Amazon S3 data lake scenario for downstream serialization. For more information about source table headers, see [Replicating source table headers using expressions](#).

To include headers, add the following transformation rule in the JSON editor in table mapping. This rule adds a new transact-id column with the LSN to all tables that the task migrates. For more information, see [Specifying table selection and transformations rules using JSON](#).

```
{
  "rule-type": "transformation",
  "rule-id": "2",
  "rule-name": "2",
  "rule-target": "column",
  "object-locator": {
    "schema-name": "%",
    "table-name": "%"
  },
  "rule-action": "add-column",
  "value": "transact_id",
  "expression": "$AR_H_STREAM_POSITION",
  "data-type": {
    "type": "string",
    "length": 50
  }
}
```

#### Note

Mapping rules are applied at the task level. You need to add a mapping rule to each task that replicates data to your data lake.

## LOB settings

Use the **sys** schema to identify the LOB columns in the tables of the Sales schema.

```
SELECT s.name AS SchemaName,  
t.name AS TableName,  
c.name AS ColumnName,  
y.name AS DataType  
FROM sys.tables AS t  
INNER JOIN sys.schemas AS s ON s.schema_id = t.schema_id  
INNER JOIN sys.columns AS c ON t.object_id = c.object_id  
INNER JOIN sys.types AS y ON y.user_type_id = c.user_type_id  
WHERE (c.user_type_id in (34,35,99,129,130,241,256) OR (c.user_type_id in (165,167,231)  
AND c.max_length = -1))  
AND s.name = 'Sales'  
ORDER BY t.name;
```

The Sales.Store table includes one LOB column. Use the following query to identify the size of the largest LOB in the migrated tables.

```
select max(datalength(Demographics)) as "Size in Bytes" from Sales.Store
```

The size of the largest LOB is 1,000 bytes. Because of that, we will leave the default value for LOB Max Size, which is 32 KB. If the size of the largest LOB is more than 32 KB, it is recommended to factor in LOB growth over time, include some buffer, and set that as the LOB Max Size value.

## Other task settings

Choose **Enable CloudWatch Logs** to upload the AWS DMS task execution log to Amazon CloudWatch. You can use these logs to troubleshoot issues because they include error and warning messages, start and end times of the run, configuration issues, and so on. Changes to the task logging setting, such as enabling debug or trace can also be helpful to diagnose performance issues.

### Note

CloudWatch log usage is charged at standard rates. For more information, see [Amazon CloudWatch pricing](#).

For **Target table preparation mode**, choose one of the following options: Do nothing, truncate, and Drop. Use Truncate in data pipelines where the downstream systems rely on a

fresh dump of clean data and do not rely on historical data. In this walkthrough, we choose **Do nothing** because we want to control the retention of files from previous runs.

For **Maximum number of tables to load in parallel**, enter the number of parallel threads that AWS DMS initiates during full load. You can increase this value to improve the full load performance and minimize the load time when you have numerous tables.

 **Note**

Increasing this parameter induces additional load on the source database, replication instance, and target database.

## Create an AWS DMS Task

After you completed all settings configurations, you can create an AWS DMS database migration task.

To create a database migration task, do the following:

1. Open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose **Database migration tasks**, and then choose **Create task**.
3. On the **Create database migration task** page, enter the following information.

| For This Parameter       | Do This                                                                           |
|--------------------------|-----------------------------------------------------------------------------------|
| Task identifier          | Enter <b>AdventureWorks-to-S3-data-lake</b> .                                     |
| Replication instance     | Choose <b>datalake-migration-ri</b> (the value that you configured on Step 1).    |
| Source database endpoint | Choose <b>datalake-source-db</b> (the value that you configured on Step 3).       |
| Target database endpoint | Choose <b>adventure-works-datalake</b> (the value that you configured on Step 4). |
| Migration type           | Choose <b>Migrate existing data and replicate ongoing changes</b> .               |

| For This Parameter                           | Do This                                                 |
|----------------------------------------------|---------------------------------------------------------|
| Editing mode                                 | Choose <b>Wizard</b> .                                  |
| Custom CDC stop mode for source transactions | Choose <b>Disable custom CDC stop mode</b> .            |
| Target table preparation mode                | Choose <b>Do nothing</b> .                              |
| Stop task after full load completes          | Choose <b>Don't stop</b> .                              |
| Include LOB columns in replication           | Choose <b>Limited LOB mode</b> .                        |
| Maximum LOB size (KB)                        | Enter <b>32</b> .                                       |
| Enable validation                            | Turn off because Amazon S3 does not support validation. |
| Enable CloudWatch logs                       | Turn on.                                                |

4. Leave the default values in the other fields and choose **Create task**.
5. The task begins immediately. The **Database migration tasks** section shows you the status of the migration task.

## Step 7: Run the AWS DMS Task

After you created your AWS Database Migration Service (AWS DMS) task, run the task a few times to identify the full load run time and ongoing replication performance. You can validate that initial configurations work as expected. You can do this by monitoring and documenting resource utilization on the source database, replication instance, and target database. These details make up the initial baseline and help determine if you need further optimizations.

After you started the task, the full load operation starts loading tables. You can see the table load completion status in the **Table Statistics** section and the corresponding target files in the Amazon S3 bucket. Because in our case the overall number of records is less than 200,000, the full load operation finishes in less than a minute. We can increase the value of **Maximum number of tables to load in parallel**, but it will not provide any meaningful gain in this scenario.

After the AWS DMS task completes full load, the status changes to the **Load complete, replication ongoing** phase. The following image shows the updated status of the task.

DMS > Database migration tasks > adventureworks-to-s3-datalake

## adventureworks-to-s3-datalake

### Summary

|                                                   |                                |
|---------------------------------------------------|--------------------------------|
| Status                                            | Type                           |
| <span>🟢</span> Load complete, replication ongoing | Full load, ongoing replication |

During this phase, AWS DMS partitions data by the year, month, and day of generation. The following image shows the structure of folders.

Amazon S3 > adventure-works-datalake > US-WEST-DATA/ > Sales/ > SalesPerson/ > 2021/ > 11/ > 29/

## 29/

**Objects** | Properties

### Objects (9)

Objects are the fundamental entities stored in Amazon S3. You can use [Amazon S3 inventory](#) to get a list of all objects in your bucket.

Find objects by prefix

| <input type="checkbox"/> | Name                       | Type    | Last modified                           | Size    |
|--------------------------|----------------------------|---------|-----------------------------------------|---------|
| <input type="checkbox"/> | 20211129-024736004.parquet | parquet | November 28, 2021, 18:47:37 (UTC-08:00) | 50.6 MB |
| <input type="checkbox"/> | 20211129-025216249.parquet | parquet | November 28, 2021, 18:52:17 (UTC-08:00) | 50.5 MB |
| <input type="checkbox"/> | 20211129-030005428.parquet | parquet | November 28, 2021, 19:00:06 (UTC-08:00) | 43.6 MB |
| <input type="checkbox"/> | 20211129-032023480.parquet | parquet | November 28, 2021, 19:20:24 (UTC-08:00) | 50.5 MB |

Following, find some of the common errors and unexpected results you might see while following this walkthrough.

**Files are not written to the Amazon S3 target even though changes are visible in the table statistics section of the c console**

- This happens due to the target endpoint configuration. After you set `CdcMaxBatchInterval=3600` and `CdcMinFileSize=64000`, AWS DMS waits for an hour or for the file size to reach 64 MB before writing data to Amazon S3.
- To write the output to Amazon S3 sooner, reduce `CdcMaxBatchInterval` to a smaller value. Alternatively, you can stop and resume the task. This will force Amazon S3 to flush events to Amazon S3 disregarding the extra connection attributes settings. Using these options means that the size of CDC files will be much smaller than the expected 64 MB.

### Parquet file sizes are less than 64 MB despite setting `CdcMinFileSize=64000`

AWS DMS creates 64 MB files in memory. When this data is encoded as Parquet the resulting file size is smaller. The file sizes vary based on the level of compression possible.

### AWS DMS captures only inserts and deletes and does not migrate update records to the target

You can see the following warning in the task logs:

```
00008570: 2021-12-07T19:52:52 [SOURCE_CAPTURE ]W: MS-REPLICATION is not enabled for
table '[Sales].[SalesPerson]'. Therefore, UPDATE changes to it will not be captured.
If you want UPDATE changes to be captured, either define a Primary Key for the table
(if missing) or enable Microsoft CDC instead. (sqlserver_log_utils.c:1292)
```

This log message indicates MS-Replication. However, for Amazon RDS for SQL Server you can use MS-CDC. This error occurs when you have not turned on MS-CDC for the table. For more information, see [the section called “Step 2: Configure a Source Amazon RDS for SQL Server Database”](#).

In this walkthrough, we covered most prerequisites that help avoid configuration related errors. If you observe issues when running the task, see [Troubleshooting migration tasks](#), [Best practices](#), or reach out to AWS Support for further assistance.

After you completed the migration, validate that your data migrated successfully and delete the cloud resources that you created.

## Migrating an Oracle Database to PostgreSQL

Using this walkthrough, you can learn how to migrate an Oracle database to a PostgreSQL database using AWS Database Migration Service (AWS DMS) and the AWS Schema Conversion Tool (AWS SCT).

AWS DMS migrates your data from your Oracle source into your PostgreSQL target. AWS DMS also captures data manipulation language (DML) and [supported data definition language \(DDL\)](#) changes that happen on your source database and applies these changes to your target database. This way, AWS DMS keeps your source and target databases in sync with each other. To facilitate the data migration, AWS SCT creates the migrated schemas on the target database, including the tables and primary key indexes on the target if necessary.

AWS DMS doesn't migrate your secondary indexes, sequences, default values, stored procedures, triggers, synonyms, views, and other schema objects not specifically related to data migration. To migrate these objects to your PostgreSQL target, use AWS SCT.

## Topics

- [Prerequisites for migrating from an Oracle database to PostgreSQL](#)
- [Step-by-step Oracle database to PostgreSQL migration walkthrough](#)
- [Rolling Back the Migration](#)
- [Oracle database migration to PostgreSQL troubleshooting](#)

## Prerequisites for migrating from an Oracle database to PostgreSQL

The following prerequisites are required to complete this walkthrough:

- Understand Amazon Relational Database Service (Amazon RDS), the applicable database technologies, and SQL.
- Create a user with AWS Identity and Access Management (IAM) credentials that allows you to launch Amazon RDS and AWS Database Migration Service (AWS DMS) instances in your AWS Region. For information about IAM credentials, see [Setting up for Amazon RDS](#).
- Understand the Amazon Virtual Private Cloud (Amazon VPC) service and security groups. For information about using Amazon VPC with Amazon RDS, see [Amazon Virtual Private Cloud \(VPCs\) and Amazon RDS](#). For information about Amazon RDS security groups, see [Amazon RDS Security Groups](#).
- Understand the supported features and limitations of AWS DMS. For information about AWS DMS, see <https://docs.aws.amazon.com/dms/latest/userguide/Welcome.html>.
- Understand the supported data type conversion options for Oracle and PostgreSQL. For information about data types for Oracle as a source, see [Using an Oracle database as a source](#). For information about data types for PostgreSQL as a target, see [Using a PostgreSQL Database as a Target](#).



- Size your target PostgreSQL database host. DBAs should be aware of the load profile of the current source Oracle database host. Consider CPU, memory, and IOPS. With RDS, you can size up the target database host, or reduce it, after the migration. If this is the first time you are migrating to PostgreSQL, then we recommend that you have extra capacity to account for performance issues and tuning opportunities.
- Audit your source Oracle database. For each schema and all the objects under each schema, determine if any of the objects are no longer being used. Deprecate these objects on the source Oracle database, because there's no need to migrate them if they are not being used.
- If load capacity permits, then get the max size (kb) for each LOB type on the source database, and keep this information for later.
- If possible, move columns with BLOB, CLOB, NCLOB, LONG, LONG RAW, and XMLTYPE to Amazon S3, Dynamo DB, or another data store. Doing so simplifies your source Oracle database for an easier migration. It will also lower the capacity requirements for the target PostgreSQL database.

For more information about AWS DMS, see [Getting started with Database Migration Service](#).

To estimate what it will cost to run this walkthrough on AWS, you can use the AWS Pricing Calculator. For more information, see <https://calculator.aws/>.

To avoid additional charges, delete all resources after you complete the walkthrough.

## Step-by-step Oracle database to PostgreSQL migration walkthrough

The following steps provide instructions for migrating an Oracle database to a PostgreSQL database. These steps assume that you have already prepared your source database as described in [Prerequisites](#).

### Topics


- [Step 1: Install the SQL Drivers and AWS Schema Conversion Tool on Your Local Computer](#)
- [Step 2: Configure Your Oracle Source Database](#)
- [Step 3: Configure Your PostgreSQL Target Database](#)
- [Step 4: Use AWS SCT to Convert the Oracle Schema to PostgreSQL](#)
- [Step 5: Create an AWS DMS Replication Instance](#)
- [Step 6: Create AWS DMS Source and Target Endpoints](#)

- [Step 7: Create and Run Your AWS DMS Migration Task](#)
- [Step 8: Cut Over to PostgreSQL](#)

## Step 1: Install the SQL Drivers and AWS Schema Conversion Tool on Your Local Computer

To install the SQL drivers and the AWS Schema Conversion Tool (AWS SCT) on your local computer, do the following:

1. Download the JDBC driver for your Oracle database release. For more information, go to <https://www.oracle.com/jdbc>.
2. Download the PostgreSQL driver ([postgresql-42.2.19.jar](#)).
3. Install AWS SCT and the required JDBC drivers.
  - a. Download AWS SCT from [Installing, verifying, and updating the Schema Conversion Tool](#).
  - b. Launch AWS SCT.
  - c. In AWS SCT, choose **Global settings** from **Settings**.
  - d. In **Global settings**, choose **Driver**, and then choose **Browse** for **Oracle driver path**. Locate the JDBC Oracle driver and choose **OK**.
  - e. Choose **Browse** for **PostgreSQL driver path**. Locate the JDBC PostgreSQL driver and choose **OK**.



f. Choose **OK** to close the dialog box.

## Step 2: Configure Your Oracle Source Database

To use Oracle as a source for AWS Database Migration Service (AWS DMS), you must first ensure that ARCHIVELOG MODE is on to provide information to LogMiner. AWS DMS uses LogMiner to read information from the archive logs so that AWS DMS can capture changes.

For AWS DMS to read this information, make sure the archive logs are retained on the database server as long as AWS DMS requires them. If you configure your task to begin capturing changes immediately, then you should only need to retain archive logs for a little longer than the duration of the longest running transaction. Retaining archive logs for 24 hours is usually sufficient. If you configure your task to begin from a point in time in the past, then archive logs must be available from that time forward. For more specific instructions about enabling ARCHIVELOG MODE and ensuring log retention for your Oracle database, see the [Oracle documentation](#).

To capture change data, AWS DMS requires supplemental logging to be enabled on your source database. Minimal supplemental logging must be enabled at the database level. AWS DMS also requires that identification key logging be enabled. This option causes the database to place all

columns of a row's primary key in the redo log file whenever a row containing a primary key is updated. This result occurs even if no value in the primary key has changed. You can set this option at the database or table level.

1. Create or configure a database account to be used by AWS DMS. We recommend that you use an account with the minimal privileges required by AWS DMS for your AWS DMS connection. AWS DMS requires the following privileges.

```
CREATE SESSION
SELECT ANY TRANSACTION
SELECT on V_$ARCHIVED_LOG
SELECT on V_$LOG
SELECT on V_$LOGFILE
SELECT on V_$DATABASE
SELECT on V_$THREAD
SELECT on V_$PARAMETER
SELECT on V_$NLS_PARAMETERS
SELECT on V_$TIMEZONE_NAMES
SELECT on V_$TRANSACTION
SELECT on ALL_INDEXES
SELECT on ALL_OBJECTS
SELECT on ALL_TABLES
SELECT on ALL_USERS
SELECT on ALL_CATALOG
SELECT on ALL_CONSTRAINTS
SELECT on ALL_CONS_COLUMNS
SELECT on ALL_TAB_COLS
SELECT on ALL_IND_COLUMNS
SELECT on ALL_LOG_GROUPS
SELECT on SYS.DBA_REGISTRY
SELECT on SYS.OBJ$
SELECT on DBA_TABLESPACES
SELECT on ALL_TAB_PARTITIONS
SELECT on ALL_ENCRYPTED_COLUMNS
* SELECT on all tables migrated
```

If you want to capture and apply changes (CDC), then you also need the following privileges.

```
EXECUTE on DBMS_LOGMNR
EXECUTE on DBMS_LOGMNR_D
SELECT on V_$LOGMNR_LOGS
SELECT on V_$LOGMNR_CONTENTS
```

```
LOGMINING /* For Oracle 12c and higher. */  
* ALTER for any table being replicated (if you want to add supplemental logging)
```

For Oracle versions before 11.2.0.3, you need the following privileges.

```
SELECT on DBA_OBJECTS /* versions before 11.2.0.3 */  
SELECT on ALL_VIEWS (required if views are exposed)
```

2. If your Oracle database is an Amazon RDS database, then connect to it as an administrative user, and run the following command to ensure that archive logs are retained on your RDS source for 24 hours:

```
exec rdsadmin.rdsadmin_util.set_configuration('archive_log_retention_hours',24);
```

If your Oracle source is an Amazon RDS database, it will be placed in ARCHIVELOG MODE if, and only if, you enable backups.

3. Run the following command to turn on supplemental logging at the database level, which AWS DMS requires:

- In Oracle SQL:

```
ALTER DATABASE ADD SUPPLEMENTAL LOG DATA;
```

- In RDS:

```
exec rdsadmin.rdsadmin_util.alter_supplemental_logging('ADD');
```

4. Use the following command to enable identification key supplemental logging at the database level. AWS DMS requires supplemental key logging at the database level. The exception is if you allow AWS DMS to automatically add supplemental logging as needed or enable key-level supplemental logging at the table level:

- In Oracle SQL:

```
ALTER DATABASE ADD SUPPLEMENTAL LOG DATA (PRIMARY KEY) COLUMNS;
```

- In RDS:

```
exec rdsadmin.rdsadmin_util.alter_supplemental_logging('ADD','PRIMARY KEY');
```

Your source database incurs a small bit of overhead when key level supplemental logging is enabled. Therefore, if you are migrating only a subset of your tables, then you might want to enable key level supplemental logging at the table level.

5. To turn on key level supplemental logging at the table level, use the following command.

```
ALTER TABLE table_name ADD SUPPLEMENTAL LOG DATA (PRIMARY KEY) COLUMNS;
```

If a table doesn't have a primary key, then you have two options.

- You can add supplemental logging on all columns involved in the first unique index on the table (sorted by index name).

To add supplemental logging on a subset of columns in a table, such as those involved in a unique index, run the following command.

```
ALTER TABLE table_name  
  ADD SUPPLEMENTAL LOG GROUP example_log_group (column_list) ALWAYS;
```

- You can add supplemental logging on all columns of the table.

To add supplemental logging on all columns of a table, run the following command.

```
ALTER TABLE table_name ADD SUPPLEMENTAL LOG DATA (ALL) COLUMNS;
```

6. Create a user for AWS SCT.

```
CREATE USER oracle_sct_user IDENTIFIED BY password;  
  
GRANT CONNECT TO oracle_sct_user;  
GRANT SELECT_CATALOG_ROLE TO oracle_sct_user;  
GRANT SELECT ANY DICTIONARY TO oracle_sct_user;
```

### Step 3: Configure Your PostgreSQL Target Database

1. If the schemas you are migrating do not exist on the PostgreSQL database, then create the schemas.
2. Create the AWS DMS user to connect to your target database, and grant Superuser or the necessary individual privileges (or use the master username for RDS).

```
CREATE USER postgresql_dms_user WITH PASSWORD 'password';  
ALTER USER postgresql_dms_user WITH SUPERUSER;
```

### 3. Create a user for AWS SCT.

```
CREATE USER postgresql_sct_user WITH PASSWORD 'password';  
  
GRANT CONNECT ON DATABASE database_name TO postgresql_sct_user;  
GRANT USAGE ON SCHEMA schema_name TO postgresql_sct_user;  
GRANT SELECT ON ALL TABLES IN SCHEMA schema_name TO postgresql_sct_user;  
GRANT ALL ON ALL SEQUENCES IN SCHEMA schema_name TO postgresql_sct_user;
```

## Step 4: Use AWS SCT to Convert the Oracle Schema to PostgreSQL

Before you migrate data to PostgreSQL, you convert the Oracle schema to a PostgreSQL schema. [This video covers all the steps of this process.](#)

To convert an Oracle schema to a PostgreSQL schema using AWS Schema Conversion Tool (AWS SCT), do the following:

1. Launch AWS SCT. In AWS SCT, choose **File**, then choose **New Project**. Create a new project named AWS Schema Conversion Tool Oracle to PostgreSQL, specify the **Location** of the project folder, and then choose **OK**.
2. Choose **Add source** to add a source Oracle database to your project, then choose **Oracle**, and choose **Next**.
3. Enter the following information, and then choose **Test Connection**.

| For This Parameter | Do This                                                                 |
|--------------------|-------------------------------------------------------------------------|
| Connection name    | Enter Oracle. AWS SCT displays this name in the tree in the left panel. |
| Type               | Choose <b>SID</b> .                                                     |
| Server name        | Enter the server name.                                                  |

| For This Parameter | Do This                                            |
|--------------------|----------------------------------------------------|
| <b>Server port</b> | Enter the Oracle port number. The default is 1521. |
| <b>Oracle SID</b>  | Enter the database SID.                            |
| <b>User name</b>   | Enter the Oracle admin username.                   |
| <b>Password</b>    | Enter the password for the admin user.             |

 Add source ×

**CONNECTION** SSL

Specify parameters for new connections to Oracle ?

Connection name  i

Type

Server name

Server port

Service name

User name

Password

☐ Use SSL

☐ Store password

- Choose **OK** to close the alert box, then choose **Connect** to close the dialog box and to connect to the Oracle DB instance.



5. Choose **Add target** to add a target PostgreSQL database to your project, then choose **Amazon RDS for PostgreSQL**, and choose **Next**.
6. Enter the following information and then choose **Test Connection**.

| Parameter              | Description                                                                                  |
|------------------------|----------------------------------------------------------------------------------------------|
| <b>Connection name</b> | Enter Amazon RDS for PostgreSQL . AWS SCT displays this name in the tree in the right panel. |
| <b>Server name</b>     | Enter the server name.                                                                       |
| <b>Server port</b>     | Enter the PostgreSQL port number. The default is 5432.                                       |
| <b>Database</b>        | Enter the database name.                                                                     |
| <b>User name</b>       | Enter the PostgreSQL admin username.                                                         |
| <b>Password</b>        | Enter the password for the admin user.                                                       |

7. Choose **OK** to close the alert box, then choose **Connect** to connect to the Amazon RDS for PostgreSQL DB instance.
8. In the tree in the left panel, select the schema to migrate. In the tree in the right panel, select your target Amazon RDS for PostgreSQL database. Choose **Create mapping**. For more information, see [Creating mapping rules](#) in the *Schema Conversion Tool User Guide*.

▼ Server mappings

| <input type="checkbox"/> | Source path                      | Target path                                         |  |
|--------------------------|----------------------------------|-----------------------------------------------------|--|
| <input type="checkbox"/> | //Servers/Oracle/Schemas/CHINOOK | //Servers/Amazon RDS for PostgreSQL/Schemas/chinook |  |
| <input type="checkbox"/> |                                  |                                                     |  |
| <input type="checkbox"/> |                                  |                                                     |  |
| <input type="checkbox"/> |                                  |                                                     |  |
| <input type="checkbox"/> |                                  |                                                     |  |
| <input type="checkbox"/> |                                  |                                                     |  |
| <input type="checkbox"/> |                                  |                                                     |  |
| <input type="checkbox"/> |                                  |                                                     |  |
| <input type="checkbox"/> |                                  |                                                     |  |

Create mapping
 Delete selected mappings
 Main view

9. Choose **Main view**. In the tree in the left panel, right-click the schema to migrate and choose **Convert schema**.

10. Choose **Yes** for the confirmation message. AWS SCT analyzes the schema, creates a database migration assessment report, and converts your schema to the target database format.

11. Choose **Assessment Report View** from the menu to check the database migration assessment report. The report breaks down by each object type and by how much manual change is needed to convert it successfully.

Generally, packages, procedures, and functions are more likely to have some issues to resolve because they contain the most custom PL/SQL code. AWS SCT also provides hints about how to fix these objects.

12. Choose the **Action Items** tab.

The **Action Items** tab shows each issue for each object that requires attention.

For each conversion issue, you can complete one of the following actions:

- Modify the objects on the source Oracle database so that AWS SCT can convert the objects to the target Amazon RDS for PostgreSQL database.
  - i. Modify the objects on the source Oracle database.
  - ii. Repeat the previous steps to convert the schema and check the assessment report.
  - iii. If necessary, repeat this process until there are no conversion issues.

- iv. Choose **Main View** from the menu. Open the context (right-click) menu for the target Amazon RDS for PostgreSQL schema, and choose **Apply to database** to apply the schema changes to the Amazon RDS for PostgreSQL database, and confirm that you want to apply the schema changes.
- Instead of modifying the source schema, modify scripts that AWS SCT generates before applying the scripts on the target Amazon RDS for PostgreSQL database.
  - i. Choose **Main View** from the menu. Open the context (right-click) menu for the target Amazon RDS for PostgreSQL schema name, and choose **Save as SQL**. Next, choose a name and destination for the script.
  - ii. In the script, modify the objects to correct conversion issues.

You can also exclude foreign key constraints, triggers, and secondary indexes from the script because they can cause problems during the migration. After the migration is complete, you can create these objects on the Amazon RDS for PostgreSQL database.

- iii. Run the script on the target Amazon RDS for PostgreSQL database.

For more information, see [Converting Database Schema to Amazon RDS](#).

### 13(Optional) Use AWS SCT to create migration rules.

- a. Choose **Mapping view** and then choose **New migration rule**.
- b. Create additional migration transformation rules that are required based on the action items.
- c. Save the migration rules.
- d. Choose **Export script for DMS** to export a JSON format of all the transformations that the AWS DMS task will use. Choose **Save**.

## Step 5: Create an AWS DMS Replication Instance

After validating the schema structure between source and target databases, continue with the core part of this walkthrough, which is the data migration. The following illustration shows a high-level view of the migration process.



An AWS DMS replication instance performs the actual data migration between source and target. The replication instance also caches the transaction logs during the migration. How much CPU and memory capacity a replication instance has influences the overall time required for the migration.

1. Sign in to the AWS Management Console, and select AWS DMS at <https://console.aws.amazon.com/dms/v2/>. Next, choose **Create Migration**. If you are signed in as an AWS Identity and Access Management (IAM) user, then you must have the appropriate permissions to access AWS DMS. For more information about the permissions required, see [IAM Permissions](#).
2. Choose **Next** to start a database migration from the console's Welcome page.
3. On the **Create replication instance** page, specify your replication instance information.

| Parameter             | Description                                                                                                                                                                                                                                              |
|-----------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Name</b>           | Select a name for your replication instance. If you will be using multiple replication servers or sharing an account, then choose a name that will help you quickly differentiate between the different servers.                                         |
| <b>Description</b>    | Enter a brief description.                                                                                                                                                                                                                               |
| <b>Instance class</b> | Select the type of replication server to create. Each size and type of instance class will have increasing CPU, memory, and I/O capacity. Generally, the t2 instances are for lower load tasks, and the c4 instances are for higher load and more tasks. |
| <b>VPC</b>            | Choose the VPC in which your replication instance will be launched. If possible, select the same VPC in which either your source or target database resides (or both).                                                                                   |
| <b>Multi-AZ</b>       | When <b>Yes</b> is selected, AWS DMS creates a second replication server in a different                                                                                                                                                                  |

| Parameter                  | Description                                                                                                                                                                          |
|----------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                            | Availability Zone for failover if there is a problem with the primary replication server.                                                                                            |
| <b>Publicly accessible</b> | If either your source or target database resides outside of the VPC in which your replication server resides, then you must make your replication server policy publicly accessible. |

4. For the **Advanced** section, specify the following information.

| Parameter                       | Description                                                                                                                                                                                                                                                                                |
|---------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Allocated storage (GB)</b>   | Amount of storage on the replication server for the AWS DMS task logs, including historical tasks logs. AWS DMS also uses disk storage to cache certain data while it replicates it from the source to the target. Additionally, more storage generally enables better IOPS on the server. |
| <b>Replication Subnet Group</b> | If you are running in a Multi-AZ configuration, then you will need at least two subnet groups.                                                                                                                                                                                             |
| <b>Availability zone</b>        | Generally, performance is better if you locate your primary replication server in the same Availability Zone as your target database.                                                                                                                                                      |
| <b>VPC Security Group(s)</b>    | Security groups enable you to control ingress and egress to your VPC. AWS DMS allows you to associate one or more security groups with the VPC in which your replication server is launched.                                                                                               |
| <b>KMS key</b>                  | With AWS DMS, all data is encrypted at rest using a KMS encryption key. By default, AWS                                                                                                                                                                                                    |

| Parameter | Description                                                                                                       |
|-----------|-------------------------------------------------------------------------------------------------------------------|
|           | DMS will create a new encryption key for your replication server. However, you may choose to use an existing key. |

For information about the KMS key, see [Setting an Encryption Key and Specifying KMS Permissions](#).

5. Click **Next**.

## Step 6: Create AWS DMS Source and Target Endpoints

While your replication instance is being created, you can specify the source and target database endpoints using the [AWS Management Console](#). However, you can only test connectivity after the replication instance has been created, because the replication instance is used in the connection.

1. Sign in to the AWS Management Console, open the [AWS DMS console](#), and then choose **Endpoints**.
2. Specify your connection information for the source Oracle database and the target PostgreSQL database. The following table describes the source settings.

| Parameter                  | Description                                                                        |
|----------------------------|------------------------------------------------------------------------------------|
| <b>Endpoint Identifier</b> | Enter a name, such as <code>Orasource</code> .                                     |
| <b>Source Engine</b>       | Choose <b>oracle</b> .                                                             |
| <b>Server name</b>         | Provide the Oracle DB instance server name.                                        |
| <b>Port</b>                | The port of the database. The default for Oracle is 1521.                          |
| <b>SSL mode</b>            | Choose an SSL mode if you want to enable encryption for your connection's traffic. |
| <b>Username</b>            | The user you want to use to connect to the source database.                        |

| Parameter | Description                        |
|-----------|------------------------------------|
| Password  | Provide the password for the user. |
| SID       | Provide the Oracle database name.  |

The following table describes the advanced source settings.

| Parameter                          | Description                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
|------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Extra connection attributes</b> | <p>Extra parameters that you can set in an endpoint to add functionality or change the behavior of AWS DMS. Some of the most common and convenient parameters to set for an Oracle source database are the following. Separate multiple entries from each other by using a semi-colon (;).</p> <ul style="list-style-type: none"> <li>• <code>addSupplementalLogging</code> - This parameter automatically configures supplemental logging when set to Y.</li> <li>• <code>useLogminerReader</code> - By default, AWS DMS uses LogMiner on the Oracle database to capture all of the changes on the source database. The other mode is called Binary Reader. When using Binary Reader instead of LogMiner, AWS DMS copies the archived redo log from the source Oracle database to the replication server and reads the entire log in order to capture changes. The Binary Reader option is recommended if you are using ASM since it has performance advantages over LogMiner on ASM. If your source database is 12c, then the Binary Reader option is currently the only way to capture CDC changes in Oracle for LOB objects.</li> </ul> <p>To use LogMiner, enter the following:<br/> <code>useLogminerReader=Y</code></p> <p>To use Binary Reader, enter the following:<br/> <code>useLogminerReader=N; useBfile=Y</code></p> |



| Parameter | Description                                                                     |
|-----------|---------------------------------------------------------------------------------|
| KMS key   | Enter the KMS key if you choose to encrypt your replication instance's storage. |

For information about extra connection attributes, see [Using Extra Connection Attributes](#).

The following table describes the target settings.

| Parameter           | Description                                                   |
|---------------------|---------------------------------------------------------------|
| Endpoint Identifier | Enter a name, such as Postgrestarget .                        |
| Target Engine       | Choose <b>postgres</b> .                                      |
| Servename           | Provide the PostgreSQL DB instance server name.               |
| Port                | The port of the database. The default for PostgreSQL is 5432. |
| SSL mode            | Choose <b>None</b> .                                          |
| Username            | The user you want to use to connect to the target database.   |
| Password            | Provide the password for the PostgreSQL DB instance.          |

The following is an example of the completed page.

## Connect source and target database endpoints

✔ Replication instance created successfully.

Your database endpoint can be on-premise, in EC2, RDS or in the cloud. Define the connection details below. It is recommended that you test your endpoint connections here to avoid errors later.

| Source database connection details                                   | Target database connection details                                        |
|----------------------------------------------------------------------|---------------------------------------------------------------------------|
| <b>Endpoint identifier*</b> <input type="text" value="Orasource"/> ⓘ | <b>Endpoint identifier*</b> <input type="text" value="Postgrestarget"/> ⓘ |
| <b>Source engine*</b> <span>oracle ▾</span> ⓘ                        | <b>Target engine*</b> <span>postgres ▾</span> ⓘ                           |
| <b>Server name*</b> <input type="text"/>                             | <b>Server name*</b> <input type="text"/>                                  |
| <b>Port*</b> <span>1521 ▾</span> ⓘ                                   | <b>Port*</b> <span>5432 ▾</span> ⓘ                                        |
| <b>SSL mode*</b> <span>none ▾</span> ⓘ                               | <b>SSL mode*</b> <span>none ▾</span> ⓘ                                    |
| <b>User name*</b> <input type="text" value="oraadmin"/> ⓘ            | <b>User name*</b> <input type="text" value="postgresadmin"/> ⓘ            |
| <b>Password*</b> <input type="password" value="....."/> ⓘ            | <b>Password*</b> <input type="password" value="....."/> ⓘ                 |
| <b>SID*</b> <input type="text" value="ORCL"/>                        | <b>Database name*</b> <input type="text" value="postgres"/>               |
| ▶ <b>Advanced</b>                                                    | ▶ <b>Advanced</b>                                                         |
| <input type="button" value="Run test"/>                              | <input type="button" value="Run test"/>                                   |

3. After the endpoints and replication instance have been created, test each endpoint connection by choosing **Run test** for the source and target endpoints.
4. Drop foreign key constraints and triggers on the target database.

During the full load process, AWS DMS does not load tables in any particular order, so it may load the child table data before parent table data. As a result, foreign key constraints might be violated if they are enabled. Also, if triggers are present on the target database, then it may change data loaded by AWS DMS in unexpected ways.

5. If you do not have one, then generate a script that enables the foreign key constraints and triggers.

Later, when you want to add them to your migrated database, you can just run this script.

6. (Optional) Drop secondary indexes on the target database.

Secondary indexes (as with all indexes) can slow down the full load of data into tables since they need to be maintained and updated during the loading process. Dropping them can improve the performance of your full load process. If you drop the indexes, then you will need to add them back later after the full load is complete.

7. Choose **Next**.

## Step 7: Create and Run Your AWS DMS Migration Task

Using an AWS DMS task, you can specify which schema to migrate and the type of migration. You can migrate existing data, migrate existing data and replicate ongoing changes, or replicate data changes only. This walkthrough migrates existing data and replicates ongoing changes.

1. On the **Create Task** page, specify the task options. The following table describes the settings.

| Parameter            | Description                                                                                                                                 |
|----------------------|---------------------------------------------------------------------------------------------------------------------------------------------|
| Task name            | Enter a name for the migration task.                                                                                                        |
| Task description     | Enter a description for the task.                                                                                                           |
| Source endpoint      | Shows the Oracle source endpoint.<br><br>If you have more than one endpoint in the account, then choose the correct endpoint from the list. |
| Target endpoint      | Shows the PostgreSQL target endpoint.                                                                                                       |
| Replication instance | Shows the AWS DMS replication instance.                                                                                                     |
| Migration type       | Choose <b>Migrate existing data and replicate ongoing changes</b> .                                                                         |
| Start task on create | Select this option.                                                                                                                         |

The page should look like the following:

The screenshot shows the 'Create task' form in the AWS DMS console. At the top, there's a title 'Create task' and a subtitle: 'A task can contain one or more table mappings which define what data is moved from the source to the target. If a table does not exist on the target, it can be created automatically.' The form contains several fields: 'Task name\*' with the value 'MigrateSchemaToPostgres', 'Task description\*' with the value 'Migrate a schema from Oracle to PostgreSQL', 'Source endpoint' with the value 'orasource', 'Target endpoint' with the value 'postgresrarget', 'Replication instance' with the value 'oracle2postgresql', and 'Migration type\*' with a dropdown menu showing 'Migrate existing data and replicate ongoing changes'. Below these fields is a yellow warning box with text about supplemental logging for Oracle databases. At the bottom, there is a checkbox labeled 'Start task on create' which is checked.

**Create task**

A task can contain one or more table mappings which define what data is moved from the source to the target. If a table does not exist on the target, it can be created automatically.

Task name\* MigrateSchemaToPostgres ⓘ

Task description\* Migrate a schema from Oracle to PostgreSQL ⓘ

Source endpoint orasource

Target endpoint postgresrarget

Replication instance oracle2postgresql

Migration type\* Migrate existing data and replicate ongoing changes ⓘ

Your source database is Oracle. Replicating ongoing changes requires supplemental logging to be turned on.

Please ensure your archive logs are retained on the server for a sufficient amount of time, (24 hours is usually enough.) To set your archive log retention on RDS databases you can use the following command: `exec rdsadmin.rdsadmin_util.set_configuration('archive_log_retention_hours', 24);`

Start task on create ☒

2. Under **Task Settings**, choose **Do nothing** or **Truncate** for **Target table preparation mode**, because you have already created the tables using the AWS Schema Conversion Tool.

If the Oracle database has LOBs, then for **Include LOB columns in replication**, select **Full LOB mode** if you want to replicate the entire LOB for all tables. Select **Limited LOB mode** if you want to replicate the LOBs only up to a certain size. You specify the size of the LOB to migrate in **Max LOB size (kb)**.

It is best to select **Enable logging**. If you enable logging, then you can see any errors or warnings that the task encounters, and you can troubleshoot those issues.

Task Settings

Target table preparation mode\*

☒ Do nothing
 ☐ Drop tables on target
 ☐ Truncate

Include LOB columns in replication\*

☐ Don't include LOB columns
 ☐ Full LOB mode
 ☒ Limited LOB mode

Max LOB size (kb)\*

32

Enable logging

☒

Advanced Settings

3. Leave the Advanced settings at their default values.
4. Choose **Table mappings**, and select the **JSON** tab. Next, select **Enable JSON editing**, and enter the table mappings you saved in the last step in [Step 4: Convert the Oracle Schema to PostgreSQL](#).

The following is an example of mappings that convert schema names and table names to lowercase.

```
{
  "rules": [
    {
      "rule-type": "transformation",
      "rule-id": "100000",
      "rule-name": "Default Lowercase Table Rule",
      "rule-action": "convert-lowercase",
      "rule-target": "table",
      "object-locator": {
        "schema-name": "%",
        "table-name": "%"
      }
    },
    {
      "rule-type": "transformation",
      "rule-id": "100001",
      "rule-name": "Default Lowercase Schema Rule",
```


```

    "rule-action": "convert-lowercase",
    "rule-target": "schema",
    "object-locator": {
      "schema-name": "%"
    }
  }
]
}

```

5. Choose **Create task**. The task will begin immediately.

The Tasks section shows you the status of the migration task.



You can monitor your task if you chose **Enable logging** when you set up your task. You can then view the CloudWatch metrics by doing the following:

1. On the navigation pane, choose **Tasks**.
2. Choose your migration task.
3. Choose the **Task monitoring** tab, and monitor the task in progress on that tab.

When the full load is complete and cached changes are applied, the task will stop on its own.

4. On the target PostgreSQL database, enable foreign key constraints and triggers using the script you saved previously.
5. On the target PostgreSQL database, re-create the secondary indexes if you removed them previously.
6. In the AWS DMS console, start the AWS DMS task by clicking **Start/Resume** for the task.

The AWS DMS task keeps the target PostgreSQL database up-to-date with source database changes. AWS DMS will keep all of the tables in the task up-to-date until it is time to implement the application migration. The latency will be zero, or close to zero, when the target has caught up to the source.

## Step 8: Cut Over to PostgreSQL

To move connections from your Oracle database to your PostgreSQL database, do the following:

1. End all Oracle database dependencies and activities, such as running scripts and client connections.

The following query should return no results:

```
SELECT MACHINE, COUNT(*) FROM V$SESSION GROUP BY MACHINE;
```

2. List any remaining sessions, and kill them.

```
SELECT SID, SERIAL#, STATUS FROM V$SESSION;  
  
ALTER SYSTEM KILL 'sid, serial_number' IMMEDIATE;
```

3. Shut down all listeners on the Oracle database.
4. (Optional) Turn off automated jobs on the Oracle database. For your production database, check that this operation doesn't influence the business logic.

```
ALTER SYSTEM SET JOB_QUEUE_PROCESSES=0
```

5. (Optional) Turn off time monitoring on queue messages on the Oracle database. For your production database, check that this operation doesn't influence the business logic.

```
ALTER SYSTEM SET AQ_TM_PROCESSES=0
```

6. Let the AWS DMS task apply the final changes from the Oracle database on the PostgreSQL database.

```
ALTER SYSTEM CHECKPOINT;
```

7. In the AWS DMS console, stop the AWS DMS task by clicking **Stop** for the task, and confirm that you want to stop the task.
8. (Optional) Set up a rollback.

You can optionally set up a rollback task, in case you run into a show stopping issue, by creating a task going in the opposite direction. Because all tables should be in sync between both databases, you only need to set up a CDC task. Therefore, you do not have to disable any foreign

key constraints. Now that the source and target databases are reversed, you must follow the instructions in the following sections:

- [Using a PostgreSQL Database as a Source](#)
- [Using an Oracle Database as a Target](#)
  - a. Disable triggers on the source Oracle database.

```
SELECT 'ALTER TRIGGER' || owner || '.' || trigger_name || 'DISABLE;'
FROM DBA_TRIGGERS WHERE OWNER = 'schema_name';
```

You do not have to disable the foreign key constraints. During the CDC process, foreign key constraints are updated in the same order as they are updated by application users.

- b. Create a new CDC-only AWS DMS task with the endpoints reversed (source PostgreSQL endpoint and target Oracle endpoint database). See [Step 7: Create and Run Your Migration Task](#).

For the rollback task, set **Migration type** to **Replicate data changes only** and **Target table preparation mode** to **Do nothing**.

- c. Start the AWS DMS task to enable you to push changes back to the original source Oracle database from the new PostgreSQL database if rollback is necessary.

9. Connect to the PostgreSQL database, and enable triggers.

```
ALTER TABLE table_name ENABLE TRIGGER ALL;
```

10. If you set up a rollback, then complete the rollback setup.

- a. Start the application services on new target PostgreSQL database (including scripts, client software, and so on).
- b. Add CloudWatch monitoring on your new PostgreSQL database. For more information, see [Monitoring Amazon RDS](#).

## Rolling Back the Migration

If there are major issues with the migration that cannot be resolved in a timely manner, you can roll back the migration. These steps assume that you have already prepared for the rollback as described in [Step 8: Cut Over to PostgreSQL](#).

### 1. Stop all application services on the target PostgreSQL database.



2. Let the AWS DMS task replicate remaining changes back to the source Oracle database.
3. Stop the PostgreSQL to Oracle AWS DMS task.
4. Start all applications back on the source Oracle database.

## Oracle database migration to PostgreSQL troubleshooting

The two most common problem areas when working with Oracle as a source and PostgreSQL as a target are: supplemental logging and case sensitivity.

- Supplemental logging – With Oracle, in order to replicate change data, supplemental logging must be enabled. However, if you enable supplemental logging at the database level, it sometimes still needs to be enabled when new tables are created. The best remedy for this is to allow AWS DMS to enable supplemental logging for you by using the extra connection attribute:

```
addSupplementalLogging=Y
```

- Case sensitivity: Oracle is case-insensitive (unless you use quotes around your object names). However, text appears in the upper case. Thus, AWS DMS defaults to naming your target objects in the upper case. In most cases, you'll want to use transformations to change schema, table, and column names to lower case.

For more tips, see [Troubleshooting migration tasks](#).

To troubleshoot issues specific to Oracle, see [Troubleshooting Oracle Specific Issues](#).

To troubleshoot PostgreSQL issues, see [Troubleshooting PostgreSQL Specific Issues](#).

## Migrating Oracle databases to Amazon Aurora MySQL with DMS Schema Conversion

This walkthrough gets you started with heterogeneous database migration from Oracle to Amazon Aurora MySQL-Compatible Edition. To automate the migration, we use the AWS DMS Schema Conversion. This service helps assess the complexity of your migration and converts source Oracle database schemas and code objects to a format compatible with MySQL. Then, you apply the converted code to your target database. This introductory exercise shows how you can use DMS Schema Conversion for this migration.

At a high level, this migration includes the following steps:

- Use the AWS Management Console to do the following:
  - Create a VPC in the Amazon VPC console.
  - Create IAM roles in the IAM console.
  - Create an Amazon S3 bucket in the Amazon S3 console.
  - Create your target Aurora MySQL database in the Amazon RDS console.
  - Store database credentials in AWS Secrets Manager.
- Use the AWS DMS console to do the following:
  - Create an instance profile for your migration project.
  - Create data providers for your source and target databases.
  - Create a migration project.
- Use DMS Schema Conversion to do the following:
  - Assess the migration complexity and review the migration action items.
  - Convert your source database.
  - Apply the converted code to your target database.

This walkthrough takes approximately three hours to complete. Make sure that you delete resources at the end of this walkthrough to avoid additional charges.

## Topics

- [Migration overview](#)
- [Prerequisites for Migrating Oracle databases to Amazon RDS for MySQL with DMS schema conversion](#)
- [Step-by-step Oracle database to Amazon RDS for MySQL with DMS schema conversion migration walkthrough](#)
- [Migration from Oracle database to Amazon RDS for MySQL with DMS schema conversion next steps](#)


## Migration overview

This section provides high-level guidance for customers looking to migrate from Oracle to MySQL using DMS Schema Conversion.

DMS Schema Conversion automatically converts your source Oracle database schemas and most of the database code objects to a format compatible with MySQL. This conversion includes tables, views, stored procedures, functions, data types, synonyms, and so on. Any objects that DMS Schema Conversion can't convert automatically are clearly marked. To complete the migration, you can convert these objects manually.

At a high level, DMS Schema Conversion operates with the following three components: instance profiles, data providers, and migration projects. An instance profile specifies network and security settings. A data provider stores database connection credentials. A migration project contains data providers, an instance profile, and migration rules. AWS DMS uses data providers and an instance profile to design a process that converts database schemas and code objects.

The following diagram illustrates the DMS Schema Conversion process for this walkthrough.



Start the walkthrough by [creating the required resources](#).

## Prerequisites for Migrating Oracle databases to Amazon RDS for MySQL with DMS schema conversion

The following prerequisites are also required to complete this walkthrough:

- Familiarity with the AWS Management Console, AWS Database Migration Service, and SQL.
- A user with AWS Identity and Access Management (IAM) credentials. Make sure that you can use these credentials to create an Amazon S3 bucket in your AWS Region.
- Basic knowledge of the Amazon Virtual Private Cloud (Amazon VPC) service and of security groups.
- An understanding of the supported features and limitations of DMS Schema Conversion. For more information, see [Schema conversion limitations](#).

We recommend that you don't use your production workloads for the migration in this walkthrough. After you get familiar with migration tools and AWS services, you can migrate your production workloads.

Make sure that you create all your AWS and DMS Schema Conversion resources in the AWS Regions that support DMS Schema Conversion. For more information, see the [list of supported Regions](#). In other Regions, you can use the AWS Schema Conversion Tool (AWS SCT). To download AWS SCT, see [Installing, verifying, and updating](#) in the *Schema Conversion Tool User Guide*.

For more information about DMS Schema Conversion, see the [user guide](#).

## Step-by-step Oracle database to Amazon RDS for MySQL with DMS schema conversion migration walkthrough

In the following sections, you can find step-by-step instructions for migrating your Oracle database to Aurora MySQL using DMS Schema Conversion.

### Topics

- [Step 1: Create AWS Resources](#)
- [Step 2: Configure Your Source Database](#)
- [Step 3: Create Your Target Aurora MySQL Database](#)
- [Step 4: Store Database Credentials in AWS Secrets Manager](#)
- [Step 5: Create an Instance Profile](#)
- [Step 6: Configure Data Providers](#)
- [Step 7: Create a Migration Project](#)
- [Step 8: Convert Database Objects](#)
- [Step 9: Edit and Apply Your Converted Code](#)

### Step 1: Create AWS Resources

In this step, you create and configure the required AWS resources for DMS Schema Conversion.

First, you create a virtual private cloud (VPC). This VPC is based on the Amazon Virtual Private Cloud (Amazon VPC) service and contains your AWS resources. Make sure that you create this VPC in one of the AWS Regions that support DMS Schema Conversion. For more information, see the [list of supported Regions](#).

## To create a VPC for DMS Schema Conversion

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. Choose your AWS Region.
3. Choose **Create VPC**.
4. On the **Create VPC** page, enter the following settings:
  - **Resources to create** — **VPC and more**
  - **Name tag auto-generation** — Choose **Auto-generate** and enter a globally unique name. For example, enter `sc-vpc`.
  - **IPv4 CIDR block** — `10.0.1.0/24`
  - **NAT gateways** — **In 1 AZ**
  - **VPC endpoints** — **None**
5. Keep the rest of the settings as they are, and then choose **Create VPC**.
6. Choose **Subnets**.
  - For **Filter by VPC**, choose `sc-vpc`.
  - Take a note of your two private subnet IDs. Private subnet IDs don't include `Public` in the name.
7. Choose **NAT gateways**.
  - Choose your **NAT gateway**.
  - Take a note of your **Elastic IP** address.

Use this VPC when you create your instance profile in [Step 5](#) and your target Aurora database in [Step 3](#).

Next, you create AWS Identity and Access Management (IAM) roles to use in your DMS Schema Conversion migration project. AWS DMS uses this IAM role to access your Amazon S3 bucket and database credentials stored in AWS Secrets Manager.

## To create an IAM role that provides access to your Amazon S3 bucket

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the navigation pane, choose **Roles**.

3. Choose **Create role**.
4. On the **Select trusted entity** page, choose **AWS service**. Choose **DMS**.
5. Choose **Next**. The **Add permissions** page opens.
6. For **Filter policies**, enter S3. Choose **AmazonS3FullAccess**.
7. Choose **Next**. The **Name, review, and create** page opens.
8. For **Role name**, enter a descriptive name. For example, enter `sc-s3-role`. Choose **Create role**.
9. On the **Roles** page, enter `sc-s3-role` for **Role name**. Choose **sc-s3-role**.
10. On the `sc-s3-role` page, choose the **Trust relationships** tab. Choose **Edit trust policy**.
11. On the **Edit trust policy** page, edit the trust relationships for the role to use the `schema-conversion.dms.amazonaws.com` service principal as the trusted entity.
12. Choose **Update trust policy**.

Use this IAM role when you create your instance profile in [Step 5](#).

### To create an IAM role that provides access to AWS Secrets Manager

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the navigation pane, choose **Roles**.
3. Choose **Create role**.
4. On the **Select trusted entity** page, choose **AWS service**. Choose **DMS**.
5. Choose **Next**. The **Add permissions** page opens.
6. For **Filter policies**, enter Secret. Choose **SecretsManagerReadWrite**.
7. Choose **Next**. The **Name, review, and create** page opens.
8. For **Role name**, enter a descriptive name. For example, enter `sc-secrets-manager-role`. Choose **Create role**.
9. On the **Roles** page, enter `sc-secrets-manager-role` for **Role name**. Choose **sc-secrets-manager-role**.
10. On the `sc-secrets-manager-role` page, choose the **Trust relationships** tab. Choose **Edit trust policy**.
11. On the **Edit trust policy** page, edit the trust relationships for the role to use `schema-conversion.dms.amazonaws.com` and your AWS DMS regional service principal as the trusted entities. This principal has the following format.

```
dms.region-name.amazonaws.com
```

Replace *region-name* with the name of your Region, such as `us-east-1`.

The following code example shows the principal for the `us-east-1` Region.

```
dms.us-east-1.amazonaws.com
```

## 12. Choose **Update trust policy**.

Use this IAM role when you create your migration project in [Step 7](#).

Next, you create an Amazon S3 bucket to use in your DMS Schema Conversion migration project. DMS Schema Conversion uses this Amazon S3 bucket to save assessment reports, SQL scripts with the converted code, and database metadata.

### To create an Amazon S3 bucket for DMS Schema Conversion

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. Choose **Create bucket**.
3. On the **Create bucket** page, select a globally unique name for your S3 bucket. For example, enter `sc-s3-bucket`.
4. For **AWS Region**, choose your Region.
5. For **Bucket Versioning**, choose **Enable**.
6. Keep the rest of the settings as they are, and then choose **Create bucket**.

Use this Amazon S3 bucket when you create your instance profile in [Step 5](#).

## Step 2: Configure Your Source Database

In this step, you configure a new database user on your source Oracle database. Also, you configure the network to set up interaction for your source database with DMS Schema Conversion.

Use the credentials of this new user in DMS Schema Conversion. We encourage not using the admin user in the DMS Schema Conversion migration project.

Make sure that you grant the following privileges to this new user to complete the migration:

- **CONNECT** — includes only the **CREATE SESSION** privilege.
- **SELECT\_CATALOG\_ROLE** — provides **SELECT** privileges on all data dictionary views for Oracle 10g users.
- **SELECT ANY DICTIONARY** — provides query access to any object in the **SYS** schema.

You can use the following code example to create a database user and grant the privileges.

```
CREATE USER user_name IDENTIFIED BY your_password;  
GRANT CONNECT TO user_name;  
GRANT SELECT_CATALOG_ROLE TO user_name;  
GRANT SELECT ANY DICTIONARY TO user_name;
```

In the preceding example, replace *user\_name* with the name of your user. Then, replace *your\_password* with a secure password.

After you configure your database user, make sure that DMS Schema Conversion can access your source Oracle database. To set up a network for DMS Schema Conversion, you can use different network configurations. These configurations depend on the settings of your source database and your network. For more information about available options, see [Setting up a network for DMS Schema Conversion](#).

In this walkthrough, you configure a Site-to-Site VPN connection using a virtual private gateway.

### To configure a Site-to-Site VPN connection

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. Choose your AWS Region.
3. Create a customer gateway.
  - In the navigation pane, choose **Customer gateways**, and then **Create customer gateway**.
  - For **Name tag**, enter a name for your customer gateway.
  - For **BGP ASN**, enter a Border Gateway Protocol (BGP) Autonomous System Number (ASN) for your customer gateway.
  - For **IP address**, enter the static, internet-routable IP address for your customer gateway device.
  - For **Certificate ARN**, choose the Amazon Resource Name of the private certificate.
  - For **Device**, enter a name for the device that hosts this customer gateway.



#### 4. Create a virtual private gateway.

- In the navigation pane, choose **Virtual private gateways**, and then **Create virtual private gateway**.
- For **Name tag**, enter a name for your virtual private gateway.
- For **Autonomous System Number (ASN)**, choose **Amazon default ASN**.
- Choose **Create virtual private gateway**.
- Select the virtual private gateway you created, choose **Actions**, and then **Attach to VPC**.
- Under **Available VPCs**, select your VPC from the list and choose **Attach to VPC**.

#### 5. Configure route propagation in your route table.

- In the navigation pane, choose **Route tables**, and then select the route table that is associated with your subnet. By default, this is the main route table for the VPC.
- On the **Route propagation** tab in the details pane, choose **Edit route propagation**.
- Select the virtual private gateway that you created before, and then choose **Save**.

#### 6. Add rules to your security group.

- In the navigation pane, choose **Security groups**, and then select the default security group for your VPC.
- On the **Inbound** tab in the details pane, add rules that allow inbound SSH, RDP, and ICMP access from your network.
- Choose **Save**.

#### 7. Create a Site-to-Site VPN connection.

- In the navigation pane, choose **Site-to-Site VPN connections**, and then **Create VPN connection**.
- For **Name tag**, enter a name for your Site-to-Site VPN connection.
- For **Target gateway type**, choose either **Virtual private gateway**.
- For **Customer gateway**, select **Existing**.
- For **Customer gateway ID**, choose the customer gateway that you created before.
- Select the routing option. If your customer gateway device supports BGP, then choose **Dynamic (requires BGP)**. Alternatively, choose **Static** and specify IP prefixes for the private network of your Site-to-Site VPN connection.
- For **Outside IP address type**, keep the default option.
- Choose **Create VPN connection**.

#### 8. Download the configuration file.

- In the navigation pane, choose **Site-to-Site VPN connections**, and then **Download configuration**.
  - Select the vendor, platform, software, and IKE version that correspond to your customer gateway device. If your device isn't listed, choose **Generic**.
  - Choose **Download**.
9. Use the sample configuration file to configure your customer gateway device.

### Step 3: Create Your Target Aurora MySQL Database

In this step, you create a new Aurora MySQL database to use as a migration target for DMS Schema Conversion. Also, you configure a new database user on your target Aurora MySQL database.

If you already created the target database, skip this step and proceed with the configuration of your database user.

#### To create an Aurora MySQL database for DMS Schema Conversion

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. Choose your AWS Region.
3. Choose **Create database**.
4. For **Engine type**, choose **Amazon Aurora**.
5. For **Edition**, choose **Amazon Aurora MySQL-Compatible Edition**.
6. For **Templates**, choose **Dev/Test**.
7. For **DB cluster identifier**, enter a unique name for your MySQL database.
8. For **Master password** and **Confirm master password**, enter a secure password that includes at least 8 printable characters.
9. For **Virtual private cloud (VPC)** under **Connectivity**, choose **sc-vpc**. You created this VPC in [Step 1](#).
- 10 For **Public access**, choose **Yes**.
- 11 Keep the rest of the settings as they are, and then choose **Create database**.

After you create your Aurora MySQL database, configure a new database user. Then, use the credentials of this user in DMS Schema Conversion. We encourage not using the admin user in the DMS Schema Conversion migration project.

To configure your target database user, create a new user and grant the following privileges:

- CREATE ON .
- ALTER ON .
- DROP ON .
- INDEX ON .
- REFERENCES ON .
- SELECT ON .
- CREATE VIEW ON .
- SHOW VIEW ON .
- TRIGGER ON .
- CREATE ROUTINE ON .
- ALTER ROUTINE ON .
- EXECUTE ON .
- CREATE TEMPORARY TABLES ON .
- INVOKE LAMBDA ON .
- INSERT, UPDATE ON AWS\_ORACLE\_EXT.\*
- INSERT, UPDATE, DELETE ON AWS\_ORACLE\_EXT\_DATA.\*
- CREATE TEMPORARY TABLES ON AWS\_ORACLE\_EXT\_DATA.\*

You can use the following code example to create a database user and grant the privileges.

```
CREATE USER 'user_name' IDENTIFIED BY 'your_password';
GRANT CREATE ON *.* TO 'user_name';
GRANT ALTER ON *.* TO 'user_name';
GRANT DROP ON *.* TO 'user_name';
GRANT INDEX ON *.* TO 'user_name';
GRANT REFERENCES ON *.* TO 'user_name';
GRANT SELECT ON *.* TO 'user_name';
GRANT CREATE VIEW ON *.* TO 'user_name';
GRANT SHOW VIEW ON *.* TO 'user_name';
GRANT TRIGGER ON *.* TO 'user_name';
GRANT CREATE ROUTINE ON *.* TO 'user_name';
GRANT ALTER ROUTINE ON *.* TO 'user_name';
```

```
GRANT EXECUTE ON *.* TO 'user_name';  
GRANT INSERT, UPDATE ON AWS_ORACLE_EXT.* TO 'user_name';  
GRANT INSERT, UPDATE, DELETE ON AWS_ORACLE_EXT_DATA.* TO 'user_name';  
GRANT CREATE TEMPORARY TABLES ON AWS_ORACLE_EXT_DATA.* TO 'user_name';
```

In the preceding example, replace *user\_name* with the name of your user. Then, replace *your\_password* with a secure password.

In your target Amazon Aurora MySQL database, set the `lower_case_table_names` parameter to 1. Also, set the `log_bin_trust_function_creators` parameter to 1, and the `character_set_server` parameter to `latin1`.

## Step 4: Store Database Credentials in AWS Secrets Manager

To connect to your source and target databases with DMS Schema Conversion, store your database credentials in AWS Secrets Manager. Make sure that you replicate these secrets to your AWS Region.

### To store your source database credentials in AWS Secrets Manager

1. Sign in to the AWS Management Console and open the AWS Secrets Manager console at <https://console.aws.amazon.com/secretsmanager/>.
2. Choose your AWS Region.
3. Choose **Store a new secret**. The **Choose secret type** page opens.
4. For **Secret type**, choose **Credentials for other database**.
5. For **User name** and **Password**, enter the credentials of the database user that you created for your source database in [Step 2](#).
6. For **Database**, choose **Oracle**.
7. For **Server name**, **Database name**, and **Port**, enter your Oracle database connection information.
8. Choose **Next**. The **Configure secret** page opens.
9. For **Secret name**, enter `sc-oracle-secret`.
10. Choose **Next**. The **Configure rotation** page opens.
11. Choose **Next**. The **Review** page opens.
12. Choose **Store**.

## To store your target database credentials in AWS Secrets Manager

1. Sign in to the AWS Management Console and open the AWS Secrets Manager console at <https://console.aws.amazon.com/secretsmanager/>.
2. Choose your AWS Region.
3. Choose **Store a new secret**. The **Choose secret type** page opens.
4. For **Secret type**, choose **Credentials for Amazon RDS database**.
5. For **User name** and **Password**, enter the credentials of the database user that you created for your target database in [Step 3](#).
6. For **Database**, choose your Aurora MySQL DB instance.
7. Choose **Next**. The **Configure secret** page opens.
8. For **Secret name**, enter `sc-mysql-secret`.
9. Choose **Next**. The **Configure rotation** page opens.
10. Choose **Next**. The **Review** page opens.
11. Choose **Store**.

Use these secrets when you create your migration project in [Step 7](#).

## Step 5: Create an Instance Profile

Before you create an instance profile, configure a subnet group for your instance profile.

### To create a subnet group

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. In the navigation pane, choose **Subnet groups**, and then choose **Create subnet group**.
4. For **Name**, enter `PrivateSubnetGroup`.
5. For **Description**, enter A group of private subnets.
6. For **VPC**, choose `sc-vpc`. You created this VPC in [Step 1](#).
7. For **Add subnets**, choose two private subnet IDs. You noted these private subnet IDs in [Step 1](#).
8. Choose **Create subnet group**.

Before you create your migration project in DMS Schema Conversion, you set up an instance profile. An instance profile specifies network and security settings for DMS Schema Conversion.

### To create an instance profile

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. In the navigation pane, choose **Instance profiles**, and then choose **Create instance profile**.
4. For **Name**, enter a unique name for your instance profile. For example, enter `sc-instance`.
5. For **Virtual private cloud (VPC)**, choose `sc-vpc`. You created this VPC in [Step 1](#).
6. For **Subnet group**, choose the `PrivateSubnetGroup` subnet group that you created before.
7. For **S3 bucket** under **Schema conversion settings - optional**, choose an Amazon S3 bucket that you created in [Step 1](#).
8. For **IAM role**, choose the AWS Identity and Access Management (IAM) role that grants access to Amazon S3. You created this role in [Step 1](#).
9. Choose **Create instance profile**.

Use this instance profile when you create your migration project in [Step 7](#).

## Step 6: Configure Data Providers

In this step, you create data providers that describe your source and target databases. A data provider stores a data store type and the location information about your database. Data providers don't include database credentials. You store database credentials in AWS Secrets Manager. Make sure that you include data providers and database secrets in your DMS Schema Conversion migration project.

You can create only one data provider for a single database. If you try to create a second data provider for the same database, DMS Schema Conversion displays an error message. However, you can use one data provider in multiple migration projects.

### To create a data provider for your Oracle database

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.

3. In the navigation pane, choose **Data providers**, and then choose **Create data provider**.
4. For **Configuration**, choose **Enter manually**.
5. For **Name**, enter a unique name for your source data provider. For example, enter `sc-oracle`.
6. For **Engine type**, choose **Oracle**.
7. For **Server name**, enter the Domain Name Service (DNS) name or IP address of your database server.
8. For **Port**, enter the port used to connect to your database server.
9. For **Database name**, enter the name of your database.
- 10 For **Secure Socket Layer (SSL) mode**, choose **none**. Optionally, choose the type of your SSL enforcement, and provide the certificate information.
- 11 Choose **Create data provider**.

### To create a data provider for your Aurora MySQL database

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. In the navigation pane, choose **Data providers**, and then choose **Create data provider**.
4. For **Configuration**, choose **RDS database instance**.
5. For **Database from RDS**, choose your Aurora MySQL database.
6. For **Name**, enter a unique name for your target data provider. For example, enter `sc-mysql`.
7. For **Database name**, enter the name of your database.
8. For **Existing CA certificate**, choose the server certificate. If you don't have any server certificates, then for **Import certificate file**, provide the `rds-ca-2019.pem` file with your certificate.
9. Choose **Create data provider**.

Use these data providers when you create your migration project in [Step 7](#).

## Step 7: Create a Migration Project

Now you can create a migration project which is the foundation of your work with DMS Schema Conversion. A migration project describes your source and target data providers, your instance profile, and migration rules.

To create a migration project

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. Choose **Migration projects**, and then choose **Create migration project**.
4. For **Name**, enter a unique name for your migration project. For example, enter `sc-project`.
5. For **Instance profile**, choose `sc-instance`. You created this instance profile in [Step 5](#).
6. For **Source**, choose **Browse**, and then choose `sc-oracle`. You created this data provider in [Step 6](#).
7. For **Secret ID**, choose `sc-oracle-secret`. You created this secret in [Step 4](#).
8. For **IAM role**, choose `sc-secrets-manager-role`. You created this role in [Step 1](#).
9. For **Target**, choose **Browse**, and then choose `sc-mysql`. You created this data provider in [Step 6](#).
10. For **Secret ID**, choose `sc-mysql-secret`. You created this secret in [Step 4](#).
11. For **IAM role**, choose `sc-secrets-manager-role`. You created this role in [Step 1](#).
12. Choose **Create migration project**.

Use this migration project to convert your Oracle database schemas to MySQL.

## Step 8: Convert Database Objects

After you create the migration project, you can convert your Oracle database schemas to MySQL. To start working with your migration project, you launch DMS Schema Conversion.

The first launch of DMS Schema Conversion requires some setup. AWS Database Migration Service (AWS DMS) starts a schema conversion instance, which can take 10-15 minutes. This process also reads the metadata from the source and target databases. After a successful first launch, you can access DMS Schema Conversion instantly.

### To convert your source Oracle database schema with DMS Schema Conversion

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.



3. Choose **Migration projects**. The **Migration projects** page opens.
4. Choose `sc-project`, and then choose **Schema conversion**.
5. Choose **Launch schema conversion**. If you launch schema conversion for the first time, then the notification appears. Choose **Launch**. The **Schema conversion** page opens. DMS Schema Conversion displays your source database schema in the left pane in a tree-view format.
6. In the source database pane, select the check box for the schema name.
7. Choose this schema in the left pane of the migration project. DMS Schema Conversion highlights the schema name in blue and activates the **Actions** menu.
8. For **Actions**, choose **Convert schema**. The conversion dialog box appears.
9. Choose **Convert** in the dialog box to confirm your choice.

After DMS Schema Conversion completes the conversion, you can review the converted code. After you choose a database object in the left pane of your project, DMS Schema Conversion automatically displays the source converted code for this object.

DMS Schema Conversion stores the converted code in your migration project and doesn't apply these code changes to your target database. You can apply the converted code in DMS Schema Conversion. Alternatively, you can save the converted code as a SQL script, edit it, and then apply to your target database. For more information, see [Step 9](#).

In the settings of your migration project, you can customize your schema conversion view. Also, you can change conversion settings to improve the performance of converted code.

### To edit the settings of your DMS Schema Conversion migration project

1. In the AWS DMS console, choose **Migration projects**. The **Migration projects** page opens.
2. Choose your migration project. Choose **Schema conversion**, then **Launch schema conversion**.
3. Choose **Settings**. The **Settings** page opens.
4. Change the settings to customize the schema conversion view. For more information, see [Specifying migration project settings](#).
5. Change the settings to improve the performance of converted code. For more information, see [Specifying Oracle to MySQL conversion settings](#).
6. Choose **Apply**, and then choose **Schema conversion**.

After you change the settings, convert your source code again.

## Step 9: Edit and Apply Your Converted Code

After you convert your source Oracle database objects, you can review the conversion statistics. DMS Schema Conversion converts most of the database objects, but some of the objects require manual conversion.

DMS Schema Conversion displays the objects that require manual conversion in the **Action items** tab. To convert these objects, you can save the converted code as a SQL script. Then you can edit it using your code editor and apply these scripts to your target database. Alternatively, you can apply the converted code as is to your target database and make the edits later.

### To save the converted code as a SQL script

1. In the target database pane, choose the converted database schema.
2. Select the check box for the name of this schema. DMS Schema Conversion highlights the schema name in blue and activates the **Actions** menu.
3. For **Actions**, choose **Save as SQL**. The **Save** dialog box appears.
4. Choose **Save as SQL** to confirm your choice.
5. Choose **S3 bucket**. The **Amazon S3** console opens.
6. Choose **Download** to save your SQL scripts.

### To apply the converted code to your target database

1. In the target database pane, choose the converted database schema.
2. Select the check box for the name of this schema. DMS Schema Conversion highlights the schema name in blue and activates the **Actions** menu.
3. For **Actions**, choose **Apply changes**. The **Apply changes** dialog box appears.
4. Choose **Apply** to confirm your choice.

Now you have successfully converted your source Oracle database schemas to MySQL. To complete the database migration, move your data and connect your applications to the new database.

## Migration from Oracle database to Amazon RDS for MySQL with DMS schema conversion next steps

After you migrate your Oracle database to Aurora MySQL using DMS Schema Conversion, you can explore several other resources:

- Use AWS DMS to migrate your source data. For more information, see the [Database Migration Service User Guide](#).
- Use DMS Fleet Advisor to inventory your source databases and discover other candidates to move to the cloud. For more information, see the [DMS Fleet Advisor User Guide](#).
- Learn more about Aurora MySQL. For more information, see the [Amazon Aurora User Guide](#).

After you've finished using DMS Schema Conversion, clean up your resources. Amazon terminates the schema conversion instance that your migration project uses in three days after you complete the conversion. You can retrieve your converted schema and assessment report from the Amazon S3 bucket that you use for DMS Schema Conversion. However, you need to terminate other resources manually.

### To clean up your DMS Schema Conversion resources

- Sign in to the AWS Management Console and open the AWS DMS console.
- In the navigation pane, choose **Migration projects**, and then choose your migration project. Choose **Schema conversion**, and then choose **Stop schema conversion**. Choose **Delete** and confirm your choice.
- Choose **Instance profiles**, and then choose `sc-instance`. Choose **Delete** and confirm your choice.
- Choose **Data providers**, and then select `sc-oracle` and `sc-mysql`. Choose **Delete** and confirm your choice.

Also, make sure that you delete your Amazon S3 bucket, database secrets in AWS Secrets Manager, IAM roles, and virtual private cloud (VPC).

# Migrating Oracle databases to Amazon RDS for PostgreSQL with DMS Schema Conversion

This walkthrough gets you started with heterogeneous database migration from Oracle to Amazon RDS for PostgreSQL. To automate the migration, we use the AWS DMS Schema Conversion. This service helps assess the complexity of your migration and converts source Oracle database schemas and code objects to a format compatible with PostgreSQL. Then, you apply the converted code to your target database. This introductory exercise shows how you can use DMS Schema Conversion for this migration.

At a high level, this migration includes the following steps:

- Use the AWS Management Console to do the following:
  - Create a VPC in the Amazon VPC console.
  - Create IAM roles in the IAM console.
  - Create an Amazon S3 bucket in the Amazon S3 console.
  - Create your target Amazon RDS for PostgreSQL database in the Amazon RDS console.
  - Store database credentials in AWS Secrets Manager.
- Use the AWS DMS console to do the following:
  - Create an instance profile for your migration project.
  - Create data providers for your source and target databases.
  - Create a migration project.
- Use DMS Schema Conversion to do the following:
  - Assess the migration complexity and review the migration action items.
  - Convert your source database.
  - Apply the converted code to your target database.

This walkthrough takes approximately three hours to complete. Make sure that you delete resources at the end of this walkthrough to avoid additional charges.

## Topics

- [Migration overview](#)
- [Prerequisites for migrating Oracle databases to Amazon Aurora PostgreSQL with DMS schema conversion](#)

- [Step-by-step Oracle databases to Amazon Aurora PostgreSQL with DMS schema conversion migration walkthrough](#)
- [Migration from Oracle databases to Amazon Aurora PostgreSQL with DMS schema conversion next steps](#)


## Migration overview

This section provides high-level guidance for customers looking to migrate from Oracle to PostgreSQL using DMS Schema Conversion.

DMS Schema Conversion automatically converts your source Oracle database schemas and most of the database code objects to a format compatible with PostgreSQL. This conversion includes tables, views, stored procedures, functions, data types, synonyms, and so on. Any objects that DMS Schema Conversion can't convert automatically are clearly marked. To complete the migration, you can convert these objects manually.

At a high level, DMS Schema Conversion operates with the following three components: instance profiles, data providers, and migration projects. An instance profile specifies network and security settings. A data provider stores database connection credentials. A migration project contains data providers, an instance profile, and migration rules. AWS DMS uses data providers and an instance profile to design a process that converts database schemas and code objects.

The following diagram illustrates the DMS Schema Conversion process.



Start the walkthrough by [creating the required resources](#).

## Prerequisites for migrating Oracle databases to Amazon Aurora PostgreSQL with DMS schema conversion

The following prerequisites are also required to complete this walkthrough:

- Familiarity with the AWS Management Console, AWS Database Migration Service, and SQL.
- A user with AWS Identity and Access Management (IAM) credentials. Make sure that you can use these credentials to create an Amazon S3 bucket in your AWS Region.
- Basic knowledge of the Amazon Virtual Private Cloud (Amazon VPC) service and of security groups.
- An understanding of the supported features and limitations of DMS Schema Conversion. For more information, see [Schema conversion limitations](#).

We recommend that you don't use your production workloads for the migration in this walkthrough. After you get familiar with migration tools and AWS services, you can migrate your production workloads.

Make sure that you create all your AWS and DMS Schema Conversion resources in the AWS Regions that support DMS Schema Conversion. For more information, see the [list of supported Regions](#). In other Regions, you can use the AWS Schema Conversion Tool. For an example of migration from Oracle to PostgreSQL with AWS SCT, see [Use Schema Conversion Tool to Convert the Oracle Schema to PostgreSQL](#).

For more information about DMS Schema Conversion, see the [user guide](#).

## Step-by-step Oracle databases to Amazon Aurora PostgreSQL with DMS schema conversion migration walkthrough

In the following sections, you can find step-by-step instructions for migrating your Oracle database to Amazon RDS for PostgreSQL using DMS Schema Conversion.

### Topics

- [Step 1: Create AWS Resources](#)
- [Step 2: Configure Your Source Database](#)
- [Step 3: Create Your Target Amazon RDS for PostgreSQL Database](#)

- [Step 4: Store Database Credentials in AWS Secrets Manager](#)
- [Step 5: Create an Instance Profile](#)
- [Step 6: Configure Data Providers](#)
- [Step 7: Create a Migration Project](#)
- [Step 8: Convert Database Objects](#)
- [Step 9: Edit and Apply Your Converted Code](#)

## Step 1: Create AWS Resources

In this step, you create and configure the required AWS resources for DMS Schema Conversion.

First, you create a virtual private cloud (VPC). This VPC is based on the Amazon Virtual Private Cloud (Amazon VPC) service and contains your AWS resources. Make sure that you create this VPC in one of the AWS Regions that support DMS Schema Conversion. For more information, see the [list of supported Regions](#).

### To create a VPC for DMS Schema Conversion

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. Choose your AWS Region.
3. Choose **Create VPC**.
4. On the **Create VPC** page, enter the following settings:
  - **Resources to create** — **VPC and more**
  - **Name tag auto-generation** — Choose **Auto-generate** and enter a globally unique name. For example, enter `sc-vpc`.
  - **IPv4 CIDR block** — `10.0.1.0/24`
  - **NAT gateways** — **In 1 AZ**
  - **VPC endpoints** — **None**
5. Keep the rest of the settings as they are, and then choose **Create VPC**.
6. Choose **Subnets**.
  - For **Filter by VPC**, choose **sc-vpc**.
  - Take a note of your two private subnet IDs. Private subnet IDs don't include `Public` in the name.

## 7. Choose **NAT gateways**.

- Choose your **NAT gateway**.
- Take a note of your **Elastic IP** address.

Use this VPC when you create your instance profile in [Step 5](#) and your target Amazon RDS database in [Step 3](#).

Next, you create AWS Identity and Access Management (IAM) roles to use in your DMS Schema Conversion migration project. AWS DMS uses this IAM role to access your Amazon S3 bucket and database credentials stored in AWS Secrets Manager.

### To create an IAM role that provides access to your Amazon S3 bucket

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the navigation pane, choose **Roles**.
3. Choose **Create role**.
4. On the **Select trusted entity** page, choose **AWS service**. Choose **DMS**.
5. Choose **Next**. The **Add permissions** page opens.
6. For **Filter policies**, enter S3. Choose **AmazonS3FullAccess**.
7. Choose **Next**. The **Name, review, and create** page opens.
8. For **Role name**, enter a descriptive name. For example, enter `sc-s3-role`. Choose **Create role**.
9. On the **Roles** page, enter `sc-s3-role` for **Role name**. Choose **sc-s3-role**.
10. On the `sc-s3-role` page, choose the **Trust relationships** tab. Choose **Edit trust policy**.
11. On the **Edit trust policy** page, edit the trust relationships for the role to use the `schema-conversion.dms.amazonaws.com` service principal as the trusted entity.
12. Choose **Update trust policy**.

Use this IAM role when you create your instance profile in [Step 5](#).

### To create an IAM role that provides access to AWS Secrets Manager

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.



2. In the navigation pane, choose **Roles**.
3. Choose **Create role**.
4. On the **Select trusted entity** page, choose **AWS service**. Choose **DMS**.
5. Choose **Next**. The **Add permissions** page opens.
6. For **Filter policies**, enter `Secret`. Choose **SecretsManagerReadWrite**.
7. Choose **Next**. The **Name, review, and create** page opens.
8. For **Role name**, enter a descriptive name. For example, enter `sc-secrets-manager-role`. Choose **Create role**.
9. On the **Roles** page, enter `sc-secrets-manager-role` for **Role name**. Choose **sc-secrets-manager-role**.
- 10 On the **sc-secrets-manager-role** page, choose the **Trust relationships** tab. Choose **Edit trust policy**.
- 11 On the **Edit trust policy** page, edit the trust relationships for the role to use `schema-conversion.dms.amazonaws.com` and your AWS DMS regional service principal as the trusted entities. This principal has the following format.

```
dms.region-name.amazonaws.com
```

Replace *region-name* with the name of your Region, such as `us-east-1`.

The following code example shows the principal for the `us-east-1` Region.

```
dms.us-east-1.amazonaws.com
```

- 12 Choose **Update trust policy**.

Use this IAM role when you create your migration project in [Step 7](#).

Next, you create an Amazon S3 bucket to use in your DMS Schema Conversion migration project. DMS Schema Conversion uses this Amazon S3 bucket to save assessment reports, SQL scripts with the converted code, and database metadata.

### To create an Amazon S3 bucket for DMS Schema Conversion

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.

2. Choose **Create bucket**.
3. On the **Create bucket** page, select a globally unique name for your S3 bucket. For example, enter `sc-s3-bucket`.
4. For **AWS Region**, choose your Region.
5. For **Bucket Versioning**, choose **Enable**.
6. Keep the rest of the settings as they are, and then choose **Create bucket**.

Use this Amazon S3 bucket when you create your instance profile in [Step 5](#).

## Step 2: Configure Your Source Database

In this step, you configure a new database user on your source Oracle database. Also, you configure the network to set up interaction for your source database with DMS Schema Conversion.

Use the credentials of this new user in DMS Schema Conversion. We encourage not using the admin user in the DMS Schema Conversion migration project.

Make sure that you grant the following privileges to this new user to complete the migration:

- **CONNECT** — includes only the **CREATE SESSION** privilege.
- **SELECT\_CATALOG\_ROLE** — provides **SELECT** privileges on all data dictionary views for Oracle 10g users.
- **SELECT ANY DICTIONARY** — provides query access to any object in the **SYS** schema.

You can use the following code example to create a database user and grant the privileges.

```
CREATE USER user_name IDENTIFIED BY your_password;  
GRANT CONNECT TO user_name;  
GRANT SELECT_CATALOG_ROLE TO user_name;  
GRANT SELECT ANY DICTIONARY TO user_name;
```

In the preceding example, replace `user_name` with the name of your user. Then, replace `your_password` with a secure password.

After you configure your database user, make sure that DMS Schema Conversion can access your source Oracle database. To set up a network for DMS Schema Conversion, you can use different network configurations. These configurations depend on the settings of your source database and

your network. For more information about available options, see [Setting up a network for DMS Schema Conversion](#).

In this walkthrough, you configure a Site-to-Site VPN connection using a virtual private gateway.

### To configure a Site-to-Site VPN connection

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. Choose your AWS Region.
3. Create a customer gateway.
  - In the navigation pane, choose **Customer gateways**, and then **Create customer gateway**.
  - For **Name tag**, enter a name for your customer gateway.
  - For **BGP ASN**, enter a Border Gateway Protocol (BGP) Autonomous System Number (ASN) for your customer gateway.
  - For **IP address**, enter the static, internet-routable IP address for your customer gateway device.
  - For **Certificate ARN**, choose the Amazon Resource Name of the private certificate.
  - For **Device**, enter a name for the device that hosts this customer gateway.
4. Create a virtual private gateway.
  - In the navigation pane, choose **Virtual private gateways**, and then **Create virtual private gateway**.
  - For **Name tag**, enter a name for your virtual private gateway.
  - For **Autonomous System Number (ASN)**, choose **Amazon default ASN**.
  - Choose **Create virtual private gateway**.
  - Select the virtual private gateway you created, choose **Actions**, and then **Attach to VPC**.
  - Under **Available VPCs**, select your VPC from the list and choose **Attach to VPC**.
5. Configure route propagation in your route table.
  - In the navigation pane, choose **Route tables**, and then select the route table that is associated with your subnet. By default, this is the main route table for the VPC.
  - On the **Route propagation** tab in the details pane, choose **Edit route propagation**.
  - Select the virtual private gateway that you created before, and then choose **Save**.

- In the navigation pane, choose **Security groups**, and then select the default security group for your VPC.
  - On the **Inbound** tab in the details pane, add rules that allow inbound SSH, RDP, and ICMP access from your network.
  - Choose **Save**.
7. Create a Site-to-Site VPN connection.
- In the navigation pane, choose **Site-to-Site VPN connections**, and then **Create VPN connection**.
  - For **Name tag**, enter a name for your Site-to-Site VPN connection.
  - For **Target gateway type**, choose either **Virtual private gateway**.
  - For **Customer gateway**, select **Existing**.
  - For **Customer gateway ID**, choose the customer gateway that you created before.
  - Select the routing option. If your customer gateway device supports BGP, then choose **Dynamic (requires BGP)**. Alternatively, choose **Static** and specify IP prefixes for the private network of your Site-to-Site VPN connection.
  - For **Outside IP address type**, keep the default option.
  - Choose **Create VPN connection**.
8. Download the configuration file.
- In the navigation pane, choose **Site-to-Site VPN connections**, and then **Download configuration**.
  - Select the vendor, platform, software, and IKE version that correspond to your customer gateway device. If your device isn't listed, choose **Generic**.
  - Choose **Download**.
9. Use the sample configuration file to configure your customer gateway device.

### Step 3: Create Your Target Amazon RDS for PostgreSQL Database

In this step, you create a new Amazon RDS for PostgreSQL database to use as a migration target for DMS Schema Conversion. Also, you configure a new database user on your target Amazon RDS for PostgreSQL database.

If you already created the target database, skip this step and proceed with the configuration of your database user.

## To create an Amazon RDS for PostgreSQL database for DMS Schema Conversion

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. Choose your AWS Region.
3. Choose **Create database**.
4. For **Engine type**, choose **PostgreSQL**.
5. For **Templates**, choose **Free tier**.
6. For **DB instance identifier**, enter a unique name for your PostgreSQL database.
7. For **Master password** and **Confirm master password**, enter a secure password that includes at least 8 printable characters.
8. For **Virtual private cloud (VPC)** under **Connectivity**, choose `sc-vpc`. You created this VPC in [Step 1](#).
9. For **Public access**, choose **Yes**.
10. Keep the rest of the settings as they are, and then choose **Create database**.

After you create your Amazon RDS for PostgreSQL database, configure a new database user. Then, use the credentials of this user in DMS Schema Conversion. We encourage not using the admin user in the DMS Schema Conversion migration project.

To configure your target database user, create a new user and grant `CREATE ON DATABASE` and the `rds_superuser` role.

You can use the following code example to create a database user and grant the privileges.

```
CREATE ROLE user_name LOGIN PASSWORD 'your_password';
GRANT CREATE ON DATABASE db_name TO user_name;
GRANT rds_superuser TO user_name;
ALTER DATABASE db_name OWNER TO user_name;
```

In the preceding example, replace *user\_name* with the name of your user. Then, replace *your\_password* with a secure password. Finally, replace *db\_name* with the name of your target Amazon RDS for PostgreSQL database.

## Step 4: Store Database Credentials in AWS Secrets Manager

To connect to your source and target databases with DMS Schema Conversion, store your database credentials in AWS Secrets Manager. Make sure that you replicate these secrets to your AWS Region.

### To store your source database credentials in AWS Secrets Manager

1. Sign in to the AWS Management Console and open the AWS Secrets Manager console at <https://console.aws.amazon.com/secretsmanager/>.
2. Choose your AWS Region.
3. Choose **Store a new secret**. The **Choose secret type** page opens.
4. For **Secret type**, choose **Credentials for other database**.
5. For **User name** and **Password**, enter the credentials of the database user that you created for your source database in [Step 2](#).
6. For **Database**, choose **Oracle**.
7. For **Server name**, **Database name**, and **Port**, enter your Oracle database connection information.
8. Choose **Next**. The **Configure secret** page opens.
9. For **Secret name**, enter `sc-oracle-secret`.
10. Choose **Next**. The **Configure rotation** page opens.
11. Choose **Next**. The **Review** page opens.
12. Choose **Store**.

### To store your target database credentials in AWS Secrets Manager

1. Sign in to the AWS Management Console and open the AWS Secrets Manager console at <https://console.aws.amazon.com/secretsmanager/>.
2. Choose your AWS Region.
3. Choose **Store a new secret**. The **Choose secret type** page opens.
4. For **Secret type**, choose **Credentials for Amazon RDS database**.
5. For **User name** and **Password**, enter the credentials of the database user that you created for your target database in [Step 3](#).
6. For **Database**, choose your Amazon RDS for PostgreSQL DB instance.

7. Choose **Next**. The **Configure secret** page opens.
8. For **Secret name**, enter `sc-postgresql-secret`.
9. Choose **Next**. The **Configure rotation** page opens.
10. Choose **Next**. The **Review** page opens.
11. Choose **Store**.

Use these secrets when you create your migration project in [Step 7](#).

## Step 5: Create an Instance Profile

Before you create an instance profile, configure a subnet group for your instance profile.

### To create a subnet group

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. In the navigation pane, choose **Subnet groups**, and then choose **Create subnet group**.
4. For **Name**, enter `PrivateSubnetGroup`.
5. For **Description**, enter `A group of private subnets`.
6. For **VPC**, choose `sc-vpc`. You created this VPC in [Step 1](#).
7. For **Add subnets**, choose two private subnet IDs. You noted these private subnet IDs in [Step 1](#).
8. Choose **Create subnet group**.

Before you create your migration project in DMS Schema Conversion, you set up an instance profile. An instance profile specifies network and security settings for DMS Schema Conversion.

### To create an instance profile

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. In the navigation pane, choose **Instance profiles**, and then choose **Create instance profile**.
4. For **Name**, enter a unique name for your instance profile. For example, enter `sc-instance`.
5. For **Virtual private cloud (VPC)**, choose `sc-vpc`. You created this VPC in [Step 1](#).

6. For **Subnet group**, choose the PrivateSubnetGroup subnet group that you created before.
7. For **S3 bucket** under **Schema conversion settings - optional**, choose an Amazon S3 bucket that you created in [Step 1](#).
8. For **IAM role**, choose the AWS Identity and Access Management (IAM) role that grants access to Amazon S3. You created this role in [Step 1](#).
9. Choose **Create instance profile**.

Use this instance profile when you create your migration project in [Step 7](#).

## Step 6: Configure Data Providers

In this step, you create data providers that describe your source and target databases. A data provider stores a data store type and the location information about your database. Data providers don't include database credentials. You store database credentials in AWS Secrets Manager. Make sure that you include data providers and database secrets in your DMS Schema Conversion migration project.

You can create only one data provider for a single database. If you try to create a second data provider for the same database, DMS Schema Conversion displays an error message. However, you can use one data provider in multiple migration projects.

### To create a data provider for your Oracle database

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. In the navigation pane, choose **Data providers**, and then choose **Create data provider**.
4. For **Configuration**, choose **Enter manually**.
5. For **Name**, enter a unique name for your source data provider. For example, enter `sc-oracle`.
6. For **Engine type**, choose **Oracle**.
7. For **Server name**, enter the Domain Name Service (DNS) name or IP address of your database server.
8. For **Port**, enter the port used to connect to your database server.
9. For **Service ID (SID) or service name**, enter the Oracle System ID (SID). To find the Oracle SID, submit the following query to your Oracle database:



```
SELECT sys_context('userenv','instance_name') AS SID FROM dual;
```

10 For **Secure Socket Layer (SSL) mode**, choose **none**. Optionally, choose the type of your SSL enforcement, and provide the certificate information.

11 Choose **Create data provider**.

### To create a data provider for your Amazon RDS for PostgreSQL database

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. In the navigation pane, choose **Data providers**, and then choose **Create data provider**.
4. For **Configuration**, choose **RDS database instance**.
5. For **Database from RDS**, choose your Amazon RDS for PostgreSQL database.
6. For **Name**, enter a unique name for your target data provider. For example, enter `sc-postgresql`.
7. Choose **Create data provider**.

Use these data providers when you create your migration project in [Step 7](#).

## Step 7: Create a Migration Project

Now you can create a migration project which is the foundation of your work with DMS Schema Conversion. A migration project describes your source and target data providers, your instance profile, and migration rules.

To create a migration project

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. Choose **Migration projects**, and then choose **Create migration project**.
4. For **Name**, enter a unique name for your migration project. For example, enter `sc-project`.
5. For **Instance profile**, choose `sc-instance`. You created this instance profile in [Step 5](#).

6. For **Source**, choose **Browse**, and then choose `sc-oracle`. You created this data provider in [Step 6](#).
7. For **Secret ID**, choose `sc-oracle-secret`. You created this secret in [Step 4](#).
8. For **IAM role**, choose `sc-secrets-manager-role`. You created this role in [Step 1](#).
9. For **Target**, choose **Browse**, and then choose `sc-postgresql`. You created this data provider in [Step 6](#).
10. For **Secret ID**, choose `sc-postgresql-secret`. You created this secret in [Step 4](#).
11. For **IAM role**, choose `sc-secrets-manager-role`. You created this role in [Step 1](#).
12. Choose **Create migration project**.

Use this migration project to convert your Oracle database schemas to PostgreSQL.

## Step 8: Convert Database Objects

After you create the migration project, you can convert your Oracle database schemas to PostgreSQL. To start working with your migration project, you launch DMS Schema Conversion.

The first launch of DMS Schema Conversion requires some setup. AWS Database Migration Service (AWS DMS) starts a schema conversion instance, which can take 10-15 minutes. This process also reads the metadata from the source and target databases. After a successful first launch, you can access DMS Schema Conversion instantly.

### To convert your source Oracle database schema with DMS Schema Conversion

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. Choose **Migration projects**. The **Migration projects** page opens.
4. Choose `sc-project`, and then choose **Schema conversion**.
5. Choose **Launch schema conversion**. If you launch schema conversion for the first time, then the notification appears. Choose **Launch**. The **Schema conversion** page opens. DMS Schema Conversion displays your source database schema in the left pane in a tree-view format.
6. In the source database pane, select the check box for the schema name.
7. Choose this schema in the left pane of the migration project. DMS Schema Conversion highlights the schema name in blue and activates the **Actions** menu.

8. For **Actions**, choose **Convert schema**. The conversion dialog box appears.
9. Choose **Convert** in the dialog box to confirm your choice.

After DMS Schema Conversion completes the conversion, you can review the converted code. After you choose a database object in the left pane of your project, DMS Schema Conversion automatically displays the source converted code for this object.

DMS Schema Conversion stores the converted code in your migration project and doesn't apply these code changes to your target database. You can apply the converted code in DMS Schema Conversion. Alternatively, you can save the converted code as a SQL script, edit it, and then apply to your target database. For more information, see [Step 9](#).

In the settings of your migration project, you can customize your schema conversion view. Also, you can change conversion settings to improve the performance of converted code.

### To edit the settings of your DMS Schema Conversion migration project

1. In the AWS DMS console, choose **Migration projects**. The **Migration projects** page opens.
2. Choose your migration project. Choose **Schema conversion**, then **Launch schema conversion**.
3. Choose **Settings**. The **Settings** page opens.
4. Change the settings to customize the schema conversion view. For more information, see [Specifying migration project settings](#).
5. Change the settings to improve the performance of converted code. For more information, see [Specifying Oracle to PostgreSQL conversion settings](#).
6. Choose **Apply**, and then choose **Schema conversion**.

After you change the settings, convert your source code again.

## Step 9: Edit and Apply Your Converted Code

After you convert your source Oracle database objects, you can review the conversion statistics. DMS Schema Conversion converts most of the database objects, but some of the objects require manual conversion.

DMS Schema Conversion displays the objects that require manual conversion in the **Action items** tab. To convert these objects, you can save the converted code as a SQL script. Then you can edit it

using your code editor and apply these scripts to your target database. Alternatively, you can apply the converted code as is to your target database and make the edits later.

### To save the converted code as a SQL script

1. In the target database pane, choose the converted database schema.
2. Select the check box for the name of this schema. DMS Schema Conversion highlights the schema name in blue and activates the **Actions** menu.
3. For **Actions**, choose **Save as SQL**. The **Save** dialog box appears.
4. Choose **Save as SQL** to confirm your choice.
5. Choose **S3 bucket**. The **Amazon S3** console opens.
6. Choose **Download** to save your SQL scripts.

### To apply the converted code to your target database

1. In the target database pane, choose the converted database schema.
2. Select the check box for the name of this schema. DMS Schema Conversion highlights the schema name in blue and activates the **Actions** menu.
3. For **Actions**, choose **Apply changes**. The **Apply changes** dialog box appears.
4. Choose **Apply** to confirm your choice.

Now you have successfully converted your source Oracle database schemas to PostgreSQL. To complete the database migration, move your data and connect your applications to the new database.

## Migration from Oracle databases to Amazon Aurora PostgreSQL with DMS schema conversion next steps

After you migrate your Oracle database to Amazon RDS for PostgreSQL using DMS Schema Conversion, you can explore several other resources:

- Use AWS DMS to migrate your source data. For more information, see the [Database Migration Service User Guide](#).
- Use DMS Fleet Advisor to inventory your source databases and discover other candidates to move to the cloud. For more information, see the [DMS Fleet Advisor User Guide](#).

- Learn more about Amazon RDS for PostgreSQL. For more information, see the [Amazon Relational Database Service User Guide](#).

After you've finished using DMS Schema Conversion, clean up your resources. Amazon terminates the schema conversion instance that your migration project uses in three days after you complete the conversion. You can retrieve your converted schema and assessment report from the Amazon S3 bucket that you use for DMS Schema Conversion. However, you need to terminate other resources manually.

### To clean up your DMS Schema Conversion resources

- Sign in to the AWS Management Console and open the AWS DMS console.
- In the navigation pane, choose **Migration projects**, and then choose your migration project. Choose **Schema conversion**, and then choose **Stop schema conversion**. Choose **Delete** and confirm your choice.
- Choose **Instance profiles**, and then choose `sc-instance`. Choose **Delete** and confirm your choice.
- Choose **Data providers**, and then select `sc-oracle` and `sc-postgresql`. Choose **Delete** and confirm your choice.

Also, make sure that you delete your Amazon S3 bucket, database secrets in AWS Secrets Manager, IAM roles, and virtual private cloud (VPC).

## Migrating SQL Server databases to Amazon Aurora PostgreSQL with DMS Schema Conversion

This walkthrough gets you started with heterogeneous database migration from Microsoft SQL Server to Amazon Aurora PostgreSQL-Compatible Edition. To automate the migration, we use the AWS DMS Schema Conversion. This service helps assess the complexity of your migration and converts source SQL Server database schemas and code objects to a format compatible with PostgreSQL. Then, you apply the converted code to your target database. This introductory exercise shows how you can use DMS Schema Conversion for this migration.

At a high level, this migration includes the following steps:

- Use the AWS Management Console to do the following:

- Create a VPC in the Amazon VPC console.
- Create IAM roles in the IAM console.
- Create an Amazon S3 bucket in the Amazon S3 console.
- Create your target Aurora PostgreSQL database in the Amazon RDS console.
- Store database credentials in AWS Secrets Manager.
- Use the AWS DMS console to do the following:
  - Create an instance profile for your migration project.
  - Create data providers for your source and target databases.
  - Create a migration project.
- Use DMS Schema Conversion to do the following:
  - Assess the migration complexity and review the migration action items.
  - Convert your source database.
  - Apply the converted code to your target database.

This walkthrough takes approximately three hours to complete. Make sure that you delete resources at the end of this walkthrough to avoid additional charges.


## Migration Overview

This section provides high-level guidance for customers looking to migrate from SQL Server to PostgreSQL using DMS Schema Conversion.

DMS Schema Conversion automatically converts your source SQL Server database schemas and most of the database code objects to a format compatible with PostgreSQL. This conversion includes tables, views, stored procedures, functions, data types, synonyms, and so on. Any objects that DMS Schema Conversion can't convert automatically are clearly marked. To complete the migration, you can convert these objects manually.

At a high level, DMS Schema Conversion operates with the following three components: instance profiles, data providers, and migration projects. An instance profile specifies network and security settings. A data provider stores database connection credentials. A migration project contains data providers, an instance profile, and migration rules. AWS DMS uses data providers and an instance profile to design a process that converts database schemas and code objects.

The following diagram illustrates the DMS Schema Conversion process for this walkthrough.



Start the walkthrough by [creating the required resources](#).

## Prerequisites for migrating SQL Server databases to Aurora PostgreSQL with DMS schema conversion

The following prerequisites are also required to complete this walkthrough:

- Familiarity with the AWS Management Console, AWS Database Migration Service, and SQL.
- A user with AWS Identity and Access Management (IAM) credentials. Make sure that you can use these credentials to create an Amazon S3 bucket in your AWS Region.
- Basic knowledge of the Amazon Virtual Private Cloud (Amazon VPC) service and of security groups.
- An understanding of the supported features and limitations of DMS Schema Conversion. For more information, see [Schema conversion limitations](#).

We recommend that you don't use your production workloads for the migration in this walkthrough. After you get familiar with migration tools and AWS services, you can migrate your production workloads.

Make sure that you create all your AWS and DMS Schema Conversion resources in the AWS Regions that support DMS Schema Conversion. For more information, see the [list of supported Regions](#). In other Regions, you can use the AWS Schema Conversion Tool (AWS SCT). To download AWS SCT, see [Installing, verifying, and updating](#) in the *Schema Conversion Tool User Guide*.

For more information about DMS Schema Conversion, see the [user guide](#).

# Step-by-step SQL Server database to Aurora PostgreSQL with DMS schema conversion migration walkthrough

In the following sections, you can find step-by-step instructions for migrating your SQL Server database to Aurora PostgreSQL using DMS Schema Conversion.

## Topics

- [Step 1: Create AWS Resources](#)
- [Step 2: Configure Your Source Database](#)
- [Step 3: Create Your Target Aurora PostgreSQL Database](#)
- [Step 4: Store Database Credentials in AWS Secrets Manager](#)
- [Step 5: Create an Instance Profile](#)
- [Step 6: Configure Data Providers](#)
- [Step 7: Create a Migration Project](#)
- [Step 8: Convert Database Objects](#)
- [Step 9: Edit and Apply Your Converted Code](#)

## Step 1: Create AWS Resources

In this step, you create and configure the required AWS resources for DMS Schema Conversion.

First, you create a virtual private cloud (VPC). This VPC is based on the Amazon Virtual Private Cloud (Amazon VPC) service and contains your AWS resources. Make sure that you create this VPC in one of the AWS Regions that support DMS Schema Conversion. For more information, see the [list of supported Regions](#).

### To create a VPC for DMS Schema Conversion

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. Choose your AWS Region.
3. Choose **Create VPC**.
4. On the **Create VPC** page, enter the following settings:
  - **Resources to create** — **VPC and more**



- **Name tag auto-generation** — Choose **Auto-generate** and enter a globally unique name. For example, enter `sc-vpc`.
  - **IPv4 CIDR block** — `10.0.1.0/24`
  - **NAT gateways** — **In 1 AZ**
  - **VPC endpoints** — **None**
5. Keep the rest of the settings as they are, and then choose **Create VPC**.
6. Choose **Subnets**.
- For **Filter by VPC**, choose **sc-vpc**.
  - Take a note of your two private subnet IDs. Private subnet IDs don't include `Public` in the name.
7. Choose **NAT gateways**.
- Choose your **NAT gateway**.
  - Take a note of your **Elastic IP** address.

Use this VPC when you create your instance profile in [Step 5](#) and your target Aurora database in [Step 3](#).

Next, you create AWS Identity and Access Management (IAM) roles to use in your DMS Schema Conversion migration project. AWS DMS uses this IAM role to access your Amazon S3 bucket and database credentials stored in AWS Secrets Manager.

### To create an IAM role that provides access to your Amazon S3 bucket

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the navigation pane, choose **Roles**.
3. Choose **Create role**.
4. On the **Select trusted entity** page, choose **AWS service**. Choose **DMS**.
5. Choose **Next**. The **Add permissions** page opens.
6. For **Filter policies**, enter `S3`. Choose **AmazonS3FullAccess**.
7. Choose **Next**. The **Name, review, and create** page opens.
8. For **Role name**, enter a descriptive name. For example, enter `sc-s3-role`. Choose **Create role**.
9. On the **Roles** page, enter `sc-s3-role` for **Role name**. Choose **sc-s3-role**.

- 10 On the `sc-s3-role` page, choose the **Trust relationships** tab. Choose **Edit trust policy**.
- 11 On the **Edit trust policy** page, edit the trust relationships for the role to use the `schema-conversion.dms.amazonaws.com` service principal as the trusted entity.
- 12 Choose **Update trust policy**.

Use this IAM role when you create your instance profile in [Step 5](#).

### To create an IAM role that provides access to AWS Secrets Manager

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the navigation pane, choose **Roles**.
3. Choose **Create role**.
4. On the **Select trusted entity** page, choose **AWS service**. Choose **DMS**.
5. Choose **Next**. The **Add permissions** page opens.
6. For **Filter policies**, enter `Secret`. Choose **SecretsManagerReadWrite**.
7. Choose **Next**. The **Name, review, and create** page opens.
8. For **Role name**, enter a descriptive name. For example, enter `sc-secrets-manager-role`. Choose **Create role**.
9. On the **Roles** page, enter `sc-secrets-manager-role` for **Role name**. Choose **sc-secrets-manager-role**.
- 10 On the `sc-secrets-manager-role` page, choose the **Trust relationships** tab. Choose **Edit trust policy**.
- 11 On the **Edit trust policy** page, edit the trust relationships for the role to use `schema-conversion.dms.amazonaws.com` and your AWS DMS regional service principal as the trusted entities. This principal has the following format.

```
dms.region-name.amazonaws.com
```

Replace *region-name* with the name of your Region, such as `us-east-1`.

The following code example shows the principal for the `us-east-1` Region.

```
dms.us-east-1.amazonaws.com
```

## 12. Choose **Update trust policy**.

Use this IAM role when you create your migration project in [Step 7](#).

Next, you create an Amazon S3 bucket to use in your DMS Schema Conversion migration project. DMS Schema Conversion uses this Amazon S3 bucket to save assessment reports, SQL scripts with the converted code, and database metadata.

### To create an Amazon S3 bucket for DMS Schema Conversion

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. Choose **Create bucket**.
3. On the **Create bucket** page, select a globally unique name for your S3 bucket. For example, enter `sc-s3-bucket`.
4. For **AWS Region**, choose your Region.
5. For **Bucket Versioning**, choose **Enable**.
6. Keep the rest of the settings as they are, and then choose **Create bucket**.

Use this Amazon S3 bucket when you create your instance profile in [Step 5](#).

## Step 2: Configure Your Source Database

In this step, you configure a new database user on your source SQL Server database. Also, you configure the network to set up interaction for your source database with DMS Schema Conversion.

Use the credentials of this new user in DMS Schema Conversion. We encourage not using the admin user in the DMS Schema Conversion migration project.

Make sure that you grant the following privileges to this new user to complete the migration:

- **VIEW DEFINITION** — makes it possible for users that have public access to see object definitions.
- **VIEW DATABASE STATE** — makes it possible for users to check the features of the SQL Server Enterprise edition.

You can use the following code example to create a database user and grant the privileges.

```
USE db_name
CREATE USER user_name FOR LOGIN user_name
GRANT VIEW DEFINITION TO user_name
GRANT VIEW DATABASE STATE TO user_name
```

In the preceding example, replace *user\_name* with the name of your user. Then, replace *db\_name* with a name of your database.

Repeat the grant for each database whose schema you are converting.

Then, make sure that you grant the following privileges on the master database:

- VIEW SERVER STATE — makes it possible for users to collect server settings and configuration.
- VIEW ANY DEFINITION — makes it possible for users to view data providers.

You can use the following code example to create a database user and grant the privileges.

```
USE master
GRANT VIEW SERVER STATE TO user_name
GRANT VIEW ANY DEFINITION TO user_name
```

In the preceding example, replace *user\_name* with the name of your user.

After you configure your database user, make sure that DMS Schema Conversion can access your source SQL Server database. To set up a network for DMS Schema Conversion, you can use different network configurations. These configurations depend on the settings of your source database and your network. For more information about available options, see [Setting up a network for DMS Schema Conversion](#).

In this walkthrough, you configure a Site-to-Site VPN connection using a virtual private gateway.

### To configure a Site-to-Site VPN connection

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. Choose your AWS Region.
3. Create a customer gateway.
  - In the navigation pane, choose **Customer gateways**, and then **Create customer gateway**.

- For **Name tag**, enter a name for your customer gateway.
  - For **BGP ASN**, enter a Border Gateway Protocol (BGP) Autonomous System Number (ASN) for your customer gateway.
  - For **IP address**, enter the static, internet-routable IP address for your customer gateway device.
  - For **Certificate ARN**, choose the Amazon Resource Name of the private certificate.
  - For **Device**, enter a name for the device that hosts this customer gateway.
4. Create a virtual private gateway.
- In the navigation pane, choose **Virtual private gateways**, and then **Create virtual private gateway**.
  - For **Name tag**, enter a name for your virtual private gateway.
  - For **Autonomous System Number (ASN)**, choose **Amazon default ASN**.
  - Choose **Create virtual private gateway**.
  - Select the virtual private gateway you created, choose **Actions**, and then **Attach to VPC**.
  - Under **Available VPCs**, select your VPC from the list and choose **Attach to VPC**.
5. Configure route propagation in your route table.
- In the navigation pane, choose **Route tables**, and then select the route table that is associated with your subnet. By default, this is the main route table for the VPC.
  - On the **Route propagation** tab in the details pane, choose **Edit route propagation**.
  - Select the virtual private gateway that you created before, and then choose **Save**.
6. Add rules to your security group.
- In the navigation pane, choose **Security groups**, and then select the default security group for your VPC.
  - On the **Inbound** tab in the details pane, add rules that allow inbound SSH, RDP, and ICMP access from your network.
  - Choose **Save**.
7. Create a Site-to-Site VPN connection.
- In the navigation pane, choose **Site-to-Site VPN connections**, and then **Create VPN connection**.
  - For **Name tag**, enter a name for your Site-to-Site VPN connection.
  - For **Target gateway type**, choose either **Virtual private gateway**.

- For **Customer gateway**, select **Existing**.
  - For **Customer gateway ID**, choose the customer gateway that you created before.
  - Select the routing option. If your customer gateway device supports BGP, then choose **Dynamic (requires BGP)**. Alternatively, choose **Static** and specify IP prefixes for the private network of your Site-to-Site VPN connection.
  - For **Outside IP address type**, keep the default option.
  - Choose **Create VPN connection**.
8. Download the configuration file.
- In the navigation pane, choose **Site-to-Site VPN connections**, and then **Download configuration**.
  - Select the vendor, platform, software, and IKE version that correspond to your customer gateway device. If your device isn't listed, choose **Generic**.
  - Choose **Download**.
9. Use the sample configuration file to configure your customer gateway device.

### Step 3: Create Your Target Aurora PostgreSQL Database

In this step, you create a new Aurora PostgreSQL database to use as a migration target for DMS Schema Conversion. Also, you configure a new database user on your target Aurora PostgreSQL database.

If you already created the target database, skip this step and proceed with the configuration of your database user.

#### To create an Aurora PostgreSQL database for DMS Schema Conversion

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. Choose your AWS Region.
3. Choose **Create database**.
4. For **Engine type**, choose **Amazon Aurora**.
5. For **Edition**, choose **Amazon Aurora PostgreSQL-Compatible Edition**.
6. For **Templates**, choose **Dev/Test**.
7. For **DB cluster identifier**, enter a unique name for your PostgreSQL database.

8. For **Master password** and **Confirm master password**, enter a secure password that includes at least 8 printable characters.
9. For **Virtual private cloud (VPC)** under **Connectivity**, choose `sc-vpc`. You created this VPC in [Step 1](#).
10. For **Public access**, choose **Yes**.
11. Keep the rest of the settings as they are, and then choose **Create database**.

After you create your Aurora PostgreSQL database, configure a new database user. Then, use the credentials of this user in DMS Schema Conversion. We encourage not using the admin user in the DMS Schema Conversion migration project.

To configure your target database user, create a new user and grant the `CREATE ON DATABASE` and the `rds_superuser` role.

You can use the following code example to create a database user and grant the privileges.

```
CREATE ROLE user_name LOGIN PASSWORD your_password;  
GRANT CREATE ON DATABASE db_name TO user_name;  
GRANT rds_superuser TO user_name;  
ALTER DATABASE db_name OWNER TO user_name;
```

In the preceding example, replace *user\_name* with the name of your user. Then, replace *your\_password* with a secure password. Finally, replace *db\_name* with the name of your target Aurora PostgreSQL database.

## Step 4: Store Database Credentials in AWS Secrets Manager

To connect to your source and target databases with DMS Schema Conversion, store your database credentials in AWS Secrets Manager. Make sure that you replicate these secrets to your AWS Region.

### To store your source database credentials in AWS Secrets Manager

1. Sign in to the AWS Management Console and open the AWS Secrets Manager console at <https://console.aws.amazon.com/secretsmanager/>.
2. Choose your AWS Region.
3. Choose **Store a new secret**. The **Choose secret type** page opens.
4. For **Secret type**, choose **Credentials for other database**.

5. For **User name** and **Password**, enter the credentials of the database user that you created for your source database in [Step 2](#).
6. For **Database**, choose **SQL Server**.
7. For **Server name**, **Database name**, and **Port**, enter your SQL Server database connection information.
8. Choose **Next**. The **Configure secret** page opens.
9. For **Secret name**, enter `sc-sql-server-secret`.
10. Choose **Next**. The **Configure rotation** page opens.
11. Choose **Next**. The **Review** page opens.
12. Choose **Store**.

### To store your target database credentials in AWS Secrets Manager

1. Sign in to the AWS Management Console and open the AWS Secrets Manager console at <https://console.aws.amazon.com/secretsmanager/>.
2. Choose your AWS Region.
3. Choose **Store a new secret**. The **Choose secret type** page opens.
4. For **Secret type**, choose **Credentials for Amazon RDS database**.
5. For **User name** and **Password**, enter the credentials of the database user that you created for your target database in [Step 3](#).
6. For **Database**, choose your Aurora PostgreSQL DB instance.
7. Choose **Next**. The **Configure secret** page opens.
8. For **Secret name**, enter `sc-postgresql-secret`.
9. Choose **Next**. The **Configure rotation** page opens.
10. Choose **Next**. The **Review** page opens.
11. Choose **Store**.

Use these secrets when you create your migration project in [Step 7](#).

## Step 5: Create an Instance Profile

Before you create an instance profile, configure a subnet group for your instance profile.

### To create a subnet group



1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. In the navigation pane, choose **Subnet groups**, and then choose **Create subnet group**.
4. For **Name**, enter PrivateSubnetGroup.
5. For **Description**, enter A group of private subnets.
6. For **VPC**, choose sc-vpc. You created this VPC in [Step 1](#).
7. For **Add subnets**, choose two private subnet IDs. You noted these private subnet IDs in [Step 1](#).
8. Choose **Create subnet group**.

Before you create your migration project in DMS Schema Conversion, you set up an instance profile. An instance profile specifies network and security settings for DMS Schema Conversion.

### To create an instance profile

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. In the navigation pane, choose **Instance profiles**, and then choose **Create instance profile**.
4. For **Name**, enter a unique name for your instance profile. For example, enter sc-instance.
5. For **Virtual private cloud (VPC)**, choose sc-vpc. You created this VPC in [Step 1](#).
6. For **Subnet group**, choose the PrivateSubnetGroup subnet group that you created before.
7. For **S3 bucket** under **Schema conversion settings - optional**, choose an Amazon S3 bucket that you created in [Step 1](#).
8. For **IAM role**, choose the AWS Identity and Access Management (IAM) role that grants access to Amazon S3. You created this role in [Step 1](#).
9. Choose **Create instance profile**.

Use this instance profile when you create your migration project in [Step 7](#).

## Step 6: Configure Data Providers

In this step, you create data providers that describe your source and target databases. A data provider stores a data store type and the location information about your database. Data providers

don't include database credentials. You store database credentials in AWS Secrets Manager. Make sure that you include data providers and database secrets in your DMS Schema Conversion migration project.

You can create only one data provider for a single database. If you try to create a second data provider for the same database, DMS Schema Conversion displays an error message. However, you can use one data provider in multiple migration projects.

### To create a data provider for your SQL Server database

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. In the navigation pane, choose **Data providers**, and then choose **Create data provider**.
4. For **Configuration**, choose **Enter manually**.
5. For **Name**, enter a unique name for your source data provider. For example, enter `sc-sql-server`.
6. For **Engine type**, choose **Microsoft SQL Server**.
7. For **Server name**, enter the Domain Name Service (DNS) name or IP address of your database server.
8. For **Port**, enter the port used to connect to your database server.
9. For **Database name**, enter the name of your database.
10. For **Secure Socket Layer (SSL) mode**, choose **none**. Optionally, choose the type of your SSL enforcement, and provide the certificate information.
11. Choose **Create data provider**.

### To create a data provider for your Aurora PostgreSQL database

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. In the navigation pane, choose **Data providers**, and then choose **Create data provider**.
4. For **Configuration**, choose **RDS database instance**.
5. For **Database from RDS**, choose your Aurora PostgreSQL database.

6. For **Name**, enter a unique name for your target data provider. For example, enter `sc-postgresql`.
7. For **Database name**, enter the name of your database.
8. For **Existing CA certificate**, choose the server certificate. If you don't have any server certificates, then for **Import certificate file**, provide the `rds-ca-2019.pem` file with your certificate.
9. Choose **Create data provider**.

Use these data providers when you create your migration project in [Step 7](#).

## Step 7: Create a Migration Project

Now you can create a migration project which is the foundation of your work with DMS Schema Conversion. A migration project describes your source and target data providers, your instance profile, and migration rules.

To create a migration project

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. Choose **Migration projects**, and then choose **Create migration project**.
4. For **Name**, enter a unique name for your migration project. For example, enter `sc-project`.
5. For **Instance profile**, choose `sc-instance`. You created this instance profile in [Step 5](#).
6. For **Source**, choose **Browse**, and then choose `sc-sql-server`. You created this data provider in [Step 6](#).
7. For **Secret ID**, choose `sc-sql-server-secret`. You created this secret in [Step 4](#).
8. For **IAM role**, choose `sc-secrets-manager-role`. You created this role in [Step 1](#).
9. For **Target**, choose **Browse**, and then choose `sc-postgresql`. You created this data provider in [Step 6](#).
10. For **Secret ID**, choose `sc-postgresql-secret`. You created this secret in [Step 4](#).
11. For **IAM role**, choose `sc-secrets-manager-role`. You created this role in [Step 1](#).
12. Choose **Create migration project**.

Use this migration project to convert your SQL Server database schemas to PostgreSQL.

## Step 8: Convert Database Objects

After you create the migration project, you can convert your SQL Server database schemas to PostgreSQL. To start working with your migration project, you launch DMS Schema Conversion.

The first launch of DMS Schema Conversion requires some setup. AWS Database Migration Service (AWS DMS) starts a schema conversion instance, which can take 10-15 minutes. This process also reads the metadata from the source and target databases. After a successful first launch, you can access DMS Schema Conversion instantly.

### To convert your source SQL Server database schema with DMS Schema Conversion

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. Choose **Migration projects**. The **Migration projects** page opens.
4. Choose `sc-project`, and then choose **Schema conversion**.
5. Choose **Launch schema conversion**. If you launch schema conversion for the first time, then the notification appears. Choose **Launch**. The **Schema conversion** page opens. DMS Schema Conversion displays your source database schema in the left pane in a tree-view format.
6. In the source database pane, select the check box for the schema name.
7. Choose this schema in the left pane of the migration project. DMS Schema Conversion highlights the schema name in blue and activates the **Actions** menu.
8. For **Actions**, choose **Convert schema**. The conversion dialog box appears.
9. Choose **Convert** in the dialog box to confirm your choice.

After DMS Schema Conversion completes the conversion, you can review the converted code. After you choose a database object in the left pane of your project, DMS Schema Conversion automatically displays the source converted code for this object.

DMS Schema Conversion stores the converted code in your migration project and doesn't apply these code changes to your target database. You can apply the converted code in DMS Schema Conversion. Alternatively, you can save the converted code as a SQL script, edit it, and then apply to your target database. For more information, see [Step 9](#).

In the settings of your migration project, you can customize your schema conversion view. Also, you can change conversion settings to improve the performance of converted code.

## To edit the settings of your DMS Schema Conversion migration project

1. In the AWS DMS console, choose **Migration projects**. The **Migration projects** page opens.
2. Choose your migration project. Choose **Schema conversion**, then **Launch schema conversion**.
3. Choose **Settings**. The **Settings** page opens.
4. Change the settings to customize the schema conversion view. For more information, see [Specifying migration project settings](#).
5. Change the settings to improve the performance of converted code. For more information, see [Specifying SQL Server to PostgreSQL conversion settings](#).
6. Choose **Apply**, and then choose **Schema conversion**.

After you change the settings, convert your source code again.

## Step 9: Edit and Apply Your Converted Code

After you convert your source SQL Server database objects, you can review the conversion statistics. DMS Schema Conversion converts most of the database objects, but some of the objects require manual conversion.

DMS Schema Conversion displays the objects that require manual conversion in the **Action items** tab. To convert these objects, you can save the converted code as a SQL script. Then you can edit it using your code editor and apply these scripts to your target database. Alternatively, you can apply the converted code as is to your target database and make the edits later.

### To save the converted code as a SQL script

1. In the target database pane, choose the converted database schema.
2. Select the check box for the name of this schema. DMS Schema Conversion highlights the schema name in blue and activates the **Actions** menu.
3. For **Actions**, choose **Save as SQL**. The **Save** dialog box appears.
4. Choose **Save as SQL** to confirm your choice.
5. Choose **S3 bucket**. The **Amazon S3** console opens.
6. Choose **Download** to save your SQL scripts.

### To apply the converted code to your target database

1. In the target database pane, choose the converted database schema.
2. Select the check box for the name of this schema. DMS Schema Conversion highlights the schema name in blue and activates the **Actions** menu.
3. For **Actions**, choose **Apply changes**. The **Apply changes** dialog box appears.
4. Choose **Apply** to confirm your choice.

Now you have successfully converted your source SQL Server database schemas to PostgreSQL. To complete the database migration, move your data and connect your applications to the new database.

## Migration from SQL Server databases to Aurora PostgreSQL with DMS schema conversion next steps

After you migrate your SQL Server database to Aurora PostgreSQL using DMS Schema Conversion, you can explore several other resources:

- Use AWS DMS to migrate your source data. For more information, see the [Database Migration Service User Guide](#).
- Use DMS Fleet Advisor to inventory your source databases and discover other candidates to move to the cloud. For more information, see the [DMS Fleet Advisor User Guide](#).
- Learn more about Aurora PostgreSQL. For more information, see the [Amazon Aurora User Guide](#).

After you've finished using DMS Schema Conversion, clean up your resources. Amazon terminates the schema conversion instance that your migration project uses in three days after you complete the conversion. You can retrieve your converted schema and assessment report from the Amazon S3 bucket that you use for DMS Schema Conversion. However, you need to terminate other resources manually.

### To clean up your DMS Schema Conversion resources

- Sign in to the AWS Management Console and open the AWS DMS console.
- In the navigation pane, choose **Migration projects**, and then choose your migration project. Choose **Schema conversion**, and then choose **Stop schema conversion**. Choose **Delete** and confirm your choice.
- Choose **Instance profiles**, and then choose `sc-instance`. Choose **Delete** and confirm your choice.

- Choose **Data providers**, and then select `sc-sql-server` and `sc-postgresql`. Choose **Delete** and confirm your choice.

Also, make sure that you delete your Amazon S3 bucket, database secrets in AWS Secrets Manager, IAM roles, and virtual private cloud (VPC).

## Migrating SQL Server databases to Amazon RDS for MySQL with DMS Schema Conversion

This walkthrough gets you started with heterogeneous database migration from Microsoft SQL Server to Amazon RDS for MySQL. To automate the migration, we use the AWS DMS Schema Conversion. This service helps assess the complexity of your migration and converts source SQL Server database schemas and code objects to a format compatible with MySQL. Then, you apply the converted code to your target database. This introductory exercise shows how you can use DMS Schema Conversion for this migration.

At a high level, this migration includes the following steps:

- Use the AWS Management Console to do the following:
  - Create a VPC in the Amazon VPC console.
  - Create IAM roles in the IAM console.
  - Create an Amazon S3 bucket in the Amazon S3 console.
  - Create your target Amazon RDS for MySQL database in the Amazon RDS console.
  - Store database credentials in AWS Secrets Manager.
- Use the AWS DMS console to do the following:
  - Create an instance profile for your migration project.
  - Create data providers for your source and target databases.
  - Create a migration project.
- Use DMS Schema Conversion to do the following:
  - Assess the migration complexity and review the migration action items.
  - Convert your source database.
  - Apply the converted code to your target database.

This walkthrough takes approximately three hours to complete. Make sure that you delete resources at the end of this walkthrough to avoid additional charges.

## Topics

- [Migration overview](#)
- [Prerequisites for migrating SQL Server databases to Amazon RDS for MySQL with DMS schema conversion](#)
- [Step-by-step SQL Server databases to Amazon RDS for MySQL with DMS schema conversion](#)
- [Migration from SQL Server databases to Amazon RDS for MySQL with DMS schema conversion next steps](#)

## Migration overview


This section provides high-level guidance for customers looking to migrate from SQL Server to MySQL using DMS Schema Conversion.

DMS Schema Conversion automatically converts your source SQL Server database schemas and most of the database code objects to a format compatible with MySQL. This conversion includes tables, views, stored procedures, functions, data types, synonyms, and so on. Any objects that DMS Schema Conversion can't convert automatically are clearly marked. To complete the migration, you can convert these objects manually.

At a high level, DMS Schema Conversion operates with the following three components: instance profiles, data providers, and migration projects. An instance profile specifies network and security settings. A data provider stores database connection credentials. A migration project contains data providers, an instance profile, and migration rules. AWS DMS uses data providers and an instance profile to design a process that converts database schemas and code objects.

The following diagram illustrates the DMS Schema Conversion process for this walkthrough.





Start the walkthrough by [creating the required resources](#).

## Prerequisites for migrating SQL Server databases to Amazon RDS for MySQL with DMS schema conversion

The following prerequisites are also required to complete this walkthrough:

- Familiarity with the AWS Management Console, AWS Database Migration Service, and SQL.
- A user with AWS Identity and Access Management (IAM) credentials. Make sure that you can use these credentials to create an Amazon S3 bucket in your AWS Region.
- Basic knowledge of the Amazon Virtual Private Cloud (Amazon VPC) service and of security groups.
- An understanding of the supported features and limitations of DMS Schema Conversion. For more information, see [Schema conversion limitations](#).

We recommend that you don't use your production workloads for the migration in this walkthrough. After you get familiar with migration tools and AWS services, you can migrate your production workloads.

Make sure that you create all your AWS and DMS Schema Conversion resources in the AWS Regions that support DMS Schema Conversion. For more information, see the [list of supported Regions](#). In other Regions, you can use the AWS Schema Conversion Tool (AWS SCT). For an example of migration from SQL Server to MySQL with AWS SCT, see [Use Schema Conversion Tool to Convert the SQL Server Schema to MySQL](#). To download AWS SCT, see [Installing, verifying, and updating](#) in the *Schema Conversion Tool User Guide*.

For more information about DMS Schema Conversion, see the [user guide](#).

# Step-by-step SQL Server databases to Amazon RDS for MySQL with DMS schema conversion

In the following sections, you can find step-by-step instructions for migrating your SQL Server database to Amazon RDS for MySQL using DMS Schema Conversion.

## Topics

- [Step 1: Create AWS Resources](#)
- [Step 2: Configure Your Source Database](#)
- [Step 3: Create Your Target Amazon RDS for MySQL Database](#)
- [Step 4: Store Database Credentials in AWS Secrets Manager](#)
- [Step 5: Create an Instance Profile](#)
- [Step 6: Configure Data Providers](#)
- [Step 7: Create a Migration Project](#)
- [Step 8: Convert Database Objects](#)
- [Step 9: Edit and Apply Your Converted Code](#)

## Step 1: Create AWS Resources

In this step, you create and configure the required AWS resources for DMS Schema Conversion.

First, you create a virtual private cloud (VPC). This VPC is based on the Amazon Virtual Private Cloud (Amazon VPC) service and contains your AWS resources. Make sure that you create this VPC in one of the AWS Regions that support DMS Schema Conversion. For more information, see the [list of supported Regions](#).

### To create a VPC for DMS Schema Conversion

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. Choose your AWS Region.
3. Choose **Create VPC**.
4. On the **Create VPC** page, enter the following settings:
  - **Resources to create** — **VPC and more**

- **Name tag auto-generation** — Choose **Auto-generate** and enter a globally unique name. For example, enter `sc-vpc`.
  - **IPv4 CIDR block** — `10.0.1.0/24`
  - **NAT gateways** — **In 1 AZ**
  - **VPC endpoints** — **None**
5. Keep the rest of the settings as they are, and then choose **Create VPC**.
6. Choose **Subnets**.
- For **Filter by VPC**, choose **sc-vpc**.
  - Take a note of your two private subnet IDs. Private subnet IDs don't include `Public` in the name.
7. Choose **NAT gateways**.
- Choose your **NAT gateway**.
  - Take a note of your **Elastic IP** address.

Use this VPC when you create your instance profile in [Step 5](#) and your target Amazon RDS database in [Step 3](#).

Next, you create AWS Identity and Access Management (IAM) roles to use in your DMS Schema Conversion migration project. AWS DMS uses this IAM role to access your Amazon S3 bucket and database credentials stored in AWS Secrets Manager.

### To create an IAM role that provides access to your Amazon S3 bucket

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the navigation pane, choose **Roles**.
3. Choose **Create role**.
4. On the **Select trusted entity** page, choose **AWS service**. Choose **DMS**.
5. Choose **Next**. The **Add permissions** page opens.
6. For **Filter policies**, enter `S3`. Choose **AmazonS3FullAccess**.
7. Choose **Next**. The **Name, review, and create** page opens.
8. For **Role name**, enter a descriptive name. For example, enter `sc-s3-role`. Choose **Create role**.
9. On the **Roles** page, enter `sc-s3-role` for **Role name**. Choose **sc-s3-role**.

- 10 On the `sc-s3-role` page, choose the **Trust relationships** tab. Choose **Edit trust policy**.
- 11 On the **Edit trust policy** page, edit the trust relationships for the role to use the `schema-conversion.dms.amazonaws.com` service principal as the trusted entity.
- 12 Choose **Update trust policy**.

Use this IAM role when you create your instance profile in [Step 5](#).

### To create an IAM role that provides access to AWS Secrets Manager

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the navigation pane, choose **Roles**.
3. Choose **Create role**.
4. On the **Select trusted entity** page, choose **AWS service**. Choose **DMS**.
5. Choose **Next**. The **Add permissions** page opens.
6. For **Filter policies**, enter `Secret`. Choose **SecretsManagerReadWrite**.
7. Choose **Next**. The **Name, review, and create** page opens.
8. For **Role name**, enter a descriptive name. For example, enter `sc-secrets-manager-role`. Choose **Create role**.
9. On the **Roles** page, enter `sc-secrets-manager-role` for **Role name**. Choose **sc-secrets-manager-role**.
- 10 On the `sc-secrets-manager-role` page, choose the **Trust relationships** tab. Choose **Edit trust policy**.
- 11 On the **Edit trust policy** page, edit the trust relationships for the role to use `schema-conversion.dms.amazonaws.com` and your AWS DMS regional service principal as the trusted entities. This principal has the following format.

```
dms.region-name.amazonaws.com
```

Replace *region-name* with the name of your Region, such as `us-east-1`.

The following code example shows the principal for the `us-east-1` Region.

```
dms.us-east-1.amazonaws.com
```

## 12. Choose **Update trust policy**.

Use this IAM role when you create your migration project in [Step 7](#).

Next, you create an Amazon S3 bucket to use in your DMS Schema Conversion migration project. DMS Schema Conversion uses this Amazon S3 bucket to save assessment reports, SQL scripts with the converted code, and database metadata.

### To create an Amazon S3 bucket for DMS Schema Conversion

1. Sign in to the AWS Management Console and open the Amazon S3 console at <https://console.aws.amazon.com/s3/>.
2. Choose **Create bucket**.
3. On the **Create bucket** page, select a globally unique name for your S3 bucket. For example, enter `sc-s3-bucket`.
4. For **AWS Region**, choose your Region.
5. For **Bucket Versioning**, choose **Enable**.
6. Keep the rest of the settings as they are, and then choose **Create bucket**.

Use this Amazon S3 bucket when you create your instance profile in [Step 5](#).

## Step 2: Configure Your Source Database

In this step, you configure a new database user on your source SQL Server database. Also, you configure the network to set up interaction for your source database with DMS Schema Conversion.

Use the credentials of this new user in DMS Schema Conversion. We encourage not using the admin user in the DMS Schema Conversion migration project.

Make sure that you grant the following privileges to this new user to complete the migration:

- **VIEW DEFINITION** — makes it possible for users that have public access to see object definitions.
- **VIEW DATABASE STATE** — makes it possible for users to check the features of the SQL Server Enterprise edition.

You can use the following code example to create a database user and grant the privileges.

```
USE db_name
CREATE USER user_name FOR LOGIN user_name
GRANT VIEW DEFINITION TO user_name
GRANT VIEW DATABASE STATE TO user_name
```

In the preceding example, replace *user\_name* with the name of your user. Then, replace *db\_name* with a name of your database.

Repeat the grant for each database whose schema you are converting.

Then, make sure that you grant the following privileges on the master database:

- VIEW SERVER STATE — makes it possible for users to collect server settings and configuration.
- VIEW ANY DEFINITION — makes it possible for users to view data providers.

You can use the following code example to create a database user and grant the privileges.

```
USE master
GRANT VIEW SERVER STATE TO user_name
GRANT VIEW ANY DEFINITION TO user_name
```

In the preceding example, replace *user\_name* with the name of your user.

After you configure your database user, make sure that DMS Schema Conversion can access your source SQL Server database. To set up a network for DMS Schema Conversion, you can use different network configurations. These configurations depend on the settings of your source database and your network. For more information about available options, see [Setting up a network for DMS Schema Conversion](#).

In this walkthrough, you configure a Site-to-Site VPN connection using a virtual private gateway.

### To configure a Site-to-Site VPN connection

1. Sign in to the AWS Management Console and open the Amazon VPC console at <https://console.aws.amazon.com/vpc/>.
2. Choose your AWS Region.
3. Create a customer gateway.
  - In the navigation pane, choose **Customer gateways**, and then **Create customer gateway**.

- For **Name tag**, enter a name for your customer gateway.
  - For **BGP ASN**, enter a Border Gateway Protocol (BGP) Autonomous System Number (ASN) for your customer gateway.
  - For **IP address**, enter the static, internet-routable IP address for your customer gateway device.
  - For **Certificate ARN**, choose the Amazon Resource Name of the private certificate.
  - For **Device**, enter a name for the device that hosts this customer gateway.
4. Create a virtual private gateway.
- In the navigation pane, choose **Virtual private gateways**, and then **Create virtual private gateway**.
  - For **Name tag**, enter a name for your virtual private gateway.
  - For **Autonomous System Number (ASN)**, choose **Amazon default ASN**.
  - Choose **Create virtual private gateway**.
  - Select the virtual private gateway you created, choose **Actions**, and then **Attach to VPC**.
  - Under **Available VPCs**, select your VPC from the list and choose **Attach to VPC**.
5. Configure route propagation in your route table.
- In the navigation pane, choose **Route tables**, and then select the route table that is associated with your subnet. By default, this is the main route table for the VPC.
  - On the **Route propagation** tab in the details pane, choose **Edit route propagation**.
  - Select the virtual private gateway that you created before, and then choose **Save**.
6. Add rules to your security group.
- In the navigation pane, choose **Security groups**, and then select the default security group for your VPC.
  - On the **Inbound** tab in the details pane, add rules that allow inbound SSH, RDP, and ICMP access from your network.
  - Choose **Save**.
7. Create a Site-to-Site VPN connection.
- In the navigation pane, choose **Site-to-Site VPN connections**, and then **Create VPN connection**.
  - For **Name tag**, enter a name for your Site-to-Site VPN connection.
  - For **Target gateway type**, choose either **Virtual private gateway**.

- For **Customer gateway**, select **Existing**.
  - For **Customer gateway ID**, choose the customer gateway that you created before.
  - Select the routing option. If your customer gateway device supports BGP, then choose **Dynamic (requires BGP)**. Alternatively, choose **Static** and specify IP prefixes for the private network of your Site-to-Site VPN connection.
  - For **Outside IP address type**, keep the default option.
  - Choose **Create VPN connection**.
8. Download the configuration file.
- In the navigation pane, choose **Site-to-Site VPN connections**, and then **Download configuration**.
  - Select the vendor, platform, software, and IKE version that correspond to your customer gateway device. If your device isn't listed, choose **Generic**.
  - Choose **Download**.
9. Use the sample configuration file to configure your customer gateway device.

### Step 3: Create Your Target Amazon RDS for MySQL Database

In this step, you create a new Amazon RDS for MySQL database to use as a migration target for DMS Schema Conversion. Also, you configure a new database user on your target Amazon RDS for MySQL database.

If you already created the target database, skip this step and proceed with the configuration of your database user.

#### To create an Amazon RDS for MySQL database for DMS Schema Conversion

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. Choose your AWS Region.
3. Choose **Create database**.
4. For **Engine type**, choose **MySQL**.
5. For **Templates**, choose **Free tier**.
6. For **DB instance identifier**, enter a unique name for your MySQL database.



7. For **Master password** and **Confirm master password**, enter a secure password that includes at least 8 printable characters.
8. For **Virtual private cloud (VPC)** under **Connectivity**, choose `sc-vpc`. You created this VPC in [Step 1](#).
9. For **Public access**, choose **Yes**.
10. Keep the rest of the settings as they are, and then choose **Create database**.

After you create your Amazon RDS for MySQL database, configure a new database user. Then, use the credentials of this user in DMS Schema Conversion. We encourage not using the admin user in the DMS Schema Conversion migration project.

To configure your target database user, create a new user and grant the following privileges:

- CREATE ON .
- ALTER ON .
- DROP ON .
- INDEX ON .
- REFERENCES ON .
- SELECT ON .
- CREATE VIEW ON .
- SHOW VIEW ON .
- TRIGGER ON .
- CREATE ROUTINE ON .
- ALTER ROUTINE ON .
- EXECUTE ON .
- CREATE TEMPORARY TABLES ON .
- INVOKE LAMBDA ON .
- INSERT, UPDATE ON AWS\_SQLSERVER\_EXT.\*
- INSERT, UPDATE, DELETE ON AWS\_SQLSERVER\_EXT\_DATA.\*
- CREATE TEMPORARY TABLES ON AWS\_SQLSERVER\_EXT\_DATA.\*

You can use the following code example to create a database user and grant the privileges.

```
CREATE USER 'user_name' IDENTIFIED BY 'your_password';
GRANT CREATE ON *.* TO 'user_name';
GRANT ALTER ON *.* TO 'user_name';
GRANT DROP ON *.* TO 'user_name';
GRANT INDEX ON *.* TO 'user_name';
GRANT REFERENCES ON *.* TO 'user_name';
GRANT SELECT ON *.* TO 'user_name';
GRANT CREATE VIEW ON *.* TO 'user_name';
GRANT SHOW VIEW ON *.* TO 'user_name';
GRANT TRIGGER ON *.* TO 'user_name';
GRANT CREATE ROUTINE ON *.* TO 'user_name';
GRANT ALTER ROUTINE ON *.* TO 'user_name';
GRANT EXECUTE ON *.* TO 'user_name';
GRANT INSERT, UPDATE ON AWS_SQLSERVER_EXT.* TO 'user_name';
GRANT INSERT, UPDATE, DELETE ON AWS_SQLSERVER_EXT_DATA.* TO 'user_name';
GRANT CREATE TEMPORARY TABLES ON AWS_SQLSERVER_EXT_DATA.* TO 'user_name';
```

In the preceding example, replace *user\_name* with the name of your user. Then, replace *your\_password* with a secure password.

## Step 4: Store Database Credentials in AWS Secrets Manager

To connect to your source and target databases with DMS Schema Conversion, store your database credentials in AWS Secrets Manager. Make sure that you replicate these secrets to your AWS Region.

### To store your source database credentials in AWS Secrets Manager

1. Sign in to the AWS Management Console and open the AWS Secrets Manager console at <https://console.aws.amazon.com/secretsmanager/>.
2. Choose your AWS Region.
3. Choose **Store a new secret**. The **Choose secret type** page opens.
4. For **Secret type**, choose **Credentials for other database**.
5. For **User name** and **Password**, enter the credentials of the database user that you created for your source database in [Step 2](#).
6. For **Database**, choose **SQL Server**.
7. For **Server name**, **Database name**, and **Port**, enter your SQL Server database connection information.

8. Choose **Next**. The **Configure secret** page opens.
9. For **Secret name**, enter `sc-sql-server-secret`.
10. Choose **Next**. The **Configure rotation** page opens.
11. Choose **Next**. The **Review** page opens.
12. Choose **Store**.

### To store your target database credentials in AWS Secrets Manager

1. Sign in to the AWS Management Console and open the AWS Secrets Manager console at <https://console.aws.amazon.com/secretsmanager/>.
2. Choose your AWS Region.
3. Choose **Store a new secret**. The **Choose secret type** page opens.
4. For **Secret type**, choose **Credentials for Amazon RDS database**.
5. For **User name** and **Password**, enter the credentials of the database user that you created for your target database in [Step 3](#).
6. For **Database**, choose your Amazon RDS for MySQL DB instance.
7. Choose **Next**. The **Configure secret** page opens.
8. For **Secret name**, enter `sc-mysql-secret`.
9. Choose **Next**. The **Configure rotation** page opens.
10. Choose **Next**. The **Review** page opens.
11. Choose **Store**.

Use these secrets when you create your migration project in [Step 7](#).

## Step 5: Create an Instance Profile

Before you create an instance profile, configure a subnet group for your instance profile.

### To create a subnet group

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.

3. In the navigation pane, choose **Subnet groups**, and then choose **Create subnet group**.
4. For **Name**, enter PrivateSubnetGroup.
5. For **Description**, enter A group of private subnets.
6. For **VPC**, choose sc-vpc. You created this VPC in [Step 1](#).
7. For **Add subnets**, choose two private subnet IDs. You noted these private subnet IDs in [Step 1](#).
8. Choose **Create subnet group**.

Before you create your migration project in DMS Schema Conversion, you set up an instance profile. An instance profile specifies network and security settings for DMS Schema Conversion.

### To create an instance profile

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. In the navigation pane, choose **Instance profiles**, and then choose **Create instance profile**.
4. For **Name**, enter a unique name for your instance profile. For example, enter sc-instance.
5. For **Virtual private cloud (VPC)**, choose sc-vpc. You created this VPC in [Step 1](#).
6. For **Subnet group**, choose the PrivateSubnetGroup subnet group that you created before.
7. For **S3 bucket** under **Schema conversion settings - optional**, choose an Amazon S3 bucket that you created in [Step 1](#).
8. For **IAM role**, choose the AWS Identity and Access Management (IAM) role that grants access to Amazon S3. You created this role in [Step 1](#).
9. Choose **Create instance profile**.

Use this instance profile when you create your migration project in [Step 7](#).

## Step 6: Configure Data Providers

In this step, you create data providers that describe your source and target databases. A data provider stores a data store type and the location information about your database. Data providers don't include database credentials. You store database credentials in AWS Secrets Manager. Make sure that you include data providers and database secrets in your DMS Schema Conversion migration project.

You can create only one data provider for a single database. If you try to create a second data provider for the same database, DMS Schema Conversion displays an error message. However, you can use one data provider in multiple migration projects.

### To create a data provider for your SQL Server database

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. In the navigation pane, choose **Data providers**, and then choose **Create data provider**.
4. For **Configuration**, choose **Enter manually**.
5. For **Name**, enter a unique name for your source data provider. For example, enter `sc-sql-server`.
6. For **Engine type**, choose **Microsoft SQL Server**.
7. For **Server name**, enter the Domain Name Service (DNS) name or IP address of your database server.
8. For **Port**, enter the port used to connect to your database server.
9. For **Database name**, enter the name of your database.
10. For **Secure Socket Layer (SSL) mode**, choose **none**. Optionally, choose the type of your SSL enforcement, and provide the certificate information.
11. Choose **Create data provider**.

### To create a data provider for your Amazon RDS for MySQL database

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. In the navigation pane, choose **Data providers**, and then choose **Create data provider**.
4. For **Configuration**, choose **RDS database instance**.
5. For **Database from RDS**, choose your Amazon RDS for MySQL database.
6. For **Name**, enter a unique name for your target data provider. For example, enter `sc-mysql`.
7. Choose **Create data provider**.

Use these data providers when you create your migration project in [Step 7](#).

## Step 7: Create a Migration Project

Now you can create a migration project which is the foundation of your work with DMS Schema Conversion. A migration project describes your source and target data providers, your instance profile, and migration rules.

To create a migration project

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. Choose **Migration projects**, and then choose **Create migration project**.
4. For **Name**, enter a unique name for your migration project. For example, enter `sc-project`.
5. For **Instance profile**, choose `sc-instance`. You created this instance profile in [Step 5](#).
6. For **Source**, choose **Browse**, and then choose `sc-sql-server`. You created this data provider in [Step 6](#).
7. For **Secret ID**, choose `sc-sql-server-secret`. You created this secret in [Step 4](#).
8. For **IAM role**, choose `sc-secrets-manager-role`. You created this role in [Step 1](#).
9. For **Target**, choose **Browse**, and then choose `sc-mysql`. You created this data provider in [Step 6](#).
10. For **Secret ID**, choose `sc-mysql-secret`. You created this secret in [Step 4](#).
11. For **IAM role**, choose `sc-secrets-manager-role`. You created this role in [Step 1](#).
12. Choose **Create migration project**.

Use this migration project to convert your SQL Server database schemas to MySQL.

## Step 8: Convert Database Objects

After you create the migration project, you can convert your SQL Server database schemas to MySQL. To start working with your migration project, you launch DMS Schema Conversion.

The first launch of DMS Schema Conversion requires some setup. AWS Database Migration Service (AWS DMS) starts a schema conversion instance, which can take 10-15 minutes. This process also reads the metadata from the source and target databases. After a successful first launch, you can access DMS Schema Conversion instantly.

### To convert your source SQL Server database schema with DMS Schema Conversion

1. Sign in to the AWS Management Console and open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. Choose your AWS Region.
3. Choose **Migration projects**. The **Migration projects** page opens.
4. Choose `sc-project`, and then choose **Schema conversion**.
5. Choose **Launch schema conversion**. If you launch schema conversion for the first time, then the notification appears. Choose **Launch**. The **Schema conversion** page opens. DMS Schema Conversion displays your source database schema in the left pane in a tree-view format.
6. In the source database pane, select the check box for the schema name.
7. Choose this schema in the left pane of the migration project. DMS Schema Conversion highlights the schema name in blue and activates the **Actions** menu.
8. For **Actions**, choose **Convert schema**. The conversion dialog box appears.
9. Choose **Convert** in the dialog box to confirm your choice.

After DMS Schema Conversion completes the conversion, you can review the converted code. After you choose a database object in the left pane of your project, DMS Schema Conversion automatically displays the source converted code for this object.

DMS Schema Conversion stores the converted code in your migration project and doesn't apply these code changes to your target database. You can apply the converted code in DMS Schema Conversion. Alternatively, you can save the converted code as a SQL script, edit it, and then apply to your target database. For more information, see [Step 9](#).

In the settings of your migration project, you can customize your schema conversion view. Also, you can change conversion settings to improve the performance of converted code.

### To edit the settings of your DMS Schema Conversion migration project

1. In the AWS DMS console, choose **Migration projects**. The **Migration projects** page opens.
2. Choose your migration project. Choose **Schema conversion**, then **Launch schema conversion**.
3. Choose **Settings**. The **Settings** page opens.
4. Change the settings to customize the schema conversion view. For more information, see [Specifying migration project settings](#).
5. Change the settings to improve the performance of converted code. For more information, see [Specifying SQL Server to MySQL conversion settings](#).

## 6. Choose **Apply**, and then choose **Schema conversion**.

After you change the settings, convert your source code again.

## Step 9: Edit and Apply Your Converted Code

After you convert your source SQL Server database objects, you can review the conversion statistics. DMS Schema Conversion converts most of the database objects, but some of the objects require manual conversion.

DMS Schema Conversion displays the objects that require manual conversion in the **Action items** tab. To convert these objects, you can save the converted code as a SQL script. Then you can edit it using your code editor and apply these scripts to your target database. Alternatively, you can apply the converted code as is to your target database and make the edits later.

### To save the converted code as a SQL script

1. In the target database pane, choose the converted database schema.
2. Select the check box for the name of this schema. DMS Schema Conversion highlights the schema name in blue and activates the **Actions** menu.
3. For **Actions**, choose **Save as SQL**. The **Save** dialog box appears.
4. Choose **Save as SQL** to confirm your choice.
5. Choose **S3 bucket**. The **Amazon S3** console opens.
6. Choose **Download** to save your SQL scripts.

### To apply the converted code to your target database

1. In the target database pane, choose the converted database schema.
2. Select the check box for the name of this schema. DMS Schema Conversion highlights the schema name in blue and activates the **Actions** menu.
3. For **Actions**, choose **Apply changes**. The **Apply changes** dialog box appears.
4. Choose **Apply** to confirm your choice.

Now you have successfully converted your source SQL Server database schemas to MySQL. To complete the database migration, move your data and connect your applications to the new database.



## Migration from SQL Server databases to Amazon RDS for MySQL with DMS schema conversion next steps

After you migrate your SQL Server database to Amazon RDS for MySQL using DMS Schema Conversion, you can explore several other resources:

- Use AWS DMS to migrate your source data. For more information, see the [Database Migration Service User Guide](#).
- Use DMS Fleet Advisor to inventory your source databases and discover other candidates to move to the cloud. For more information, see the [DMS Fleet Advisor User Guide](#).
- Learn more about Amazon RDS for MySQL. For more information, see the [Amazon Relational Database Service User Guide](#).

After you've finished using DMS Schema Conversion, clean up your resources. Amazon terminates the schema conversion instance that your migration project uses in three days after you complete the conversion. You can retrieve your converted schema and assessment report from the Amazon S3 bucket that you use for DMS Schema Conversion. However, you need to terminate other resources manually.

### To clean up your DMS Schema Conversion resources

- Sign in to the AWS Management Console and open the AWS DMS console.
- In the navigation pane, choose **Migration projects**, and then choose your migration project. Choose **Schema conversion**, and then choose **Stop schema conversion**. Choose **Delete** and confirm your choice.
- Choose **Instance profiles**, and then choose `sc-instance`. Choose **Delete** and confirm your choice.
- Choose **Data providers**, and then select `sc-sql-server` and `sc-mysql`. Choose **Delete** and confirm your choice.

Also, make sure that you delete your Amazon S3 bucket, database secrets in AWS Secrets Manager, IAM roles, and virtual private cloud (VPC).

# Migrating an Amazon RDS for Oracle Database to Amazon Redshift

This walkthrough gets you started with heterogeneous database migration from Amazon RDS for Oracle to Amazon Redshift using AWS Database Migration Service (AWS DMS) and the AWS Schema Conversion Tool (AWS SCT). This introductory exercise doesn't cover all scenarios but provides you with a good understanding of the steps involved in such a migration.

It is important to understand that AWS DMS and AWS SCT are two different tools and serve different needs. They don't interact with each other in the migration process. At a high level, the steps involved in this migration are the following:

1. Using AWS SCT to do the following:
  - Run the conversion report for Oracle to Amazon Redshift to identify the issues, limitations, and actions required for the schema conversion.
  - Generate the schema scripts and apply them on the target before performing the data load by using AWS DMS. AWS SCT performs the necessary code conversion for objects like procedures and views.
2. Identify and implement solutions to the issues reported by AWS SCT.
3. Disable foreign keys or any other constraints that might impact the AWS DMS data load.
4. AWS DMS loads the data from source to target using the Full Load approach. Although AWS DMS is capable of creating objects in the target as part of the load, it follows a minimalistic approach to efficiently migrate the data so that it doesn't copy the entire schema structure from source to target.
5. Perform postmigration activities such as creating additional indexes, enabling foreign keys, and making the necessary changes in the application to point to the new database.

This walkthrough uses a custom AWS CloudFormation template to create RDS DB instances for Oracle and Amazon Redshift. It then uses a SQL command script to install a sample schema and data onto the RDS Oracle DB instance that you then migrate to Amazon Redshift.

This walkthrough takes approximately two hours to complete. Be sure to follow the instructions to delete resources at the end of this walkthrough to avoid additional charges.

To estimate what it will cost to run this walkthrough on AWS, you can use the AWS Pricing Calculator. For more information, see <https://calculator.aws/>.

## Topics

- [Prerequisites for migrating from Amazon RDS for Oracle to Amazon Redshift](#)
- [Migration architecture for migrating from Amazon RDS for Oracle to Amazon Redshift](#)
- [Step-by-step Amazon RDS for Oracle to Amazon Redshift migration walkthrough](#)
- [Migration from Amazon RDS for Oracle to Amazon Redshift next steps](#)

## Prerequisites for migrating from Amazon RDS for Oracle to Amazon Redshift

The following prerequisites are also required to complete this walkthrough:

- Familiarity with Amazon RDS, Amazon Redshift, the applicable database technologies, and SQL.
- The custom scripts that include creating the tables to be migrated and SQL queries for confirming the migration, as listed following:
  - `Oracle_Redshift_For_DMSDemo.template` — an AWS CloudFormation template.
  - `Oraclesalesstarschema.sql` — SQL statements to build the **SH** schema.

These scripts are available at the following link: [dms-sbs-RDSOracle2Redshift.zip](#) .

Each step in the walkthrough also contains a link to download the file involved or includes the exact query in the step.

- A user with AWS Identity and Access Management (IAM) credentials that allow you to launch Amazon RDS, AWS Database Migration Service (AWS DMS) instances, and Amazon Redshift clusters in your AWS Region. For information about IAM credentials, see [Setting up for Amazon RDS](#).
- Basic knowledge of the Amazon Virtual Private Cloud (Amazon VPC) service and of security groups. For information about using Amazon VPC with Amazon RDS, see [Virtual Private Clouds \(VPCs\) and Amazon RDS](#). For information about Amazon RDS security groups, see [Amazon RDS Security Groups](#). For information about using Amazon Redshift in a VPC, see [Managing Clusters in an Amazon Virtual Private Cloud \(VPC\)](#).
- An understanding of the supported features and limitations of AWS DMS. For information about AWS DMS, see <https://docs.aws.amazon.com/dms/latest/userguide/Welcome.html>.
- Knowledge of the supported data type conversion options for Oracle and Amazon Redshift. For information about data types for Oracle as a source, see [Using an Oracle database as a source](#).

For information about data types for Amazon Redshift as a target, see [Using an Amazon Redshift Database as a Target](#).

For more information about AWS DMS, see [Getting started with Database Migration Service](#).

## Migration architecture for migrating from Amazon RDS for Oracle to Amazon Redshift

This walkthrough uses AWS CloudFormation to create a simple network topology for database migration that includes the source database, the replication instance, and the target database in the same VPC. For more information about AWS CloudFormation, see the [AWS CloudFormation documentation](#).


We provision the AWS resources that are required for this AWS DMS walkthrough through AWS CloudFormation. These resources include a VPC and Amazon RDS instance for Oracle and an Amazon Redshift cluster. We provision through AWS CloudFormation because it simplifies the process, so we can concentrate on tasks related to data migration. When you create a stack from the AWS CloudFormation template, it provisions the following resources:

- A VPC with CIDR (10.0.0.0/24) with two public subnets in your region, DBSubnet1 at the address 10.0.0.0/26 in Availability Zone (AZ) 1 and DBSubnet2 at the address 10.0.0.64/26, in AZ 12.
- A DB subnet group that includes DBSubnet1 and DBSubnet2.
- Oracle RDS Standard Edition Two with these deployment options:
  - License Included
  - Single-AZ setup
  - db.m3.medium or equivalent instance class
  - Port 1521
  - Default option and parameter groups
- Amazon Redshift cluster with these deployment options:
  - dc1.large
  - Port 5439
  - Default parameter group
- A security group with ingress access from your computer or 0.0.0.0/0 (access from anywhere) based on the input parameter

We have designed the AWS CloudFormation template to require few inputs from the user. It provisions the necessary AWS resources with minimum recommended configurations. However, if you want to change some of the configurations and parameters, such as the VPC CIDR block and Amazon RDS instance types, feel free to update the template.

We use the [AWS Management Console](#) to provision the AWS DMS resources, such as the replication instance, endpoints, and tasks. You install client tools such as SQL Workbench/J and the AWS Schema Conversion Tool (AWS SCT) on your local computer to connect to the Amazon RDS instances.

Following is an illustration of the migration architecture for this walkthrough.



## Step-by-step Amazon RDS for Oracle to Amazon Redshift migration walkthrough

In the following sections, you can find step-by-step instructions for migrating an Amazon RDS for Oracle database to Amazon Redshift. These steps assume that you have already prepared your source database as described in preceding sections.

### Topics

- [Step 1: Launch the RDS Instances in a VPC by Using the AWS CloudFormation Template](#)
- [Step 2: Install the SQL Tools and AWS Schema Conversion Tool on Your Local Computer](#)
- [Step 3: Test Connectivity to the Oracle DB Instance and Create the Sample Schema](#)
- [Step 4: Test the Connectivity to the Amazon Redshift Database](#)
- [Step 5: Use AWS SCT to Convert the Oracle Schema to Amazon Redshift](#)
- [Step 6: Validate the Schema Conversion](#)
- [Step 7: Create an AWS DMS Replication Instance](#)
- [Step 8: Create AWS DMS Source and Target Endpoints](#)
- [Step 9: Create and Run Your AWS DMS Migration Task](#)
- [Step 10: Verify That Your Data Migration Completed Successfully](#)
- [Step 11: Delete Walkthrough Resources](#)

## Step 1: Launch the RDS Instances in a VPC by Using the AWS CloudFormation Template

Before you begin, you'll need to download an AWS CloudFormation template. Follow these instructions:

1. Download the following archive to your computer: <http://docs.aws.amazon.com/dms/latest/sbs/samples/dms-sbs-RDSOracle2Redshift.zip>
2. Extract the AWS CloudFormation template (`Oracle_Redshift_For_DMSDemo.template`) from the archive.
3. Copy and paste the `Oracle_Redshift_For_DMSDemo.template` file into your current directory.

Now you need to provision the necessary AWS resources for this walkthrough.

1. Sign in to the AWS Management Console and open the AWS CloudFormation console at <https://console.aws.amazon.com/cloudformation>.
2. Choose **Create stack**.
3. On the **Select Template** page, choose **Upload a template to Amazon S3**.
4. Click **Choose File**, and then choose the `Oracle_Redshift_For_DMSDemo.template` file that you extracted from the `dms-sbs-RDSOracle2Redshift.zip` archive.

5. Choose **Next**. On the **Specify Details** page, provide parameter values as shown following.

| Parameter                 | Action                                                                                                                                                                                                                                                                                                                                                                       |
|---------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Stack Name</b>         | Enter OracletoRedshiftDWusingDMS .                                                                                                                                                                                                                                                                                                                                           |
| <b>OracleDBName</b>       | Provide a unique name for your database. The name should begin with a letter. The default is ORCL.                                                                                                                                                                                                                                                                           |
| <b>OracleDBUsername</b>   | Specify the admin (DBA) user for managing the Oracle instance. The default is oraadmin.                                                                                                                                                                                                                                                                                      |
| <b>OracleDBPassword</b>   | Provide the password for the admin user. The default is oraadmin123                                                                                                                                                                                                                                                                                                          |
| <b>RedshiftDBName</b>     | Provide any unique name for your database. The name should begin with a letter. The default is test.                                                                                                                                                                                                                                                                         |
| <b>RedshiftDBUsername</b> | Provide the password for the master user. The default is Redshift#123 .                                                                                                                                                                                                                                                                                                      |
| <b>ClientIP</b>           | Specify the IP address in CIDR (x.x.x.x/32) format for your local computer. You can get your IP address from <a href="https://whatismyip.org">whatismyip.org</a> . Your RDS instances' security group will allow ingress to this IP address. The default is access from anywhere (0.0.0.0/0), which is not recommended; you should use your IP address for this walkthrough. |

## Create stack

Select Template  
Specify Details  
Options  
Review

### Specify Details

Specify a stack name and parameter values. You can use or change the default parameter values, which are defined in the AWS CloudFormation template. [Learn more.](#)

Stack name OracletoRedshiftDWusingDMS

### Parameters

#### Source Oracle Database Configuration

OracleDBName ORCL Enter Oracle Database name

OracleDBUsername oraadmin Enter database Admin username for Oracle

OracleDBPassword \*\*\*\*\* Enter password for Oracle Admin user

#### Target Redshift Database Configuration

RedshiftDBName test Enter Redshift DB name

RedshiftDBUsername masteruser Enter master user name for Redshift

RedshiftDBPassword \*\*\*\*\* Enter master user password for Redshift

#### Enter IP address for DB Security group Configuration

ClientIP 0.0.0.0/0

The IP address range that can be used to connect to the RDS instances from your local machine. It must be a valid IP CIDR range of the form x.x.x.x/x. Pls get your address using [checkip.amazonaws.com](http://checkip.amazonaws.com) or [whatsmyip.org](http://whatsmyip.org)

Cancel

Previous

Next

6. Choose **Next**. On the **Options** page, choose **Next**.

7. On the **Review** page, review the details, and if they are correct choose **Create**.



Create stack

Select Template

### Specify Details

Options

## Review

## Review

## Template

|               |                                                                                                                                                                                                                                                                                 |
|---------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Template URL  | <a href="https://t">https://t</a>                                                                                                                                                                                                                                               |
| Description   | This CloudFormation sample template OracleDWtoRedshift_DMS creates a Oracle and Redshift instances in a VPC which can be used to test the datawarehouse migration using AWS DMS service. You will be billed for the AWS resources used if you create a stack from this template |
| Estimate cost | Link is not available                                                                                                                                                                                                                                                           |

### Details

Stack name OracletoRedshiftDWusingDMS

### Source Oracle Database Configuration

```
OracleDBName      ORCL
OracleDBUsername   oraadmin
OracleDBPassword   .....
```

### Target Redshift Database Configuration

```
RedshiftDBName      test
RedshiftDBUsername   masteruser
RedshiftDBPassword   .....
```

Enter IP address for DB Security group Configuration

ClientIP 0.0.0.0/0

## Options

## Tags

No tags provided

### Advanced


|                     |      |
|---------------------|------|
| Notification        |      |
| Timeout             | none |
| Rollback on failure | Yes  |

Cancel

[Previous](#)

Create

8. AWS can take about 20 minutes or more to create the stack with an Amazon RDS for Oracle instance and an Amazon Redshift cluster.



9. After the stack is created, select the **OracletoRedshiftDWusingDMS** stack, and then choose the **Outputs** view. Record the JDBC connection strings, **OracleJDBCConnectionString** and **RedshiftJDBCConnectionString**, for use later in this walkthrough to connect to the Oracle and Amazon Redshift databases.


## Step 2: Install the SQL Tools and AWS Schema Conversion Tool on Your Local Computer

Next, you need to install a SQL client and AWS SCT on your local computer.


This walkthrough assumes you will use the SQL Workbench/J client to connect to the RDS instances for migration validation.

1. Download SQL Workbench/J from [the SQL Workbench/J website](#), and then install it on your local computer. This SQL client is free, open-source, and DBMS-independent.
2. Download the JDBC driver for your Oracle database release. For more information, go to <https://www.oracle.com/jdbc>.
3. Download the Amazon Redshift driver file, `RedshiftJDBC41-1.1.17.1017.jar`, as described following.

- a. Find the Amazon S3 URL to the file in [Previous JDBC Driver Versions](#) of the *Amazon Redshift Cluster Management Guide*.
  - b. Download the driver as described in [Download the Amazon Redshift JDBC Driver](#) of the same guide.
4. Using SQL Workbench/J, configure JDBC drivers for Oracle and Amazon Redshift to set up connectivity, as described following.
- a. In SQL Workbench/J, choose **File**, then choose **Manage Drivers**.
  - b. From the list of drivers, choose **Oracle**.
  - c. Choose the **Open** icon, then choose the `ojdbc7.jar` file that you downloaded in the previous step. Choose **OK**.




- d. From the list of drivers, choose **Redshift**.
- e. Choose the **Open** icon, then choose the Amazon Redshift JDBC driver that you downloaded in the previous step. Choose **OK**.



Next, install AWS SCT and the required JDBC drivers.

1. Download AWS SCT from [Installing, verifying, and updating the Schema Conversion Tool](#).
2. Follow the instructions to install AWS SCT.
3. Launch AWS SCT.
4. In AWS SCT, choose **Global settings** from **Settings**.
5. Choose **Settings, Global settings**, then choose **Drivers**, and then choose **Browse** for **Oracle driver path**. Locate the Oracle JDBC driver and choose **OK**.
6. Choose **Browse** for **Amazon Redshift driver path**. Locate the Amazon Redshift JDBC driver and choose **OK**. Choose **OK** to close the dialog box.



### Step 3: Test Connectivity to the Oracle DB Instance and Create the Sample Schema


After the AWS CloudFormation stack has been created, test the connection to the Oracle DB instance by using SQL Workbench/J and then create the HR sample schema.

1. In SQL Workbench/J, choose **File**, then choose **Connect window**. Create a new connection profile using the following information.

| Parameter               | Action                                                  |
|-------------------------|---------------------------------------------------------|
| <b>New profile name</b> | Enter <code>RDSOracleConnection</code> .                |
| <b>Driver</b>           | Choose <code>Oracle (oracle.jdbc.OracleDriver)</code> . |

| Parameter | Action                                                                                                                                     |
|-----------|--------------------------------------------------------------------------------------------------------------------------------------------|
| URL       | Use the <b>OracleJDBCConnectionString</b> value you recorded when you examined the output details of the DMSdemo stack in a previous step. |
| Username  | Enter oraadmin.                                                                                                                            |
| Password  | Enter oraadmin123 .                                                                                                                        |

- Test the connection by choosing **Test**. Choose **OK** to close the dialog box, then choose **OK** to create the connection profile.



 **Note**

If your connection is unsuccessful, ensure that the IP address you assigned when creating the AWS CloudFormation template is the one you are attempting to connect from. This issue is the most common one when trying to connect to an instance.

3. Create the **SH** schema you will use for migration using a custom `Oraclesalesstarschema.sql` SQL script. To obtain this script, do the following:
  - Download the following archive to your computer: <http://docs.aws.amazon.com/dms/latest/sbs/samples/dms-sbs-RDSOracle2Redshift.zip>
  - Extract the `Oraclesalesstarschema.sql` SQL script from the archive.
  - Copy and paste the `Oraclesalesstarschema.sql` file into your current directory.
    - a. Open the SQL script in a text editor. Copy the entire script.
    - b. In SQL Workbench/J, paste the SQL script in the Default.wksp window showing **Statement 1**.
    - c. Choose **SQL**, then choose **Execute All**.

The screenshot shows the SQL Workbench/J interface with the following SQL statements in the editor:

```

1553 ALTER TABLE "SH"."SALES" MODIFY ("CHANNEL_ID" NOT NULL ENABLE);
1554 ALTER TABLE "SH"."SALES" MODIFY ("TIME_ID" NOT NULL ENABLE);
1555 ALTER TABLE "SH"."SALES" MODIFY ("CUST_ID" NOT NULL ENABLE);
1556 ALTER TABLE "SH"."SALES" MODIFY ("PROD_ID" NOT NULL ENABLE);
1557 -----
1558 -- Ref Constraints for Table SALES
1559 -----
1560
1561 ALTER TABLE "SH"."SALES" ADD CONSTRAINT "SALES_CHANNEL_FK" FOREIGN KEY ("CHANNEL_ID")
1562 REFERENCES "SH"."CHANNELS" ("CHANNEL_ID") ENABLE NOVALIDATE;
1563 ALTER TABLE "SH"."SALES" ADD CONSTRAINT "SALES_CUSTOMER_FK" FOREIGN KEY ("CUST_ID")
1564 REFERENCES "SH"."CUSTOMERS" ("CUST_ID") ENABLE NOVALIDATE;
1565 ALTER TABLE "SH"."SALES" ADD CONSTRAINT "SALES_PRODUCT_FK" FOREIGN KEY ("PROD_ID")
1566 REFERENCES "SH"."PRODUCTS" ("PROD_ID") ENABLE NOVALIDATE;
1567 ALTER TABLE "SH"."SALES" ADD CONSTRAINT "SALES_PROMO_FK" FOREIGN KEY ("PROMO_ID")
1568 REFERENCES "SH"."PROMOTIONS" ("PROMO_ID") ENABLE NOVALIDATE;
1569 -----
1570
1571 -- Collect Statistics of all Tables in SH
1572 -----
1573

```

The Messages pane at the bottom shows the execution results:

```

Execution time: 0.05s
Statement 1293 of 1294 finished

0 rows affected
Table SH.SALES analyzed

Execution time: 0.22s
Statement 1294 of 1294 finished

```

The status bar at the bottom indicates: Ready, if you are | L:1567 C:1 | 1m 40s | Timeout: 0 | Max. Rows: 0

4. Verify the object types and count in **SH** Schema were created successfully by running the following SQL query.

```

Select OBJECT_TYPE, COUNT(*) from dba_OBJECTS where owner='SH'
GROUP BY OBJECT_TYPE;

```

The results of this query should be similar to the following.

| OBJECT_TYPE     | COUNT(*) |
|-----------------|----------|
| INDEX PARTITION | 40       |
| TABLE PARTITION | 8        |
| TABLE           | 5        |
| INDEX           | 15       |

5. Verify the total number of tables and number of rows for each table by running the following SQL query.



```
Select table_name, num_rows from dba_tables where owner='SH' order by 1;
```

The results of this query should be similar to the following.

| TABLE_NAME | NUM_ROWS |
|------------|----------|
| CHANNELS   | 5        |
| CUSTOMERS  | 8        |
| PRODUCTS   | 66       |
| PROMOTIONS | 503      |
| SALES      | 553      |

6. Verify the integrity in tables. Check the number of sales made in different channels by running the following SQL query.

```
Select b.channel_desc,count(*) from SH.SALES a,SH.CHANNELS b where
a.channel_id=b.channel_id
group by b.channel_desc
order by 1;
```

The results of this query should be similar to the following.

| CHANNEL_DESC | COUNT(*) |
|--------------|----------|
| Direct Sales | 710      |
| Internet     | 52       |
| Partners     | 344      |

#### Note

The preceding examples are representative of validation queries. When you perform actual migrations, you should develop similar queries to validate the schema and the data integrity.


## Step 4: Test the Connectivity to the Amazon Redshift Database

Next, test your connection to your Amazon Redshift database.

1. In SQL Workbench/J, choose **File**, then choose **Connect window**. Choose the **Create a new connection profile** icon. Connect to the Amazon Redshift database in SQL Workbench/J by using the information shown following.

| Parameter               | Action                                                                                                                                        |
|-------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------|
| <b>New profile name</b> | Enter RedshiftConnection .                                                                                                                    |
| <b>Driver</b>           | Choose Redshift (com.amazon.redshift.jdbc42.Driver) .                                                                                         |
| <b>URL</b>              | Use the <b>RedshiftJDBCCConnectionString</b> value you recorded when you examined the output details of the DMSdemo stack in a previous step. |
| <b>Username</b>         | Enter redshiftadmin .                                                                                                                         |
| <b>Password</b>         | Enter Redshift#123 .                                                                                                                          |

2. Test the connection by choosing **Test**. Choose **OK** to close the dialog box, then choose **OK** to create the connection profile.



**Note**

If your connection is unsuccessful, ensure that the IP address you assigned when creating the AWS CloudFormation template is the one you are attempting to connect from. This issue is the most common one when trying to connect to an instance.

3. Verify your connectivity to the Amazon Redshift DB instance by running a sample SQL command, such as `select current_date;`.

## Step 5: Use AWS SCT to Convert the Oracle Schema to Amazon Redshift

Before you migrate data to Amazon Redshift, you convert the Oracle schema to an Amazon Redshift schema. [This video covers all the steps of this process.](#)

To convert an Oracle schema to an Amazon Redshift schema using AWS Schema Conversion Tool (AWS SCT), do the following:

1. Launch AWS SCT. In AWS SCT, choose **File**, then choose **New Project**. Create a new project named `DWSchemaMigrationDemoProject`, specify the **Location** of the project folder, and then choose **OK**.
2. Choose **Add source** to add a source Oracle database to your project, then choose **Oracle**, and choose **Next**.
3. Enter the following information, and then choose **Test Connection**.

| Parameter       | Action                                                                                                                                                                                                                                           |
|-----------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Connection name | Enter <code>Oracle DW</code> . AWS SCT displays this name in the tree in the left panel.                                                                                                                                                         |
| Type            | Choose <b>SID</b> .                                                                                                                                                                                                                              |
| Server name     | Use the <b>OracleJDBCConnectionString</b> value you used to connect to the Oracle DB instance, but remove the JDBC prefix information and the port and database name suffix. For example, a sample connection string you use with SQL Workbench/ |

| Parameter   | Action                                                                                                                                                                                                                                                                 |
|-------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|             | J might be "jdbc:oracle:thin:@abc12345678.cqi87654abc.us-west-2.rds.amazonaws.com:1521:ORCL" . For AWS SCT <b>Server name</b> , you remove "jdbc:oracle:thin:@" and ":1521:ORCL" and use just the server name: "abc12345678.cqi87654abc.us-west-2.rds.amazonaws.com" . |
| Server port | Enter 1521.                                                                                                                                                                                                                                                            |
| Oracle SID  | Enter ORCL.                                                                                                                                                                                                                                                            |
| User name   | Enter oraadmin.                                                                                                                                                                                                                                                        |
| Password    | Enter oraadmin123 .                                                                                                                                                                                                                                                    |

**Add source**

**CONNECTION** SSL

Specify parameters for new connections to Oracle

Connection name: Oracle DW

Type: SID

Server name: abc12345678.cqi87654abc.us-west-2.rds.amazonaws.com

Server port: 1521

Oracle SID: ORCL

User name: oraadmin

Password: .....

☐ Use SSL

☐ Store password

Test connection Previous Connect

4. Choose **OK** to close the alert box, then choose **Connect** to close the dialog box and to connect to the Oracle DB instance.
5. Choose **Add target** to add a target Amazon Redshift database to your project, then choose **Amazon Redshift**, and choose **Next**.
6. Enter the following information and then choose **Test Connection**.

| Parameter              | Action                                                                            |
|------------------------|-----------------------------------------------------------------------------------|
| <b>Connection name</b> | Enter Amazon Redshift. AWS SCT displays this name in the tree in the right panel. |
| <b>Type</b>            | Choose <b>SID</b> .                                                               |

| Parameter    | Action                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                       |
|--------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Server name  | Use the <b>RedshiftJDBCConnectionString</b> value you used to connect to the Amazon Redshift cluster, but remove the JDBC prefix information and the port suffix. For example, a sample connection string you use with SQL Workbench/J might be "jdbc:redshift://oracletoreshiftdwusingdms-redshiftcluster-abc123567.abc87654321.us-west-2.redshift.amazonaws.com:5439/test". For AWS SCT <b>Server name</b> , you remove "jdbc:redshift://" and :5439/test" to use just the server name: "oracletoreshiftdwusingdms-redshiftcluster-abc123567.abc87654321.us-west-2.redshift.amazonaws.com" |
| Server port  | Enter 5439.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
| User name    | Enter redshiftadmin .                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |
| Password     | Enter Redshift#123 .                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| Use AWS Glue | Turn off this option.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |

- Choose **OK** to close the alert box, then choose **Connect** to connect to the Amazon Redshift DB instance.
- In the tree in the left panel, select only the **SH** schema. In the tree in the right panel, select your target Amazon Redshift database. Choose **Create mapping**.

▼ Server mappings

| <input type="checkbox"/> | Source path                    | Target path                          |  |
|--------------------------|--------------------------------|--------------------------------------|--|
| <input type="checkbox"/> | //Servers/Oracle DW/Schemas/SH | //Servers/Amazon Redshift/Schemas/sh |  |
| <input type="checkbox"/> |                                |                                      |  |
| <input type="checkbox"/> |                                |                                      |  |
| <input type="checkbox"/> |                                |                                      |  |
| <input type="checkbox"/> |                                |                                      |  |
| <input type="checkbox"/> |                                |                                      |  |
| <input type="checkbox"/> |                                |                                      |  |
| <input type="checkbox"/> |                                |                                      |  |
| <input type="checkbox"/> |                                |                                      |  |

Create mapping
 Delete selected mappings
 Main view

## 9. Choose **Main view**.

10 In the tree in the left panel, right-click the **SH** schema and choose **Collect Statistics**. AWS SCT analyzes the source data to recommend the best keys for the target Amazon Redshift database. For more information, see [Collecting or Uploading Statistics](#).

### Note

If the **SH** schema does not appear in the list, choose **Actions**, then choose **Refresh from Database**.

11 In the tree in the left panel, right-click the **SH** schema and choose **Create report**. AWS SCT analyzes the **SH** schema and creates a database migration assessment report for the conversion to Amazon Redshift.

12 Check the report and the action items it suggests. The report discusses the type of objects that can be converted by using AWS SCT, along with potential migration issues and actions to resolve these issues. For this walkthrough, you should see something like the following:


### Database objects with conversion actions for Amazon Redshift

Of the total 25 database storage object(s) and 49 database code object(s) in the source database, we identified 16 (64%) database storage object(s) and 39 (80%) database code object(s) that can be converted to Amazon Redshift automatically or with minimal changes. We found 1 encrypted object(s).

9 (36%) database storage object(s) require 13 medium and 2 complex user action(s) to complete the conversion.

10 (20%) database code object(s) require 3 medium and 7 complex user action(s) to complete the conversion.

**Figure: Conversion statistics for database storage objects**



13 Review the report summary. To save the report, choose either **Save to CSV** or **Save to PDF**.

14 Choose the **Action Items** tab. The report discusses the type of objects that can be converted by using AWS SCT, along with potential migration issues and actions to resolve these issues.

15 In the tree in the left panel, right-click the **SH** schema and choose **Convert schema**.

16 Choose **Yes** for the confirmation message. AWS SCT then converts your schema to the target database format.

#### **Note**

The choice of the Amazon Redshift sort keys and distribution keys is critical for optimal performance. You can use key management in AWS SCT to customize the choice of keys. For this walkthrough, we use the defaults recommended by AWS SCT. For more information, see [Optimizing Amazon Redshift](#).

17 In the tree in the right panel, choose the converted **sh** schema, and then choose **Apply to database** to apply the schema scripts to the target Amazon Redshift instance.

18 In the tree in the right panel, choose the **sh** schema, and then choose **Refresh from Database** to refresh from the target database.

The database schema has now been converted and imported from source to target.



## Step 6: Validate the Schema Conversion

To validate the schema conversion, you compare the objects found in the Oracle and Amazon Redshift databases using SQL Workbench/J.

1. In SQL Workbench/J, choose **File**, then choose **Connect window**. Choose the **RedshiftConnection** you created in an earlier step. Choose **OK**.
2. Run the following script to verify the number of object types and count in **SH** schema in the target Amazon Redshift database. These values should match the number of objects in the source Oracle database.

```
SELECT 'TABLE' AS OBJECT_TYPE,
       TABLE_NAME AS OBJECT_NAME,
       TABLE_SCHEMA AS OBJECT_SCHEMA
FROM information_schema.TABLES
WHERE TABLE_TYPE = 'BASE TABLE'
AND    OBJECT_SCHEMA = 'sh';
```

The output from this query should be similar to the following.

| object_type | object_name | object_schema |
|-------------|-------------|---------------|
| TABLE       | channels    | sh            |
| TABLE       | customers   | sh            |
| TABLE       | products    | sh            |
| TABLE       | promotions  | sh            |
| TABLE       | sales       | sh            |

3. Verify the sort and distributions keys that are created in the Amazon Redshift cluster by using the following query.

```
set search_path to '$user', 'public', 'sh';

SELECT tablename,
       "column",
       TYPE,
       encoding,
       distkey,
       sortkey,
       "notnull"
FROM pg_table_def
```


```
WHERE (distkey = TRUE OR sortkey <> 0);
```

The results of the query reflect the distribution key (distkey) and sort key (sortkey) choices made by using AWS SCT key management.

| tablename                                 | column              | type                        | encoding | distkey | sortkey | notnull |
|-------------------------------------------|---------------------|-----------------------------|----------|---------|---------|---------|
| -----+-----+-----+-----+-----+-----+----- |                     |                             |          |         |         |         |
| channels                                  | channel_id          | numeric(38,18)              | none     | true    |         |         |
| 1                                         | true                |                             |          |         |         |         |
| customers                                 | cust_id             | numeric(38,18)              | none     | false   |         |         |
| 4                                         | true                |                             |          |         |         |         |
| customers                                 | cust_gender         | character(2)                | none     | false   |         |         |
| 1                                         | true                |                             |          |         |         |         |
| customers                                 | cust_year_of_birth  | smallint                    | none     | false   |         |         |
| 3                                         | true                |                             |          |         |         |         |
| customers                                 | cust_marital_status | character varying(40)       | none     | false   |         |         |
| 2                                         | false               |                             |          |         |         |         |
| products                                  | prod_id             | integer                     | none     | true    |         |         |
| 4                                         | true                |                             |          |         |         |         |
| products                                  | prod_subcategory    | character varying(100)      | none     | false   |         |         |
| 3                                         | true                |                             |          |         |         |         |
| products                                  | prod_category       | character varying(100)      | none     | false   |         |         |
| 2                                         | true                |                             |          |         |         |         |
| products                                  | prod_status         | character varying(40)       | none     | false   |         |         |
| 1                                         | true                |                             |          |         |         |         |
| promotions                                | promo_id            | integer                     | none     | true    |         |         |
| 1                                         | true                |                             |          |         |         |         |
| sales                                     | prod_id             | numeric(38,18)              | none     | false   |         |         |
| 4                                         | true                |                             |          |         |         |         |
| sales                                     | cust_id             | numeric(38,18)              | none     | false   |         |         |
| 3                                         | true                |                             |          |         |         |         |
| sales                                     | time_id             | timestamp without time zone | none     | true    |         |         |
| 1                                         | true                |                             |          |         |         |         |
| sales                                     | channel_id          | numeric(38,18)              | none     | false   |         |         |
| 2                                         | true                |                             |          |         |         |         |
| sales                                     | promo_id            | numeric(38,18)              | none     | false   |         |         |
| 5                                         | true                |                             |          |         |         |         |

## Step 7: Create an AWS DMS Replication Instance

After we validate the schema structure between source and target databases, as described preceding, we proceed to the core part of this walkthrough, which is the data migration. The following illustration shows a high-level view of the migration process.



A DMS replication instance performs the actual data migration between source and target. The replication instance also caches the transaction logs during the migration. How much CPU and memory capacity a replication instance has influences the overall time required for the migration.

To create an AWS DMS replication instance, do the following:

1. Sign in to the AWS Management Console, open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>, and choose **Create Migration**. If you are signed in as an AWS Identity and Access Management (IAM) user, you must have the appropriate permissions to access AWS DMS. For more information about the permissions required, see [IAM Permissions](#).
2. Choose **Create migration** to start a database migration.
3. On the **Welcome** page, choose **Next**.
4. On the **Create replication instance** page, specify your replication instance information as shown following.

| Parameter             | Action                                                                                              |
|-----------------------|-----------------------------------------------------------------------------------------------------|
| <b>Name</b>           | Enter <code>DMSdemo-repserver</code> .                                                              |
| <b>Description</b>    | Enter a brief description, such as <code>DMS demo replication server</code> .                       |
| <b>Instance class</b> | Choose <b>dms.t2.medium</b> . This instance class is large enough to migrate a small set of tables. |

| Parameter           | Action                                                                                                            |
|---------------------|-------------------------------------------------------------------------------------------------------------------|
| VPC                 | Choose <code>OracleToRedshiftusingDMS</code> , which is the VPC that was created by the AWS CloudFormation stack. |
| Multi-AZ            | Choose No.                                                                                                        |
| Publicly accessible | Leave this item selected.                                                                                         |

5. For the **Advanced** section, leave the default settings as they are, and choose **Next**.

## Step 8: Create AWS DMS Source and Target Endpoints

While your replication instance is being created, you can specify the source and target database endpoints using the [AWS Management Console](#). However, you can only test connectivity after the replication instance has been created, because the replication instance is used in the connection.

1. Specify your connection information for the source Oracle database and the target Amazon Redshift database. The following table describes the source settings.

| Parameter           | Action                                                                                                                                                                 |
|---------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Endpoint Identifier | Enter <code>OraSource</code> (the Amazon RDS for Oracle endpoint).                                                                                                     |
| Source Engine       | Choose <b>oracle</b> .                                                                                                                                                 |
| Server name         | Provide the Oracle DB instance name. This name is the <b>Server name</b> value that you used for AWS SCT, such as "abc123567.abc87654321.us-west-2.rds.amazonaws.com". |
| Port                | Enter 1521.                                                                                                                                                            |
| SSL mode            | Choose <b>None</b> .                                                                                                                                                   |
| Username            | Enter <code>oraadmin</code> .                                                                                                                                          |

| Parameter | Action              |
|-----------|---------------------|
| Password  | Enter oraadmin123 . |
| SID       | Enter ORCL.         |

The following table describes the target settings.

| Parameter           | Action                                                                                                                                                                                                                             |
|---------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Endpoint Identifier | Enter Redshifttarget (the Amazon Redshift endpoint).                                                                                                                                                                               |
| Target Engine       | Choose <b>redshift</b> .                                                                                                                                                                                                           |
| Servename           | Provide the Amazon Redshift DB instance name. This name is the <b>Server name</b> value that you used for AWS SCT, such as "oracletooredshiftdwusingdms-redshiftcluster-abc123567.abc87654321.us-west-2.redshift.amazonaws.com" .. |
| Port                | Enter 5439.                                                                                                                                                                                                                        |
| SSL mode            | Choose <b>None</b> .                                                                                                                                                                                                               |
| Username            | Enter redshiftadmin .                                                                                                                                                                                                              |
| Password            | Enter Redshift#123 .                                                                                                                                                                                                               |
| Database name       | Enter test.                                                                                                                                                                                                                        |

The completed page should look like the following.



Replication instance created successfully.

Your database endpoint can be on-premise, in EC2, RDS or in the cloud. Define the connection details below. It is recommended that you test your endpoint connections here to avoid errors later.

#### Source database connection details

Endpoint  
identifier\*



Source engine\*



Server name\*

Port\*

SSL mode\*



User name\*

Password\*

SID\*

► Advanced

#### Target database connection details

Endpoint  
identifier\*



Target engine\*



Server name\*

Port\*

SSL mode\*



User name\*

Password\*

Database name\*

► Advanced

2. Wait for the status to say **Replication instance created successfully..**
3. To test the source and target connections, choose **Run Test** for the source and target connections.
4. Choose **Next**.

## Step 9: Create and Run Your AWS DMS Migration Task

Using an AWS DMS task, you can specify what schema to migrate and the type of migration. You can migrate existing data, migrate existing data and replicate ongoing changes, or replicate data changes only. This walkthrough migrates existing data only.

1. On the **Create Task** page, specify the task options. The following table describes the settings.

| Parameter                   | Action                                                                                              |
|-----------------------------|-----------------------------------------------------------------------------------------------------|
| <b>Task name</b>            | Enter <code>migrateSHschema</code> .                                                                |
| <b>Replication instance</b> | Shows <code>DMSdemo-repserver</code> (the AWS DMS replication instance created in an earlier step). |
| <b>Source endpoint</b>      | Shows <code>orasource</code> (the Amazon RDS for Oracle endpoint).                                  |
| <b>Target endpoint</b>      | Shows <code>redshifttarget</code> (the Amazon Redshift endpoint).                                   |
| <b>Migration type</b>       | Choose <b>Migrate existing data</b> .                                                               |
| <b>Start task on create</b> | Choose this option.                                                                                 |

The page should look like the following.

### Create task

A task can contain one or more table mappings which define what data is moved from the source to the target. If a table does not exist on the target, it can be created automatically.

Task name\*  ⓘ

Replication instance\*  ▼

Source endpoint\*  ▼

Target endpoint\*  ▼

Migration type\*  ▼ ⓘ

Start task on create ☒

2. On the **Task Settings** section, specify the settings as shown in the following table.

| Parameter                          | Action                           |
|------------------------------------|----------------------------------|
| Target table preparation mode      | Choose <b>Do nothing</b> .       |
| Include LOB columns in replication | Choose <b>Limited LOB mode</b> . |
| Max LOB size (kb)                  | Accept the default (32).         |

The section should look like the following.

▼ Task Settings

Target table preparation mode\* ☐ Do nothing ☒ Drop tables on target ☐ Truncate ?

Include LOB columns in replication\* ☐ Don't include LOB columns ☒ Limited LOB mode ?

Max LOB size (kb)\*  ?

Enable logging ☐

[Advanced Settings](#)

3. In the **Selection rules** section, specify the settings as shown in the following table.

| Parameter           | Action                 |
|---------------------|------------------------|
| Schema name is      | Choose Enter a schema. |
| Schema name is like | Enter SH%.             |
| Table name is like  | Enter %.               |
| Action              | Choose Include.        |



The section should look like the following:

▼ Table mappings

GuidedJSON

Selection rules ⓘ

At least one selection rule with an include action is required. Once you have one or more selection rules, you can add transformation rules.

Where ⓘ

Schema name is

Enter a schema

Schema name is like

SH%

Table name is like

%

Use % as a wildcard.

Action

Include ⓘ

Filter ⓘ

Add column filter

Add selection rule

4. Choose **Add selection rule**.
5. Choose **Create task**. The task begins immediately. The **Tasks** section shows you the status of the migration task.

DMS

Dashboard

Get started

Tasks

Endpoints

Certificates

Replication instances

Subnet groups

Create taskModifyStart/ResumeStopDelete

Filter: Filter X

|                                     | ID             | Status   | Source    | Target         | Type      | Complete %   | Elapsed time | Tables loaded |
|-------------------------------------|----------------|----------|-----------|----------------|-----------|--------------|--------------|---------------|
| <input type="checkbox"/>            | migrateschema2 | Creating | orasource | redshifttarget | Full Load | <div>0</div> |              | 0             |
| <input checked="" type="checkbox"/> | migrateschema  | Ready    | orasource | redshifttarget | Full Load | <div>0</div> | 0m           | 0             |

Step 10: Verify That Your Data Migration Completed Successfully

When the migration task completes, you can compare your task results with the expected results.

1. On the navigation pane, choose **Tasks**.

2. Choose your migration task (migrateSHschema).
3. Choose the **Table statistics** tab, shown following.

[Create task](#)
[Modify](#)
[Start/Resume](#)
[Stop](#)
[Delete](#)

Filter:  X

| <input type="checkbox"/>            | ID              | Status    | Source    | Target         | Type      | Complete %                | Elapsed time | Tables loaded |
|-------------------------------------|-----------------|-----------|-----------|----------------|-----------|---------------------------|--------------|---------------|
| <input checked="" type="checkbox"/> | migrateshschema | Modifying | orasource | redshifftarget | Full Load | <div><div>100</div></div> | 1m           | 5             |

migrateshschema

[Overview](#)
[Task monitoring](#)
[Table statistics](#)
[Logs](#)

Filter:  X

| Schema | Table      | State           | Inserts | Deletes | Updates | DDLs | Full Load Rows | Total | Last updated      |
|--------|------------|-----------------|---------|---------|---------|------|----------------|-------|-------------------|
| SH     | CHANNELS   | Table completed | 0       | 0       | 0       | 0    | 5              | 5     | 12/26/16, 9:00 PM |
| SH     | CUSTOMERS  | Table completed | 0       | 0       | 0       | 0    | 8              | 8     | 12/26/16, 9:00 PM |
| SH     | PRODUCTS   | Table completed | 0       | 0       | 0       | 0    | 66             | 66    | 12/26/16, 9:00 PM |
| SH     | PROMOTIONS | Table completed | 0       | 0       | 0       | 0    | 503            | 503   | 12/26/16, 9:00 PM |
| SH     | SALES      | Table completed | 0       | 0       | 0       | 0    | 1,106          | 1,106 | 12/26/16, 9:00 PM |

4. Connect to the Amazon Redshift instance by using SQL Workbench/J, and then check whether the database tables were successfully migrated from Oracle to Amazon Redshift by running the SQL script shown following.

```
select "table", tbl_rows
from svv_table_info
where
SCHEMA = 'sh'
order by 1;
```

Your results should look similar to the following.

```
table      | tbl_rows
-----+-----
channels   |         5
```

|            |  |      |
|------------|--|------|
| customers  |  | 8    |
| products   |  | 66   |
| promotions |  | 503  |
| sales      |  | 1106 |

5. To verify whether the output for tables and number of rows from the preceding query matches what is expected for RDS Oracle, compare your results with those in previous steps.
6. Run the following query to check the relationship in tables; this query checks the departments with employees greater than 10.

```
Select b.channel_desc,count(*) from SH.SALES a,SH.CHANNELS b where
a.channel_id=b.channel_id
group by b.channel_desc
order by 1;
```

The output from this query should be similar to the following.

| channel_desc |  | count |
|--------------|--|-------|
| -----+       |  | ----- |
| Direct Sales |  | 355   |
| Internet     |  | 26    |
| Partners     |  | 172   |

7. Verify column compression encoding.

DMS uses an Amazon Redshift COPY operation to load data. By default, the COPY command applies automatic compression whenever loading to an empty target table. The sample data for this walkthrough is not large enough for automatic compression to be applied. When you migrate larger data sets, COPY will apply automatic compression.

For more details about automatic compression on Amazon Redshift tables, see [Loading Tables with Automatic Compression](#).

To view compression encodings, run the following query.

```
SELECT *
FROM pg_table_def
WHERE schemaname = 'sh';
```

Now you have successfully completed a database migration from an Amazon RDS for Oracle DB instance to Amazon Redshift.

## Step 11: Delete Walkthrough Resources

After you have completed this walkthrough, perform the following steps to avoid being charged further for AWS resources used in the walkthrough. It's necessary that you do the steps in order, because some resources cannot be deleted if they have a dependency upon another resource.

To delete AWS DMS resources, do the following:

1. On the navigation pane, choose **Tasks**, choose your migration task (migratehrschema), and then choose **Delete**.
2. On the navigation pane, choose **Endpoints**, choose the Oracle source endpoint (orasource), and then choose **Delete**.
3. Choose the Amazon Redshift target endpoint (redshifttarget), and then choose **Delete**.
4. On the navigation pane, choose **Replication instances**, choose the replication instance (DMSdemo-repserver), and then choose **Delete**.

Next, you must delete your AWS CloudFormation stack, DMSdemo. Do the following:

1. Sign in to the AWS Management Console and open the AWS CloudFormation console at <https://console.aws.amazon.com/cloudformation>.

If you are signed in as an IAM user, you must have the appropriate permissions to access AWS CloudFormation.

2. Choose your AWS CloudFormation stack, OracletoRedshiftDWusingDMS.
3. For **Actions**, choose **Delete stack**.

The status of the stack changes to DELETE\_IN\_PROGRESS while AWS CloudFormation cleans up the resources associated with the OracletoRedshiftDWusingDMS stack. When AWS CloudFormation is finished cleaning up resources, it removes the stack from the list.

## Migration from Amazon RDS for Oracle to Amazon Redshift next steps

You can explore several other features of AWS DMS that were not included in this walkthrough, including the following:

- The AWS DMS change data capture (CDC) feature, for ongoing replication of data.
- Transformation actions that let you specify and apply transformations to the selected schema or table as part of the migration process.

For more information, see [Getting started with Database Migration Service](#).

## Migrating a BigQuery Project to Amazon Redshift

This walkthrough gets you started with heterogeneous database migration from BigQuery to Amazon Redshift. To automate the migration, we use the AWS Schema Conversion Tool (AWS SCT) that runs on Windows. This introductory exercise provides you with a good understanding of the steps involved in such a migration.

At a high level, the steps involved in this migration are the following:

- Use the Google Cloud management console to do the following:
  - Create a service account, which AWS SCT can use to connect to your source BigQuery project.
  - Create a Cloud Storage bucket to store your source data during migration.
- Use the AWS Management Console to do the following:
  - Create an Amazon Redshift cluster.
  - Create an Amazon Simple Storage Service (Amazon S3) bucket.
- Use AWS SCT to convert source database schemas and apply converted code to your target database.
- Use data extraction agents to migrate data.

To see all the steps of the migration process, [watch this video](#).

This walkthrough takes approximately three hours to complete. Make sure that you delete resources at the end of this walkthrough to avoid additional charges.


### Topics

- [Migration overview](#)
- [Prerequisites for migrating a BigQuery project to Amazon Redshift](#)
- [Step-by-Step BigQuery project to Amazon Redshift migration walkthrough](#)
- [Migration from a BigQuery project to Amazon Redshift next steps](#)

## Migration overview

This section provides high-level guidance for customers looking for a way to migrate from BigQuery to Amazon Redshift. After you complete this introductory exercise, understand the migration process, and become familiar with migration automation tools, plan the migration of your production workloads.

The following illustration demonstrates the migration architecture for this walkthrough.



First, you create a service account to connect to your BigQuery project. Then you create an Amazon Redshift database, as well as the buckets in Cloud Storage and Amazon S3. After this setup, you use AWS SCT to convert source database schemas and apply them to your target database. Finally, you install and configure a data extraction agent to migrate data, upload it to your S3 bucket, and then copy to Amazon Redshift. For big datasets, you can use several data extraction agents to increase the speed of data migration.

To connect to BigQuery, AWS SCT uses the Application Layer Transport Security (ALTS) authentication and encryption system. To connect to Amazon S3 and Amazon Redshift, AWS SCT uses the HTTPS and SSL protocols.

## Migration strategy

For BigQuery to Amazon Redshift migrations, you can use the following typical migration approach.

### 1. Future State Architecture Design

This step defines the architecture of your new system in the target environment. This architecture includes databases, applications, scripts, and so on.

### 2. Database Schema Conversion

You can use AWS SCT to automate the conversion of your source database to Amazon Redshift. For more information, see [Convert Database Schemas](#).

### 3. Application Conversion or Remediation

After you migrate your data storage, make sure that you update your applications. You can use AWS SCT to convert SQL queries in your application code. For more information, see [Converting SQL code in your applications](#).

### 4. Scripts, ETL, Reports Conversion

In addition to applications, make sure that you update all other components of your source system. These include business intelligence reports, extract, transform, and load (ETL) processes, and other scripts.

### 5. Integration with Third-Party Applications

Your applications usually connect to other applications or monitoring tools. Your migration from BigQuery to Amazon Redshift affects these dependencies.

### 6. Data Migration

You can use AWS SCT to manage a data extraction agent that migrates data from BigQuery to Amazon Redshift. For more information, see [Data Extraction Agents](#).

### 7. Testing and Bug Fixing

Migration touches all the stored procedures and functions and affects substantial parts of the application code. For this reason, good testing is required both at the unit and system functional level.

### 8. Performance Tuning

Because of database platform differences and syntax, certain constructs or combinations of data objects can perform differently on the new platform. The performance tuning part of the migration resolves any bottlenecks.

## 9. Setup, DevOps, Integration, Deployment, and Security

Take the opportunity to embrace infrastructure as code for the migration. Make sure that you also focus on the application security. Finally, plan the cutover.

## Security in the AWS Cloud

Cloud security at AWS is the highest priority. As an AWS customer, you benefit from a data center and network architecture that is built to meet the requirements of the most security-sensitive organizations.

AWS is responsible for protecting the global infrastructure that runs all of the AWS Cloud. You are responsible for maintaining control over your content that is hosted on this infrastructure. For more information, see [Shared Responsibility Model](#).

Amazon Redshift protects data with AWS encryption solutions, along with all default security controls within AWS services. Your data is encrypted at rest and in transit. Amazon Redshift automatically integrates with AWS Key Management Service (AWS KMS) for key management. AWS KMS uses envelope encryption. For more information, see [Data protection in Amazon Redshift](#).

Access to Amazon Redshift requires credentials that AWS can use to authenticate your requests. Those credentials must have permissions to access AWS resources, such as an Amazon Redshift cluster. You can use AWS Identity and Access Management (IAM) to secure your data by controlling who can access your Amazon Redshift cluster. For more information, see [Identity and access management in Amazon Redshift](#).

## Data types mapping

Amazon Redshift supports all BigQuery data types. The following table shows the data type mappings that AWS SCT uses by default. Users can set up migration rules in AWS SCT to change the data type of columns. For more information, see [Creating migration rules](#).

| BigQuery data type | Amazon Redshift data type |
|--------------------|---------------------------|
| BOOLEAN            | BOOLEAN                   |



| BigQuery data type | Amazon Redshift data type                                                                     |
|--------------------|-----------------------------------------------------------------------------------------------|
| BYTES(L)           | BINARY VARYING(L)                                                                             |
| BYTES              | BINARY VARYING(1024000)                                                                       |
| DATE               | DATE                                                                                          |
| DATETIME           | TIMESTAMP WITHOUT TIME ZONE                                                                   |
| GEOGRAPHY          | GEOGRAPHY                                                                                     |
| INTERVAL           | CHARACTER VARYING(256)                                                                        |
| JSON               | SUPER                                                                                         |
| INTEGER            | BIGINT                                                                                        |
| NUMERIC(p,s)       | NUMERIC(p,s)                                                                                  |
| NUMERIC            | NUMERIC(38,9)                                                                                 |
| BIGNUMERIC         | NUMERIC(38,9)                                                                                 |
| BIGNUMERIC(p,s)    | NUMERIC(p,s) if <i>p</i> is less than or equal to 38 or <i>s</i> is less than or equal to 37. |
| BIGNUMERIC(p,s)    | CHARACTER VARYING(256) if <i>p</i> is more than 38 or <i>s</i> is more than 37.               |
| FLOAT              | DOUBLE PRECISION                                                                              |
| STRING(L)          | CHARACTER VARYING(L) if <i>L</i> is less than 65,535.                                         |
| STRING             | CHARACTER VARYING(65535)                                                                      |
| STRUCT             | SUPER                                                                                         |
| TIME               | TIME WITHOUT TIME ZONE                                                                        |
| TIMESTAMP          | TIMESTAMP WITHOUT TIME ZONE                                                                   |

## Limitations

You can use AWS SCT to automatically convert a majority of your BigQuery code and storage objects. These objects include datasets, tables, views, stored procedures, functions, data types, and so on. However, AWS SCT has some limitations when using BigQuery as a source.

For example, AWS SCT can't convert subqueries in analytic functions, as well as geography, statistical aggregate, or some of the string functions. You can find the full list of limitations in the AWS SCT user guide. For more information, see [Limitations on using BigQuery as a source](#).

## Prerequisites for migrating a BigQuery project to Amazon Redshift

The following prerequisites are also required to complete this walkthrough:

- Familiarity with the AWS Management Console, AWS SCT, Amazon Redshift, Google Cloud management console, and SQL.
- An AWS user with AWS Identity and Access Management (IAM) credentials. Make sure that you can use these credentials to launch Amazon Redshift clusters and create an Amazon S3 bucket in your AWS Region. For more information, see [Creating an IAM role for your Amazon Redshift cluster](#) and [Create an IAM user for your Amazon S3 bucket](#).
- Basic knowledge of the Amazon Virtual Private Cloud (Amazon VPC) service and of security groups. For information about using Amazon Redshift in a VPC, see [Managing clusters in a VPC](#).
- An understanding of the supported features and limitations on using BigQuery as a source for AWS SCT. For more information, see [Using BigQuery as a source](#) in AWS SCT User Guide.
- An AWS service profile in AWS SCT with access to the S3 bucket. For more information, see [Storing service profiles](#) in AWS SCT User Guide.

For more information about the AWS Schema Conversion Tool, see [https://docs.aws.amazon.com/SchemaConversionTool/latest/userguide/CHAP\\_Welcome.html](https://docs.aws.amazon.com/SchemaConversionTool/latest/userguide/CHAP_Welcome.html).

For this migration walkthrough, we expect that you are familiar with BigQuery. You can use your BigQuery project for this migration, or create a new one. For example, you can use one of the public BigQuery datasets that are available in [Google Cloud Marketplace](#).

We recommend that you don't use your production workloads for the migration in this walkthrough. After you get familiar with migration tools and AWS services, you can migrate your production workloads.

# Step-by-Step BigQuery project to Amazon Redshift migration walkthrough

In the following sections, you can find step-by-step instructions for migrating your BigQuery project to Amazon Redshift. These steps assume that you have already prepared your source and target databases as described in preceding sections.

## Topics

- [Step 1: Create a BigQuery Service Account Key File](#)
- [Step 2: Create an Amazon Redshift Cluster](#)
- [Step 3: Create Buckets to Store Your Temporary Data](#)
- [Step 4: Install AWS SCT on Your Local Computer](#)
- [Step 5: Create an AWS SCT Project](#)
- [Step 6: Convert Database Schemas](#)
- [Step 7: Install and Configure Data Extraction Agents](#)
- [Step 8: Run Your Migration Task](#)
- [Step 9: Delete Walkthrough Resources](#)

## Step 1: Create a BigQuery Service Account Key File

You can connect to BigQuery with a user account or a service account. A *service account* is a special kind of account designed to be used by applications or compute workloads, rather than a person.

Service accounts don't have passwords and use a unique email address for identification. You can associate each service account with a service account key, which is a public or private RSA key pair. In this walkthrough, we use a service account key in AWS SCT to access your BigQuery project.

### To create a BigQuery service account key

1. Sign in to the [Google Cloud management console](#).
2. Make sure that you have API enabled on your [BigQuery API](#) page. If you don't see **API Enabled**, choose **Enable**.
3. On the [Service accounts](#) page, choose your BigQuery project, and then choose **Create service account**.

4. On the **Service account details** page, enter a descriptive value for **Service account name**. Choose **Create and continue**. The **Grant this service account access to the project** page opens.
5. For **Select a role**, choose **BigQuery**, and then choose **BigQuery Admin**. AWS SCT uses permissions to manage all resources within the project to load your BigQuery metadata in the migration project.
6. Choose **Add another role**. For **Select a role**, choose **Cloud Storage**, and then choose **Storage Admin**. AWS SCT uses full control of data objects and buckets to extract your data from BigQuery and then load it into Amazon Redshift.
7. Choose **Continue**, and then choose **Done**.
8. On the [Service accounts](#) page, choose the service account that you created.
9. Choose **Keys, Add key, Create new key**.
10. Choose **JSON**, and then choose **Create**. Choose the folder to save your private key or check the default folder for downloads in your browser.

## Step 2: Create an Amazon Redshift Cluster

To store your data in the AWS cloud, you can use your existing Amazon Redshift cluster or create a new one. You don't need to create any tables because AWS SCT automates this process.

If you don't plan to migrate data as part of this walkthrough, you can skip this step. To see how AWS SCT converts your database code objects, use a virtual Amazon Redshift target in your project. For more information, see [Using virtual targets](#).

### To create an Amazon Redshift cluster

1. Sign in to the AWS Management Console and open the Amazon Redshift console at <https://console.aws.amazon.com/redshift/>.
2. On the navigation menu, choose **Clusters**.
3. Choose **Create cluster**.
4. For **Cluster identifier**, enter the unique name of your Amazon Redshift cluster.
5. Choose **Free trial**.
6. For **Admin user name**, enter the login for the admin user of your Amazon Redshift cluster.
7. For **Admin user password**, enter the password for the admin user.
8. Choose **Create cluster**.

After you create your Amazon Redshift database, configure a new database user. Then, use the credentials of this user in AWS SCT to access your Amazon Redshift cluster. We don't recommend you to use the admin user for the migration.

Make sure that you grant the following privileges to this new user to complete the migration:

- `CREATE ON DATABASE` — allows to create new schemas in the database.
- `GRANT USAGE ON LANGUAGE` — allows to create new functions and procedures in the database.
- `GRANT SELECT ON ALL TABLES IN SCHEMA pg_catalog` — provides the user with system information about the Amazon Redshift cluster.
- `GRANT SELECT ON pg_class_info` — provides the user with information about tables distribution style.

You can use the following code example to create a database user and grant the privileges.

```
CREATE USER user_name PASSWORD your_password;  
GRANT CREATE ON DATABASE db_name TO user_name;  
GRANT USAGE ON LANGUAGE plpythonu TO user_name;  
GRANT USAGE ON LANGUAGE plpgsql TO user_name;  
GRANT SELECT ON ALL TABLES IN SCHEMA pg_catalog TO user_name;  
GRANT SELECT ON pg_class_info TO user_name;  
GRANT SELECT ON sys_serverless_usage TO user_name;  
GRANT SELECT ON pg_database_info TO user_name;  
GRANT SELECT ON pg_statistic TO user_name;
```

In the preceding example, replace *user\_name* with the name of your user. Then, replace *db\_name* with the name of your target Amazon Redshift database. Finally, replace *your\_password* with a secure password.

### Step 3: Create Buckets to Store Your Temporary Data

Data migration from BigQuery to Amazon Redshift includes the following steps:

1. Export data from BigQuery to a Cloud Storage bucket.
2. Extract data from a Cloud Storage bucket.
3. Upload data to an Amazon Simple Storage Service (Amazon S3) bucket.
4. Copy data from an S3 bucket to Amazon Redshift.

You need all four steps because you can't access data directly in BigQuery and you can't upload data directly to Amazon Redshift. Because of these limitations, you need to create buckets to store your data during migration.

### To create a Cloud Storage bucket

1. Sign in to the [Google Cloud management console](#).
2. Open the [Cloud Storage Browser](#) page.
3. Choose **Create bucket**.
4. For **Name your bucket**, enter a name for your Cloud Storage bucket.
5. On the **Choose where to store your data** page, choose **Region** for **Location type** and then choose your region for **Location**.
6. Leave the default values for other options, and choose **Create**.

### To create an Amazon S3 bucket

1. Sign in to the AWS Management Console and open the Amazon Redshift console at <https://console.aws.amazon.com/s3/>.
2. Choose **Create bucket**.
3. For **Bucket name**, enter the globally unique name of your Amazon S3 bucket.
4. For **AWS Region**, choose the AWS Region where you want the bucket to reside. Choose a Region close to you to minimize latency and costs.
5. Leave the default values for other options, and choose **Create bucket**.

## Step 4: Install AWS SCT on Your Local Computer

In this step, you install and configure the AWS Schema Conversion Tool. In this walkthrough, we run AWS SCT and the data extraction agent on Windows. However, you can use AWS SCT and data extraction agents on other supported operating systems. For more information, see [Installing the schema conversion tool](#) and [Installing extraction agents](#).

### To install AWS SCT

1. Download the compressed file that contains AWS SCT installer for Microsoft Windows from <https://s3.amazonaws.com/publicsctdownload/Windows/aws-schema-conversion-tool-1.0.latest.zip>.

2. Extract AWS SCT installer file.
3. Run AWS SCT installer file that you extracted in the previous step.
4. Choose **Next**, accept the terms of the License Agreement, and choose **Next** again.
5. Enter the path to the folder where you want to install AWS SCT, and choose **Next**.
6. Choose **Install**.
7. Choose **Finish** to close the installation wizard.

Now you can run AWS SCT. Before you create a new project, make sure that you add the path to an Amazon Redshift JDBC driver in the application settings. You don't need a JDBC driver to connect to your BigQuery project.

### To configure driver settings in AWS SCT

1. Download an Amazon Redshift JDBC driver version 2.1.0.9 or later from <https://docs.aws.amazon.com/redshift/latest/mgmt/jdbc20-download-driver.html>.
2. Extract the JDBC driver from the compressed file that you downloaded.
3. Open AWS SCT, and choose **Global settings** from **Settings**.
4. Choose **Drivers**.
5. For **Amazon Redshift driver path**, choose Browse and choose the `redshift-jdbc42-2.1.0.9.jar` file that you extracted.
6. Choose **Apply**, and then choose **OK** to close the settings window.

To access AWS services such as Amazon S3 from AWS SCT, you configure an AWS service profile. An AWS service profile is a set of AWS credentials that includes your AWS access key, AWS secret access key, AWS Region, and Amazon S3 bucket.

### To create an AWS service profile in AWS SCT

1. Open AWS SCT, and choose **Global settings** from **Settings**.
2. Choose **AWS service profiles**.
3. Choose **Add a new AWS service profile**.
4. For **Profile name**, enter a descriptive name for your profile.
5. For **AWS access key**, enter your AWS access key.

6. For **AWS secret key**, enter your AWS secret access key. For more information about AWS access keys, see [Programmatic access](#).
7. For **Region**, choose the AWS Region where you created your Amazon S3 bucket in the previous step.
8. For **Amazon S3 bucket folder**, choose the Amazon S3 bucket that you created in the previous step.
9. Choose **Apply**, and then choose **OK** to close the settings window.

## Step 5: Create an AWS SCT Project

After you configure AWS SCT, create a new migration project.

1. In AWS SCT, choose **File**, then choose **New Project**.
2. For **Project name**, enter a descriptive name of your project, and then choose **OK**.
3. Choose **Add source** to add a source BigQuery data warehouse to your project, then choose **BigQuery**, and choose **Next**.
4. For **Connection name**, enter a name for your source data warehouse. AWS SCT displays this name in the tree in the left panel.
5. For **Key path**, choose **Browse** and then choose the BigQuery service account key file that you created in step 1.
6. Choose **Connect** to close the dialog box and to connect to your BigQuery data warehouse.
7. Choose **Add target** to add a target Amazon Redshift database to your project, then choose **Amazon Redshift**, and choose **Next**.
8. If you store your database credentials in AWS Secrets Manager, choose your secret and then choose **Populate**. For more information, see [Using Secrets Manager](#).

If you don't use Secrets Manager, enter your database credentials manually.

- For **Connection name**, enter a name for your target data warehouse. AWS SCT displays this name in the tree in the right panel.
- For **Server name**, enter the server name of the Amazon Redshift cluster that you created in step 2. You can copy the server name as **JDBC URL** in the **General information** for your Amazon Redshift cluster. Remove `jdbc:redshift://` from the URL that you copied.
- For **Server port**, enter 5439.
- For **User name**, enter the name of the user that you created in step 2.



- For **Password**, enter the password for the user that you created in step 2.
9. Turn off **Use AWS Glue** and choose **Connect**.
  10. In the tree in the left panel, choose your BigQuery dataset. In the tree in the right panel, choose your target Amazon Redshift database. Choose **Create mapping**. You can add multiple mapping rules a single AWS SCT project. For more information about mapping rules, see [Creating mapping rules](#).
  11. Choose **Main view**.

## Step 6: Convert Database Schemas

After you create a new AWS SCT project, convert your source database schemas and apply converted code to your target database.

1. In the tree in the left panel, choose your source dataset. Open the context (right-click) menu, and choose **Convert schema**.
2. Choose **Yes** for the confirmation message. AWS SCT then converts your schema to the target database format.
3. AWS SCT also generates the assessment report. This report includes database objects that require manual conversion. To view this report, choose **View**, and then choose **Assessment report view**.
4. On the **Action items** tab, AWS SCT provides you with the recommended actions for each conversion issue.
5. Check the report and make changes in your source or converted code where necessary. You can optionally save the report as a .CSV or .PDF file for later analysis.
6. Choose **Action Items**, and review any recommendations that you see.
7. In the tree in the right panel, choose the converted schema. Open the context (right-click) menu, and choose **Apply to database** to apply the schema scripts to the target Amazon Redshift cluster.

## Step 7: Install and Configure Data Extraction Agents

AWS SCT uses a data extraction agent to migrate data from BigQuery to Amazon Redshift. The .zip file that you downloaded to install AWS SCT, includes the extraction agent installer file. In this walkthrough, we install the data extraction agent on Windows. However, you can install data

extraction agents on Red Hat Enterprise Linux or Ubuntu. For more information, see [Installing extraction agents](#).

### To install and configure a data extraction agent

1. Find the `aws-schema-conversion-tool-extractor-2.0.1.<version>.msi` file in the `agents` folder. The number of the `<version>` in the file name depends on the version of AWS SCT that you use. To migrate data from BigQuery to Amazon Redshift, make sure that you use an extraction agent version 665 or higher.
2. Run the file.
3. Choose **Next**, accept the terms of the License Agreement, and choose **Next** again.
4. Enter the path to the folder where you want to install the data extraction agent, and choose **Next**.
5. Choose **Install**.
6. On Windows, the data extraction agent installer launches the configuration wizard in the command prompt window. On Linux, run the `sct-extractor-setup.sh` file from the location where you installed the agent.
7. For **Listening port**, enter 8192. This is the default value. You can choose another port.
8. For **Add a source vendor**, enter no. You don't need to configure the data extraction agent to work with your BigQuery data warehouse because you don't need a driver to connect to BigQuery.
9. For **Add the Amazon Redshift driver**, enter yes and then enter the path to the Amazon Redshift JDBC driver that you downloaded in [Step 4](#).
- 10 For **Working folder**, enter the folder where the data extraction agent can store its data. Choose the project folder and make sure that you don't need admin rights to write data to this folder.
- 11 For **Enable SSL communication**, enter no. Then enter yes to confirm your choice. In this walkthrough, we don't use SSL to connect to databases. If you use SSL, configure the agent.

## Step 8: Run Your Migration Task

After you install and configure the data extraction agent, register it in AWS SCT.

### To register a data extraction agent

1. In AWS SCT, for **View** choose **Data migration view (other)**, and then choose **Register**.
2. For **Description**, enter a name for your data extraction agent.

3. For **Host name**, enter `0.0.0.0` because you run the data extraction agent on the same computer as AWS SCT. If you install the data extraction agent on another computer, enter the IP address of this computer.
4. For **Port**, enter 8192. If you configured another listening port in the previous step, use the value that you configured.
5. Choose **Register**.

AWS SCT now can use the data extraction agent for data migration tasks.

When you migrate big tables, you can split data into virtual partitions in AWS SCT. Then AWS SCT creates subtasks for each virtual partition.

### To create virtual partitions for your table in AWS SCT

1. In the tree in the left panel, choose your source table. Open the context (right-click) menu, and choose **Add virtual partitioning**.
2. For **Partition type**, choose **Range**.
3. For **Column name**, choose the column of your table. AWS SCT partitions data based on a range of column values. This partition type creates a WHERE clause, and you provide the range of values for each partition.
4. For **Values**, enter a list of values for the partitioned column.
5. Choose **OK** to create virtual partitions for your table.

Now, you can start the data migration.

### To create and run a migration task

1. In the tree in the left panel, choose your source table. Open the context (right-click) menu, and choose **Create local task**.
2. For **Task name**, enter a descriptive name for your data migration task.
3. For **Migration mode**, choose **Extract, upload, and copy**.
4. Choose **Advanced**. For **Google CS bucket folder**, enter the name for your Cloud Storage bucket that you created in [Step 3](#).
5. Choose **Amazon S3 settings**. For **Amazon S3 bucket folder**, enter the name of your Amazon S3 bucket that you created in [Step 3](#).

## 6. Choose **Create** and then choose **Start**.

The AWS SCT data extraction agents migrates data from your BigQuery dataset to Amazon Redshift. You can manage the migration process in AWS SCT. After the data extraction agent completes the migration, check your data in Amazon Redshift. Make sure that all your source data migrated to the new target database.

## Step 9: Delete Walkthrough Resources

After you complete this step-by-step guide, make sure that you delete your Amazon Redshift cluster to avoid additional charges.

### To delete an Amazon Redshift cluster

1. Sign in to the AWS Management Console and open the Amazon Redshift console at <https://console.aws.amazon.com/redshift/>.
2. On the navigation menu, choose **Clusters**.
3. Choose the cluster to delete.
4. For **Actions**, choose **Delete**. The **Delete cluster** page appears.
5. Choose **Delete cluster**.

## Migration from a BigQuery project to Amazon Redshift next steps

After you migrate your BigQuery project to Amazon Redshift, you can explore several other resources:

- Get started with Amazon Redshift. For more information, see [Amazon Redshift Getting Started Guide](#).
- Consider Amazon Redshift Serverless as a migration target. For more information, see [Amazon Redshift Serverless](#).
- Learn more about [Amazon Redshift performance optimization](#).
- You can use AWS SCT command line interface (CLI) to automate database migrations with scripts. For more information, see [CLI Reference](#).

# Migrating a MySQL-Compatible Database to Amazon Aurora MySQL

If your database supports the InnoDB or MyISAM tablespaces, you have these options for migrating your data to an Amazon Aurora MySQL DB cluster:

- You can create a dump of your data using the `mysqldump` utility, and then import that data into an existing Amazon Aurora MySQL DB cluster. For more information, see [Migrating MySQL to Amazon Aurora MySQL by Using mysqldump](#).
- You can copy the source files from your database to an Amazon S3 bucket, and then restore an Amazon Aurora MySQL DB cluster from those files. This option can be considerably faster than migrating data using `mysqldump`. For more information, see [Migrating Data from an External MySQL Database to an Amazon Aurora MySQL Using Amazon S3](#).

## Migrating Data from an External MySQL Database to an Amazon Aurora MySQL Using Amazon S3

You can copy the source files from your source MySQL version 5.5, 5.6, or 5.7 database to an Amazon S3 bucket, and then restore an Amazon Aurora MySQL DB cluster from those files.

This option can be considerably faster than migrating data using `mysqldump`, because using `mysqldump` replays all of the commands to recreate the schema and data from your source database in your new Amazon Aurora MySQL DB cluster. By copying your source MySQL data files, Amazon Aurora MySQL can immediately use those files as the data for DB cluster.

### Note

Restoring an Amazon Aurora MySQL DB cluster from backup files in an Amazon S3 bucket is not supported for the Asia Pacific (Mumbai) region.

Amazon Aurora MySQL does not restore everything from your database. You should save the database schema and values for the following items from your source MySQL or MariaDB database and add them to your restored Amazon Aurora MySQL DB cluster after it has been created.

- User accounts
- Functions

- Stored procedures
- Time zone information. Time zone information is loaded from the local operating system of your Amazon Aurora MySQL DB cluster.

## Prerequisites

Before you can copy your data to an Amazon S3 bucket and restore a DB cluster from those files, you must do the following:

- Install Percona XtraBackup on your local server.
- Permit Amazon Aurora MySQL to access your Amazon S3 bucket on your behalf.

## Installing Percona XtraBackup

Amazon Aurora MySQL can restore a DB cluster from files that were created using Percona XtraBackup. You can install Percona XtraBackup from the Percona website at <https://www.percona.com/software/mysql-database/percona-xtrabackup>.

## Required Permissions

To migrate your MySQL data to an Amazon Aurora MySQL DB cluster, several permissions are required:

- The user that is requesting that Amazon RDS create a new cluster from an Amazon S3 bucket must have permission to list the buckets for your user. You grant the user this permission using an AWS Identity and Access Management (IAM) policy.
- Amazon RDS requires permission to act on your behalf to access the Amazon S3 bucket where you store the files used to create your Amazon Aurora MySQL DB cluster. You grant Amazon RDS the required permissions using an IAM service role.
- The user making the request must also have permission to list the IAM roles for your user.
- If the user making the request will create the IAM service role, or will request that Amazon RDS create the IAM service role (by using the console), then the user must have permission to create an IAM role for your user.

For example, the following IAM policy grants a user the minimum required permissions to use the console to both list IAM roles, create an IAM role, and list the S3 buckets for your user.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "iam:ListRoles",
        "iam:CreateRole",
        "iam:CreatePolicy",
        "iam:AttachRolePolicy",
        "s3:ListBucket",
        "s3:ListObjects"
      ],
      "Resource": "*"
    }
  ]
}
```

Additionally, for a user to associate an IAM role with an S3 bucket, the IAM user must have the `iam:PassRole` permission for that IAM role. This permission allows an administrator to restrict which IAM roles a user can associate with S3 buckets.

For example, the following IAM policy allows a user to associate the role named `S3Access` with an S3 bucket.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "AllowS3AccessRole",
      "Effect": "Allow",
      "Action": "iam:PassRole",
      "Resource": "arn:aws:iam::123456789012:role/S3Access"
    }
  ]
}
```

## Creating the IAM Service Role

You can have the Amazon RDS Management Console create a role for you by choosing the **Create a New Role** option (shown later in this topic). If you select this option and specify a name for the

new role, then Amazon RDS will create the IAM service role required for Amazon RDS to access your Amazon S3 bucket with the name that you supply.

As an alternative, you can manually create the role using the following procedure.

1. Sign in to the AWS Management Console and open the IAM console at <https://console.aws.amazon.com/iam/>.
2. In the left navigation pane, choose **Roles**.
3. Choose **Create New Role**, specify a value for **Role Name** for the new role, and then choose **Next Step**.
4. Under **AWS Service Roles**, find **Amazon RDS** and choose **Select**.
5. Do not select a policy to attach in the **Attach Policy** step. Instead, choose **Next Step**.
6. Review your role information, and then choose **Create Role**.
7. In the list of roles, choose the name of your newly created role. Choose the **Permissions** tab.
8. Choose **Inline Policies**. Because your new role has no policy attached, you will be prompted to create one. Click the link to create a new policy.
9. On the **Set Permissions** page, choose **Custom Policy** and then choose **Select**.
10. Enter a **Policy Name** such as S3-bucket-policy. Add the following code for **Policy Document**, replacing *<bucket name>* with the name of the S3 bucket that you are allowing access to.

As part of the policy document, you can also include a file name prefix. If you specify a prefix, then Amazon Aurora MySQL will create the DB cluster using the files in the S3 bucket that begin with the specified prefix. If you don't specify a prefix, then Amazon Aurora MySQL will create the DB cluster using all of the files in the S3 bucket.

To specify a prefix, replace *<prefix>* following with the prefix of your file names. Include the asterisk (\*) after the prefix. If you don't want to specify a prefix, specify only an asterisk.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "s3:ListBucket",
        "s3:GetBucketLocation"
      ],

```



```
        "Resource": [
            "arn:aws:s3:::<bucket name>"
        ],
    },
    {
        "Effect": "Allow",
        "Action": [
            "s3:GetObject"
        ],
        "Resource": [
            "arn:aws:s3:::<bucket name>/<prefix>*"
        ]
    }
]
```

11 Choose **Apply Policy**.

## Step 1: Backing Up Files to be Restored as a DB Cluster

To create a backup of your MySQL database files that can be restored from S3 to create an Amazon Aurora MySQL DB cluster, use the Percona Xtrabackup utility (innobackupex) to back up your database.

For example, the following command creates a backup of a MySQL database and stores the files in the /s3-restore/backup folder.

```
innobackupex --user=myuser --password=<password> --no-timestamp /s3-restore/backup
```

If you want to compress your backup into a single file (which can be split, if needed), you can use the `--stream` option to save your backup in one of the following formats:

- Gzip (.gz)
- tar (.tar)
- Percona xstream (.xstream)

For example, the following command creates a backup of your MySQL database split into multiple Gzip files. The parameter values shown are for a small test database; for your scenario, you should determine the parameter values needed.

```
innobackupex --user=myuser --password=<password> --stream=tar \  
/mydata/s3-restore/backup | split -d --bytes=512000 \  
- /mydata/s3-restore/backup3/backup.tar.gz
```

For example, the following command creates a backup of your MySQL database split into multiple tar files.

```
innobackupex --user=myuser --password=<password> --stream=tar \  
/mydata/s3-restore/backup | split -d --bytes=512000 \  
- /mydata/s3-restore/backup3/backup.tar
```

For example, the following command creates a backup of your MySQL database split into multiple xstream files.

```
innobackupex --stream=xstream \  
/mydata/s3-restore/backup | split -d --bytes=512000 \  
- /mydata/s3-restore/backup/backup.xstream
```

Amazon S3 limits the size of a file uploaded to a bucket to 5 terabytes (TB). If the backup data for your database exceeds 5 TB, then you must use the `split` command to split the backup files into multiple files that are each less than 5 TB.

Amazon Aurora MySQL does not support partial backups created using Percona Xtrabackup. You cannot use the `--include`, `--tables-file`, or `--databases` options to create a partial backup when you backup the source files for your database.

For more information, see [The innobackupex Script](#).

Amazon Aurora MySQL consumes your backup files based on the file name. Be sure to name your backup files with the appropriate file extension based on the file format—for example, `0xstream` for files stored using the Percona xstream format.

Amazon Aurora MySQL consumes your backup files in alphabetical order as well as natural number order. Always use the `split` option when you issue the `innobackupex` command to ensure that your backup files are written and named in the proper order.

## Step 2: Copying Files to an Amazon S3 Bucket

Once you have backed up your MySQL database using the Percona Xtrabackup utility, then you can copy your backup files to an Amazon S3 bucket.

For information about creating and uploading a file to an Amazon S3 bucket, see [Getting Started with Amazon Simple Storage Service](#) in the *Amazon S3 Getting Started Guide*.

### Step 3: Restoring an Aurora MySQL DB Cluster from an Amazon S3 Bucket

You can restore your backup files from your Amazon S3 bucket to a create new Amazon Aurora MySQL DB cluster by using the Amazon RDS console.

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. In the RDS Dashboard, choose **Restore Aurora MySQL DB Cluster from S3**.
3. In the **Create database by restoring from S3** page, specify the following settings in the following sections:
  - a. In the **S3 Destination** section, specify the following:

| Parameter                   | Action                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                  |
|-----------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>S3 Bucket</b>            | Select the Amazon S3 bucket where your backup files are stored.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                         |
| <b>S3 Prefix (Optional)</b> | <p>Specify a file path prefix for the files stored in your Amazon S3 bucket. The <b>S3 Bucket Prefix</b> is optional. If you don't specify a prefix, then Amazon Aurora MySQL will create the DB cluster using all of the files in the root folder of the S3 bucket. If you specify a prefix, then Amazon Aurora MySQL will create the DB cluster using the files in the S3 bucket where the full path for the file begins with the specified prefix.</p> <p>Amazon Aurora MySQL does not traverse subfolders in your S3 bucket looking for backup files. Only the files from the folder identified by the <b>S3 Bucket Prefix</b> are used. If you store your backup files in a subfolder in your S3 bucket, then you must specify</p> |

| Parameter | Action                                                                                                                                                                                                                                                                                                                                                                                                                    |
|-----------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|           | <p>a prefix that identifies the full path to the folder where the files are stored.</p> <p>For example, if you store your backup files in a subfolder of your S3 bucket named backups, and you have multiple sets of backup files, each in its own directory (gzip_backup1 , gzip_backup2 , and so on), then you would specify a prefix of backups/gzip_backup1 to restore from the files in the gzip_backup1 folder.</p> |

b. In the **Engine Options** section, specify the following:

| Parameter          | Action                                                                                                                                       |
|--------------------|----------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Engine Type</b> | Keep Amazon Aurora selected.                                                                                                                 |
| <b>Edition</b>     | Keep Amazon Aurora with MySQL compatibility selected.                                                                                        |
| <b>Version</b>     | Specify the version of the MySQL database that the backup files were created from, for example 5.7. MySQL version 5.6 and 5.7 are supported. |

c. In the **IAM role** section, specify the following:

| Parameter       | Action                                                                                                                                                                                                |
|-----------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>IAM Role</b> | Choose the IAM role that you created to authorize Amazon Aurora MySQL to access Amazon S3 on your behalf. If you have not created an IAM role, you can choose <b>Create a New Role</b> to create one. |

d. In the **Settings** section, specify the following:

| Parameter                       | Action                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               |
|---------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>DB cluster identifier</b>    | <p>Enter a name for your DB cluster. This identifier will be used in the endpoint address for the primary instance of your DB cluster.</p> <p>The DB instance identifier has the following constraints:</p> <ul style="list-style-type: none"><li>• It must contain from 1 to 63 alphanumeric characters or hyphens.</li><li>• Its first character must be a letter.</li><li>• It cannot end with a hyphen or contain two consecutive hyphens.</li><li>• It must be unique for all DB instances per user, for each region.</li></ul> |
| <b>Master Username</b>          | <p>Enter a name using alphanumeric characters that you will use as the master user name to log on to your DB cluster. The default privileges granted to the master user name account include: create, drop, references, event, alter, delete, index, insert, select, update, create temporary tables, lock tables, trigger, create view, show view, alter routine, create routine, execute, create user, process, show databases, grant option.</p>                                                                                  |
| <b>Auto generate a password</b> | Leave unchecked.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                     |

| Parameter               | Action                                                                                                                        |
|-------------------------|-------------------------------------------------------------------------------------------------------------------------------|
| <b>Master Password</b>  | Enter a password that contains from 8 to 41 printable ASCII characters (excluding /, ", and @) for your master user password. |
| <b>Confirm Password</b> | Retype the Master Password.                                                                                                   |

e. In the **DB Instance Class** section, specify the following:

| Parameter                | Action                                                                                                                                                                                                                                                                                                                                                                                                                      |
|--------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>DB Instance Class</b> | Select a DB instance class that defines the processing and memory requirements for each instance in the DB cluster. Aurora MySQL supports the <code>db.r3.large</code> , <code>db.r3.xlarge</code> , <code>db.r3.2xlarge</code> , <code>db.r3.4xlarge</code> , and <code>db.r3.8xlarge</code> DB instance classes. For more information about DB instance class options, see the <a href="#">Amazon RDS documentation</a> . |

f. In the **Availability & durability** section, specify the following:

| Parameter                  | Action                                                                                                                                                                                                             |
|----------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Multi-AZ Deployment</b> | Determine if you want to create Aurora MySQL Replicas in other Availability Zones for failover support. For more information about multiple Availability Zones, see the <a href="#">Amazon RDS documentation</a> . |

g. In the **Connectivity** section, specify the following:

| Parameter                          | Action                                                                                                                                                                                                                                                                                                |
|------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Virtual private cloud (VPC)</b> | Select the VPC that will host the DB cluster. Select <b>Create a New VPC</b> to have Amazon RDS create a VPC for you. For more information, see earlier in this topic.                                                                                                                                |
| <b>Subnet group</b>                | Select the DB subnet group to use for the DB cluster. Select <b>Create a New DB Subnet Group</b> to have Amazon RDS create a DB subnet group for you. For more information, see earlier in this topic.                                                                                                |
| <b>Public access</b>               | Select <b>Yes</b> to give the DB cluster a public IP address; otherwise, select <b>No</b> . The instances in your DB cluster can be a mix of both public and private DB instances. For more information about hiding instances from public access, see the <a href="#">Amazon RDS documentation</a> . |
| <b>VPC Security Group(s)</b>       | Select one or more VPC security groups to secure network access to the DB cluster. Select <b>Create a New VPC Security Group</b> to have Amazon RDS create a VPC security group for you. For more information, see earlier in this topic.                                                             |
| <b>Availability Zone</b>           | Determine if you want to specify a particular Availability Zone. For more information about Availability Zones, see the <a href="#">Amazon RDS documentation</a> .                                                                                                                                    |

| Parameter     | Action                                                                                                                                                                                                                                                                                                                         |
|---------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Database Port | Specify the port that applications and utilities will use to access the database. Aurora MySQL DB clusters default to the default MySQL port, 3306. The firewalls at some companies block connections to the default MySQL port. If your company firewall blocks the default port, choose another port for the new DB cluster. |

- h. In the **Database authentication** section, specify the following:

| Parameter               | Action                                         |
|-------------------------|------------------------------------------------|
| Database Authentication | Leave <b>Password authentication</b> selected. |

- i. In the **Additional configuration** section, specify the following:

| Parameter                  | Action                                                                                                                                                                                                                                        |
|----------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Initial Database Name      | Enter a name for your database of up to 8 alphanumeric characters. If you don't provide a name, Amazon RDS will not create a database on the DB cluster you are creating.                                                                     |
| DB cluster parameter Group | Select a parameter group for the cluster. Aurora MySQL has a default parameter group you can use, or you can create your own parameter group. For more information about parameter groups, see the <a href="#">Amazon RDS documentation</a> . |
| DB parameter Group         | Select a parameter group for the database.                                                                                                                                                                                                    |



| Parameter                   | Action                                                                                                                                                                                                                                                                                                        |
|-----------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Option Group                | Select an option group. Aurora MySQL has a default option group you can use, or you can create your own option group. For more information about option groups, see the <a href="#">Amazon RDS documentation</a> .                                                                                            |
| Failover Priority           | Choose a failover priority for the instance. If you don't select a value, the default is <b>tier-1</b> . This priority determines the order in which Aurora MySQL Replicas are promoted when recovering from a primary instance failure. For more information, see <a href="#">Amazon RDS documentation</a> . |
| Backup Retention Period     | Select the length of time, from 1 to 35 days, that Aurora MySQL will retain backup copies of the database. Backup copies can be used for point-in-time restores (PITR) of your database, timed down to the second.                                                                                            |
| Copy tags to snapshots      | Leave checked.                                                                                                                                                                                                                                                                                                |
| Enable Encryption           | Check the box to enable encryption at rest for this DB cluster. Leave <b>AWS KMS Key</b> set to <b>(default) aws/rds</b> . For more information, see <a href="#">Amazon RDS documentation</a> .                                                                                                               |
| Backtrack                   | Leave unchecked.                                                                                                                                                                                                                                                                                              |
| Enable Performance insights | Leave checked. Leave <b>Retention Period</b> and <b>AWS KMS Key</b> as they are.                                                                                                                                                                                                                              |
| Enable Enhanced Monitoring  | Choose <b>Yes</b> to enable gathering metrics in real time for the operating system that your DB cluster runs on. For more information, see <a href="#">Amazon RDS documentation</a> .                                                                                                                        |

| Parameter                                | Action                                                                                                                                                                                                                                                                                                                                                                                              |
|------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Granularity</b>                       | This option is only available if <b>Enable Enhanced Monitoring</b> is set to <b>Yes</b> . Set the interval, in seconds, between times at which metrics are collected for your DB cluster.                                                                                                                                                                                                           |
| <b>Monitoring role</b>                   | Leave as <b>default</b> .                                                                                                                                                                                                                                                                                                                                                                           |
| <b>Log exports</b>                       | Leave unchecked.                                                                                                                                                                                                                                                                                                                                                                                    |
| <b>Enable auto Minor Version Upgrade</b> | <p>Check this box if you want to enable your Aurora MySQL DB cluster to receive minor MySQL DB engine version upgrades automatically when they become available.</p> <p>The <b>Auto Minor Version Upgrade</b> option only applies to upgrades to MySQL minor engine versions for your Amazon Aurora MySQL DB cluster. It doesn't apply to regular patches applied to maintain system stability.</p> |
| <b>Maintenance Window</b>                | Select the weekly time range during which system maintenance can occur.                                                                                                                                                                                                                                                                                                                             |
| <b>Enable deletion protection</b>        | Leave unchecked.                                                                                                                                                                                                                                                                                                                                                                                    |

4. Choose **Launch DB Instance** to launch your Aurora MySQL DB instance, and then choose **Close** to close the wizard.

On the Amazon RDS console, the new DB instance appears in the list of DB instances. The DB instance has a status of **creating** until the DB instance is created and ready for use. When the state changes to **available**, you can connect to the primary instance for your DB cluster. Depending on the DB instance class and store allocated, it can take several minutes for the new instance to be available.

To view the newly created cluster, choose the **Clusters** view in the Amazon RDS console. For more information, see [Amazon RDS documentation](#).

Note the port and the endpoint of the cluster. Use the endpoint and port of the cluster in your JDBC and ODBC connection strings for any application that performs write or read operations.

## Migrating MySQL to Amazon Aurora MySQL by Using mysqldump

You can create a dump of your data using the `mysqldump` utility, and then import that data into an existing Amazon Aurora MySQL DB cluster.

Because Amazon Aurora MySQL is a MySQL-compatible database, you can use the `mysqldump` utility to copy data from your MySQL or MariaDB database to an existing Amazon Aurora MySQL DB cluster.

## Migrating Data from an Amazon RDS MySQL DB Instance to an Amazon Aurora MySQL DB Cluster

You can migrate (copy) data to an Amazon Aurora MySQL DB cluster from an Amazon RDS snapshot, as described following.

### Note

Because Amazon Aurora MySQL is compatible with MySQL, you can migrate data from your MySQL database by setting up replication between your MySQL database, and an Amazon Aurora MySQL DB cluster. We recommend that your MySQL database run MySQL version 5.5 or later.

## Migrating an Amazon RDS for MySQL Snapshot to Aurora MySQL

You can migrate a DB snapshot of an Amazon RDS MySQL DB instance to create an Aurora MySQL DB cluster. The new DB cluster is populated with the data from the original Amazon RDS MySQL DB instance. The DB snapshot must have been made from an Amazon RDS DB instance running MySQL 5.6.

You can migrate either a manual or automated DB snapshot. After the DB cluster is created, you can then create optional Aurora MySQL Replicas.

The general steps you must take are as follows:

1. Determine the amount of space to provision for your Amazon Aurora MySQL DB cluster. For more information, see [Amazon RDS documentation](#).
2. Use the console to create the snapshot in the region where the Amazon RDS MySQL 5.6 instance is located
3. If the DB snapshot is not in the region as your DB cluster, use the Amazon RDS console to copy the DB snapshot to that region. For information about copying a DB snapshot, see the [Amazon RDS documentation](#).
4. Use the console to migrate the DB snapshot and create an Amazon Aurora MySQL DB cluster with the same databases as the original DB instance of MySQL 5.6.

 **Warning**

Amazon RDS limits each user to one snapshot copy into each region at a time.

## How Much Space Do I Need?

When you migrate a snapshot of a MySQL DB instance into an Aurora MySQL DB cluster, Aurora MySQL uses an Amazon Elastic Block Store (Amazon EBS) volume to format the data from the snapshot before migrating it. In some cases, additional space is needed to format the data for migration. When migrating data into your DB cluster, observe the following guidelines and limitations:

- Although Amazon Aurora MySQL supports storage up to 64 TB in size, the process of migrating a snapshot into an Aurora MySQL DB cluster is limited by the size of the EBS volume of the snapshot. Thus, the maximum size for a snapshot that you can migrate is 6 TB.
- Tables that are not MyISAM tables and are not compressed can be up to 6 TB in size. If you have MyISAM tables, then Aurora MySQL must use additional space in the volume to convert the tables to be compatible with Aurora MySQL. If you have compressed tables, then Aurora MySQL must use additional space in the volume to expand these tables before storing them on the Aurora MySQL cluster volume. Because of this additional space requirement, you should ensure that none of the MyISAM and compressed tables being migrated from your MySQL DB instance exceeds 3 TB in size.

## Reducing the Amount of Space Required to Migrate Data into Amazon Aurora MySQL

You might want to modify your database schema prior to migrating it into Amazon Aurora MySQL. Such modification can be helpful in the following cases:

- You want to speed up the migration process.
- You are unsure of how much space you need to provision.
- You have attempted to migrate your data and the migration has failed due to a lack of provisioned space.

You can make the following changes to improve the process of migrating a database into Amazon Aurora MySQL.

### Important

Be sure to perform these updates on a new DB instance restored from a snapshot of a production database, rather than on a production instance. You can then migrate the data from the snapshot of your new DB instance into your Amazon Aurora MySQL DB cluster to avoid any service interruptions on your production database.

| Table Type    | Limitation or Guideline                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             |
|---------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| MyISAM tables | <p>Amazon Aurora MySQL supports InnoDB tables only. If you have MyISAM tables in your database, then those tables must be converted before being migrated into Amazon Aurora MySQL. The conversion process requires additional space for the MyISAM to InnoDB conversion during the migration procedure.</p> <p>To reduce your chances of running out of space or to speed up the migration process, convert all of your MyISAM tables to InnoDB tables before migrating them. The size of the resulting InnoDB table is equivalent to the size</p> |

| Table Type        | Limitation or Guideline                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                |
|-------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
|                   | <p>required by Amazon Aurora MySQL for that table. To convert a MyISAM table to InnoDB, run the following command:</p> <pre>alter table &lt;schema&gt;.&lt;table_name&gt; engine=innodb, algorithm=copy;</pre>                                                                                                                                                                                                                                                                                                                                                         |
| Compressed tables | <p>Amazon Aurora MySQL does not support compressed tables (that is, tables created with <code>ROW_FORMAT=COMPRESSED</code> ).</p> <p>To reduce your chances of running out of space or to speed up the migration process, expand your compressed tables by setting <code>ROW_FORMAT</code> to <code>DEFAULT</code>, <code>COMPACT</code>, <code>DYNAMIC</code>, or <code>REDUNDANT</code> . For more information, see <a href="https://dev.mysql.com/doc/refman/5.6/en/innodb-row-format.html">https://dev.mysql.com/doc/refman/5.6/en/innodb-row-format.html</a>.</p> |

You can use the following SQL script on your existing MySQL DB instance to list the tables in your database that are MyISAM tables or compressed tables.

```
-- This script examines a MySQL database for conditions that will block
-- migrating the database into an Amazon Aurora MySQL DB.
-- It needs to be run from an account that has read permission for the
-- INFORMATION_SCHEMA database.

-- Verify that this is a supported version of MySQL.

select msg as `==> Checking current version of MySQL.`
from
(
select
  'This script should be run on MySQL version 5.6. ' +
  'Earlier versions are not supported.' as msg,
  cast(substring_index(version(), '.', 1) as unsigned) * 100 +
  cast(substring_index(substring_index(version(), '.', 2), '.', -1)
```

```

        as unsigned)
    as major_minor
) as T
where major_minor <> 506;

-- List MyISAM and compressed tables. Include the table size.

select concat(TABLE_SCHEMA, '.', TABLE_NAME) as `==> MyISAM or Compressed Tables`,
round(((data_length + index_length) / 1024 / 1024), 2) "Approx size (MB)"
from INFORMATION_SCHEMA.TABLES
where
    ENGINE <> 'InnoDB'
and
(
    -- User tables
    TABLE_SCHEMA not in ('mysql', 'performance_schema',
                          'information_schema')

or
    -- Non-standard system tables
    (
        TABLE_SCHEMA = 'mysql' and TABLE_NAME not in
        (
            'columns_priv', 'db', 'event', 'func', 'general_log',
            'help_category', 'help_keyword', 'help_relation',
            'help_topic', 'host', 'ndb_binlog_index', 'plugin',
            'proc', 'procs_priv', 'proxies_priv', 'servers', 'slow_log',
            'tables_priv', 'time_zone', 'time_zone_leap_second',
            'time_zone_name', 'time_zone_transition',
            'time_zone_transition_type', 'user',
            'general_log_backup', 'slow_log_backup'
        )
    )
)
or
(
    -- Compressed tables
    ROW_FORMAT = 'Compressed'
);

```

The script produces output similar to the output in the following example. The example shows two tables that must be converted from MyISAM to InnoDB. The output also includes the approximate size of each table in megabytes (MB).

```
+-----+-----+
| ==> MyISAM or Compressed Tables | Approx size (MB) |
+-----+-----+
| test.name_table                  |          2102.25 |
| test.my_table                    |           65.25  |
+-----+-----+
2 rows in set (0.01 sec)
```

## Migrating a DB Snapshot by Using the Console

You can migrate a DB snapshot of an Amazon RDS for MySQL DB instance to create an Aurora MySQL DB cluster. The new DB cluster will be populated with the data from the original Amazon RDS for MySQL DB instance. The DB snapshot must have been made from an Amazon RDS DB instance running MySQL 5.6 and must not be encrypted. For information about creating a DB snapshot, see the [Amazon RDS documentation](#).


If the DB snapshot is not in the AWS Region where you want to locate your data, use the Amazon RDS console to copy the DB snapshot to that region. For information about copying a DB snapshot, see the [Amazon RDS documentation](#).

When you migrate the DB snapshot by using the console, the console takes the actions necessary to create both the DB cluster and the primary instance.

You can also choose for your new Aurora MySQL DB cluster to be encrypted "at rest" using an AWS Key Management Service (AWS KMS) encryption key. This option is available only for unencrypted DB snapshots.

1. Sign in to the AWS Management Console and open the Amazon RDS console at <https://console.aws.amazon.com/rds/>.
2. Choose **Snapshots**.
3. On the **Snapshots** page, choose the snapshot that you want to migrate into an Aurora MySQL DB cluster.
4. Choose **Migrate Database**.





5. Set the following values on the **Migrate Database** page:

- **DB Instance Class:** Select a DB instance class that has the required storage and capacity for your database, for example `db.r3.large`. Aurora MySQL cluster volumes automatically grow as the amount of data in your database increases, up to a maximum size of 64 terabytes (TB). So you only need to select a DB instance class that meets your current storage requirements.
- **DB Instance Identifier:** Enter a name for the DB cluster that is unique for your account in the region you selected. This identifier is used in the endpoint addresses for the instances in your DB cluster. You might choose to add some intelligence to the name, such as including the region and DB engine you selected, for example `aurora-cluster1`.

The DB instance identifier has the following constraints:

- It must contain from 1 to 63 alphanumeric characters or hyphens.
- Its first character must be a letter.
- It cannot end with a hyphen or contain two consecutive hyphens.
- It must be unique for all DB instances per user, for each AWS Region.

- **VPC:** If you have an existing VPC, then you can use that VPC with your Amazon Aurora MySQL DB cluster by selecting your VPC identifier, for example `vpc-a464d1c1`. For information about using an existing VPC, see the [Amazon RDS documentation](#).

Otherwise, you can choose to have Amazon RDS create a VPC for you by selecting **Create a new VPC**.

- **Subnet Group:** If you have an existing subnet group, then you can use that subnet group with your Amazon Aurora MySQL DB cluster by selecting your subnet group identifier, for example `gs-subnet-group1`.

Otherwise, you can choose to have Amazon RDS create a subnet group for you by selecting **Create a new subnet group**.

- **Publicly Accessible:** Select **No** to specify that instances in your DB cluster can only be accessed by resources inside of your VPC. Select **Yes** to specify that instances in your DB cluster can be accessed by resources on the public network. The default is **Yes**.

 **Note**

Your production DB cluster might not need to be in a public subnet, because only your application servers will require access to your DB cluster. If your DB cluster doesn't need to be in a public subnet, set **Publicly Accessible** to **No**.

- **Availability Zone:** Select the Availability Zone to host the primary instance for your Aurora MySQL DB cluster. To have Amazon RDS select an Availability Zone for you, select **No Preference**.
- **Database Port:** Enter the default port to be used when connecting to instances in the DB cluster. The default is 3306.

 **Note**

You might be behind a corporate firewall that doesn't allow access to default ports such as the MySQL default port, 3306. In this case, provide a port value that your corporate firewall allows. Remember that port value later when you connect to the Aurora MySQL DB cluster.

- **Enable Encryption:** Choose **Yes** for your new Aurora MySQL DB cluster to be encrypted "at rest." If you choose **Yes**, you will be required to choose an AWS KMS encryption key as the KMS key value.
- **Auto Minor Version Upgrade:** Select **Yes** if you want to enable your Aurora MySQL DB cluster to receive minor MySQL DB engine version upgrades automatically when they become available.

The **Auto Minor Version Upgrade** option only applies to upgrades to MySQL minor engine versions for your Amazon Aurora MySQL DB cluster. It doesn't apply to regular patches applied to maintain system stability.

## Instance Specifications

**Migrate to DB Engine**

**DB Instance Class**

## Settings

**DB Snapshot ID**

**DB Instance Identifier\***

## Network & Security

This instance will be created with the new Certificate Authority `rds-ca-2015`. If you are using SSL to connect to this instance, you should use the [new certificate bundle](#). Learn more [here](#)

**VPC\***

**Subnet Group**

**Publicly Accessible**

**Availability Zone**

## Database Options

**Database Port**

**Enable Encryption**


## Maintenance

**Auto Minor Version Upgrade**

Cancel

Migrate

6. Choose **Migrate** to migrate your DB snapshot.
7. Choose **Instances**, and then choose the arrow icon to show the DB cluster details and monitor the progress of the migration. On the details page, you will find the cluster endpoint used to connect to the primary instance of the DB cluster. For more information about connecting to an Amazon Aurora MySQL DB cluster, see the [Amazon RDS documentation](#).



## Migrating a MariaDB Database to Amazon RDS for MySQL or Amazon Aurora MySQL

You can migrate data from existing on-premises MariaDB or [Amazon RDS for MariaDB](#) to [Amazon Aurora](#) MySQL using [Database Migration Service](#). Amazon Aurora is a MySQL and PostgreSQL-compatible relational database built for the cloud. Amazon Aurora features a distributed, fault-tolerant, self-healing storage system that auto-scales up to 64 TB per database instance. It delivers high performance and availability with up to 15 low-latency read replicas, point-in-time recovery, and continuous backup to Amazon S3, and replication across three Availability Zones (AZs).

Some key features offered by Aurora MySQL are the following:

- High throughput with low latency
- Push-button compute scaling
- Storage autoscaling
- Custom database endpoints
- Parallel queries for faster analytics

In the following sections, we demonstrate migration from MariaDB as a source database to an Aurora MySQL database as a target using AWS DMS. At a high level, the steps involved in this migration are:

- Provision MariaDB as a source DB instance and load the data
- Provision Aurora MySQL as target DB instance
- Provision DMS replication instance and create DMS endpoints
- Create DMS task, migrate data and perform validation

For the purpose of this section, we are using the AWS CloudFormation templates for creating Amazon RDS for MariaDB, Aurora MySQL database and AWS DMS replication instance with their source and endpoints. We will be loading sample tables and data in MariaDB located on [GitHub](#).

To estimate what it will cost to run this walkthrough on AWS, you can use the AWS Pricing Calculator. For more information, see <https://calculator.aws/>.

## Topics

- [Set up MariaDB as a source database](#)
- [Set up Aurora MySQL as a target database](#)
- [Set up an AWS DMS replication instance](#)
- [Test the endpoints for MariaDB database migration](#)
- [Create a migration task for a MariaDB database](#)
- [Validate the MariaDB database migration](#)
- [Cut over for the migration from a MariaDB database](#)

## Set up MariaDB as a source database

To provision MariaDB as a source database, download [Mariadb\\_CF.zip with the YAML template](#). This AWS CloudFormation template creates an Amazon RDS for MariaDB instance with the required parameters.

1. On the [AWS Management Console](#), under **Services**, choose **CloudFormation**.
2. Choose **Create stack**, and then choose **With new resources (standard)**.
3. For **Specify template**, choose **Upload a template file**.
4. Select **Choose file**.

5. Choose the `Mariadb_CF.yaml` file, and then choose **Next**.
6. On the **Specify stack details** page, edit the predefined values as needed, and then choose **Next**:
  - **Stack name** — Enter a name for the stack.
  - **CIDR** — Enter the CIDR IP range to access the instance.
  - **DBAllocatedStorage** — Enter the database storage size in GB. The default is 20 GB.
  - **DBBackupRetentionPeriod** — The number of days to retain backups.
  - **DBInstanceClass** — Enter the instance type of the database server.
  - **DBMonitoringInterval** — Interval to publish database logs to Amazon CloudWatch.
  - **DBSubnetGroup** — Enter the DB subnet group name. For more information, see [Create a DB subnet group](#) in the *Amazon RDS User Guide*.
  - **MariaDBEngine** — Enter the MariaDB engine version. Leave empty to use default. For more information, see [MariaDB on Amazon RDS versions](#) in the *Amazon RDS User Guide*.
  - **DBMasterPassword** — Enter the master password for the DB instance.
  - **DBMasterUsername** — Enter the master user name for the DB instance.
  - **PreferredBackupWindow** — Enter the daily time range in UTC during which you want to create automated backups.
  - **PreferredMaintenanceWindow** — Enter the weekly time range in UTC during which system maintenance can occur.
  - **RDSDBName** — Enter the name of the database.
  - **RDSMultiAZ** — Choose **true** to use Amazon RDS Multi-AZ for this instance. The default value for this option is **false**. For more information, see [Multi-AZ deployments for high availability](#) in the *Amazon Relational Database Service User Guide*.
  - **VPCID** — Enter the VPC to launch your DB instance. For more information, see [Working with a DB instance in a VPC](#) in the *Amazon RDS User Guide*.

Make sure that you entered the stack name, DB subnet group name, user name, password, database name, and VPC ID.
7. On the **Configure stack options** page, for **Tags**, specify any optional tags, and then choose **Next**.
8. On the **Review** page, select **I acknowledge that AWS CloudFormation might create IAM resources**, and then choose **Next**.
9. Choose **Create stack**.

After the Amazon RDS for MariaDB instance is created, log in to MariaDB and run the following statements to create `webdb_user`, a superuser that connects to a DMS instance for migration, and grant necessary privileges.

```
CREATE USER 'webdb_user'@'%' IDENTIFIED BY '*****';
GRANT ALL ON migrate.* TO 'webdb_user'@'%' with grant option;
grant REPLICATION SLAVE ON *.* TO webdb_user;
grant REPLICATION CLIENT ON *.* TO webdb_user;
```

In this walkthrough, we created a database called *migration* and few sample tables, along with stored procedures, triggers, functions, and so on. The following query provides the list of tables in *migration* database:

```
MariaDB [(none)]> use migration

Database changed
MariaDB [migration]> show tables;
+-----+
| Tables_in_migration |
+-----+
| animal_count        |
| animals              |
| contacts             |
| seat_type            |
| sport_location       |
| sport_team           |
| sport_type           |
+-----+
7 rows in set (0.000 sec)
```

The following query returns a list of secondary indexes.

```
MariaDB [migration]> SELECT DISTINCT TABLE_NAME, INDEX_NAME, NON_UNIQUE
-> FROM INFORMATION_SCHEMA.STATISTICS
-> WHERE TABLE_SCHEMA = 'migration' and INDEX_NAME <> 'PRIMARY';
+-----+-----+-----+
| TABLE_NAME | INDEX_NAME | NON_UNIQUE |
+-----+-----+-----+
| sport_location | city_id_sport_loc | 1 |
| sport_team    | sport_team_u      | 0 |
| sport_team    | home_field_fk     | 1 |
```



```
+-----+-----+-----+
3 rows in set (0.000 sec)
```

The following query returns a list of triggers.

```
MariaDB [migration]> select TRIGGER_SCHEMA,TRIGGER_NAME
->          from information_schema.triggers
->          where TRIGGER_SCHEMA='migration';
+-----+-----+
| TRIGGER_SCHEMA | TRIGGER_NAME          |
+-----+-----+
| migration      | increment_animal      |
| migration      | contacts_after_update |
+-----+-----+
2 rows in set (0.001 sec)
```

The following query returns a list of procedures and functions.

```
MariaDB [(none)]> select routine_schema as database_name,
->          routine_name,
->          routine_type as type,
->          data_type as return_type
->          from information_schema.routines
->          where routine_schema not in ('sys', 'information_schema',
->          'mysql', 'performance_schema');
+-----+-----+-----+-----+
| database_name | routine_name          | type          | return_type |
+-----+-----+-----+-----+
| migration     | CalcValue             | FUNCTION      | int          |
| migration     | loadMLBPlayers        | PROCEDURE     |              |
| migration     | loadNFLPlayers        | PROCEDURE     |              |
+-----+-----+-----+-----+
3 rows in set (0.000 sec)
```

After all the data is loaded, use `mysqldump` to back up the database metadata. The `mysqldump` utility to dump one or more databases for backup or transfer to another database server. The dump typically contains SQL statements to create the table, populate it, or both. You can also use `mysqldump` to generate files in comma-separated value (CSV), other delimited text, or XML format.

Use the following command exports tables and index definitions:

```
$ mysqldump --no-data --no-create-db --single_transaction -u root -p migration --skip-triggers > mysql_tables_indexes.sql
```

Use following command to exports routines (stored procedures, functions, and triggers) into the `routines.sql` file:

```
$ mysqldump -u root --routines --no-create-info --no-data --no-create-db --skip-opt -p migration > routines.sql
```

The `mysqldump` utility doesn't provide the option to remove a `DEFINER` statement. Some MySQL clients provide the option to ignore the definer when creating a logical backup, but this isn't the default behavior. Use the following command in a UNIX or Linux environment to remove the `DEFINER` from `routines.sql`:

```
$ sed -i -e 's/DEFINER=`root`@`localhost`/DEFINER=`master`@`%`/g' routines.sql
```

We now have a backup of MariaDB, in two `.sql` files (`mysql_tables_indexes.sql` and `routines.sql`). We will use these files to load the table definition into an Aurora MySQL database.

After backups are completed into two `.sql` files (`mysql_tables_indexes.sql`, `routines.sql`), use these files to load the table definition into the Aurora MySQL database.

## Set up Aurora MySQL as a target database

To provision Aurora MySQL as a target database, download the [AuroraMysql\\_CF.yaml template](#). This template creates an Aurora MySQL database with required parameters.

1. On the [AWS Management Console](#), under **Services**, choose **CloudFormation**.
2. Choose **Create stack**, and then choose **With new resources (standard)**.
3. For **Specify template**, choose **Upload a template file**.
4. Select **Choose file**.
5. Choose the `AuroraMySQL.yaml` file.
6. Choose **Next**.
7. On the **Specify stack details** page, edit the predefined values as needed, and then choose **Next**:

- **Stack name** — Enter a name for the stack.
  - **CIDR** — Enter the CIDR IP range to access the instance.
  - **DBBackupRetentionPeriod** — The number of days for backup retention.
  - **DBInstanceClass** — Enter the instance type of the database server.
  - **DBMasterPassword** — Enter the master password for the DB instance.
  - **DBMasterUsername** — Enter the master user name for the DB instance.
  - **DBName** — Enter the name of the database.
  - **DBSubnetGroup** — Enter the DB subnet group.
  - **Engine** — Enter the Aurora engine version; the default is `5.7.mysql-aurora.2.03.4`.
  - **PreferredBackupWindow** — Enter the daily time range in UTC during which you want to create automated backups.
  - **PreferredMaintenanceWindow** — Enter the weekly time range in UTC during which system maintenance can occur.
  - **VPCID** — Enter the ID for the VPC to launch your DB instance in.
8. On the **Configure stack options** page, for **Tags**, specify any optional tags, and then choose **Next**.
9. On the **Review** page, choose **Next**.
10. Choose **Create stack**.

After the Aurora MySQL database is created, log in to the Aurora MySQL instance:

```
$ mysql -h mysqltrg-instance-1.xxxxxxxxxx.us-east-1.rds.amazonaws.com -u master -p
migration -P 3306
MySQL [(none)]> show databases;
+-----+
| Database                |
+-----+
| information_schema      |
| awsdms_control          |
| mysql                   |
| performance_schema      |
| source                   |
| tmp                     |
| webdb                   |
+-----+
7 rows in set (0.001 sec)
```

```
MySQL [(none)]> create database migration;
Query OK, 1 row affected (0.016 sec)

MySQL [(none)]> use migration;
Database changed

MySQL [migration]> show tables;
Empty set (0.001 sec)
```

Use `mysql_tables_indexes.sql` to create table and index structures in Aurora MySQL.

```
$ mysql -h mysqltrg-instance-1.xxxxxxxxxx.us-east-1.rds.amazonaws.com -u master -p
migration -P 3306 < mysql_tables_indexes.sql
Enter password:
$
```

After the tables and indexes are successfully created, the next step is to set up and use AWS DMS.

## Set up an AWS DMS replication instance

To provision an AWS DMS replication instance, download the [DMS\\_CF.yaml template](#).

1. On the [AWS Management Console](#), under **Services**, choose **CloudFormation**.
2. Choose **Create stack**.
3. For **Specify template**, choose **Upload a template file**.
4. Select **Choose File**.
5. Choose the `DMS_CF.yaml` file.
6. Choose **Next**.
7. On the **Specify Stack Details** page, edit the predefined values as needed, and then choose **Next**:
  - **Stack name** — Enter a name for the stack.
  - **AllocatedStorageSize** — Enter the storage size in GB. The default is 200 GB.
  - **DMSReplicationSubnetGroup** — Enter the subnet group for DMS replication.
  - **DMSSecurityGroup** — Enter the security group for DMS replication.
  - **InstanceType** — Enter the instance type.
  - **SourceDBPort** — Enter the source database port.
  - **SourceDatabaseName** — Enter the source database name.

- **SourceServerName** — Enter the IP address of the source database server.
  - **SourceUsername** — Enter the source database user name.
  - **SourcePassword** — Enter the source database password.
  - **TargetDBPort** — Enter the target database port.
  - **TargetDatabaseName** — Enter the target database name.
  - **TargetServerName** — Enter the IP address of the target database server.
  - **TargetUsername** — Enter the target database user name.
  - **TargetPassword** — Enter the target database password.
8. On the **Configure stack options** page, for **Tags**, specify any optional tags, and then choose **Next**.
  9. On the **Review** page, choose **I acknowledge that AWS CloudFormation might create IAM resources**.
  10. Choose **Create Stack**.

This AWS CloudFormation template creates a replication instance named `mariadb-mysql`. This replication instance has a source endpoint named `maria-on-prem` and a target endpoint named `mysqltrg-rds`. This target endpoint has extra connection attributes to disable foreign key constraint checks during the AWS DMS replication, as shown following.

```
ExtraConnectionAttributes : "initstmt=SET FOREIGN_KEY_CHECKS=0;parallelLoadThreads=1"
```

## Test the endpoints for MariaDB database migration

1. On the navigation pane, choose **Endpoints**.
2. Choose the source endpoint name (`maria-on-prem`) and do the following:
  - a. Choose **Test connections**.
  - b. Choose the replication instance to test (`mariadb-mysql`).
  - c. Choose **Run Test** and wait for the status to be **successful**.
3. On the navigation pane, choose **Endpoints**.
4. Choose the target endpoint name (`mysqltrg-rds`) and do the following:
  - a. Choose **Test Connections**.
  - b. Choose the replication instance to test (`mariadb-mysql`).
  - c. Choose **Run Test** and wait for the status to be **successful**.

**Note**

If **Run Test** returns a status other than **successful**, the reason for the failure is displayed. Make sure that you resolve the issue before proceeding further.

## Create a migration task for a MariaDB database

We've now verified that the replication instance can connect to both the source and target endpoints. The next step is to create a database migration task.

1. On the navigation pane, choose **Database Migration Tasks**.
2. Choose **Create Task**. Provide the specified values for the following, and then choose **Next**:
  - **Task identifier** — `maria-mysql`
  - **Replication instance** — Choose the replication instance, `mariadb-mysql`.
  - **Source database endpoint** — Choose the source database, `maria-on-prem`.
  - **Target database endpoint** — Choose the target database, `mysqltrg-rds`.
  - **Migration Type** — Choose **Migrate existing data and replicate ongoing changes** for CDC, or **Migrate existing data** for full load.
3. For **Task settings**, choose the following settings:
  - **Target table preparation mode** — Do nothing
  - **Stop task after full load completes** — Don't stop
  - **Include LOB columns in replication** — Limited LOB mode
  - **Maximum LOB size (KB)** — 32
  - **Enable validation**
  - **Enable CloudWatch logs**
4. For **Table mappings**, choose the following settings:
  - **Schema** — Choose **migration** (assuming the schema and database to be migrated appear correctly).
  - **Table name** — Enter the table name, or % to specify all the tables in the database.
  - **Action** — Enter **Include** to include specific tables, or **Exclude** to exclude specific tables.

### 5. Choose **Create Task**.

Your new AWS DMS migration task reads the data from the tables in the MariaDB source and migrates your data to the Aurora MySQL target.

You can use an AWS DMS full-load-only migration task to migrate views or a combination of tables and views. For more information, see [Specifying table selection and transformations rules](#) in the *DMS User Guide*.

## Validate the MariaDB database migration

AWS DMS performs data validation to confirm that your data successfully migrated the source database to the target. You can check the **Table statistics** page to determine the DML changes that occurred after the AWS DMS task started. During data validation, AWS DMS compares each row in the source with its corresponding row at the target, and verifies that those rows contain the same data. To accomplish this, AWS DMS issues the appropriate queries to retrieve the data.

After your data is loaded successfully, you can select your task on the AWS DMS page and choose **Table statistics** to show statistics about your migration. The following screen shot shows the **Table statistics** page and its relevant entries.

The following screenshot shows the table statics page and its relevant entries.

**Table statistics (8)**

Find schema

| <input type="checkbox"/> | Schema name ▾ | Table ▾        | Load state ▾    | Inserts ▾ | Deletes ▾ | Updates ▾ | DDLs ▾ | Full load rows ▾ | Total ▾ | Validation state ▾ |
|--------------------------|---------------|----------------|-----------------|-----------|-----------|-----------|--------|------------------|---------|--------------------|
| <input type="checkbox"/> | migration     | contacts       | Table completed | 0         | 0         | 0         | 0      | 0                | 0       | Validated          |
| <input type="checkbox"/> | migration     | seat_type      | Table completed | 0         | 0         | 0         | 0      | 6                | 6       | Validated          |
| <input type="checkbox"/> | migration     | player         | Table completed | 0         | 0         | 0         | 0      | 0                | 0       | Validated          |
| <input type="checkbox"/> | migration     | sport_type     | Table completed | 0         | 0         | 0         | 0      | 2                | 2       | Validated          |
| <input type="checkbox"/> | migration     | sport_team     | Table completed | 0         | 0         | 0         | 0      | 32               | 32      | Validated          |
| <input type="checkbox"/> | migration     | sport_location | Table completed | 0         | 0         | 0         | 0      | 30               | 30      | Validated          |
| <input type="checkbox"/> | migration     | animals        | Table completed | 0         | 0         | 0         | 0      | 0                | 0       | Validated          |
| <input type="checkbox"/> | migration     | animal_count   | Table completed | 0         | 0         | 0         | 0      | 1                | 1       | No primary Key     |

AWS DMS can validate the data between source and target engines. The **Validation state** column helps us to validate the data migration. This ensures that your data was migrated accurately from the source to the target.

## Cut over for the migration from a MariaDB database

After the data validation is complete and any problems resolved, you can load the database triggers, functions, and procedures.

To do this, use the `routines.sql` file generated from MariaDB to create the necessary routines in Aurora MySQL. The following statement loads all procedures, functions, and triggers into the Aurora MySQL database.

```
$ mysql -h mysqltrg-instance-1.xxxxxxxxxx.us-east-1.rds.amazonaws.com -u master -p
migration -P 3306 < routines.sql
```

After the routines are loaded, connect to the Aurora MySQL database to validate as shown following.

```
$ mysql -h mysqltrg-instance-1.xxxxxxxxxx.us-east-1.rds.amazonaws.com -u master -p
migration -P 3306
```

Enter password:

Reading table information for completion of table and column names

You can turn off this feature to get a quicker startup with -A

Welcome to the MariaDB monitor. Commands end with ; or \g.

Your MySQL connection id is 957

Server version: 5.6.10 MySQL Community Server (GPL)

Copyright (c) 2000, 2018, Oracle, MariaDB Corporation Ab and others.

Enter 'help;' or '\h' for help. Enter '\c' to clear the current input statement.

```
MySQL [migration]> select routine_schema as database_name,
->          routine_name,
->          routine_type as type,
->          data_type as return_type
->          from information_schema.routines
->          where routine_schema not in ('sys', 'information_schema',
->                                     'mysql', 'performance_schema');
```

```
+-----+-----+-----+-----+
| database_name | routine_name | type      | return_type |
+-----+-----+-----+-----+
| migration    | CalcValue    | FUNCTION  | int          |
| migration    | loadMLBPlayers | PROCEDURE |              |
| migration    | loadNFLPlayers | PROCEDURE |              |
+-----+-----+-----+-----+
3 rows in set (0.002 sec)
```



```
MySQL [migration]> select TRIGGER_SCHEMA, TRIGGER_NAME from information_schema.triggers
where TRIGGER_SCHEMA='migration';
+-----+-----+
| TRIGGER_SCHEMA | TRIGGER_NAME |
+-----+-----+
| migration      | increment_animal |
| migration      | contacts_after_update |
+-----+-----+
2 rows in set (0.009 sec)
```

The preceding output shows that all the procedures, triggers, and functions are loaded successfully to the Aurora MySQL database.

## Migrating from MongoDB to Amazon DocumentDB

Use the following tutorial to guide you through the process of migrating from MongoDB to Amazon DocumentDB (with MongoDB compatibility). In this tutorial, you do the following:

- Install MongoDB on an Amazon EC2 instance.
- Populate MongoDB with sample data.
- Create an AWS DMS replication instance, a source endpoint (for MongoDB), and a target endpoint (for Amazon DocumentDB).
- Run an AWS DMS task to migrate the data from the source endpoint to the target endpoint.

### Important

Before you begin, make sure to launch an Amazon DocumentDB cluster in your default virtual private cloud (VPC). For more information, see [Getting started](#) in the *Amazon DocumentDB Developer Guide*.

To estimate what it will cost to run this walkthrough on AWS, you can use the AWS Pricing Calculator. For more information, see <https://calculator.aws/>.

### Topics

- [Launch an Amazon EC2 instance for MongoDB migration](#)
- [Install and configure MongoDB community edition](#)

- [Create an AWS DMS replication instance for MongoDB migration](#)
- [Create source and target endpoints for MongoDB migration](#)
- [Create and run a MongoDB migration task](#)

## Launch an Amazon EC2 instance for MongoDB migration

For this tutorial, you launch an Amazon EC2 instance into your default VPC.

1. Open the Amazon EC2 console at <https://console.aws.amazon.com/ec2/>.
2. Choose **Launch Instance**, and do the following:
  - a. On the **Choose an Amazon Machine Image (AMI)** page, at the top of the list of AMIs, go to **Amazon Linux AMI** and choose **Select**.
  - b. On the **Choose an Instance Type** page, at the top of the list of instance types, choose **t2.micro**. Then choose **Next: Configure Instance Details**.
  - c. On the **Configure Instance Details** page, for **Network**, choose your default VPC. Then choose **Next: Add Storage**.
  - d. On the **Add Storage** page, skip this step by choosing **Next: Add Tags**.
  - e. On the **Add Tags** page, skip this step by choosing **Next: Configure Security Group**.
  - f. On the **Configure Security Group** page, do the following:
    - i. Choose **Select an existing security group**.
    - ii. In the list of security groups, choose **default**. Doing this chooses the default security group for your VPC. By default, the security group accepts inbound Secure Shell (SSH) connections on TCP port 22. If this isn't the case for your VPC, add this rule; for more information, see [What is Amazon VPC?](#) in the *Amazon VPC User Guide*.
    - iii. Choose **Next: Review and Launch**.
  - g. Review the information, and choose **Launch**.
3. In the **Select an existing key pair or create a new key pair** window, do one of the following:
  - If you don't have an Amazon EC2 key pair, choose **Create a new key pair** and follow the instructions. You are asked to download a private key file (.pem file). You need this file later when you log in to your Amazon EC2 instance.
  - If you already have an Amazon EC2 key pair, for **Select a key pair** choose your key pair from the list. You must already have the private key file (.pem file) available in order to log in to your Amazon EC2 instance.

#### 4. After you configure your key pair, choose **Launch Instances**.

In the console navigation pane, choose **EC2 Dashboard**, and then choose the instance that you launched. In the lower pane, on the **Description** tab, find the **Public DNS** location for your instance, for example: `ec2-11-22-33-44.us-west-2.compute.amazonaws.com`.

It takes a few minutes for your Amazon EC2 instance to become available.

#### 5. Use the `ssh` command to log in to your Amazon EC2 instance, as in the following example.

```
chmod 400 my-keypair.pem
ssh -i my-keypair.pem ec2-user@public-dns-name
```

Specify your private key file (.pem file) and the public DNS name of your EC2 instance. The login ID is `ec2-user`. No password is required.

For further details about connecting to your EC instance, see [Connecting to your Linux instance using SSH](#) in the *Amazon EC2 User Guide for Linux Instances*.

## Install and configure MongoDB community edition

Perform these steps on the Amazon EC2 instance that you launched in [Launch an Amazon EC2 instance](#).

1. Go to [Install MongoDB community edition on Amazon Linux](#) in the MongoDB documentation and follow the instructions there.
2. By default, the MongoDB server (`mongod`) only allows loopback connections from IP address 127.0.0.1 (localhost). To allow connections from elsewhere in your Amazon VPC, do the following:
  - a. Edit the `/etc/mongod.conf` file and look for the following lines.

```
# network interfaces
net:
  port: 27017
  bindIp: 127.0.0.1 # Enter 0.0.0.0,:: to bind to all IPv4 and IPv6 addresses or,
                  alternatively, use the net.bindIpAll setting.
```

- b. Modify the `bindIp` line so that it looks like the following.

```
bindIp: public-dns-name
```

- c. Replace *public-dns-name* with the actual public DNS name for your instance, for example `ec2-11-22-33-44.us-west-2.compute.amazonaws.com`.
- d. Save the `/etc/mongod.conf` file, and then restart `mongod`.

```
sudo service mongod restart
```

### 3. Populate your MongoDB instance with data by doing the following:

- a. Use the `wget` command to download a JSON file containing sample data.

```
wget http://media.mongodb.org/zip.js
```

- b. Use the `mongoimport` command to import the data into a new database (`zip`).

```
mongoimport --host public-dns-name:27017 --db zip --file zip.js
```

- c. After the import completes, use the mongo shell to connect to MongoDB and verify that the data was loaded successfully.

```
mongo --host public-dns-name:27017
```

- d. Replace *public-dns-name* with the actual public DNS name for your instance.
- e. At the mongo shell prompt, enter the following commands.

```
use zip
db.zip.count()

db.zip.aggregate( [
  { $group: { _id: { state: "$state", city: "$city" }, pop: { $sum: "$pop" } } },
  { $group: { _id: "$_id.state", avgCityPop: { $avg: "$pop" } } }
] )
```

The output should display the following:

- The name of the database (`zip`)
- The number of documents in the `zip` collection (29353)
- The average population for cities in each state

- f. Exit from the mongo shell and return to the command prompt by using the following command.

```
exit
```

## Create an AWS DMS replication instance for MongoDB migration

To perform replication in AWS DMS, you need a replication instance.

1. Open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. In the navigation pane, choose **Replication instances**.
3. Choose **Create replication instance** and enter the following information:
  - For **Name**, enter `mongodb2docdb`.
  - For **Description**, enter `MongoDB to Amazon DocumentDB replication instance`.
  - For **Instance class**, keep the default value.
  - For **Engine version**, keep the default value.
  - For **VPC**, choose your default VPC.
  - For **Multi-AZ**, choose **No**.
  - For **Publicly accessible**, enable this option.

When the settings are as you want them, choose **Create replication instance**.

### Note

You can begin using your replication instance when its status becomes **available**. This can take several minutes.

## Create source and target endpoints for MongoDB migration

The source endpoint is the endpoint for your MongoDB installation running on your Amazon EC2 instance.

1. Open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. In the navigation pane, choose **Endpoints**.

### 3. Choose **Create endpoint** and enter the following information:

- For **Endpoint type**, choose **Source**.
- For **Endpoint identifier**, enter a name that's easy to remember, for example `mongodb-source`.
- For **Source engine**, choose **mongodb**.
- For **Server name**, enter the public DNS name of your Amazon EC2 instance, for example `ec2-11-22-33-44.us-west-2.compute.amazonaws.com`.
- For **Port**, enter `27017`.
- For **SSL mode**, choose **none**.
- For **Authentication mode**, choose **none**.
- For **Database name**, enter `zips-db`.
- For **Authentication mechanism**, choose **default**.
- For **Metadata mode**, choose **document**.

When the settings are as you want them, choose **Create endpoint**.

Next, you create a target endpoint. This endpoint is for your Amazon DocumentDB cluster, which should already be running. For more information about launching your Amazon DocumentDB cluster, see [Getting started](#) in the *Amazon DocumentDB Developer Guide*.

#### **Important**

Before you proceed, do the following:

- Create indexes on your Amazon DocumentDB cluster before you begin migration because it can reduce the overall time and increase the speed of the migration. To extract indexes from a running MongoDB instance, you can use the [Amazon DocumentDB Index Tool](#).
- Get the master user name and password for your Amazon DocumentDB cluster.
- Get the DNS name and port number of your Amazon DocumentDB cluster, so that AWS DMS can connect to it. To determine this information, use the following AWS CLI command, replacing `cluster-id` with the name of your Amazon DocumentDB cluster.

```
aws docdb describe-db-clusters \  
  --db-cluster-identifier cluster-id \  
  --output text
```

```
--query "DBClusters[*].[Endpoint,Port]"
```

- Download a certificate bundle that Amazon DocumentDB can use to verify SSL connections. To do this, enter the following command. Here, *aws-api-domain* completes the Amazon S3 domain in your AWS Region required to access the specified S3 bucket and the `rds-combined-ca-bundle.pem` file that it provides.

```
wget https://s3.amazonaws.com/rds-downloads/rds-combined-ca-bundle.pem
```

To create a target endpoint, do the following:

1. In the navigation pane, choose **Endpoints**.
2. Choose **Create endpoint** and enter the following information:
  - For **Endpoint type**, choose **Target**.
  - For **Endpoint identifier**, enter a name that's easy to remember, for example `docdb-target`.
  - For **Target engine**, choose **docdb**.
  - For **Server name**, enter the DNS name of your Amazon DocumentDB cluster.
  - For **Port**, enter the port number of your Amazon DocumentDB cluster.
  - For **SSL mode**, choose **verify-full**.
  - For **CA certificate**, do one of the following to attach the SSL certificate to your endpoint:
    - If available, choose the existing **rds-combined-ca-bundle** certificate from the **Choose a certificate** drop down.
    - Choose **Add new CA certificate**. Then, for **Certificate identifier**, enter `rds-combined-ca-bundle`. For **Import certificate file**, choose **Choose file** and navigate to the `rds-combined-ca-bundle.pem` file that you previously downloaded. Select and open the file. Choose **Import certificate**, then choose **rds-combined-ca-bundle** from the **Choose a certificate** drop down.
  - For **User name**, enter the master user name of your Amazon DocumentDB cluster.
  - For **Password**, enter the master password of your Amazon DocumentDB cluster.
  - For **Database name**, enter `zips-db`.

When the settings are as you want them, choose **Create endpoint**.

Now that you've created the source and target endpoints, test them to ensure that they work correctly. Also, to ensure that AWS DMS can access the database objects at each endpoint, refresh the endpoints' schemas.

To test an endpoint, do the following:

1. In the navigation pane, choose **Endpoints**.
2. Choose the source endpoint (mongodb-source), and then choose **Test connection**.
3. Choose your replication instance (mongodb2docdb), and then choose **Run test**. It takes a few minutes for the test to complete, and for the **Status** to change to **successful**.

If the **Status** changes to **failed** instead, review the failure message. Correct any errors that might be present, and test the endpoint again.

 **Note**

Repeat this procedure for the target endpoint (docdb-target).

To refresh schemas, do the following:

1. In the navigation pane, choose **Endpoints**.
2. Choose the source endpoint (mongodb-source), and then choose **Refresh schemas**.
3. Choose your replication instance (mongodb2docdb), and then choose **Refresh schemas**.

 **Note**

Repeat this procedure for the target endpoint (docdb-target).

## Create and run a MongoDB migration task

You are now ready to launch an AWS DMS migration task, to migrate the zips data from MongoDB to Amazon DocumentDB.

1. Open the AWS DMS console at <https://console.aws.amazon.com/dms/v2/>.
2. In the navigation pane, choose **Database migration tasks**.



### 3. Choose **Create task** and enter the following information:

- For **Task configuration**, choose the following settings:
  - **Task identifier** — enter a name that's easy to remember, for example `my-dms-task`.
  - **Replication instance** — choose the replication instance that you created in [Create a replication instance](#).
  - **Source database endpoint** — choose the source endpoint that you created in [Create source and target endpoints](#).
  - **Target database endpoint** — choose the target endpoint that you created in [Create source and target endpoints](#).
  - **Migration type** — choose **Migrate existing data**.
- For **Task settings**, choose the following settings:
  - **Target table preparation mode** — Do nothing
  - **Include LOB columns in replication** — Limited LOB mode
  - **Maximum LOB size (KB)** — 32
  - **Enable validation**
  - **Enable CloudWatch logs**

#### Note

CloudWatch logs usage will be charged at standard rates. See [here](#) for more details.

- For **Advanced task settings**, keep all of the options at their default values.
- For **Premigration assessment**, keep the option at its default value.
- For **Start migration task** in **Migration task startup configuration**, choose **Automatically on create**.
- For **Tags**, keep all of the options at their default values.

When the settings are as you want them, choose **Create task**.

AWS DMS now begins migrating data from MongoDB to Amazon DocumentDB. The task status changes from **Starting** to **Running**. You can monitor the progress by choosing **Tasks** in the AWS DMS console. After several minutes, the status changes to **Load complete**.

**Note**

After the migration is complete, you can use the mongo shell to connect to your Amazon DocumentDB cluster and view the `zips` data. For more information, see [Access your Amazon DocumentDB cluster using the mongo shell](#) in the *Amazon DocumentDB Developer Guide*.